



US011275584B2

(12) **United States Patent**
Harthcock

(10) **Patent No.:** US 11,275,584 B2
(45) **Date of Patent:** Mar. 15, 2022

(54) **UNIVERSAL FLOATING-POINT INSTRUCTION SET ARCHITECTURE FOR COMPUTING DIRECTLY WITH DECIMAL CHARACTER SEQUENCES AND BINARY FORMATS IN ANY COMBINATION**

(71) Applicant: **Jerry D. Harthcock**, Boerne, TX (US)

(72) Inventor: **Jerry D. Harthcock**, Boerne, TX (US)

(*) Notice: Subject to any disclaimer, the term of this patent is extended or adjusted under 35 U.S.C. 154(b) by 98 days.

(21) Appl. No.: **16/943,077**

(22) Filed: **Jul. 30, 2020**

(65) **Prior Publication Data**

US 2021/0049011 A1 Feb. 18, 2021

Related U.S. Application Data

(60) Provisional application No. 62/886,570, filed on Aug. 14, 2019.

(51) **Int. Cl.**

G06F 9/30 (2018.01)

G06F 9/355 (2018.01)

G06F 9/54 (2006.01)

G06F 12/06 (2006.01)

G06F 9/32 (2018.01)

G06F 9/35 (2018.01)

G06F 9/38 (2018.01)

G06F 12/02 (2006.01)

(52) **U.S. Cl.**

CPC **G06F 9/30025** (2013.01); **G06F 9/30079** (2013.01); **G06F 9/30101** (2013.01); **G06F**

9/3557 (2013.01); **G06F 9/544** (2013.01);
G06F 12/063 (2013.01)

(58) **Field of Classification Search**

CPC G06F 9/30025; G06F 9/30079; G06F 9/30101; G06F 9/3557; G06F 9/544; G06F 12/063; G06F 9/30; G06F 9/321; G06F 9/325; G06F 9/35; G06F 9/38; G06F 9/3824; G06F 12/0292; G06F 2212/1024; G06F 2212/1048; G06F 12/0284

USPC 712/220

See application file for complete search history.

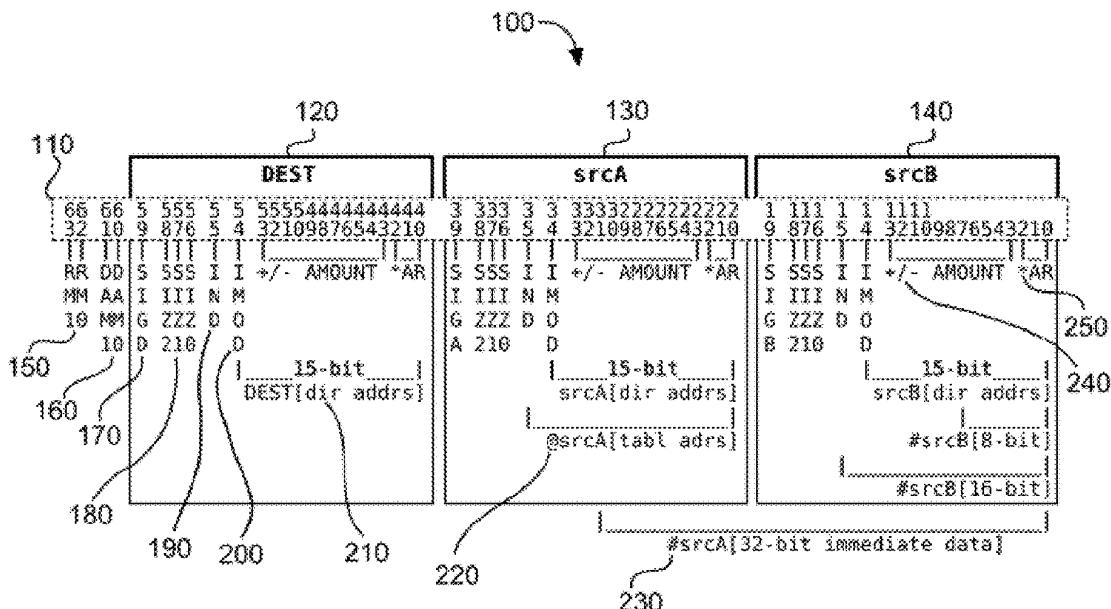
Primary Examiner — Shawn Doman

(74) *Attorney, Agent, or Firm — Steven W. Smith*

(57) **ABSTRACT**

A universal floating-point Instruction Set Architecture (ISA) implemented entirely in hardware. Using a single instruction, the universal floating-point ISA has the ability, in hardware, to compute directly with dual decimal character sequences up to IEEE 754-2008 “H=20” in length, without first having to explicitly perform a conversion-to-binary-format process in software before computing with these human-readable floating-point or integer representations. The ISA does not employ opcodes, but rather pushes and pulls “gobs” of data without the encumbering opcode fetch, decode, and execute bottleneck. Instead, the ISA employs stand-alone, memory-mapped operators, complete with their own pipeline that is completely decoupled from the processor’s primary push-pull pipeline. The ISA employs special three-port, 1024-bit wide SRAMs; a special dual asymmetric system stack; memory-mapped stand-alone hardware operators with private result buffers having simultaneously readable side-A and side-B read ports; and dual hardware H=20 convertFromDecimalCharacter conversion operators.

16 Claims, 89 Drawing Sheets



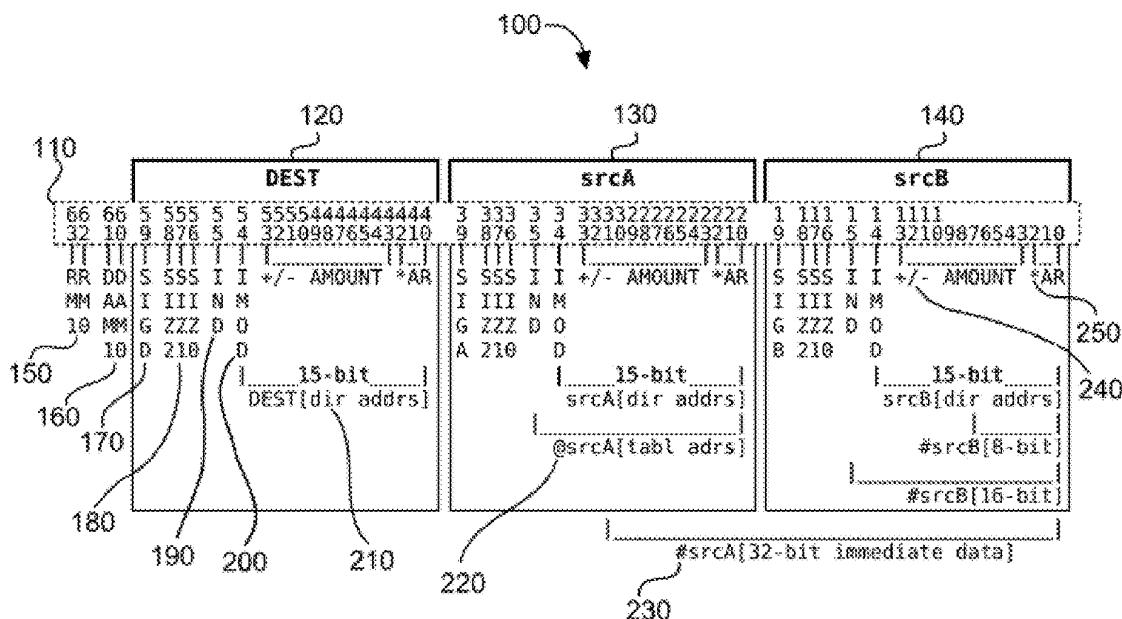


FIG. 1

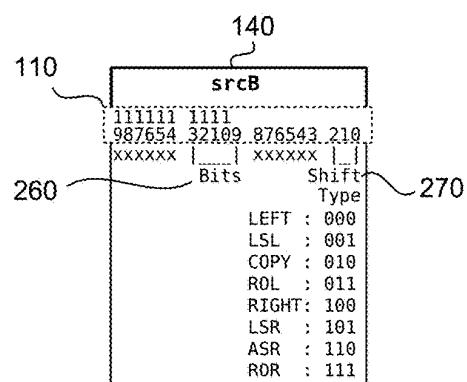


FIG. 2

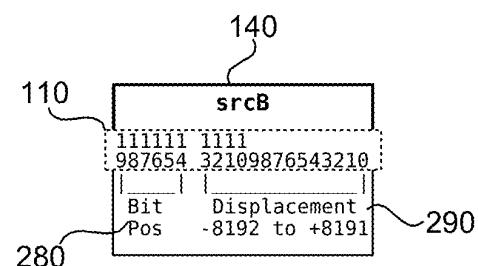


FIG. 3

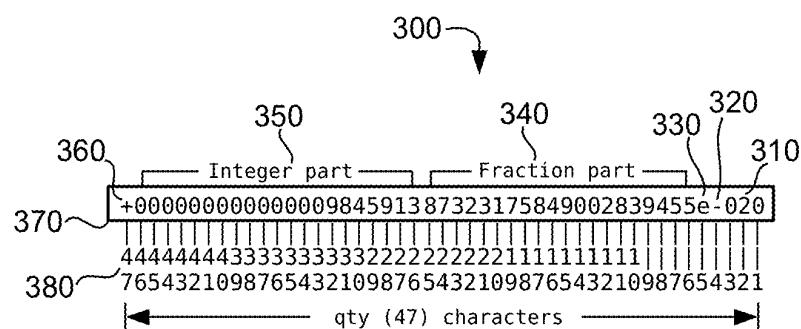
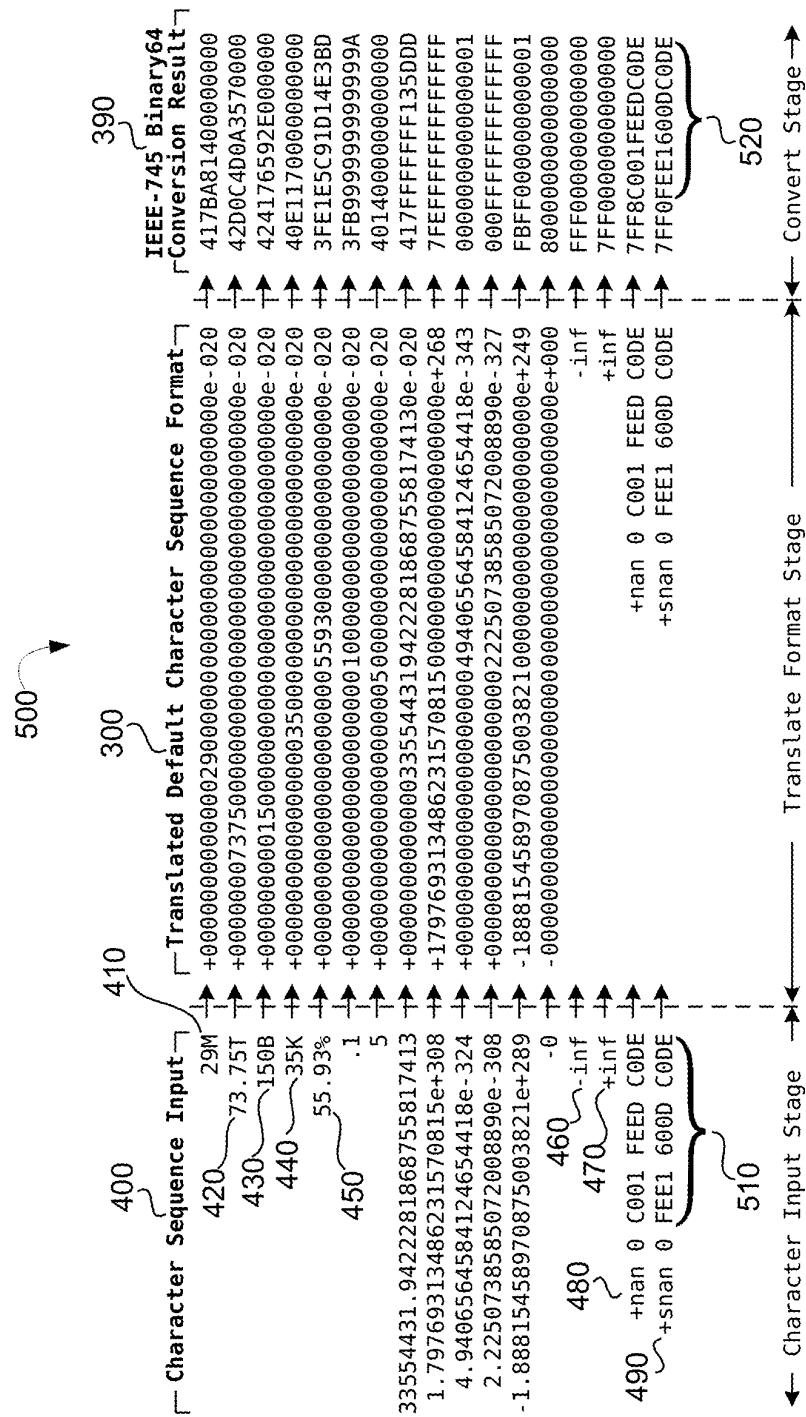


FIG. 4



5
FIG.

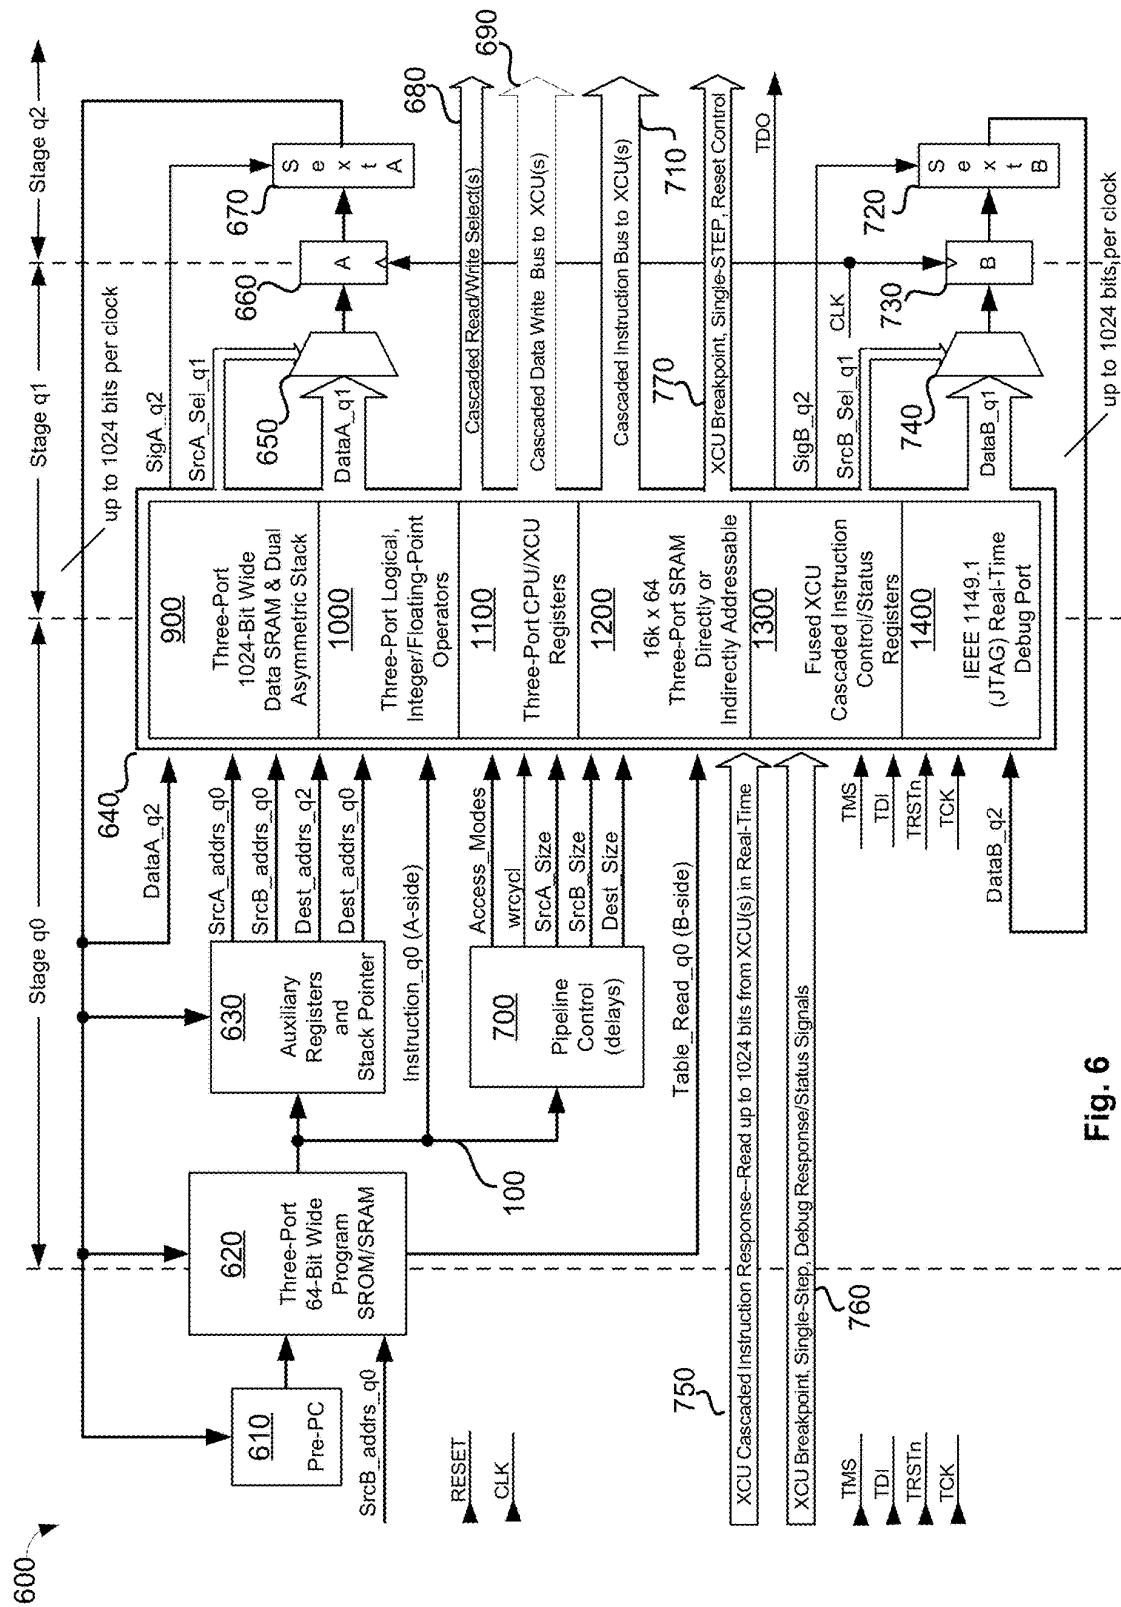


Fig. 6

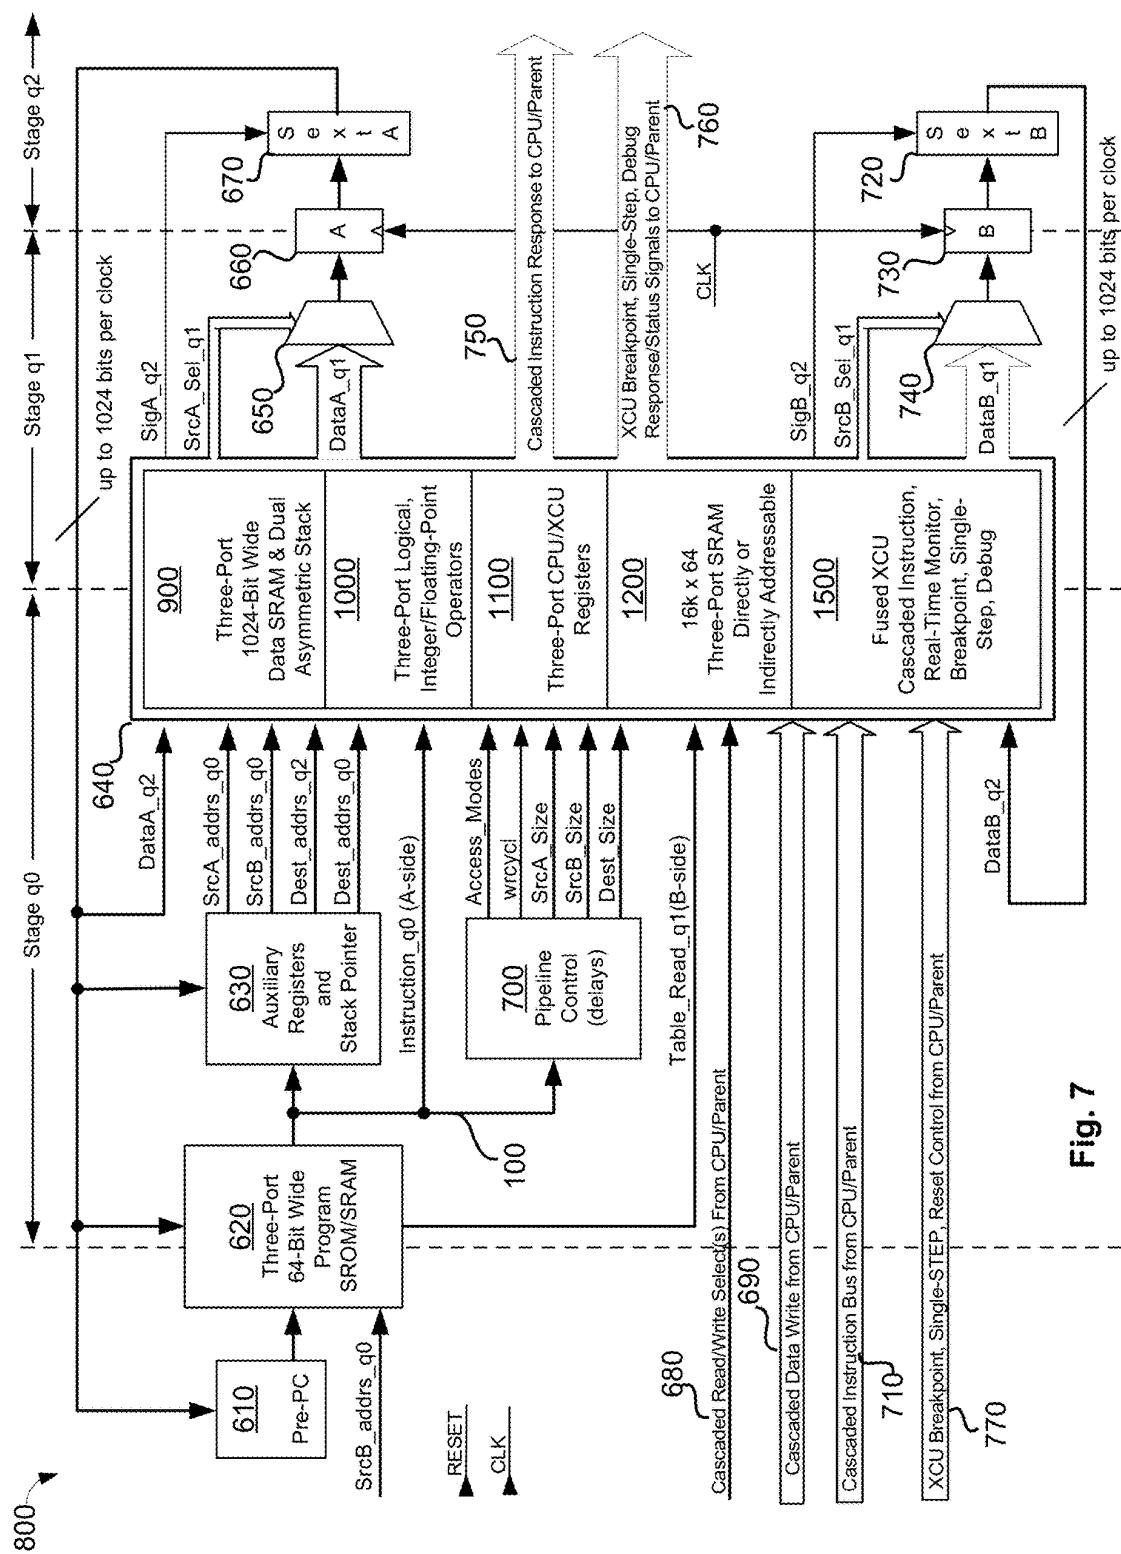
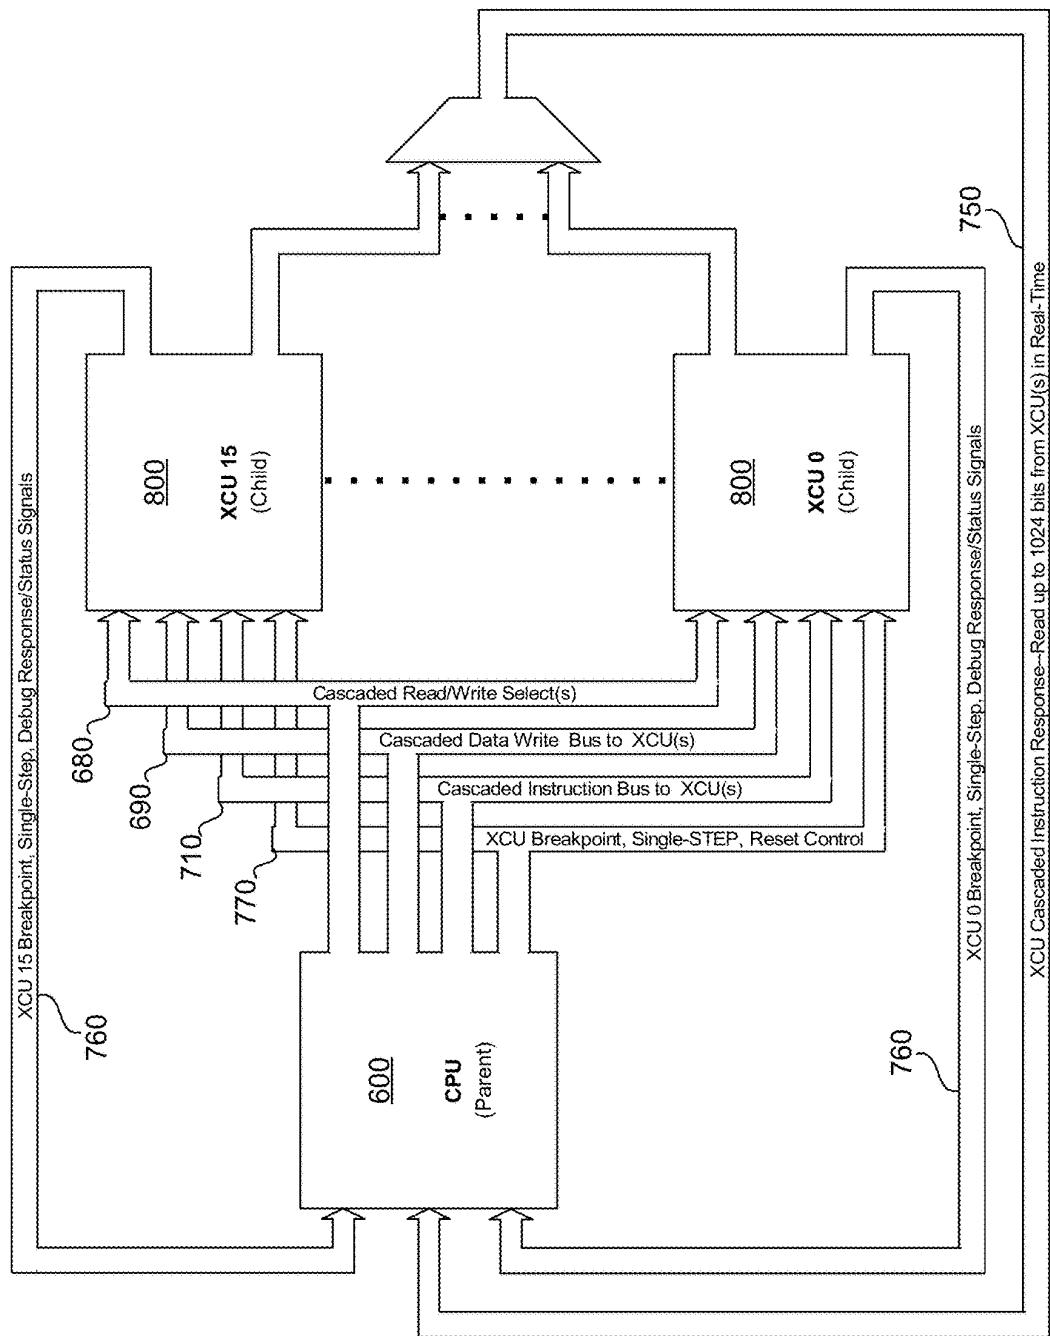


Fig. 7

**FIG. 8**

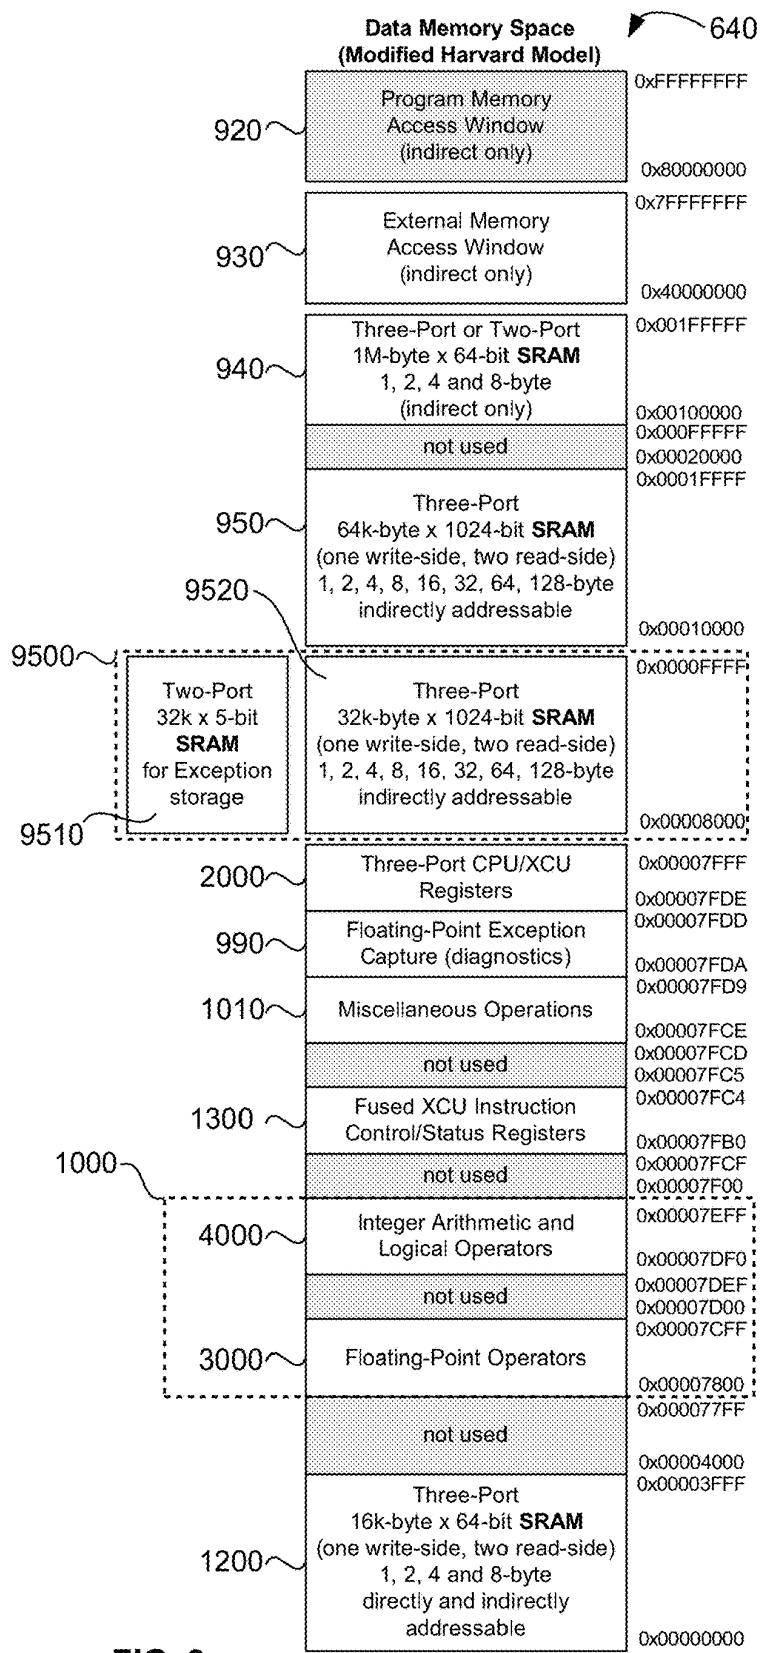


FIG. 9

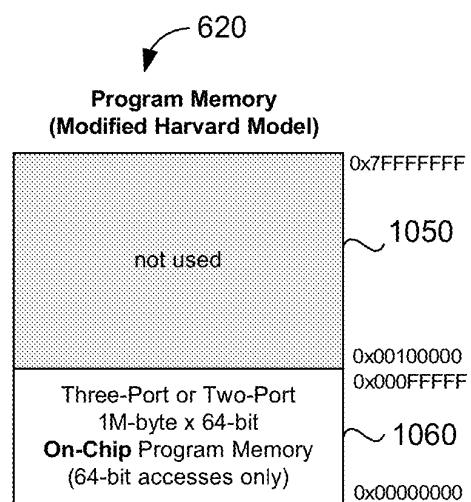


FIG. 10

2000

Direct Address	Programming Model Registers	Assembler Label	Access Modes	Reset State
0x7FFF	not used	n/a	n/a	n/a
0x7FFE	Stack Pointer	SP	R/W	0x0000FF80
0x7FFD	Auxiliary Register 6	AR6	R/W	0x00000000
0x7FFC	Auxiliary Register 5	AR5	R/W	0x00000000
0x7FFB	Auxiliary Register 4	AR4	R/W	0x00000000
0x7FFA	Auxiliary Register 3	AR3	R/W	0x00000000
0x7FF9	Auxiliary Register 2	AR2	R/W	0x00000000
0x7FF8	Auxiliary Register 1	AR1	R/W	0x00000000
0x7FF7	Auxiliary Register 0	AR0	R/W	0x00000000
0x7FF6	Branch Always Long	PCR	W	n/a
0x7FF5	Program Counter (PC)	PC	R/W	0x00000100
0x7FF4	Bit-Test and Branch If Set	PCS	W	n/a
0x7FF3	Bit-Test and Branch If Clear	PCC	W	n/a
0x7FF2	PC Copy	PC_COPY	R	0x00000100
0x7FF1	Status Register	STATUS	R/W	0x0000000040000191
0x7FF0	not used	n/a	n/a	n/a
0x7FEF	Repeat Counter	REPEAT	R/W	0x00000000
0x7FEE	Loop Counter 1	LPCNT1	R/W	0x00000000
0x7FED	Loop Counter 0	LPCNT0	R/W	0x00000000
0x7FEC	Timer	TIMER	R/W	0x00000000
0x7FEB	Real-Time Monitor Data	n/a	R	0x00000000
0x7FEA	not used	n/a	n/a	n/a
0x7FE9	not used	n/a	n/a	n/a
0x7FE8	Spare Vector 1	n/a	n/a	n/a
0x7FE7	Spare Vector 0	n/a	n/a	n/a
0x7FE6	NMI Vector	NMI_VECT	W	0x00000000
0x7FE5	IRQ Vector	IRQ_VECT	W	0x00000000
0x7FE4	Invalid Vector	INV_VECT	W	0x00000000
0x7FE3	Divide By Zero Vector	DIVx0_VECT	W	0x00000000
0x7FE2	Overflow Vector	OVFL_VECT	W	0x00000000
0x7FE1	Underflow Vector	UNFL_VECT	W	0x00000000
0x7FE0	Inexact Vector	INEXT_VECT	W	0x00000000
0x7FDF	XCU Control Register	XCUCR	R/W	0x00000000FFFF0000
0x7FDE	XCU Status Register	XCUSR	R	0xFFFF000000000000
0x7FDD	Exception Capture Register 3	CAPT_3	R	0x0000000000000000
0x7FDC	Exception Capture Register 2	CAPT_2	R	0x0000000000000000
0x7FDB	Exception Capture Register 1	CAPT_1	R	0x0000000000000000
0x7FDA	Exception Capture Register 0	CAPT_0	R	0x0000000000000000
0x7FC4	XCU Pre-emption Register	PREEMPT	W	0x0000
0x7FC3	XCU Single-Step Register	SSTEP	W	0x0000
0x7FC2	XCU Force_Break Register	FBREAK	W	0x0000
0x7FC1	XCU Force-Reset Register	FRESET	W	0x0000
0x7FC0	XCU Poke_All Register	POKEALL	W	0x0000
0x7FBX	XCU Monitor Request	MREQ	W	0x0000000000000000

FIG. 11



IEEE 754-2008 Mandated Computational Operators*	Assembler Label with 1st buffer# Last buffer#		1st Direct Address	Input Size (in bytes)	Result Size
roundToIntegralTiesToEven	RTOIE.0	RTOIE.15	0x7A10	2,4,8	2,4,8
roundToIntegralTowardZero	RTOIZ.0	RTOIZ.15	0x7A30	2,4,8	2,4,8
roundToIntegralTowardPositive	RTOIP.0	RTOIP.15	0x7A40	2,4,8	2,4,8
roundToIntegralTowardNegative	RTOIN.0	RTOIN.15	0x7A50	2,4,8	2,4,8
roundToIntegralTiesToAway	RTOIA.0	RTOIA.15	0x7A20	2,4,8	2,4,8
roundToIntegralExact	RTOIX.0	RTOIX.15	0x7A60	2,4,8	2,4,8
nextUp	nextUp.0	nextUp.15	0x7930	2,4,8	2,4,8
nextDown	nextDwn.0	nextDwn.15	0x7920	2,4,8	2,4,8
remainder	rem.0	rem.31	0x7880	2,4,8	2,4,8
minNum	minNum.0	minNum.15	0x78E0	2,4,8	2,4,8
maxNum	maxNum.0	maxNum.15	0x78F0	2,4,8	2,4,8
minNumMag	minNumMg.0	minNumMg.15	0x7900	2,4,8	2,4,8
maxNumMag	maxNumMg.0	maxNumMg.15	0x7910	2,4,8	2,4,8
scaleB	scaleB.0	scaleB.15	0x7950	2,4,8	2,4,8
logB	logB.0	logB.15	0x7940	2,4,8	2,4,8
addition	fadd.0	fadd.15	0x7800	2,4,8	2,4,8
subtraction	fsub.0	fsub.15	0x7810	2,4,8	2,4,8
multiplication	fmul.0	fmul.15	0x7820	2,4,8	2,4,8
division	fdiv.0	fdiv.15	0x7850	2,4,8	2,4,8
squareRoot	sqrt.0	sqrt.31	0x7860	2,4,8	2,4,8
fusedMultiplyAdd	fma.0	fma.15	0x7830	2,4,8	2,4,8
C Register/Accumulator	C.0	C.15	0x7840	2,4,8	16
convertFromInt	CFI.0	CFI.15	0x7960	1,2,4,8	2,4,8
convertToIntegerTiesToEven	CTOIE.0	CTOIE.15	0x7970	2,4,8	8
convertToIntegerTowardZero	CTOIZ.0	CTOIZ.15	0x7990	2,4,8	8
convertToIntegerTowardPositive	CTOIP.0	CTOIP.15	0x79A0	2,4,8	8
convertToIntegerTowardNegative	CTOIN.0	CTOIN.15	0x79B0	2,4,8	8
convertToIntegerTiesToAway	CTOIA.0	CTOIA.15	0x7980	2,4,8	8
convertToIntegerExactTiesToEven	CTOIXE.0	CTOIXE.15	0x79C0	2,4,8	8
convertToIntegerExactTowardZero	CTOIXZ.0	CTOIXZ.15	0x79E0	2,4,8	8
convertToIntegerExactTowardPositive	CTOIXP.0	CTOIXP.15	0x79F0	2,4,8	8
convertToIntegerExactTowardNegative	CTOIXN.0	CTOIXN.15	0x7A00	2,4,8	8
convertToIntegerExactTiesToAway	CTOIXA.0	CTOIXA.15	0x79D0	2,4,8	8
convertFormat	CONV.0	CONV.15	0x7A70	2,4,8	2,4,8
convertFromDecimalCharacter	cnvFDCS.0	cnvFDCS.31	0x7AC0	1 to 47	2,4,8
convertToDecimalCharacter	cnvTDCS.0	cnvTDCS.31	0x7AA0	2,4,8	47
convertFromHexCharacter	cnvFHCS.0	cnvFHCS.15	0x7A90	24	8
convertToHexCharacter	cnvTHCS.0	cnvTHCS.15	0x7A80	2,4,8	24
copy	copy.0	copy.15	0x78D0	2,4,8	2,4,8
negate	negate.0	negate.15	0x78C0	2,4,8	2,4,8
abs	abs.0	abs.15	0x78B0	2,4,8	2,4,8
copySign	copySign.0	copySign.15	0x78A0	2,4,8	2,4,8

FIG. 12A

3020

IEEE 754-2008 Mandated Comparison Predicates	Assembler Label	Direct Address	Input Size (in bytes)
compareSignalingEqual	cmpSE	0x7CFF	2,4,8
compareQuietEqual	cmpQE	0x7CFE	2,4,8
compareSignalingNotEqual	cmpSNE	0x7CFD	2,4,8
compareQuietNotEqual	cmpQNE	0x7CFC	2,4,8
compareSignalingGreater	cmpSG	0x7CFB	2,4,8
compareQuietGreater	cmpQG	0x7CFA	2,4,8
compareSignalingGreaterEqual	cmpSGE	0x7CF9	2,4,8
compareQuietGreaterEqual	cmpQGE	0x7CF8	2,4,8
compareSignalingLess	cmpSL	0x7CF7	2,4,8
compareQuietLess	cmpQL	0x7CF6	2,4,8
compareSignalingLessEqual	cmpSLE	0x7CF5	2,4,8
compareQuietLessEqual	cmpQLE	0x7CF4	2,4,8
compareSignalingNotGreater	cmpSNG	0x7CF3	2,4,8
compareQuietNotGreater	cmpQNG	0x7CF2	2,4,8
compareSignalingLessUnordered	cmpSLU	0x7CF1	2,4,8
compareQuietLessUnordered	cmpQLU	0x7CF0	2,4,8
compareSignalingNotLess	cmpSNL	0x7CEF	2,4,8
compareQuietNotLess	cmpQNL	0x7CEE	2,4,8
compareSignalingGreaterUnordered	cmpSGU	0x7CED	2,4,8
compareQuietGreaterUnordered	cmpQGU	0x7CEC	2,4,8
compareQuietUnordered	cmpQU	0x7CEB	2,4,8
compareQuietOrdered	cmpQO	0x7CEA	2,4,8

FIG. 12B

3030

IEEE 754-2008 Mandated (dual-operand) non-computational non-exceptional operators	Assembler Label	Direct Address	Input Size (in bytes)
totalOrder	tOrd	0x7CE9	2,4,8
totalOrderMag	tOrdM	0x7CE8	2,4,8

FIG. 12C

3040

IEEE 754-2008 Mandated (single-operand) non-computational non-exceptional operators	Assembler Label	Direct Address	Input Size (in bytes)
isCanonical	isCanonical	0x7CD8	2,4,8
isSignaling	isSignaling	0x7CD7	2,4,8
isNaN	isNaN	0x7CD6	2,4,8
isInfinite	isInfinite	0x7CD5	2,4,8
isSubnormal	isSubnormal	0x7CD4	2,4,8
isZero	isZero	0x7CD3	2,4,8
isFinite	isFinite	0x7CD2	2,4,8
isNormal	isNormal	0x7CD1	2,4,8
isSignMinus	isSignMinus	0x7CD0	2,4,8
Class	clas	0x7FCE	2,4,8

Fig. 12D

3050

Miscellaneous Mandated Operations	Assembler Label	Direct Address	Example Code
saveAllFlags	saveAllFlags	0x7FD9	_ _8:saveAllFlags = _8:STATUS
testSavedFlags	tstSavFlg	0x7FCE	_ _1:tstSavflg = _1:saveAllFlags
testFlags	tstFlg	0x7CE3	_ _1:tstFlg = (_1:saveAllFlags, _1:#3)
getBinaryRoundingDirection	sgtBinRnDir	0x7FD8	_ _1:sgtBinRnDir = _8:STATUS
setBinaryRoundingDirection	sgtBinRnDir	0x7FD8	_ _1:sgtBinRnDir = _1:getRndDir
saveModes	saveModes	0x7FD7	_ _1:saveModes = _8: STATUS
restoreModes	restModes	0x7FD7	_ _1:restModes = _1: saveModes
defaultModes	defModes	0x7FD7	_ _1:defModes = _1:#0
radix	radix	0x7FD5	_ _1:<dest> = _1:radix
is1985	is1985	0x0000	_ _1:<dest> = _1:is1985
is2008	is2008	0x7FD0	_ _1:<dest> = _1:is2008

FIG. 12E

3060

Implemented but not Mandated Computational Operators	Assembler Table with		1st Direct Address	Input Size (in bytes)	Result Size
	1st buffer#	Last buffer#			
SIND (integer in, SP degrees out)	SIND.0	SIND.15	0x7B50	1,2	2
COSD (integer in, SP degrees out)	COSD.0	COSD.15	0x7B40	1,2	2
TAND (integer in, SP degrees out)	TAND.0	TAND.15	0x7B30	1,2	2
COTD (integer in, SP degrees out)	COTD.0	COTD.15	0x7B20	1,2	2
LOG	LOG.0	LOG.15	0x7B10	2,4,8	2,4,8
EXP	EXP.0	EXP.15	0x7B00	2,4,8	2,4,8
POW	POW.0	POW.31	0x7AE0	2,4,8	2,4,8
POWR	POWR.0	POWR.31	0x7AE0	2,4,8	2,4,8
POWN	POWN.0	POWN.31	0x7AE0	2,4,8	2,4,8
universal FusedMultiplyAdd	UFMA.0	UFMA.31	0x7B60	1 to 47	2,4,8
universal C register/Accumulator	UC.0	UC.31	0x7B80	1 to 47	2,4,8

FIG. 12F

4000

Logical and Integer Arithmetic Operators	Assembler Table with		Direct Address	Input Size (in bytes)	Result Size
	1st buffer#	Last buffer#			
AND	and.0	and.15	0x7EF0	1,2,4,8	1,2,4,8
OR	or.0	or.15	0x7EE0	1,2,4,8	1,2,4,8
XOR	xor.0	xor.15	0x7ED0	1,2,4,8	1,2,4,8
ADD	add.0	add.15	0x7EC0	1,2,4,8	1,2,4,8
ADDC	addc.0	addc.15	0x7EB0	1,2,4,8	1,2,4,8
SUB	sub.0	sub.15	0x7EA0	1,2,4,8	1,2,4,8
SUBB	subb.0	subb.15	0x7E90	1,2,4,8	1,2,4,8
MUL	mul.0	mul.15	0x7E80	1,2,4,8	1,2,4,8
DIV	div.0	div.15	0x7E70	1,2,4,8	1,2,4,8
Shift	shft.0	shft.15	0x7E60	1,2,4,8	1,2,4,8
Maximum	max.0	max.15	0x7E50	1,2,4,8	1,2,4,8
Minimum	min.0	min.15	0x7E40	1,2,4,8	1,2,4,8
Bit Set	bset.0	bset.15	0x7E30	1,2,4,8	1,2,4,8
Bit Clear	bclr.0	bclr.15	0x7E20	1,2,4,8	1,2,4,8
Swap Endian-ness	endi.0	endi.15	0x7E10	1,2,4,8	1,2,4,8
Convert from Hex to ASCII	cnvFHTA.0	cnvFHTA.15	0x7E00	1,2,4	1,2,4,8
Convert to Hex from ASCII	cnvTHFA.0	cnvTHFA.15	0x7DF0	1,2,4,8	1,2,4
Compare	cmp	n/a	0x7FCF	1,2,4,8	1,2,4,8

FIG. 13

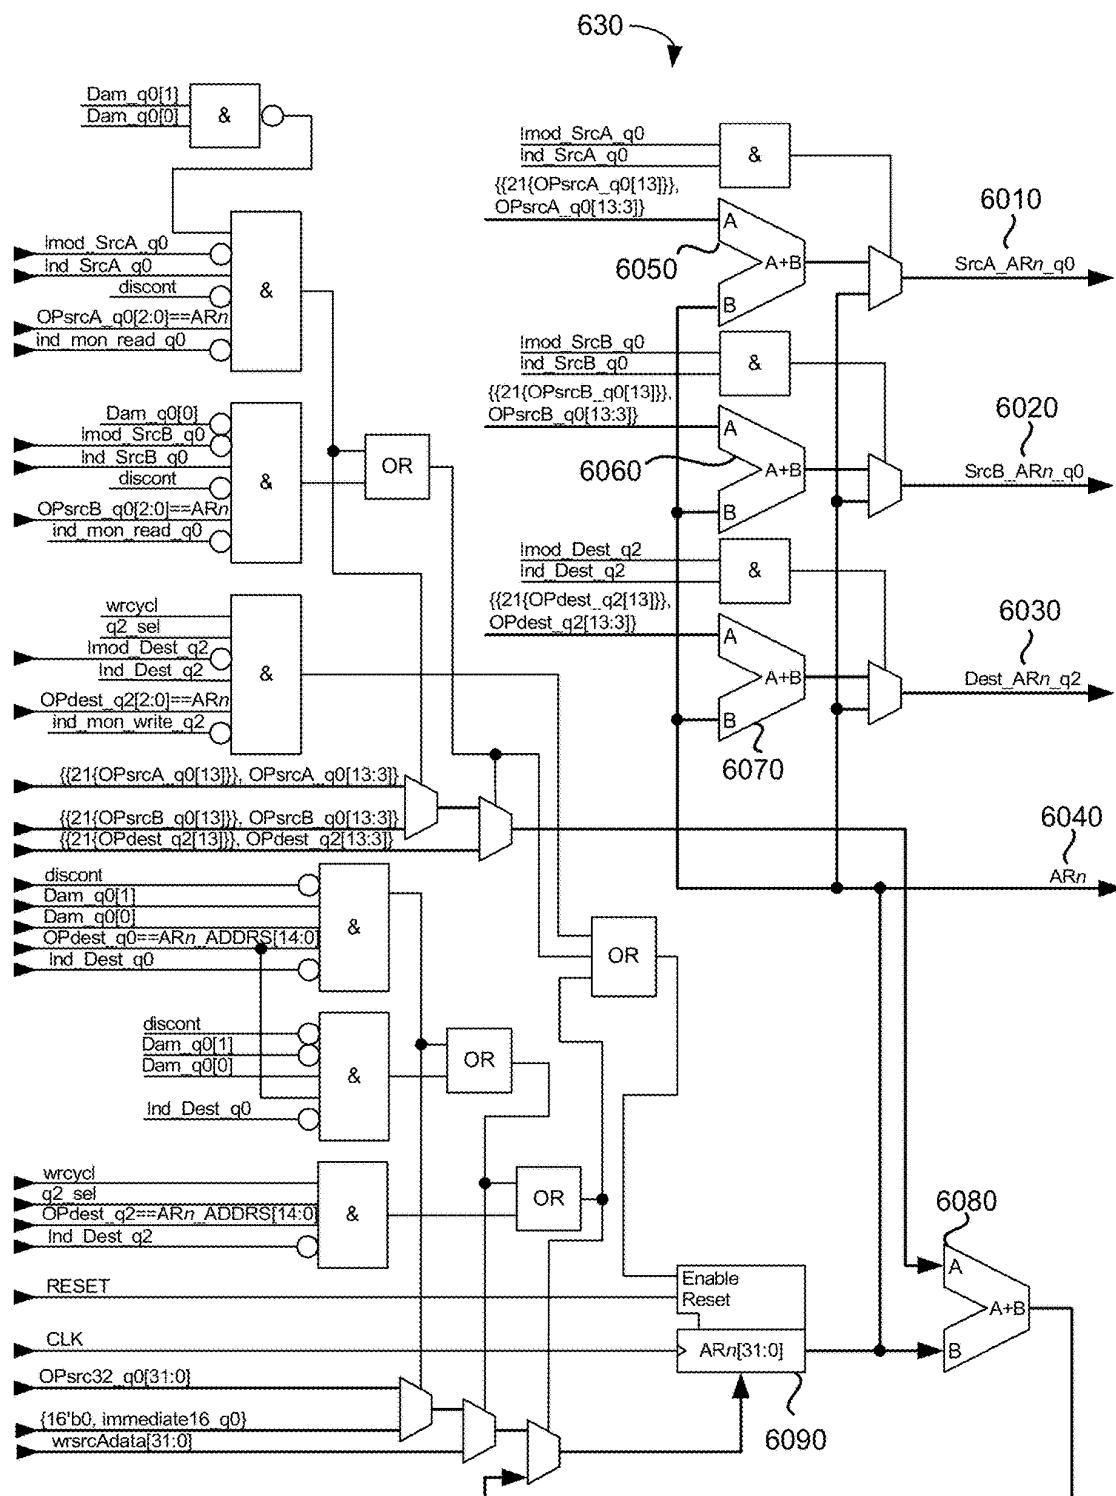


FIG. 14A

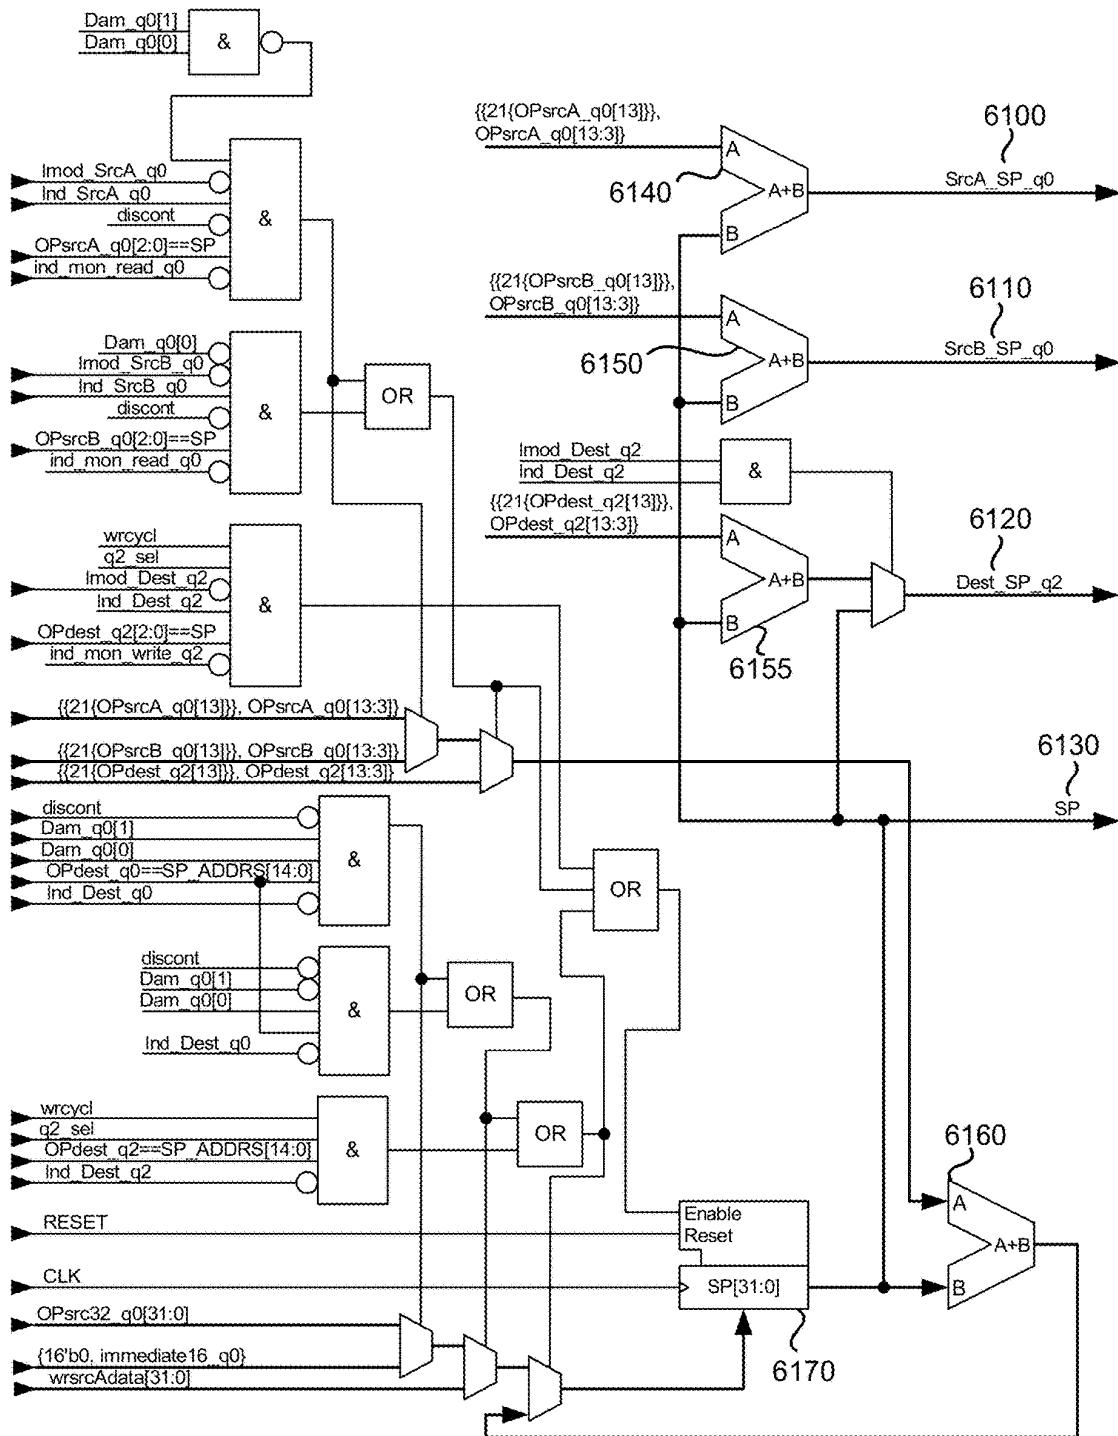


FIG. 14B

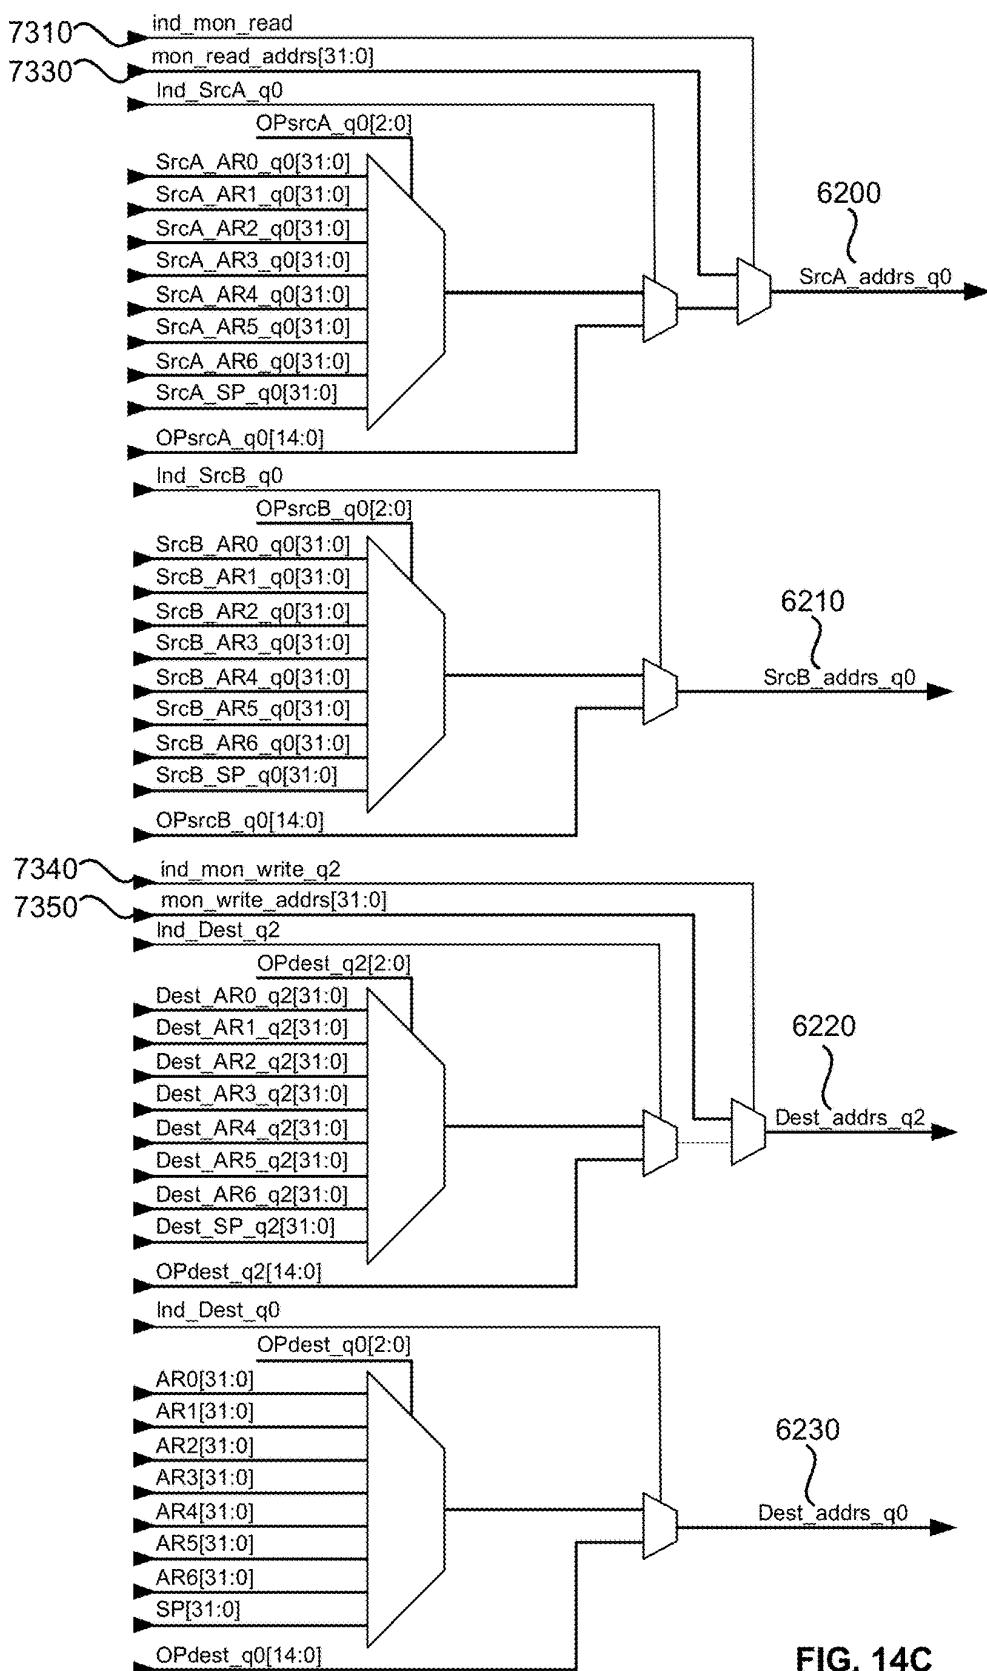


FIG. 14C

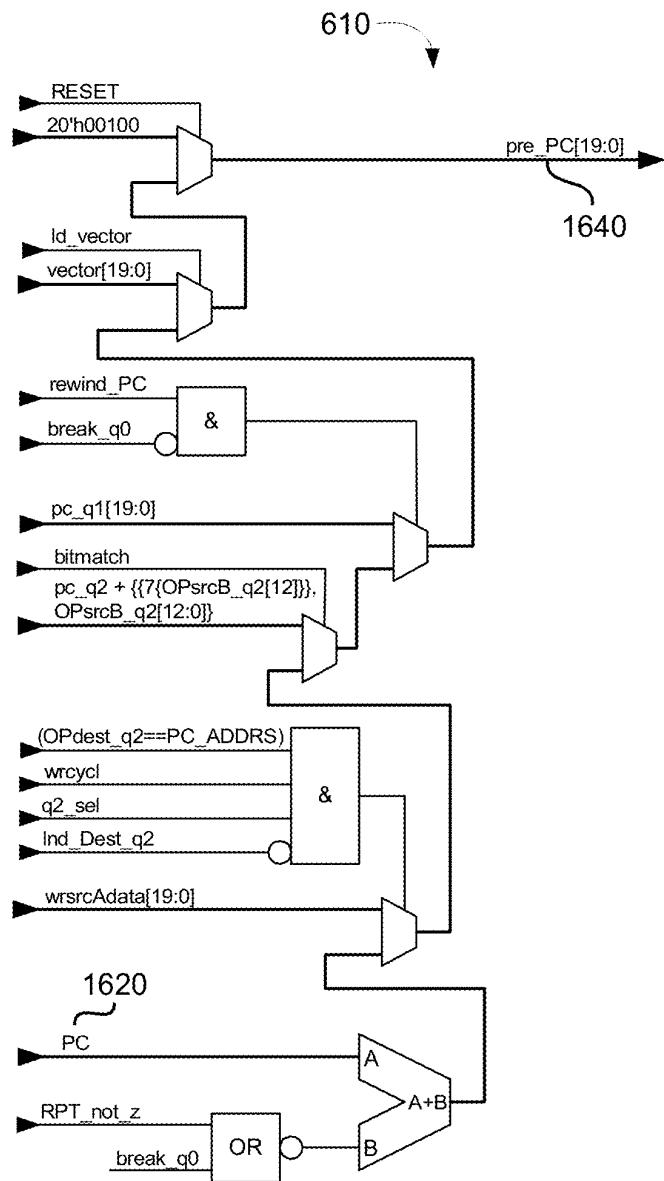


FIG. 15

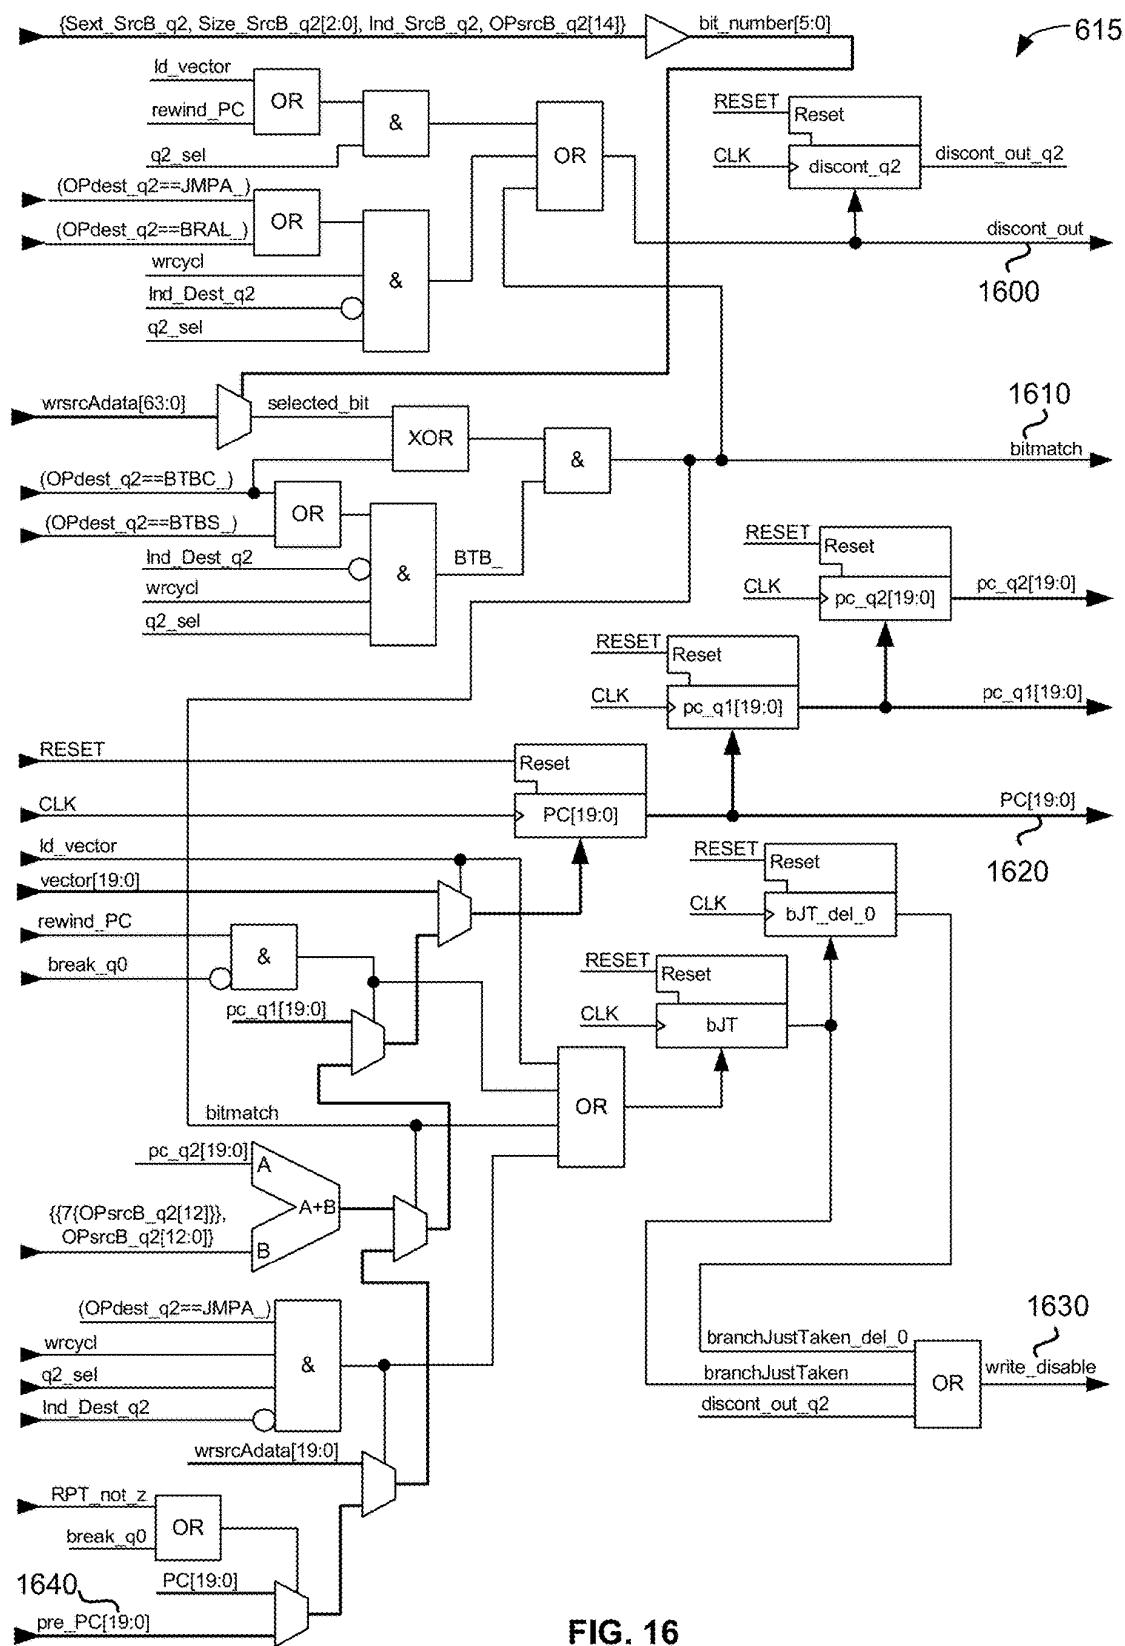


FIG. 16

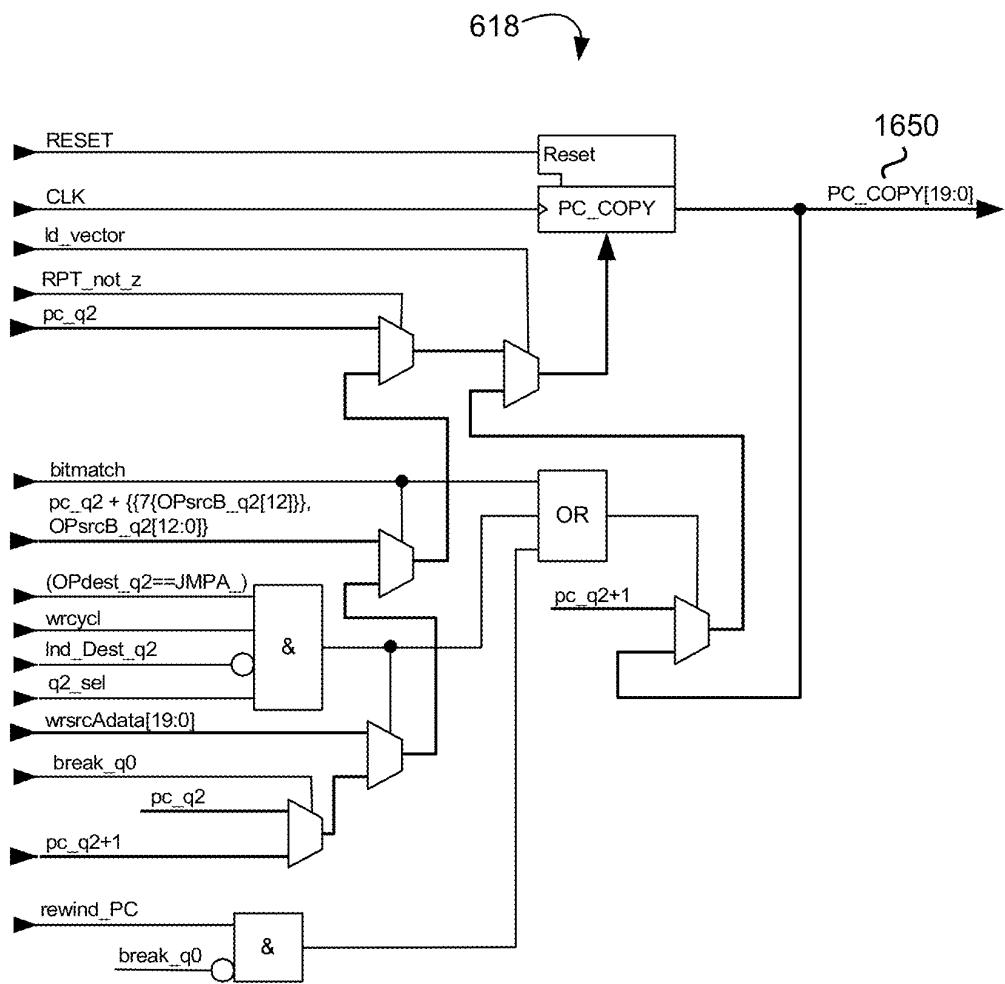


FIG. 17

5000

bit #	name	setBit Address	clearBit Address	group Bit Pos	reset State	readable	writable
63	subsAbruptUnderflow	0x7CDF	0x7CDE	7	0	yes	yes
62	substitute X	0x7CDF	0x7CDE	6	0	yes	yes
61	substitute xor(X)	0x7CDF	0x7CDE	5	0	yes	yes
60	enable subs Inexact	0x7CDF	0x7CDE	4	0	yes	yes
59	enable subs Underflow	0x7CDF	0x7CDE	3	0	yes	yes
58	enable subs Overflow	0x7CDF	0x7CDE	2	0	yes	yes
57	enable subs DivBy0	0x7CDF	0x7CDE	1	0	yes	yes
56	enable subs Invalid	0x7CDF	0x7CDE	0	0	yes	yes
55	override RM bits	0x7CD9	0x7CD9	0	0	yes	yes
54	Away	0x7FD8	0x7FD8	3	0	yes	yes
53	enable RM attribute	0x7FD8	0x7FD8	2	0	yes	yes
52	RM bit 1	0x7FD8	0x7FD8	1	0	yes	yes
51	RM bit 0	0x7FD8	0x7FD8	0	0	yes	yes
50	IEEE compare True	n/a	n/a	n/a	0	yes	yes
49	IEEE "is" True	n/a	n/a	n/a	0	yes	yes
48	a flag raised	n/a	n/a	n/a	0	yes	yes
47	total Order Mag True	n/a	n/a	n/a	0	yes	yes
46	total Order True	n/a	n/a	n/a	0	yes	yes
45	positive Infinity	n/a	n/a	n/a	0	yes	yes
44	positive Normal	n/a	n/a	n/a	0	yes	yes
43	positive subNormal	n/a	n/a	n/a	0	yes	yes
42	positive Zero	n/a	n/a	n/a	0	yes	yes
41	negative Zero	n/a	n/a	n/a	0	yes	yes
40	negative subNormal	n/a	n/a	n/a	0	yes	yes
39	negative Normal	n/a	n/a	n/a	0	yes	yes
38	negative Infinity	n/a	n/a	n/a	0	yes	yes
37	quiet NaN	n/a	n/a	n/a	0	yes	yes
36	signaling NaN	n/a	n/a	n/a	0	yes	yes
35	enAltImmInexactHandl	0x7CDB	0x7CDA	4	0	yes	yes
34	enAltImmUnderflowHandl	0x7CDB	0x7CDA	3	0	yes	yes
33	enAltImmOverflowHandl	0x7CDB	0x7CDA	2	0	yes	yes
32	enAltImmDivByZeroHandl	0x7CDB	0x7CDA	1	0	yes	yes
31	enAltImmInvalidHandl	0x7CDB	0x7CDA	0	0	yes	yes
30	razNoInexactFlag	0x7CE5	0x7CE4	4	1	yes	yes
29	razNoUnderflowFlag	0x7CE5	0x7CE4	3	0	yes	yes
28	razNoOverflowFlag	0x7CE5	0x7CE4	2	0	yes	yes
27	razNoDivByZeroFlag	0x7CE5	0x7CE4	1	0	yes	yes
26	razNoInvalidFlag	0x7CE5	0x7CE4	0	0	yes	yes
25	inexact Flag	0x7CE7	0x7CE6	4	0	yes	yes
24	underflow Flag	0x7CE7	0x7CE6	3	0	yes	yes
23	overflow Flag	0x7CE7	0x7CE6	2	0	yes	yes
22	divByZero Flag	0x7CE7	0x7CE6	1	0	yes	yes
21	invalid Flag	0x7CE7	0x7CE6	0	0	yes	yes
20	inexact Signal	0x7CE1	0x7CE0	4	0	yes	yes
19	underflow Signal	0x7CE1	0x7CE0	3	0	yes	yes
18	overflow Signal	0x7CE1	0x7CE0	2	0	yes	yes
17	divByZero Signal	0x7CE1	0x7CE0	1	0	yes	yes
16	invalid Signal	0x7CE1	0x7CE0	0	0	yes	yes

FIG. 18A

bit #	name	setBit Address	clearBit Address	group	reset Bit Pos	readable State	writable	
15	spare4	n/a	n/a	n/a	0	yes	no	
14	spare3	n/a	n/a	n/a	0	yes	no	
13	spare2	n/a	n/a	n/a	0	yes	no	
12	spare1	n/a	n/a	n/a	0	yes	no	
11	spare0	n/a	n/a	n/a	0	yes	no	
10	IRQ	n/a	n/a	n/a	0	yes	no	
9	exception Source	n/a	n/a	n/a	0	yes	no	
8	A GTE B	n/a	n/a	n/a	1	yes	no	
7	A LTE B	n/a	n/a	n/a	1	yes	no	
6	A GT B	n/a	n/a	n/a	0	yes	no	
5	interrupt Enable	0x7CDD	0x7CDC	5	0	yes	yes	
4	Done Bit	0x7CDD	0x7CDC	4	1	yes	yes	
3	V Flag	0x7CDD	0x7CDC	3	0	yes	yes	
2	N Flag	0x7CDD	0x7CDC	2	0	yes	yes	
1	C Flag	0x7CDD	0x7CDC	1	0	yes	yes	
0	Z Flag	0x7CDD	0x7CDC	0	1	yes	yes	

5130

unused

5140

5150

5160

5170

5180

5190

Processor integer arithmetic flags;
status and control bits

FIG. 18B

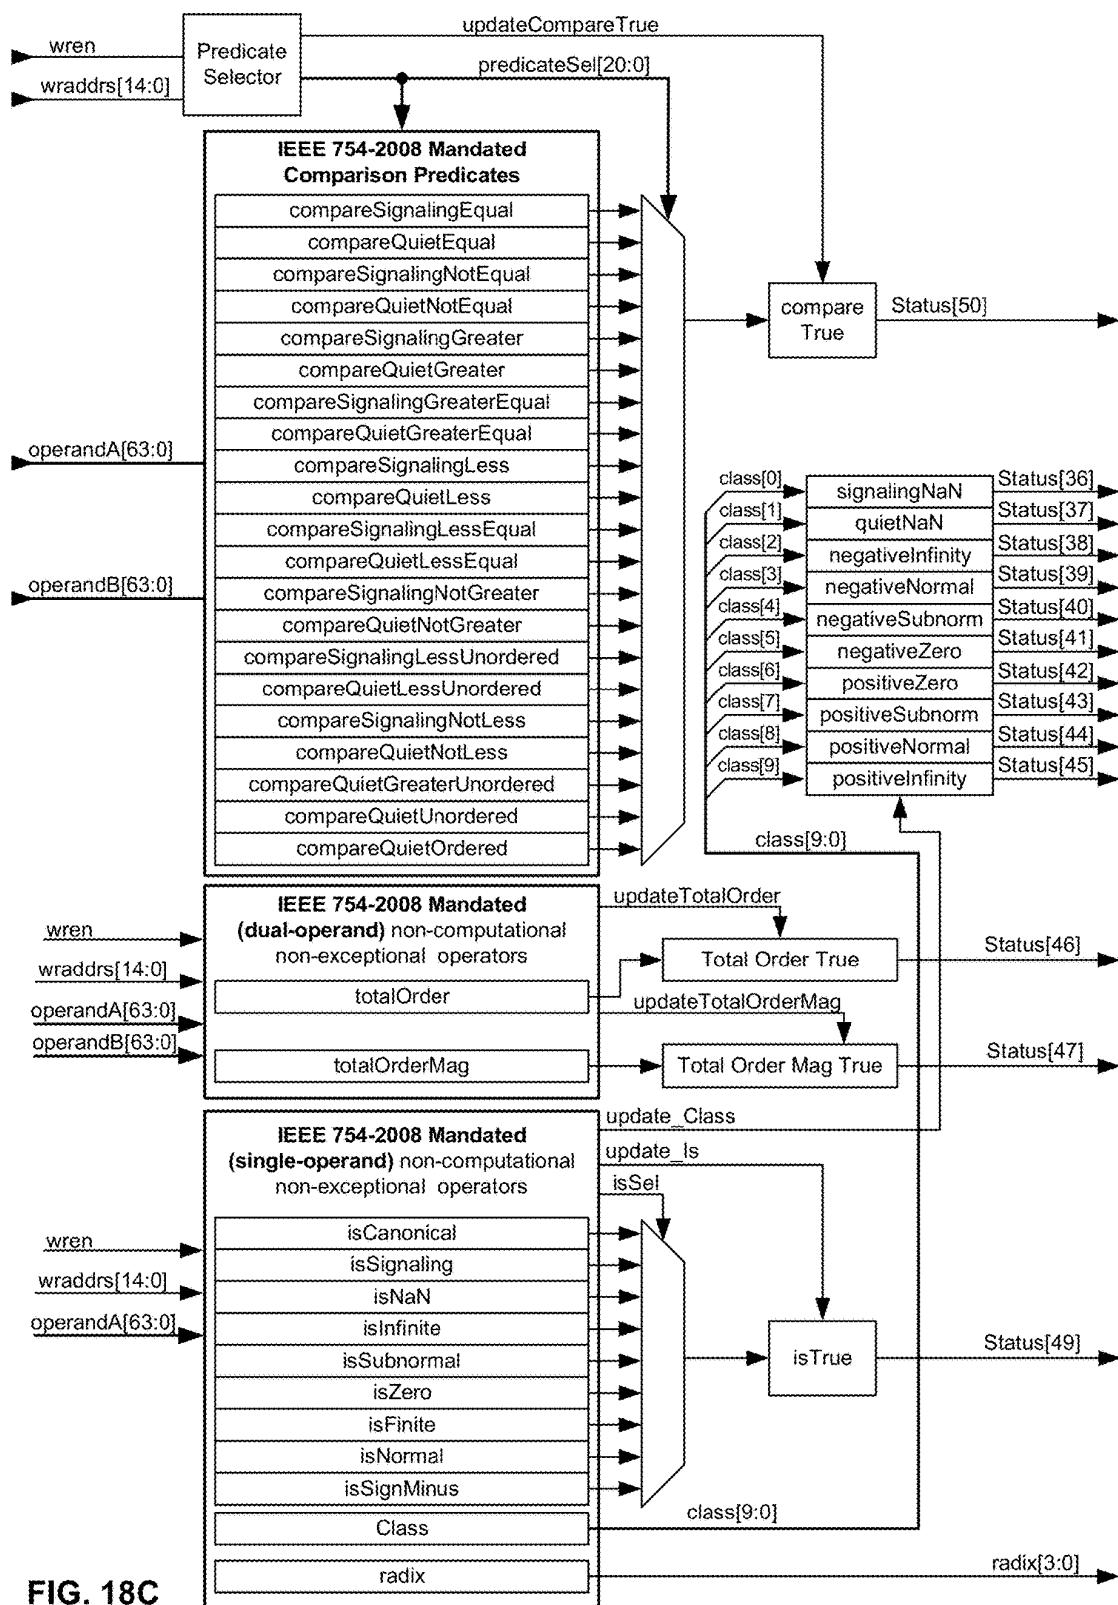


FIG. 18C

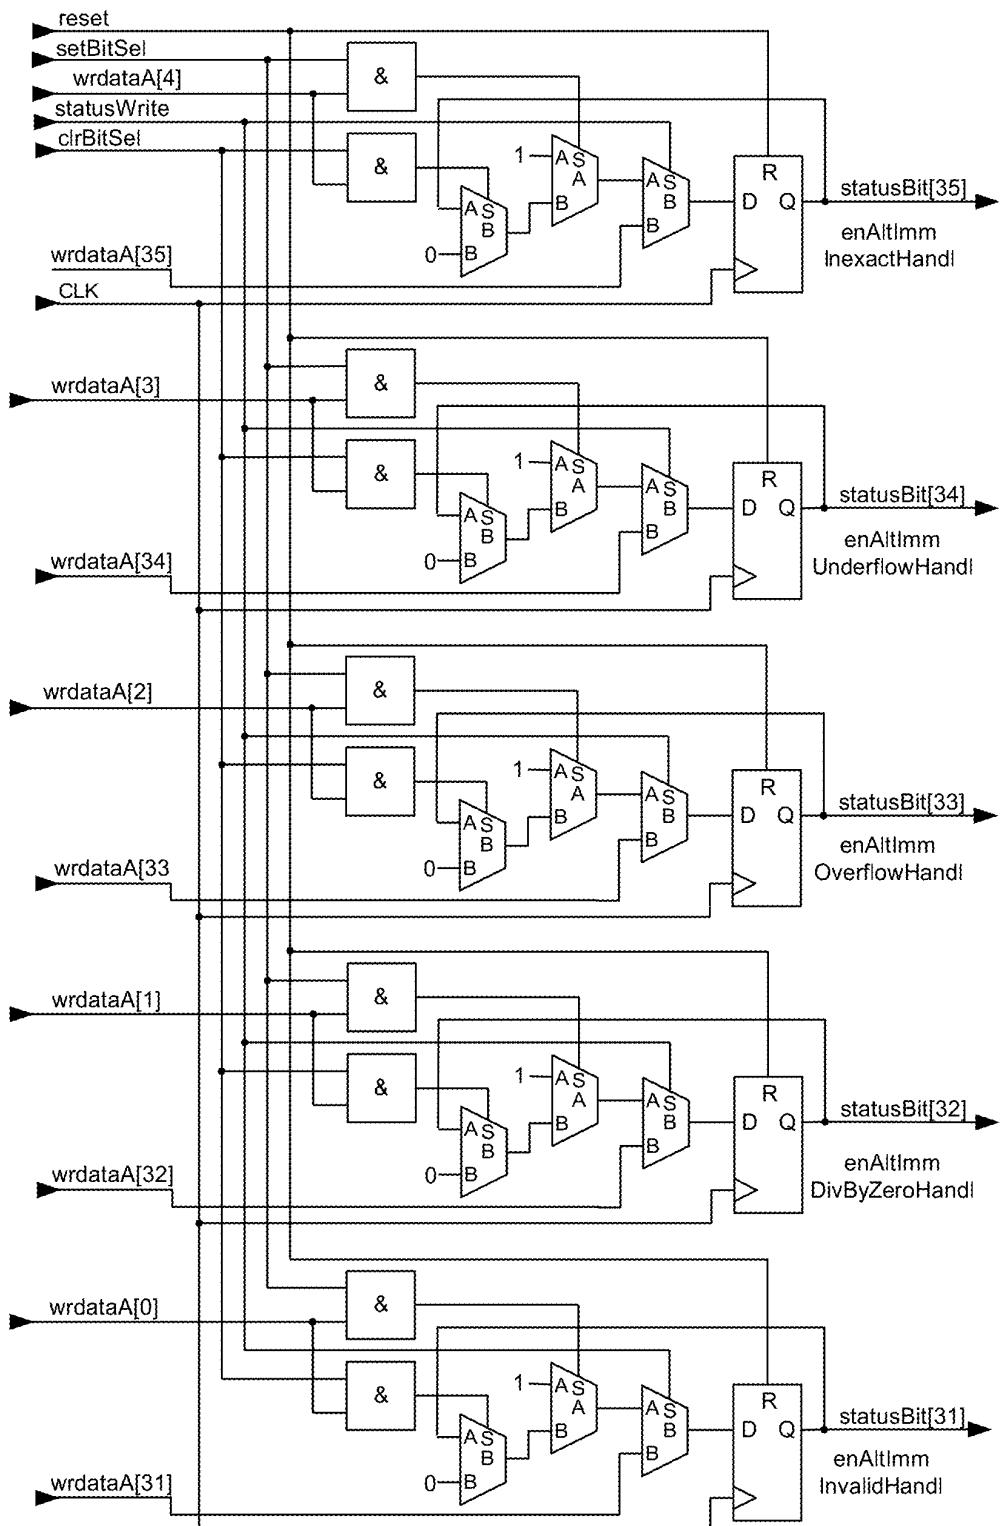


FIG. 18D

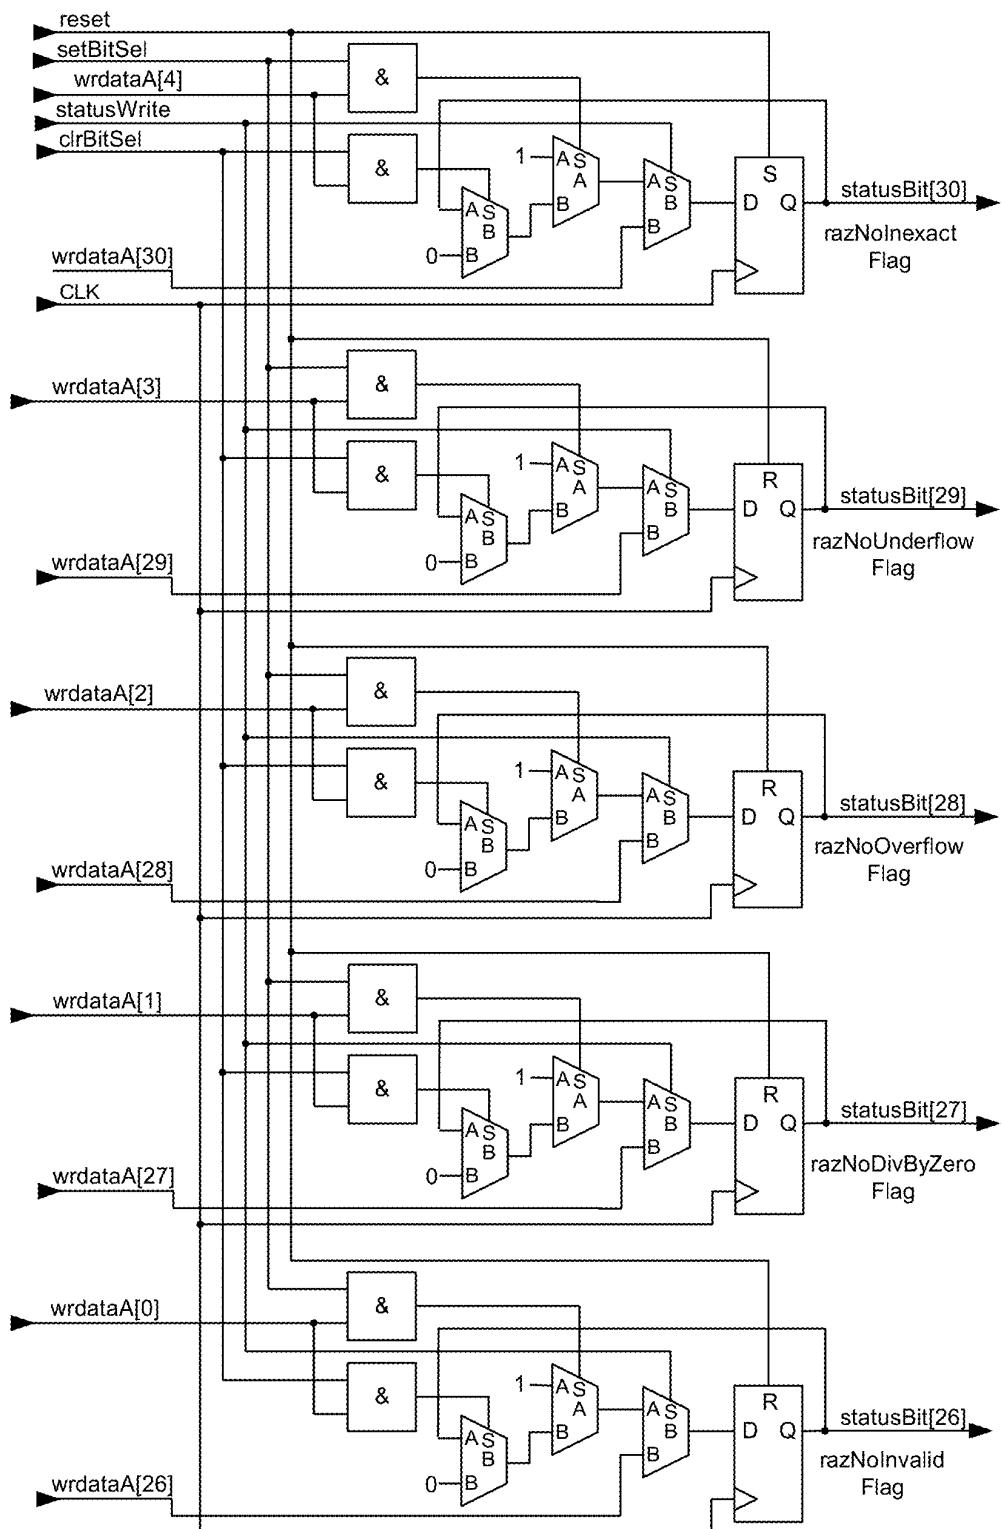


FIG. 18E

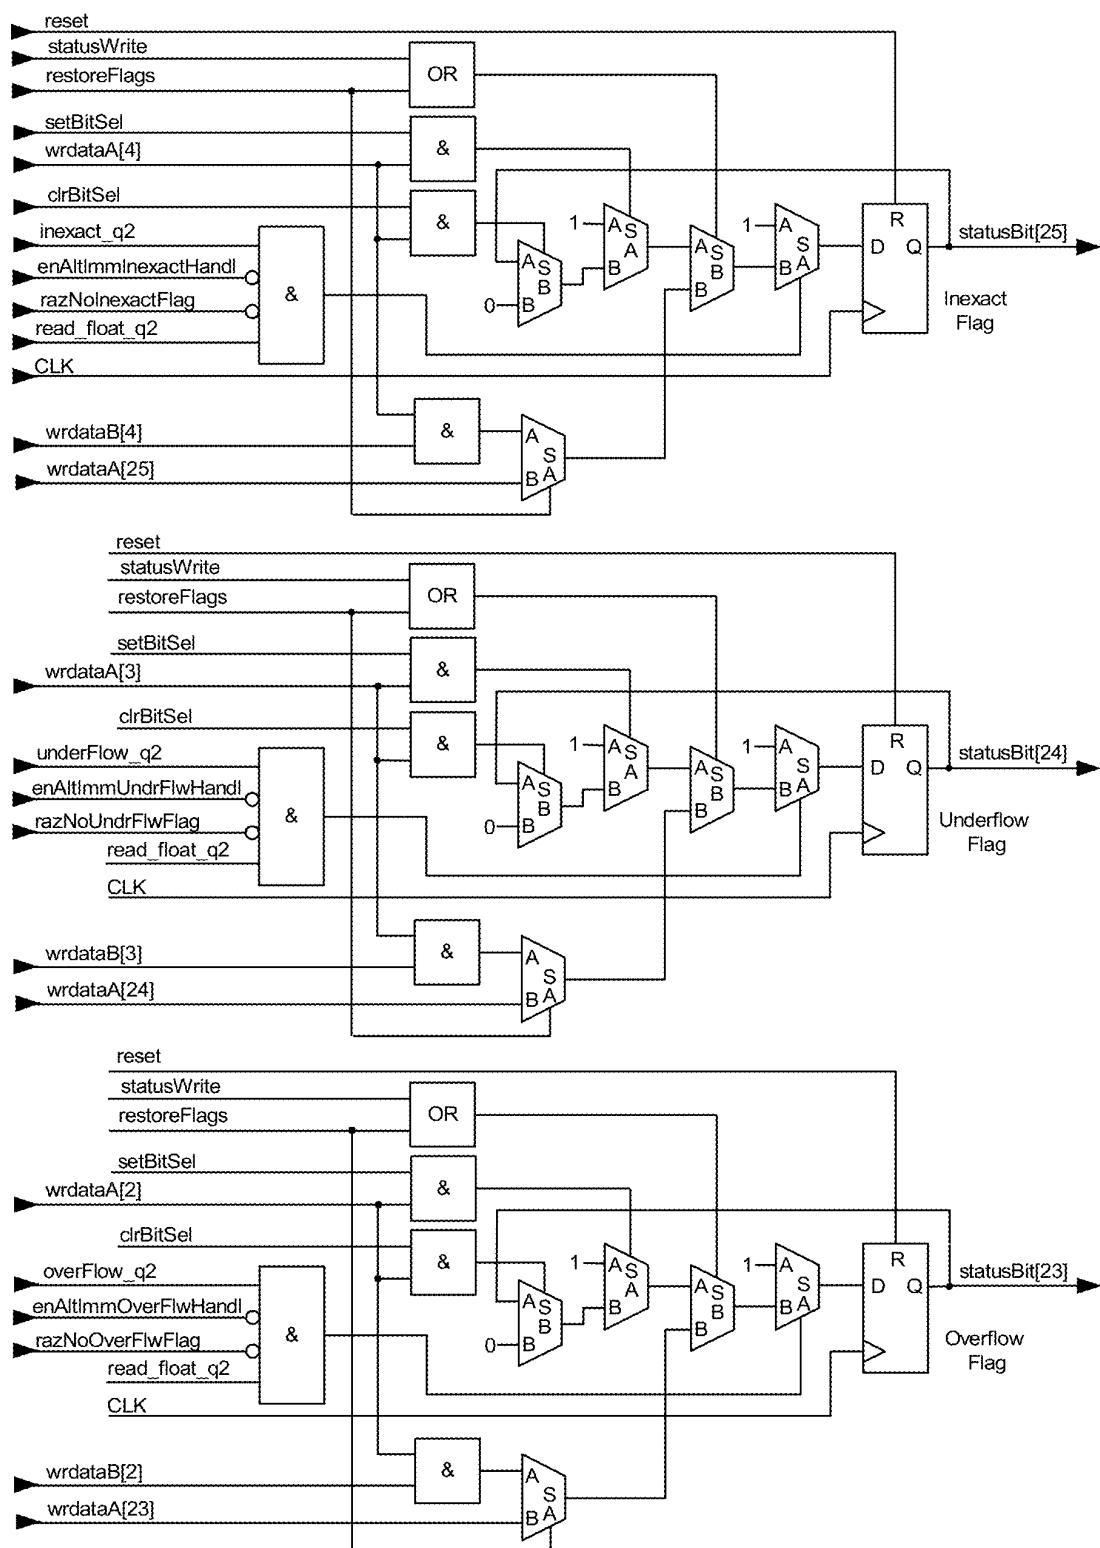


FIG. 18F

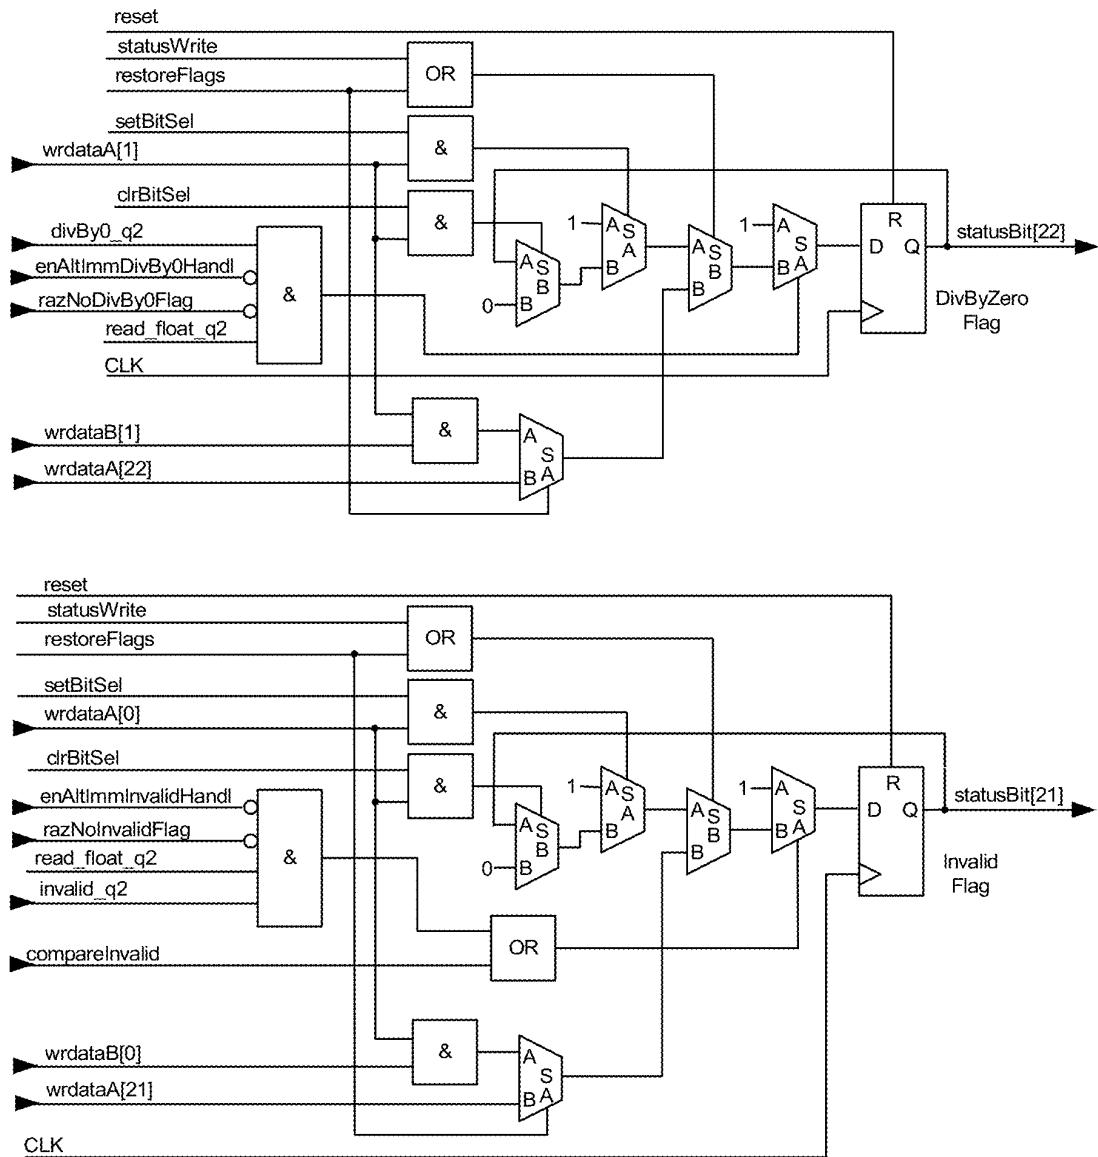


FIG. 18G

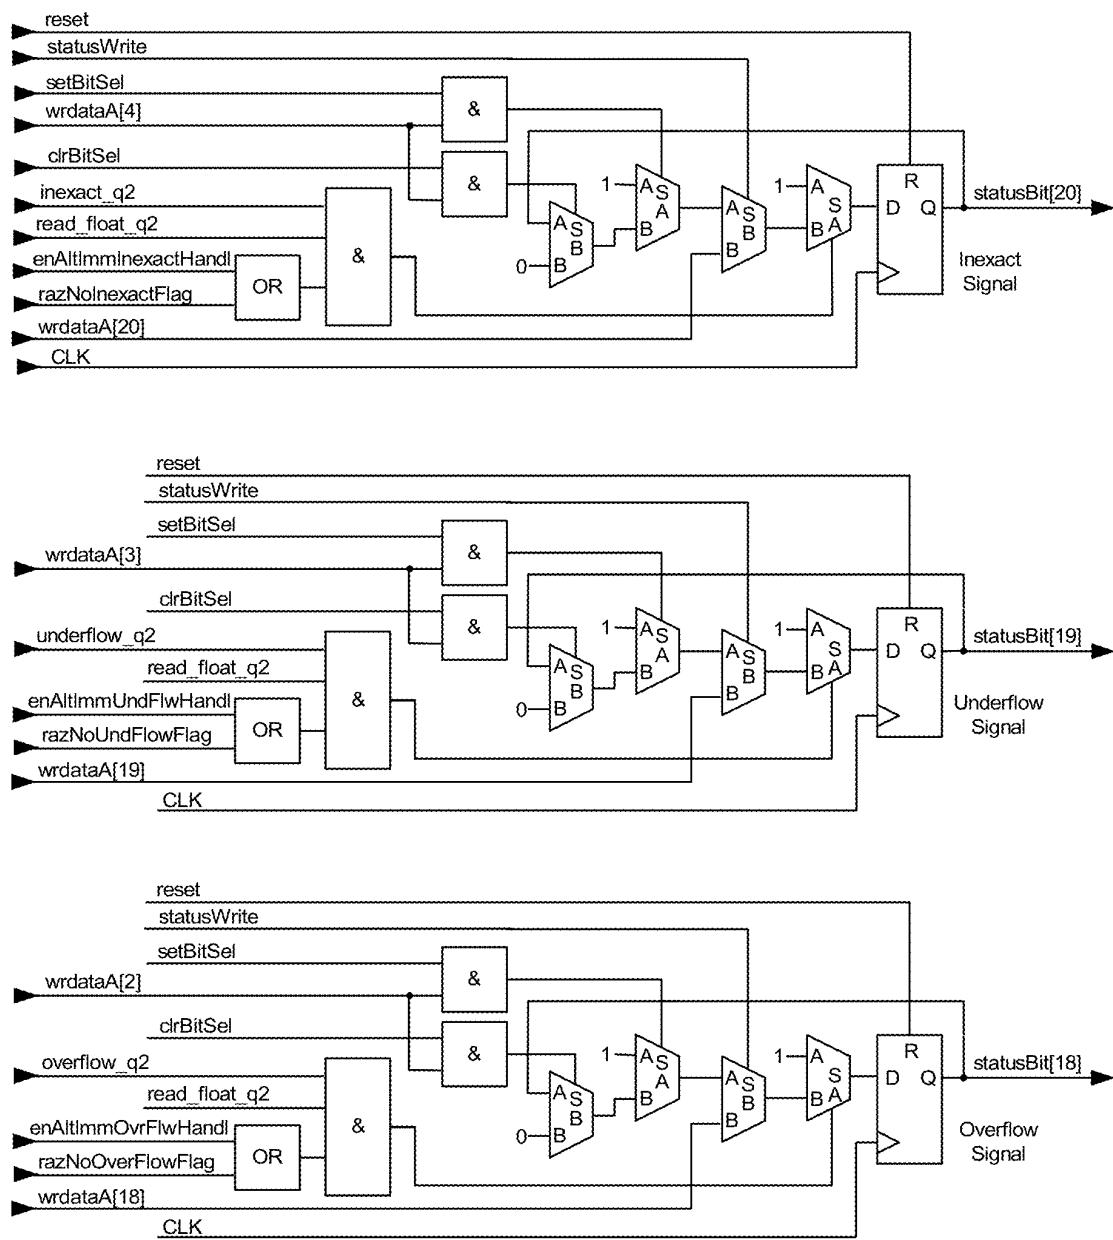


FIG. 18H

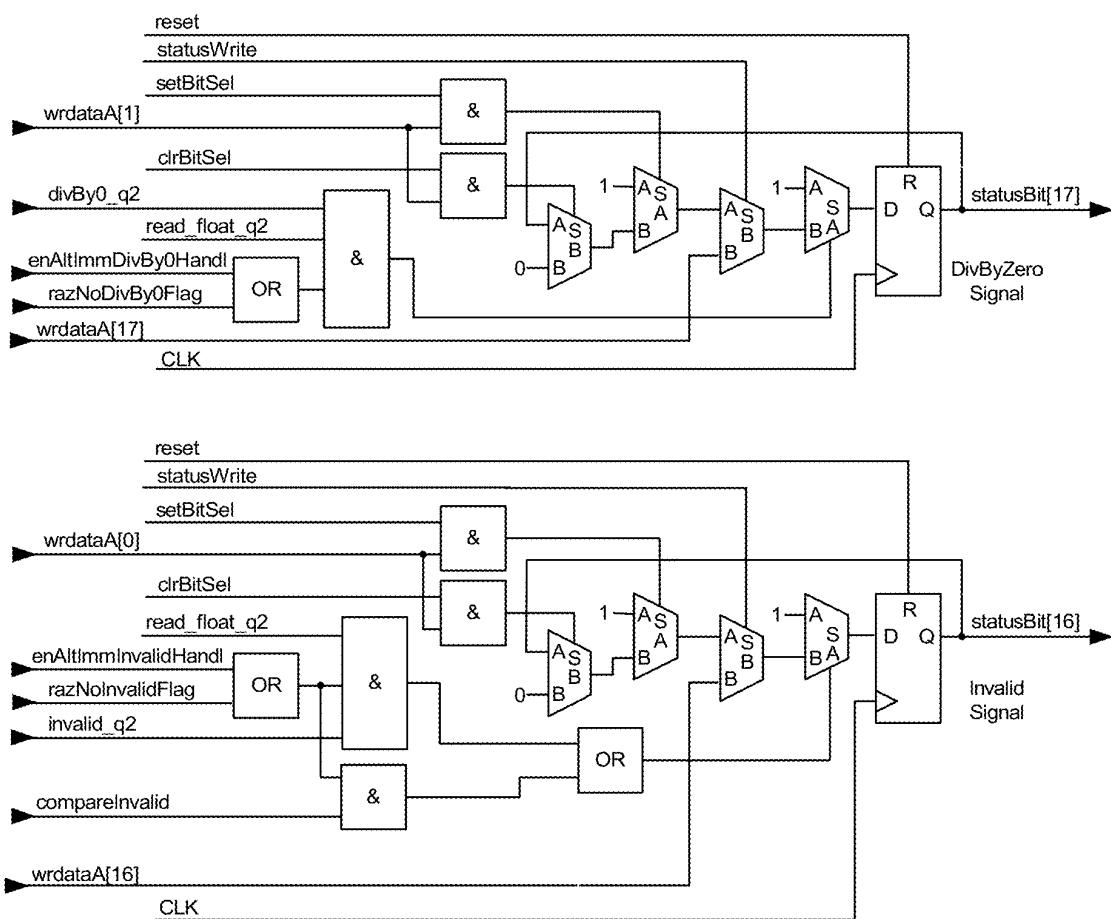


FIG. 18I

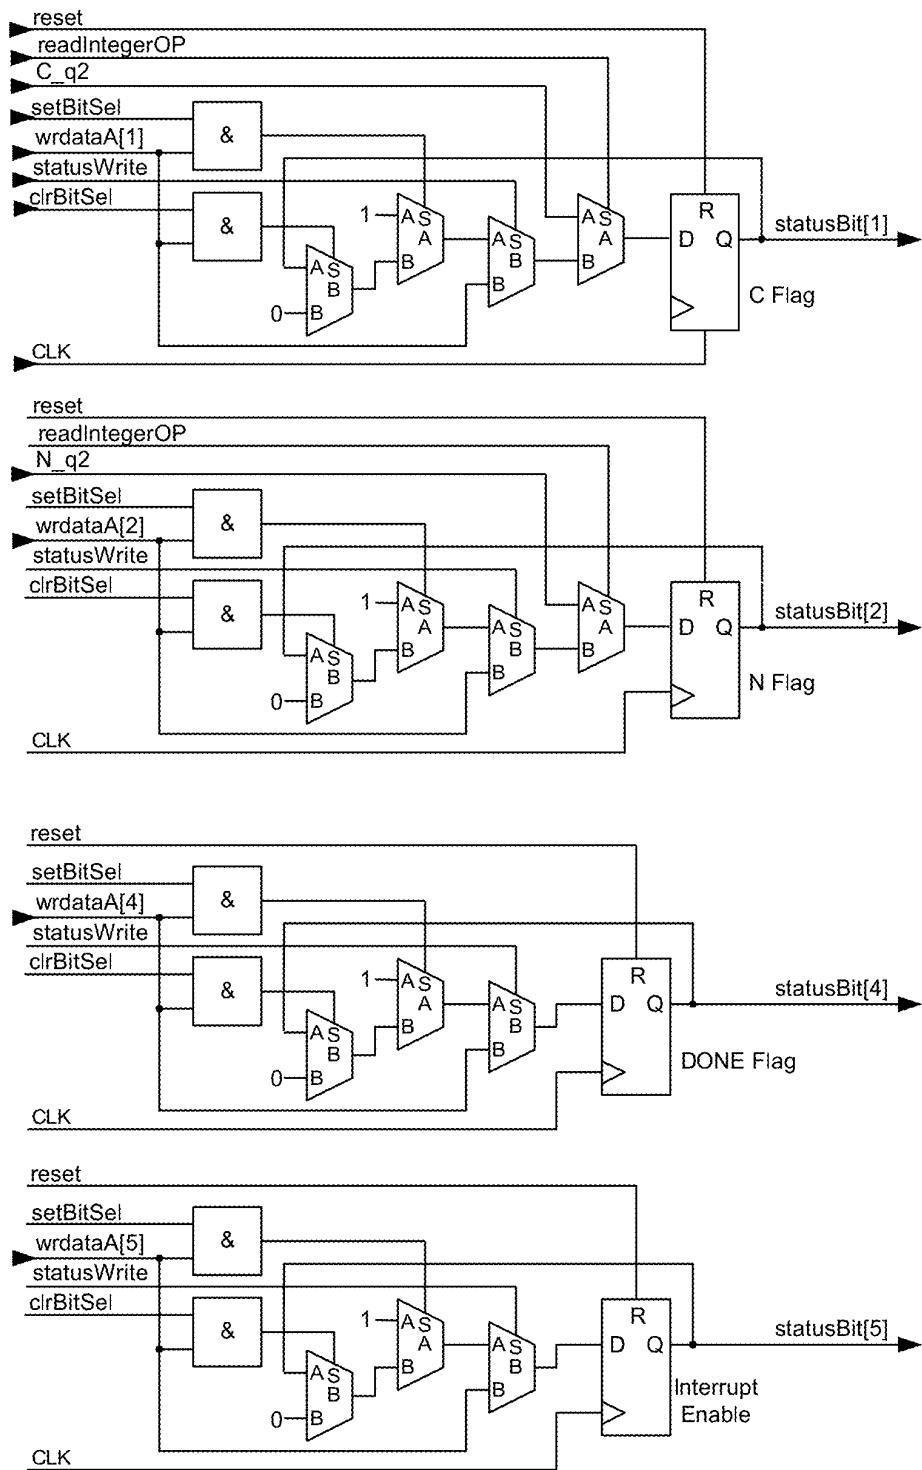


FIG. 18J

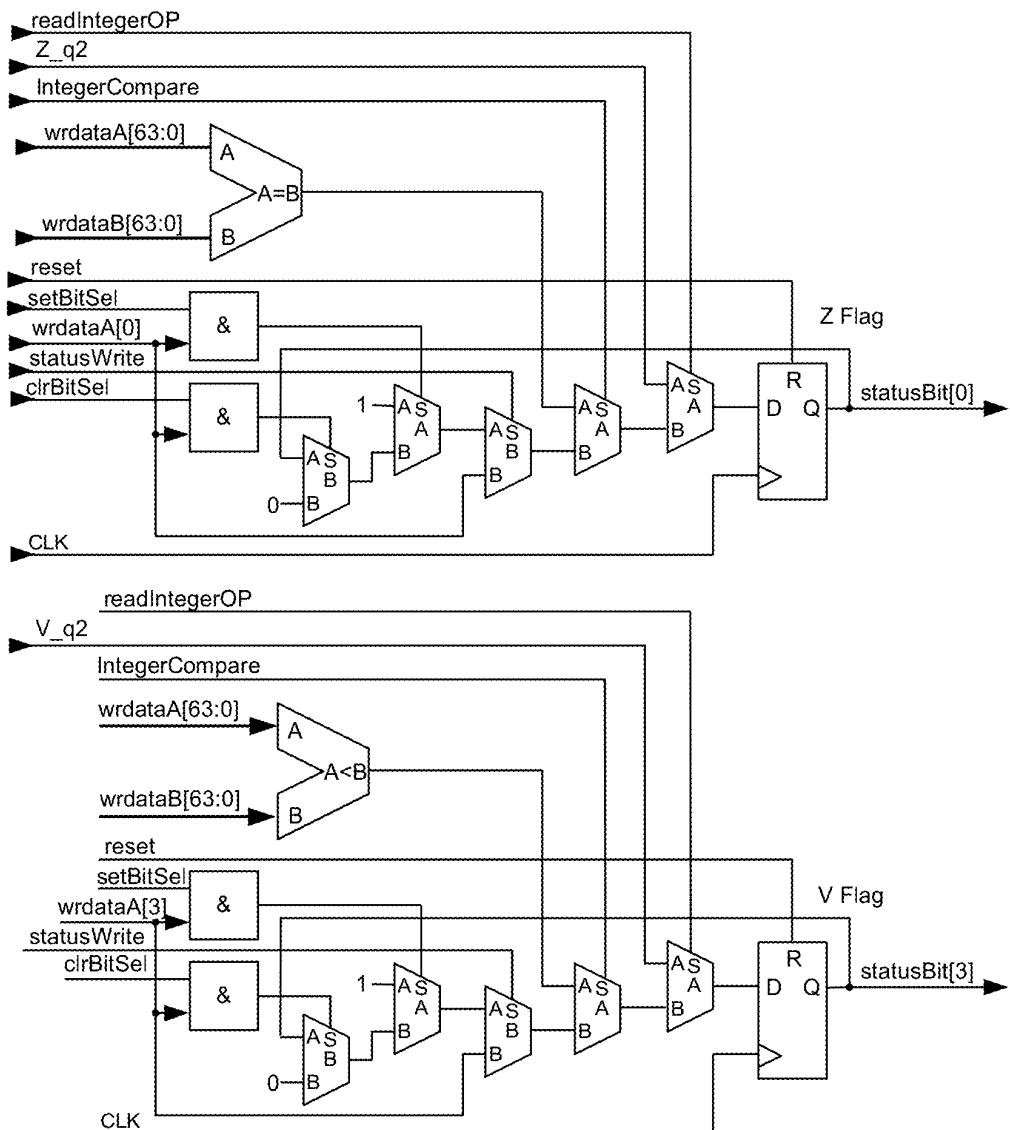


FIG. 18K

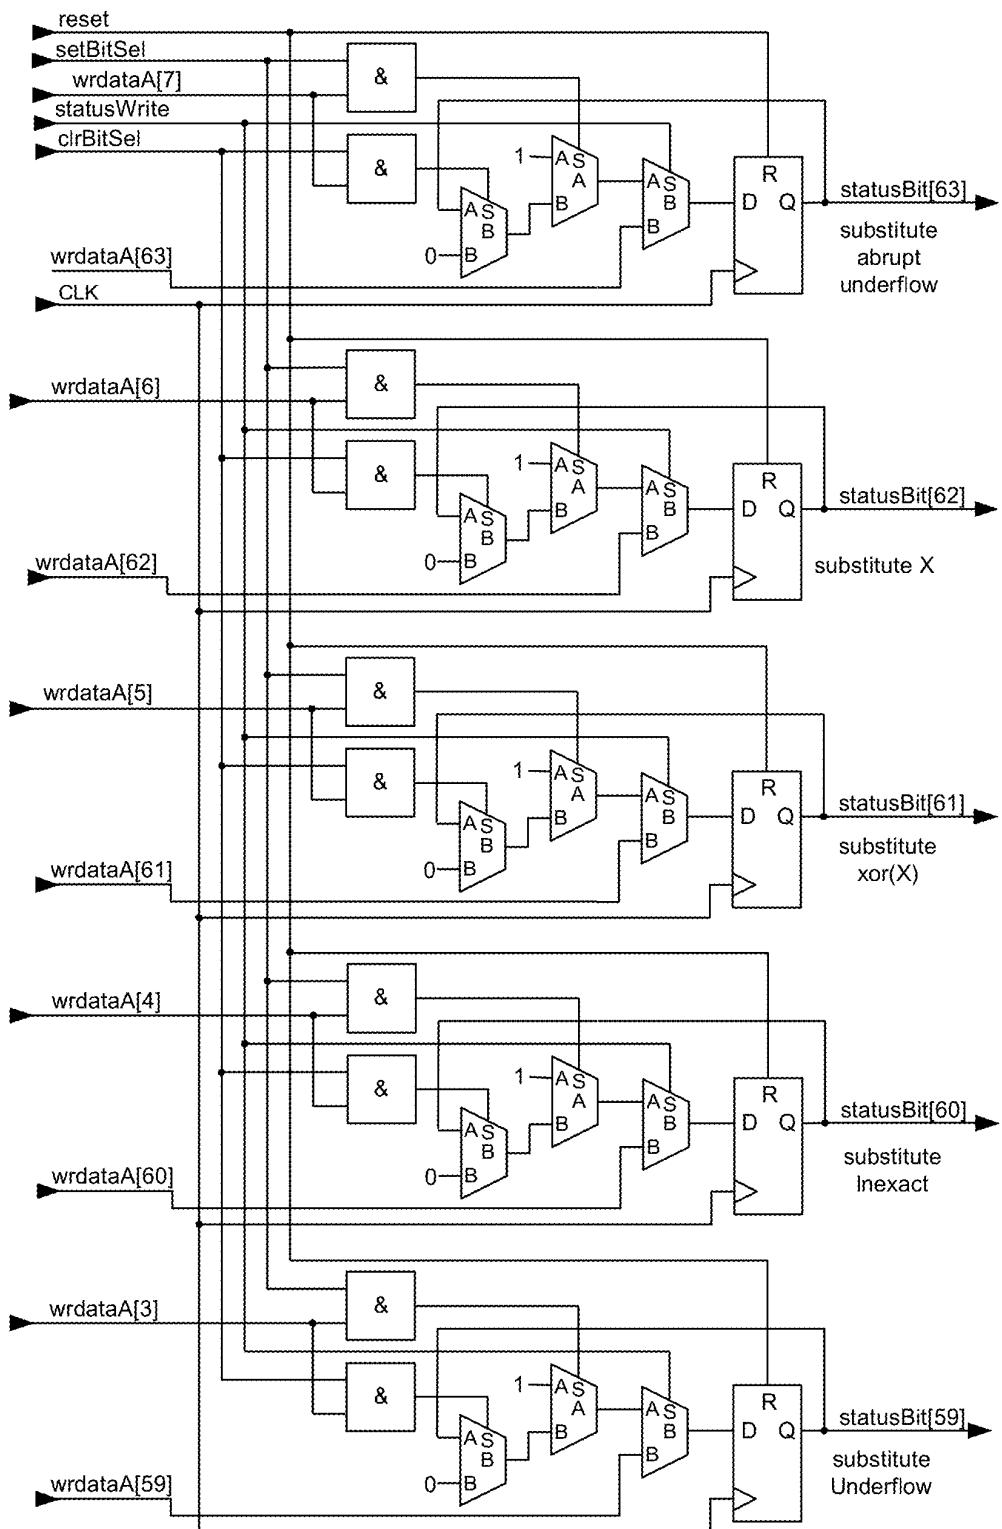


FIG. 18L

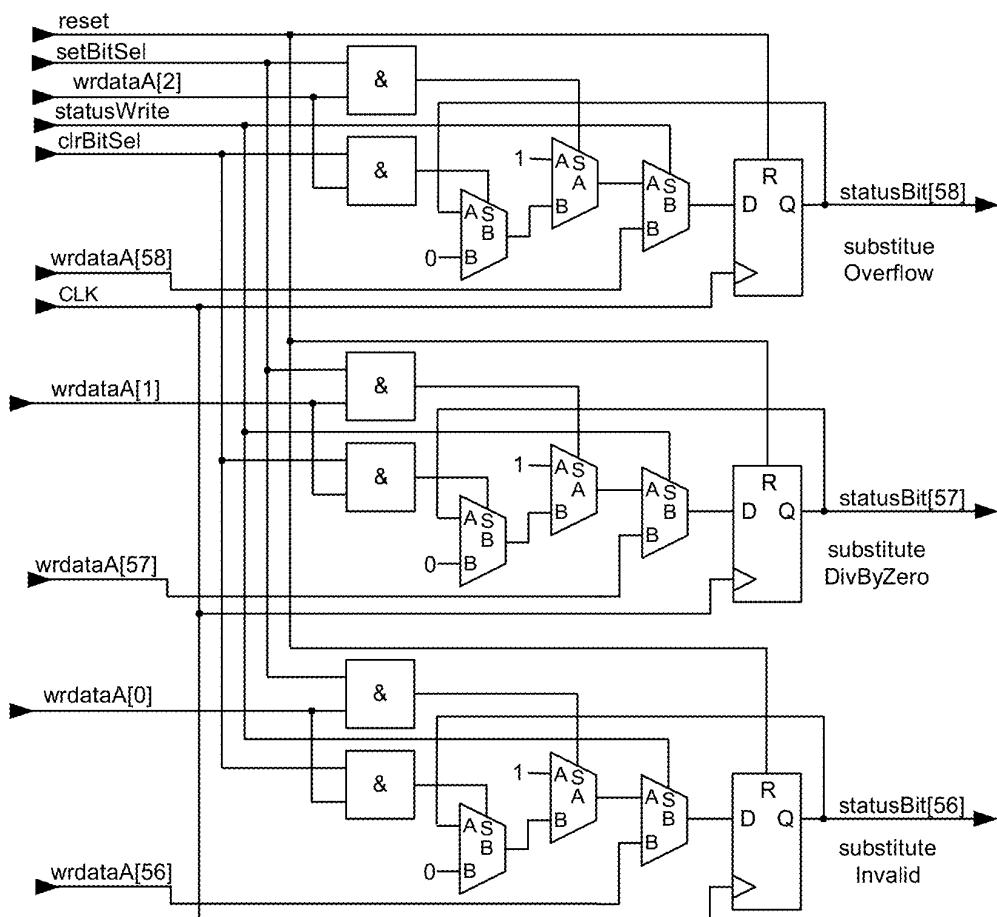


FIG. 18M

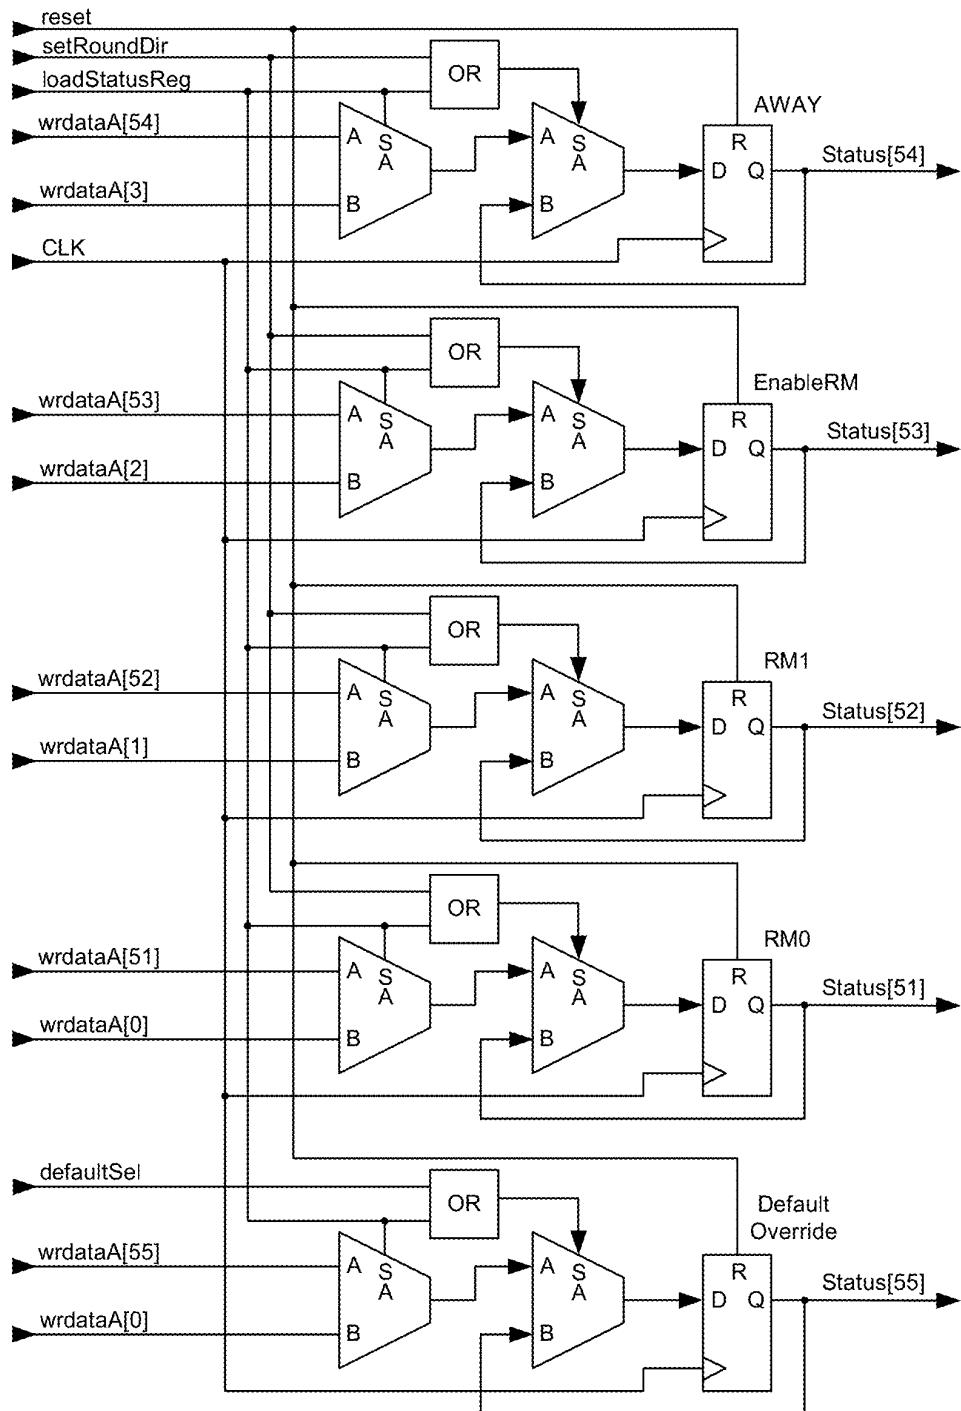


FIG. 18N

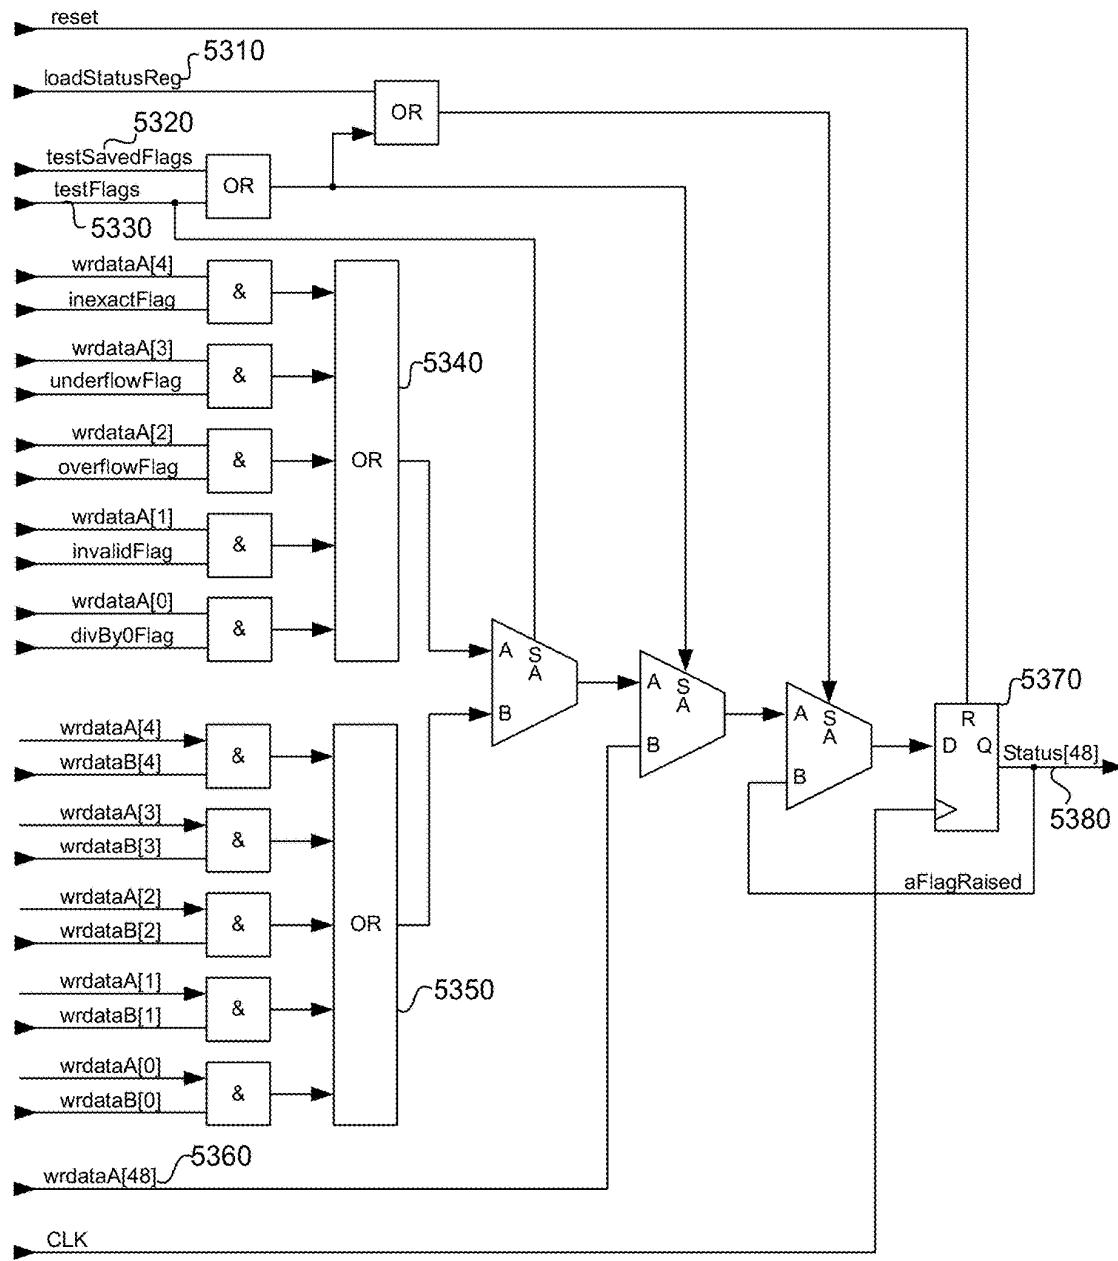


FIG. 18O

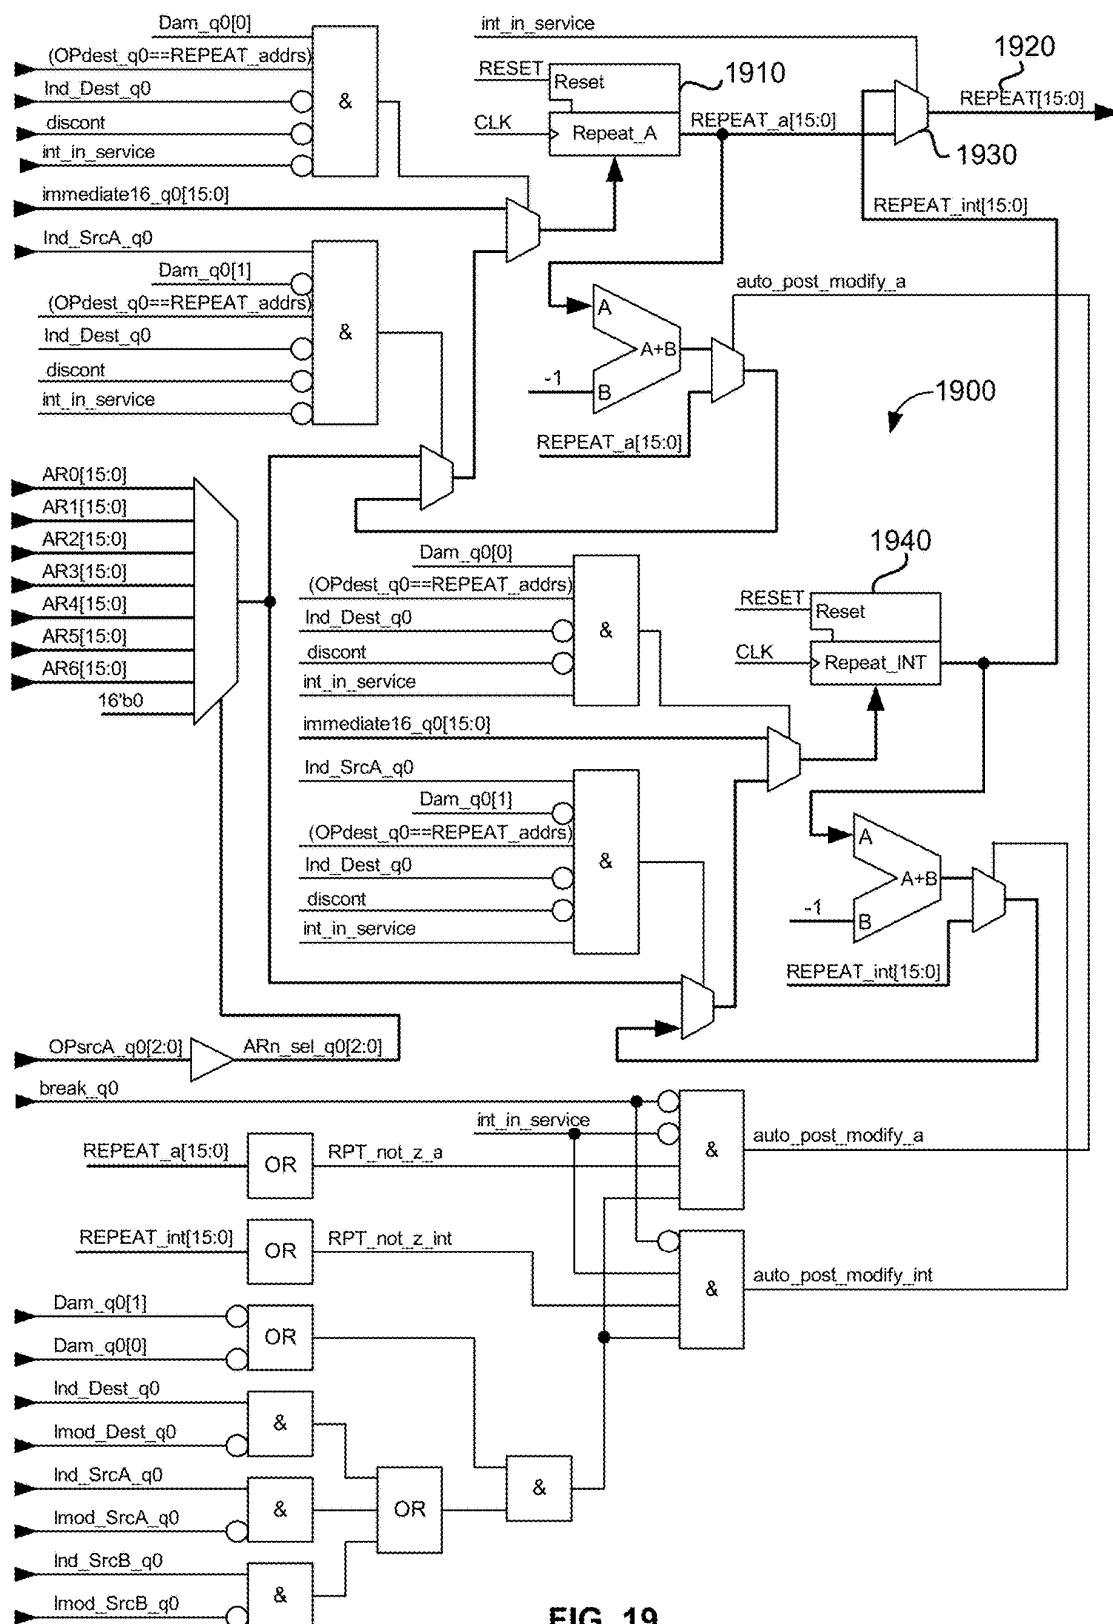


FIG. 19

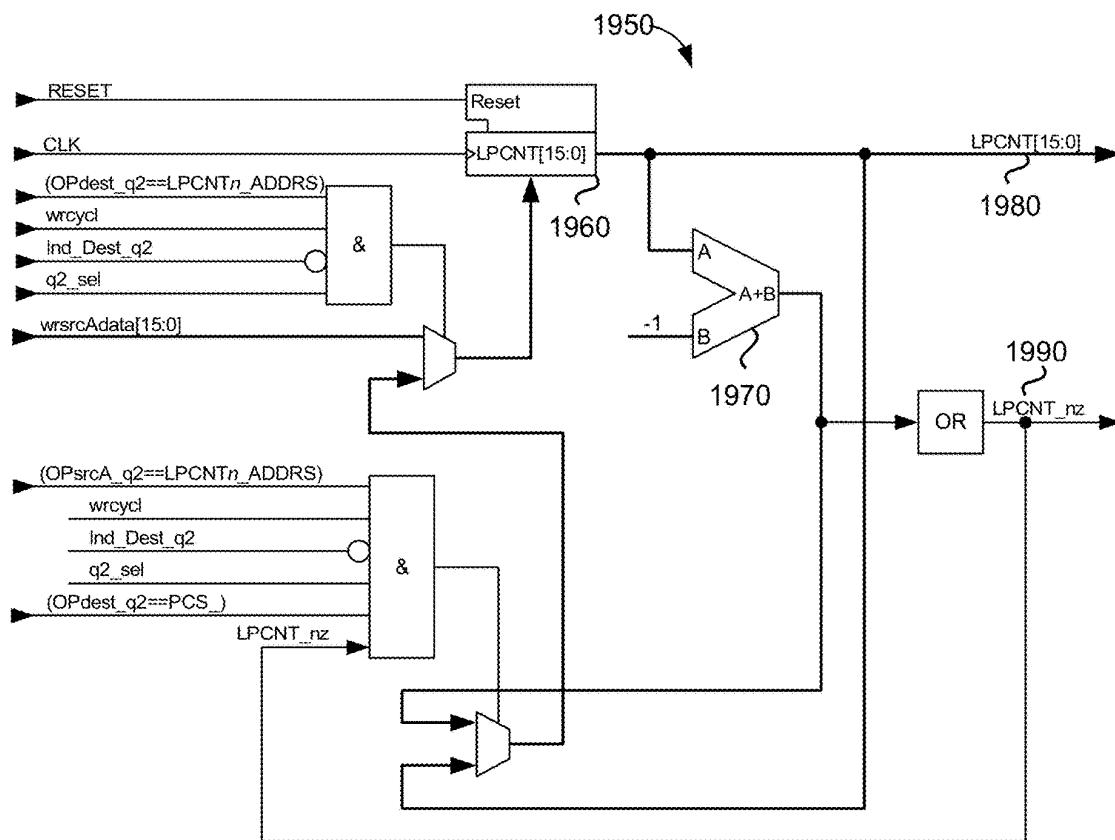


FIG. 20

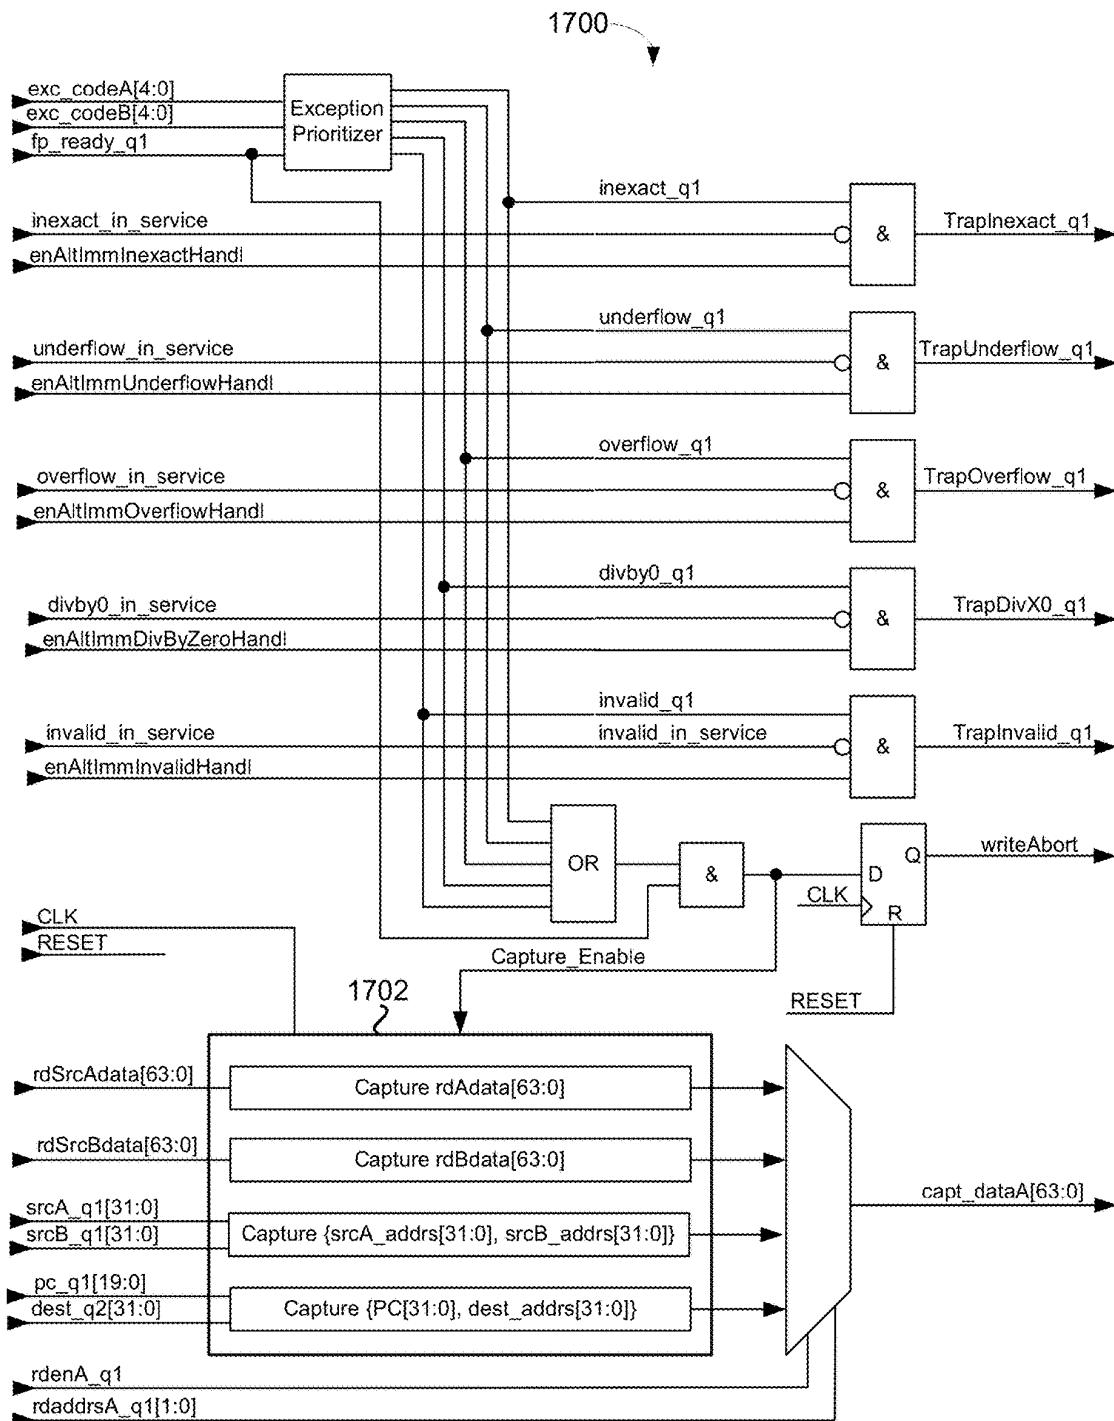


FIG. 21

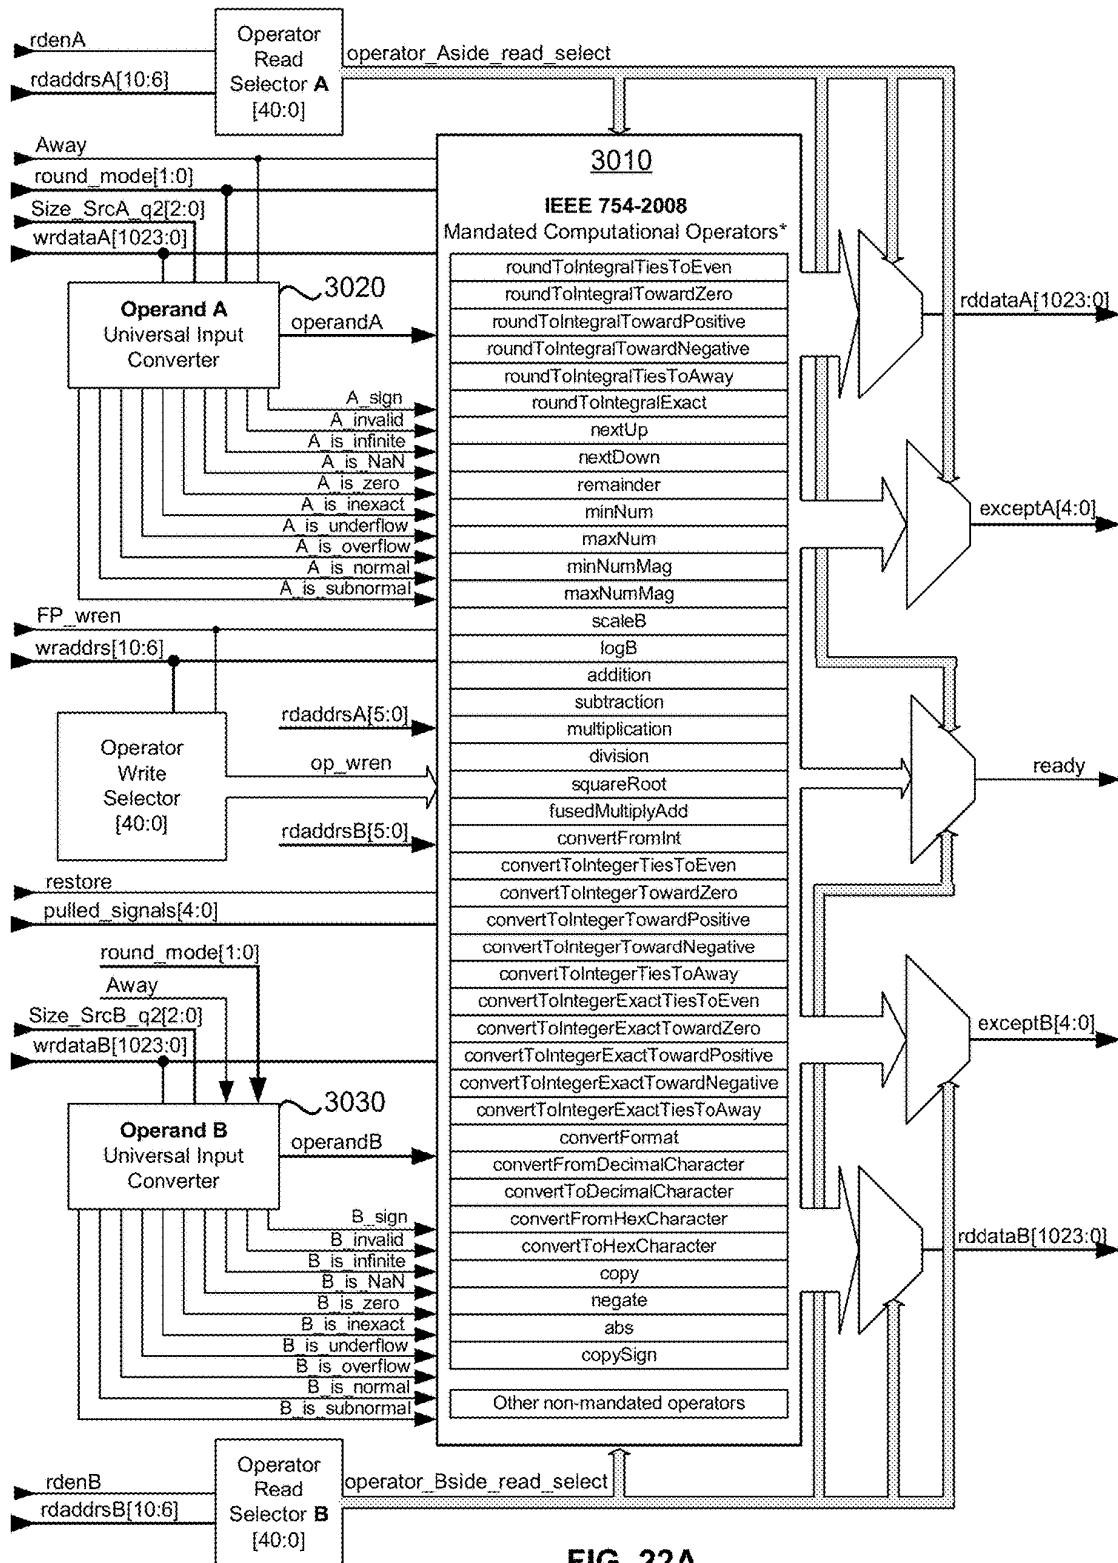


FIG. 22A

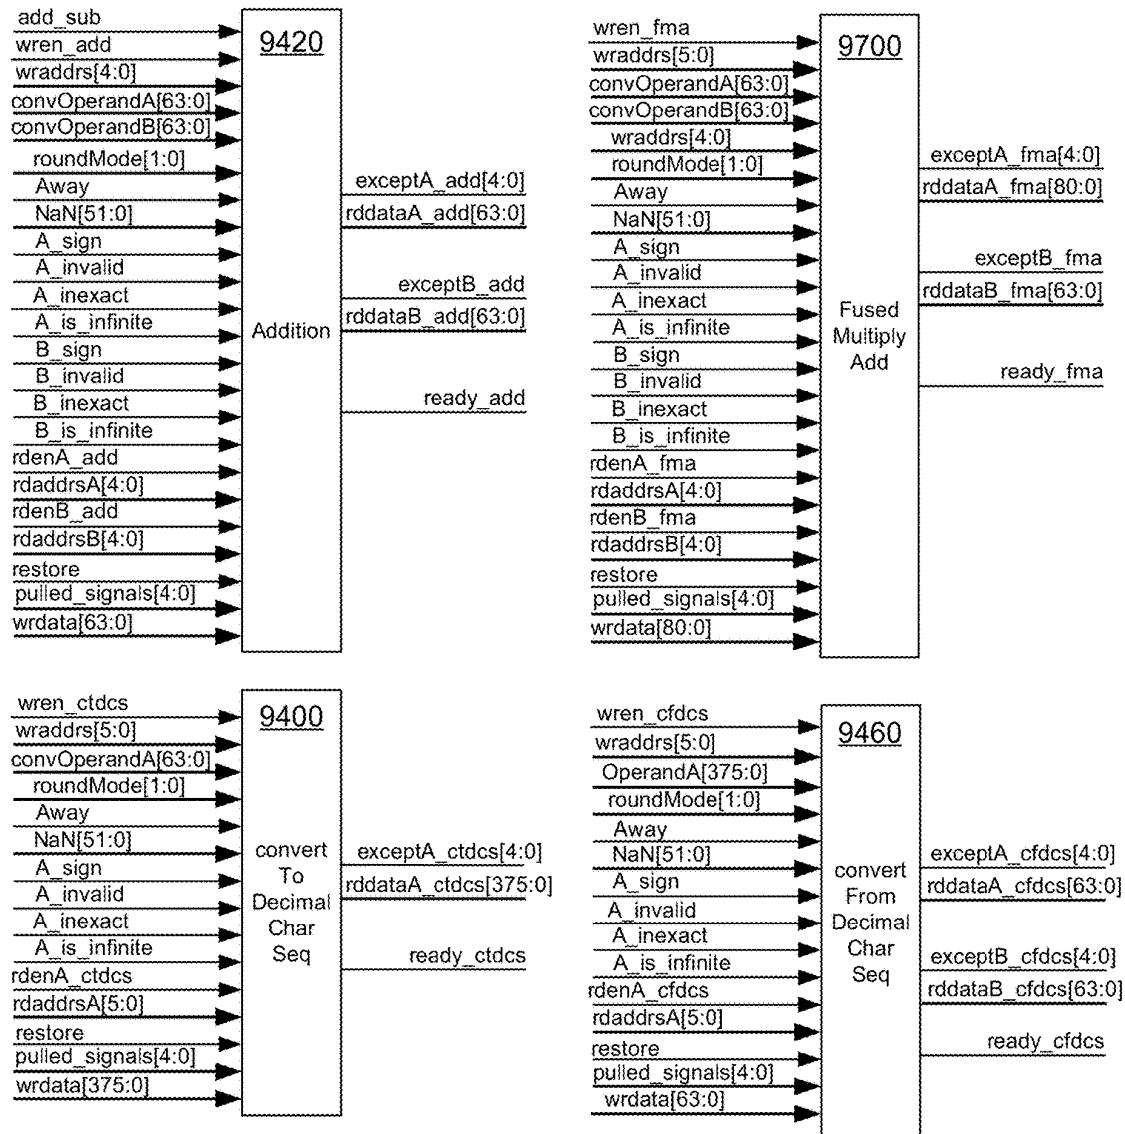


FIG. 22B

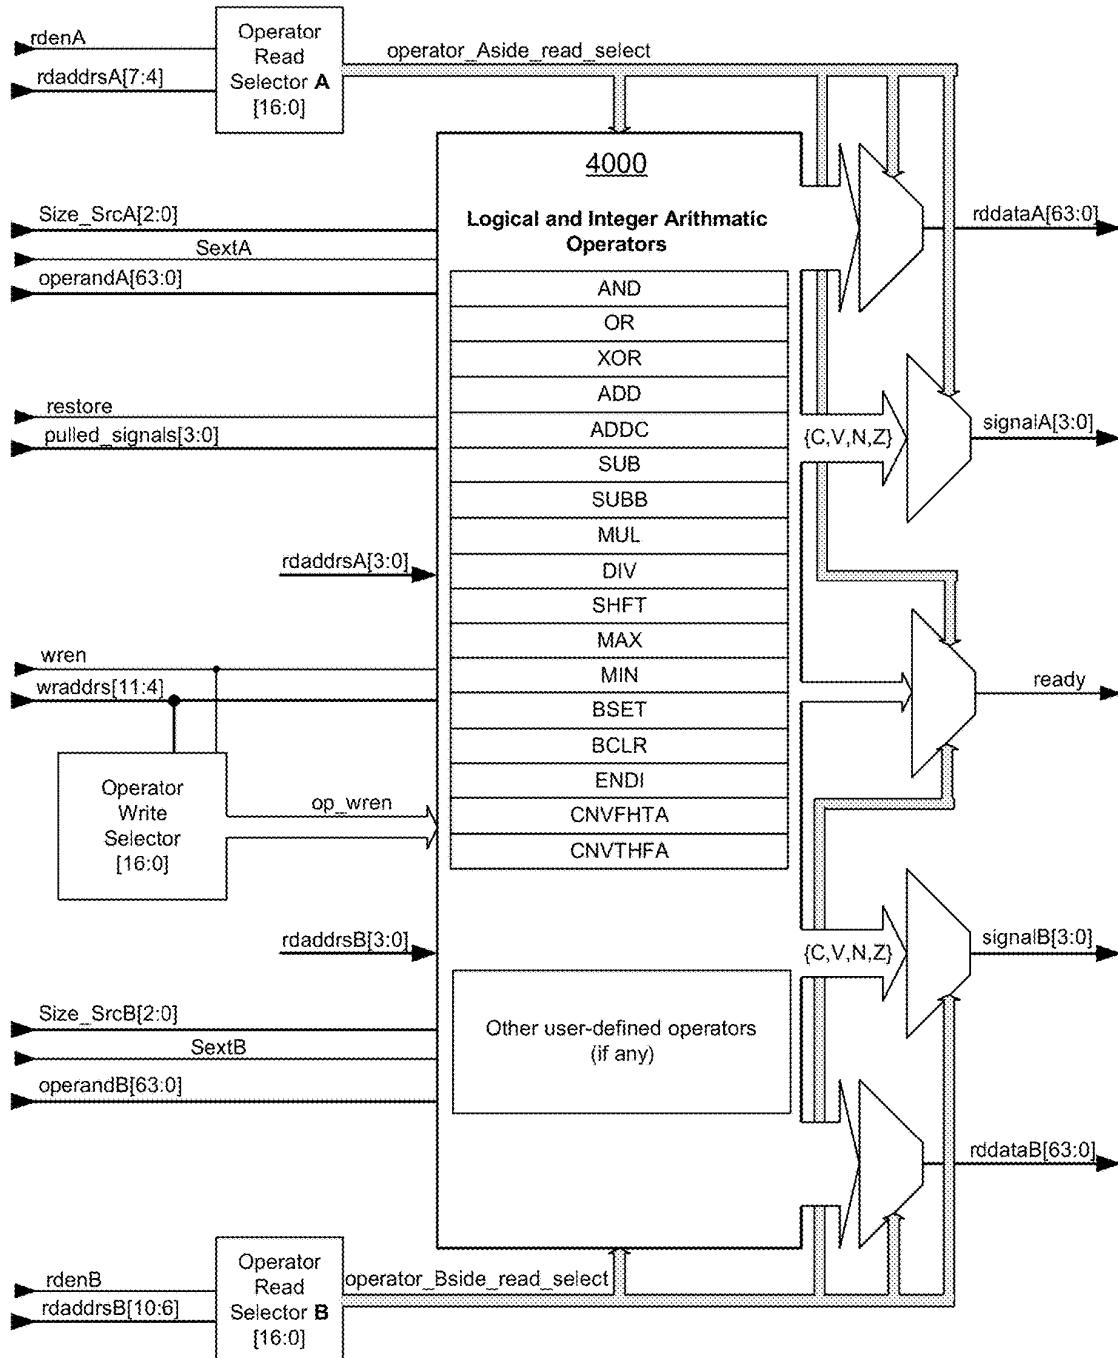


FIG. 23

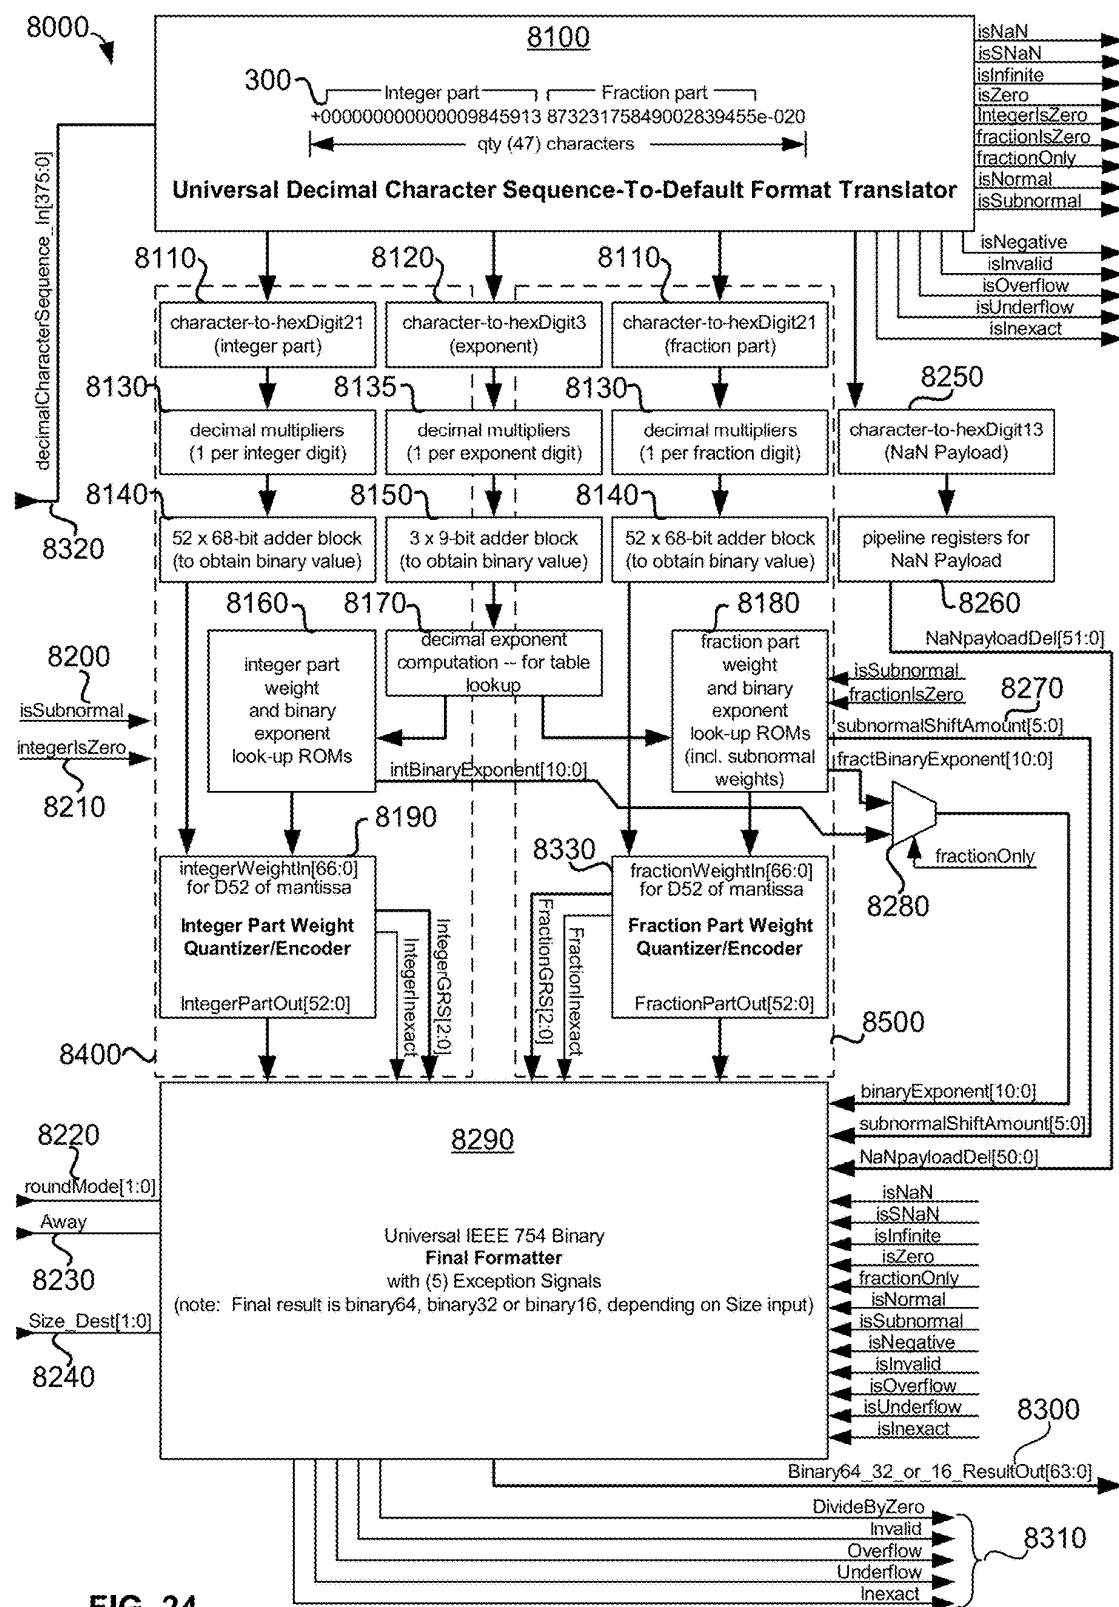


FIG. 24

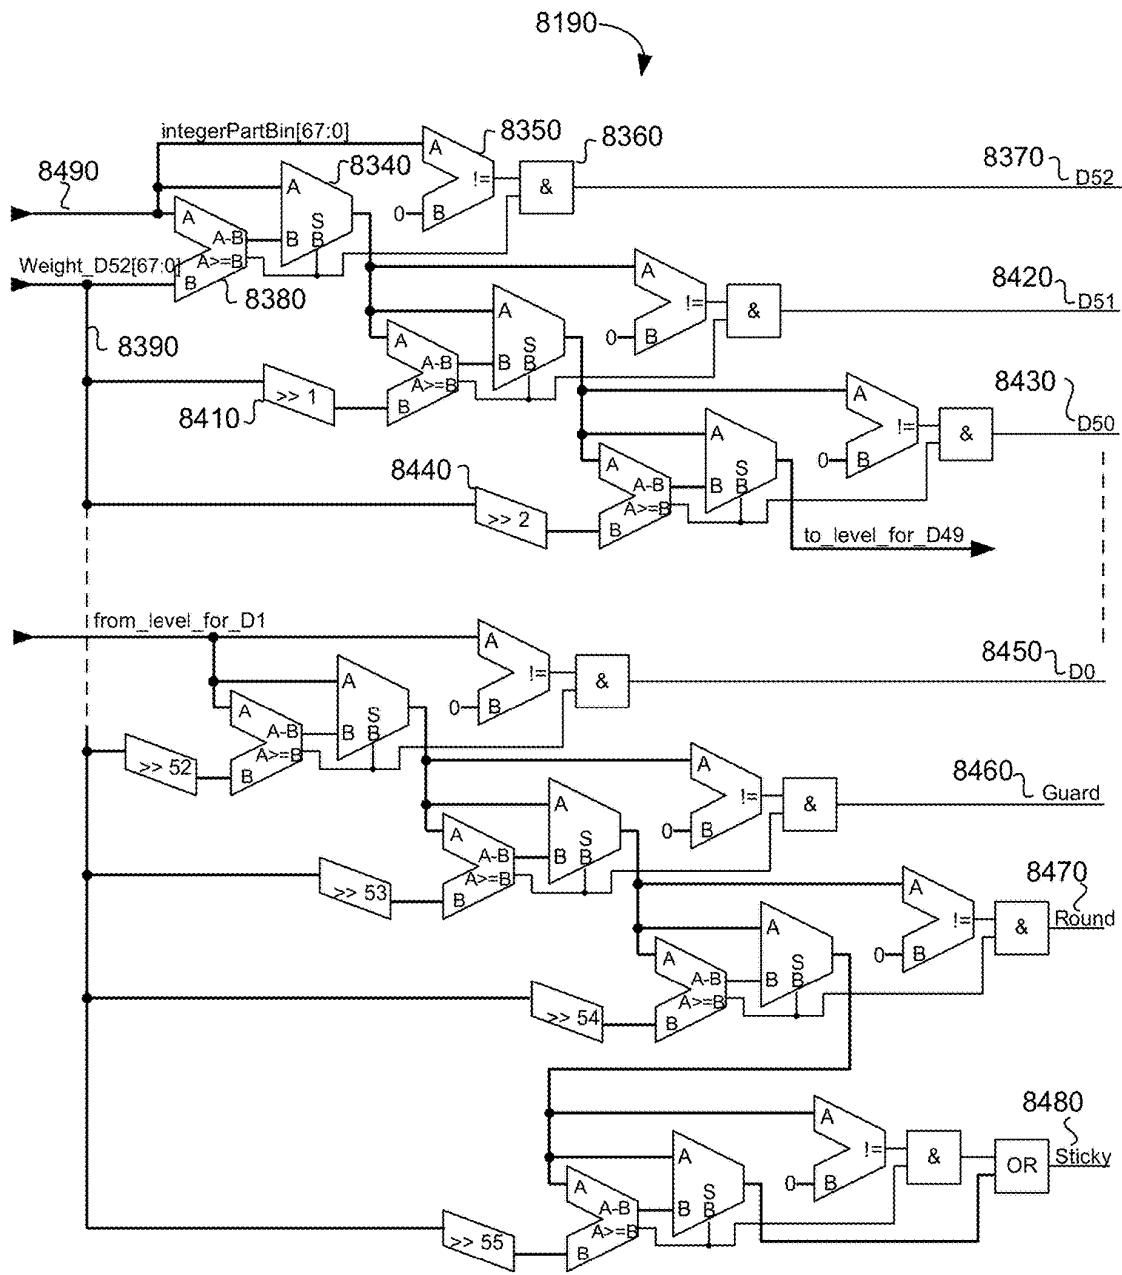


FIG. 25

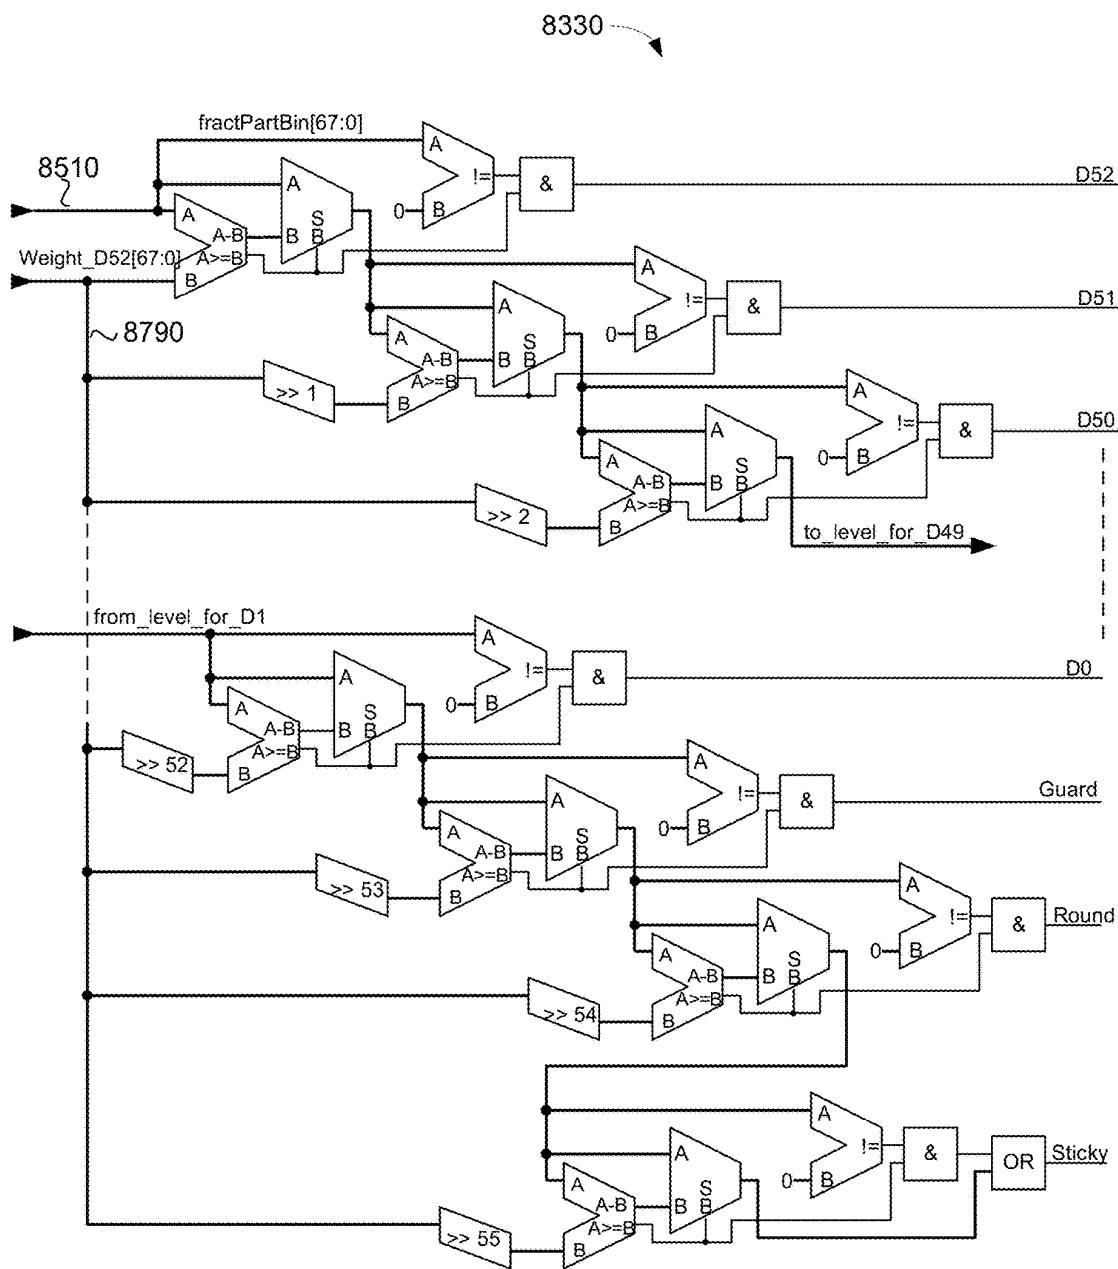


FIG. 26

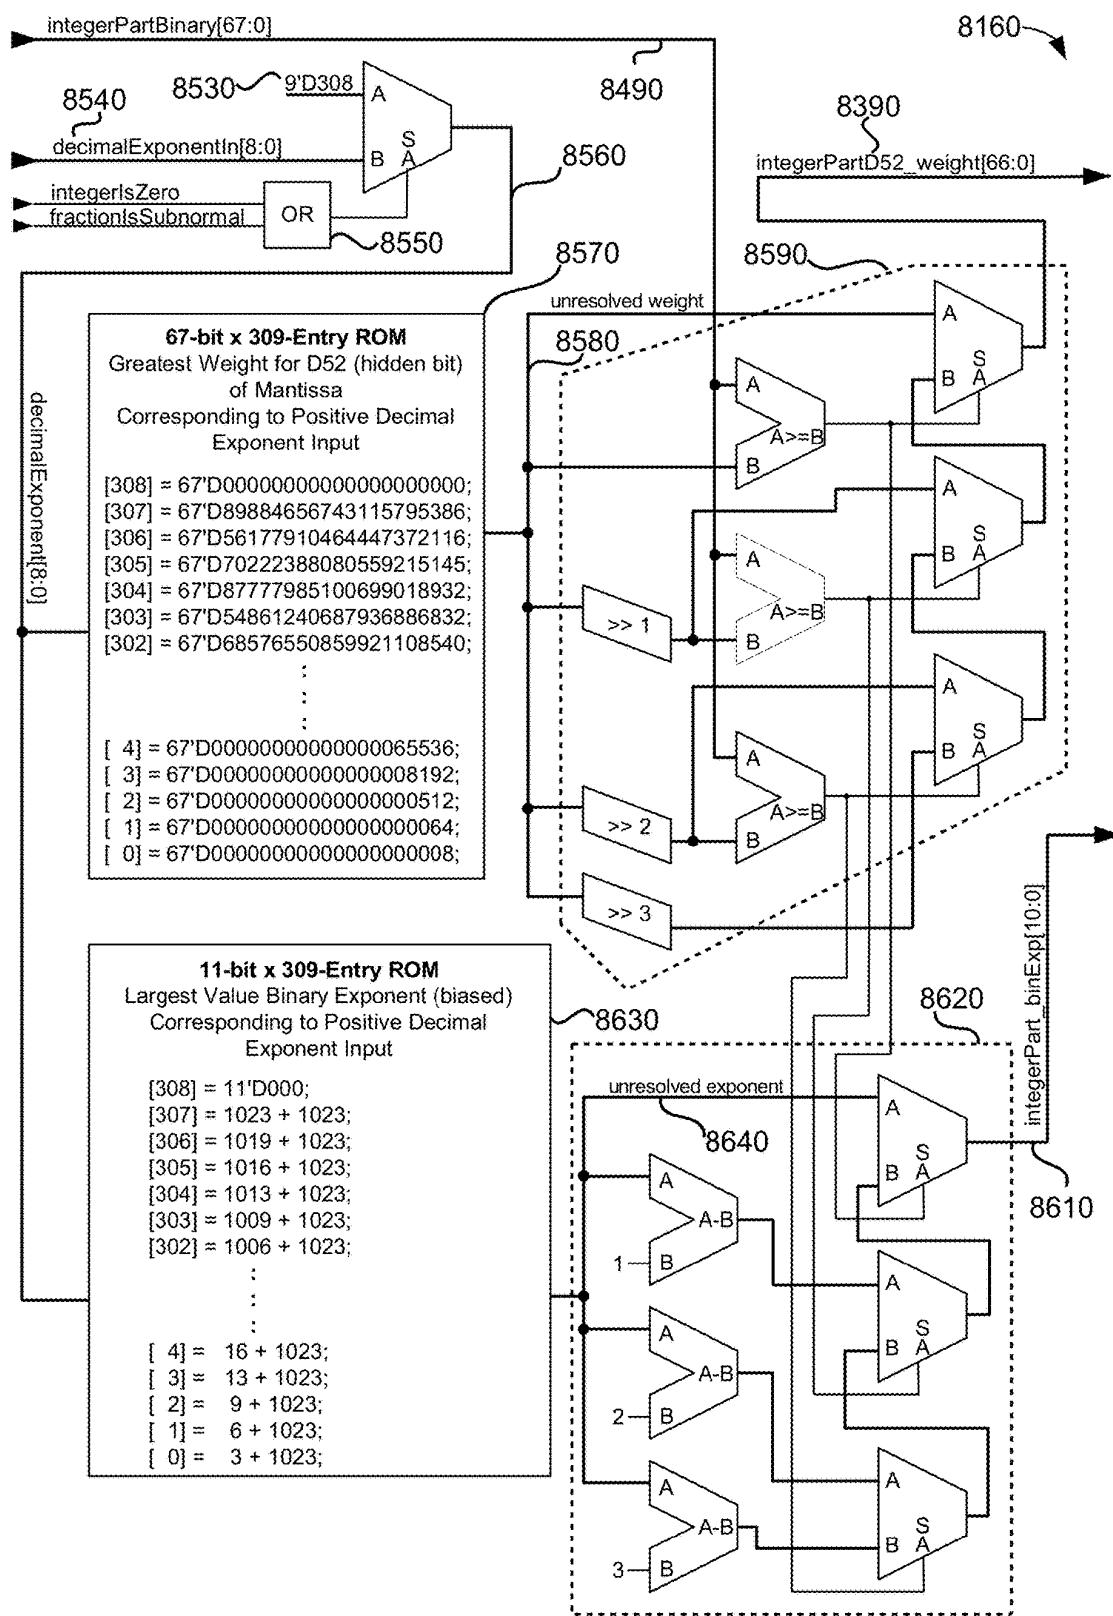


FIG. 27

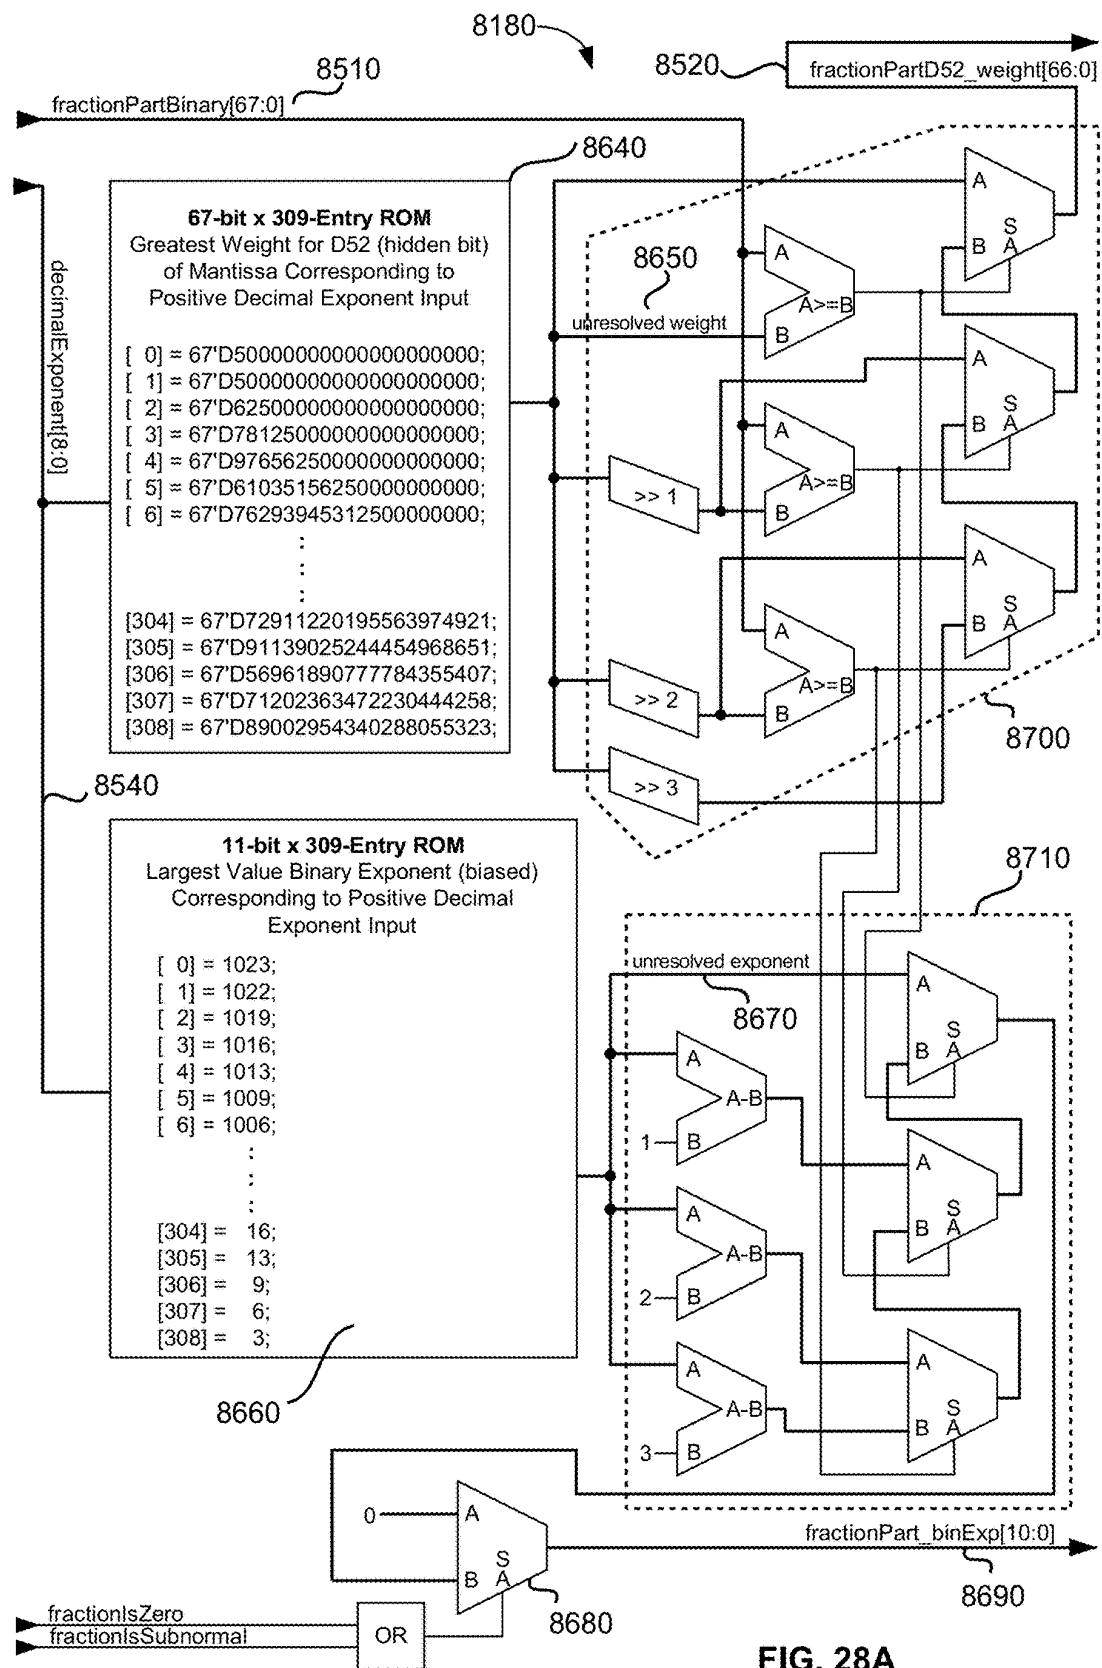
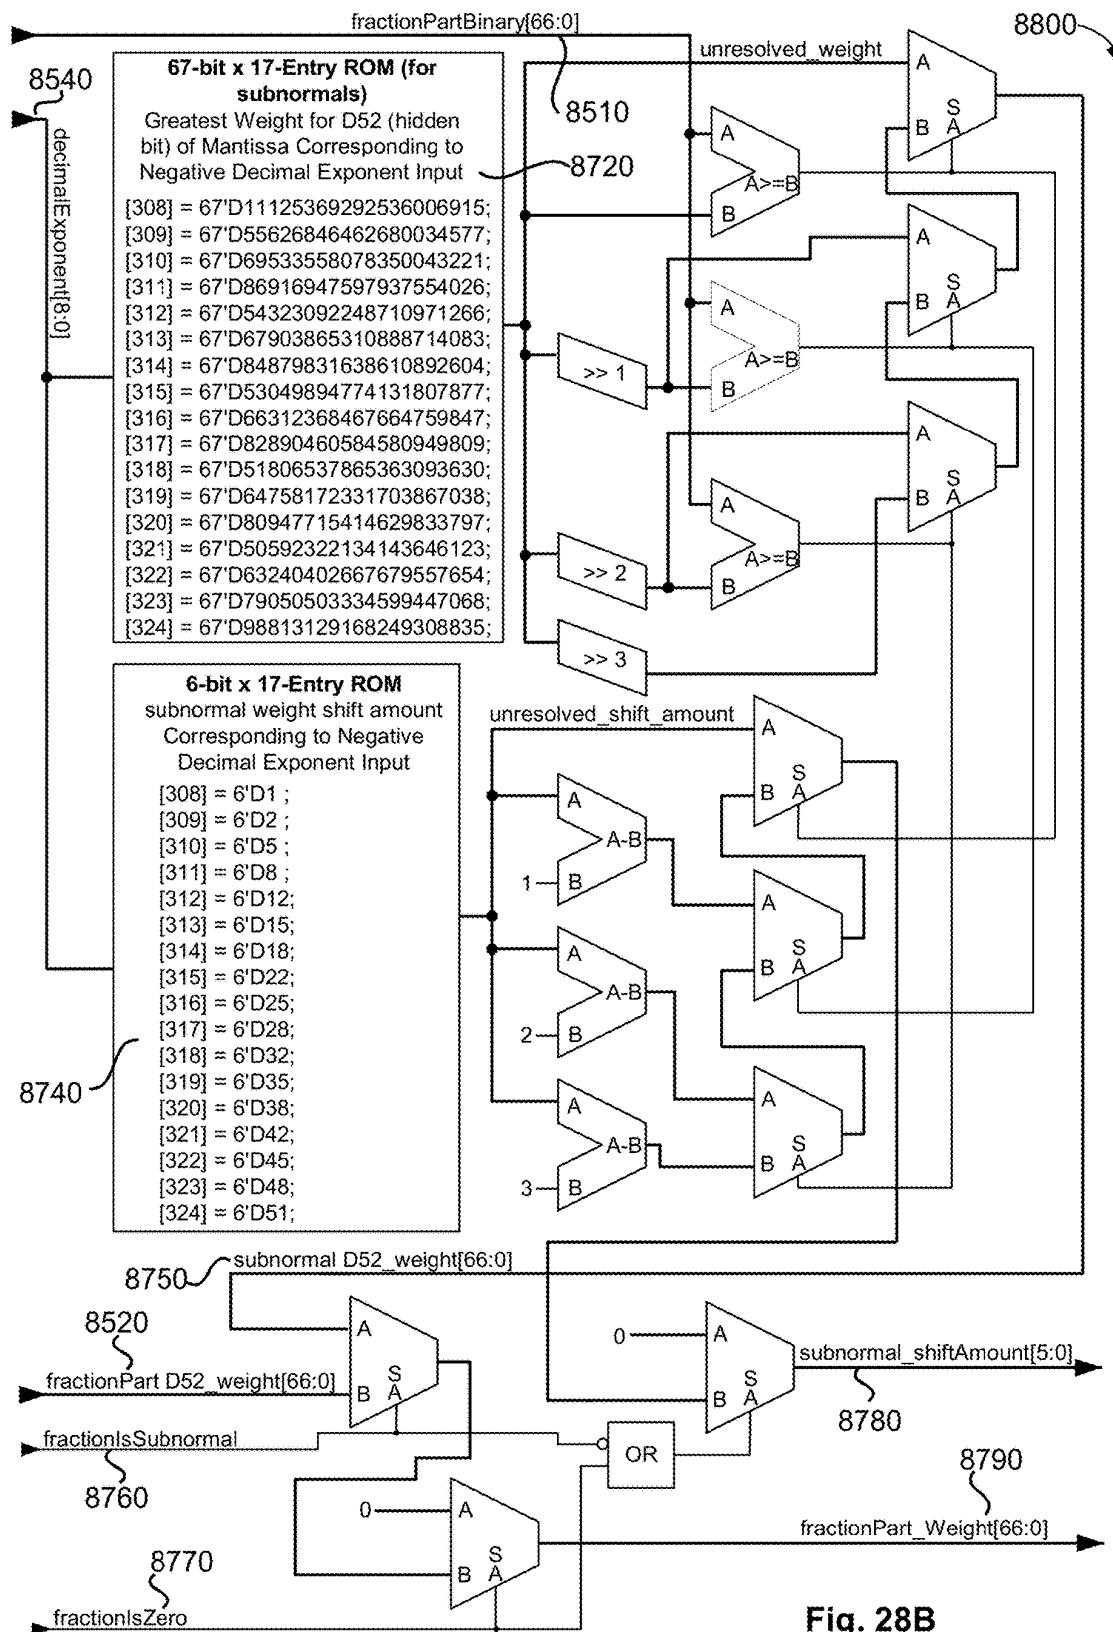


FIG. 28A



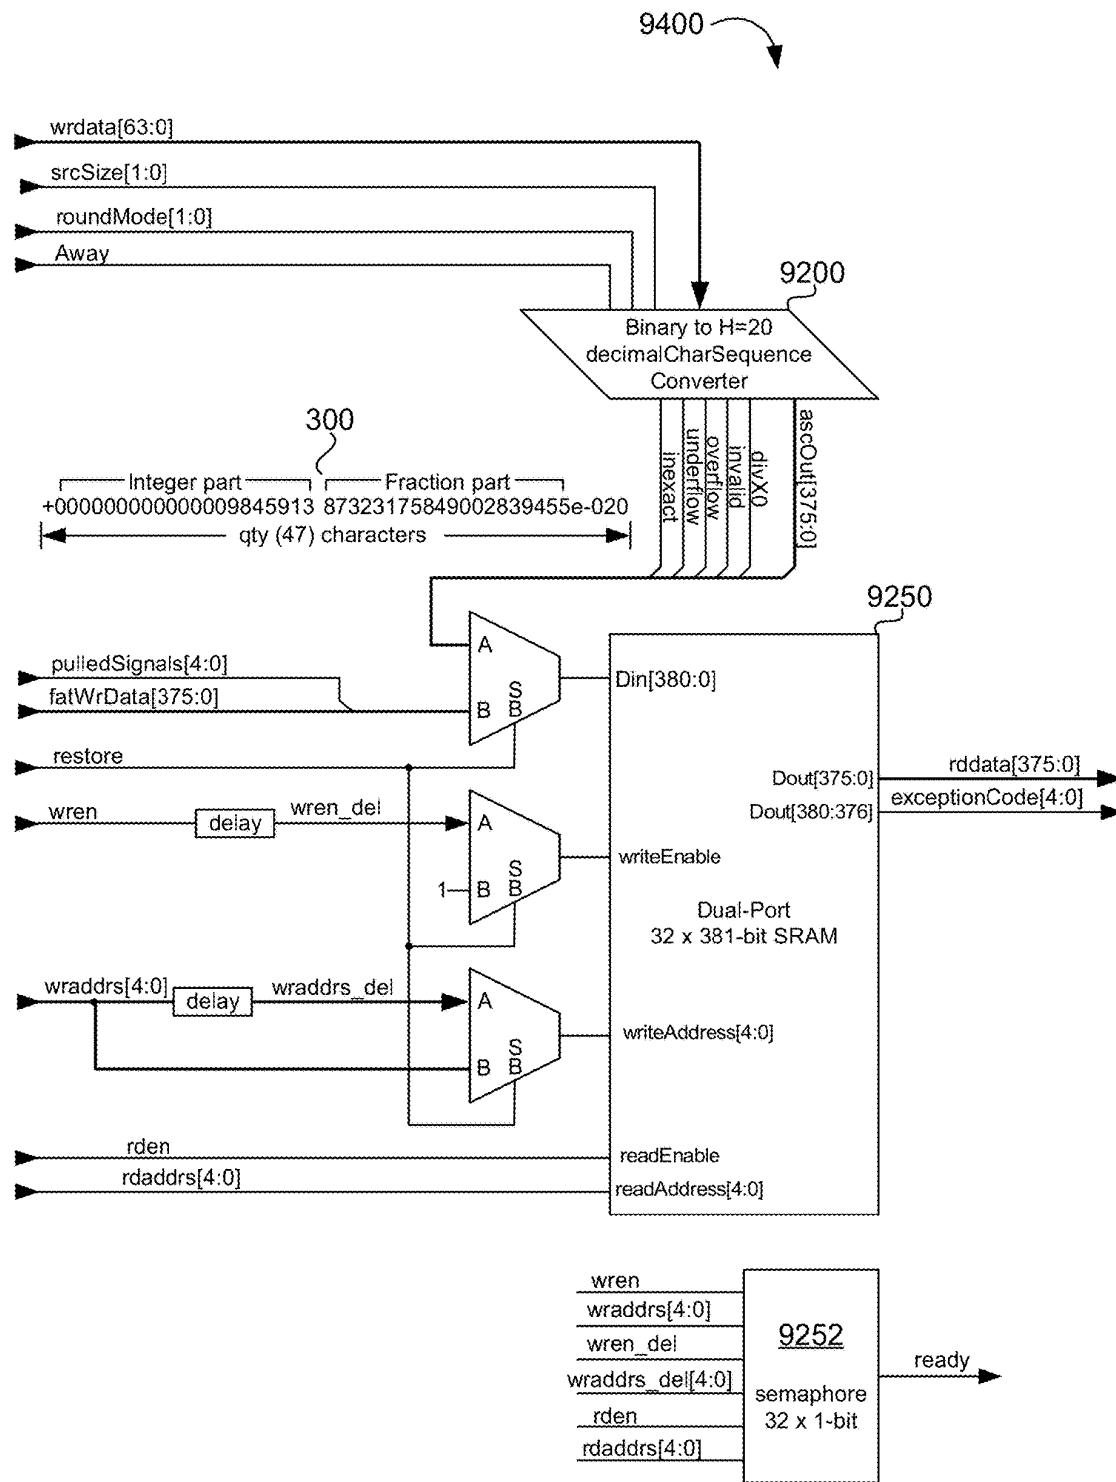


FIG. 29

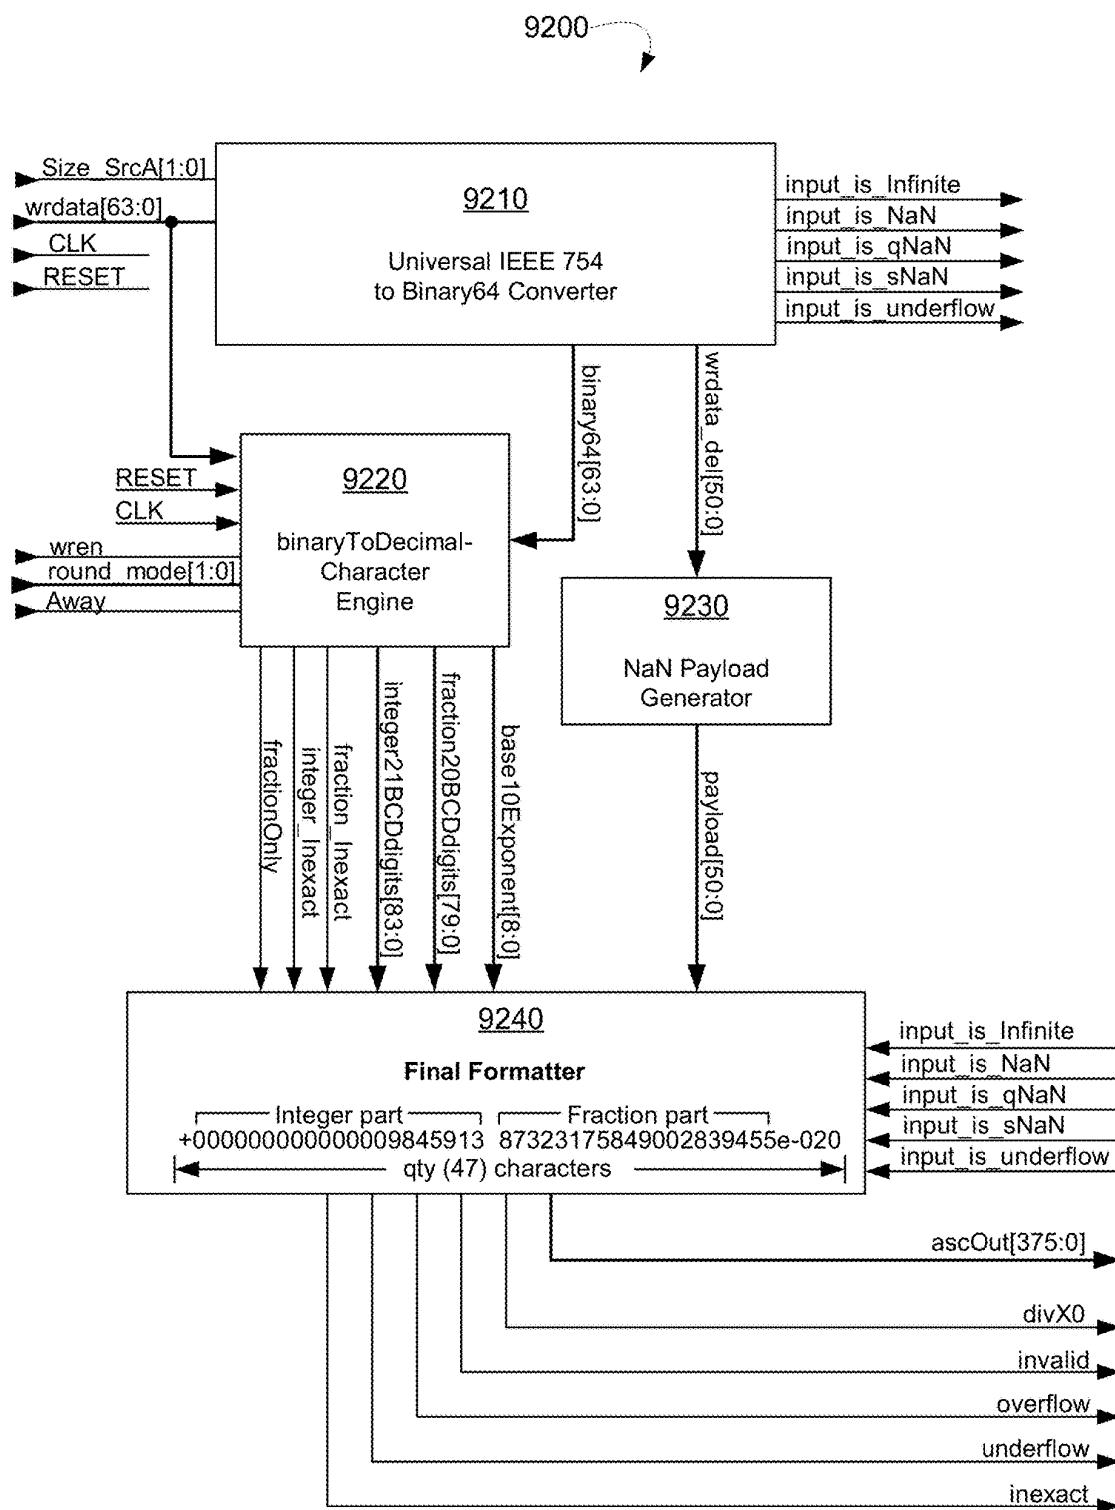


FIG. 30

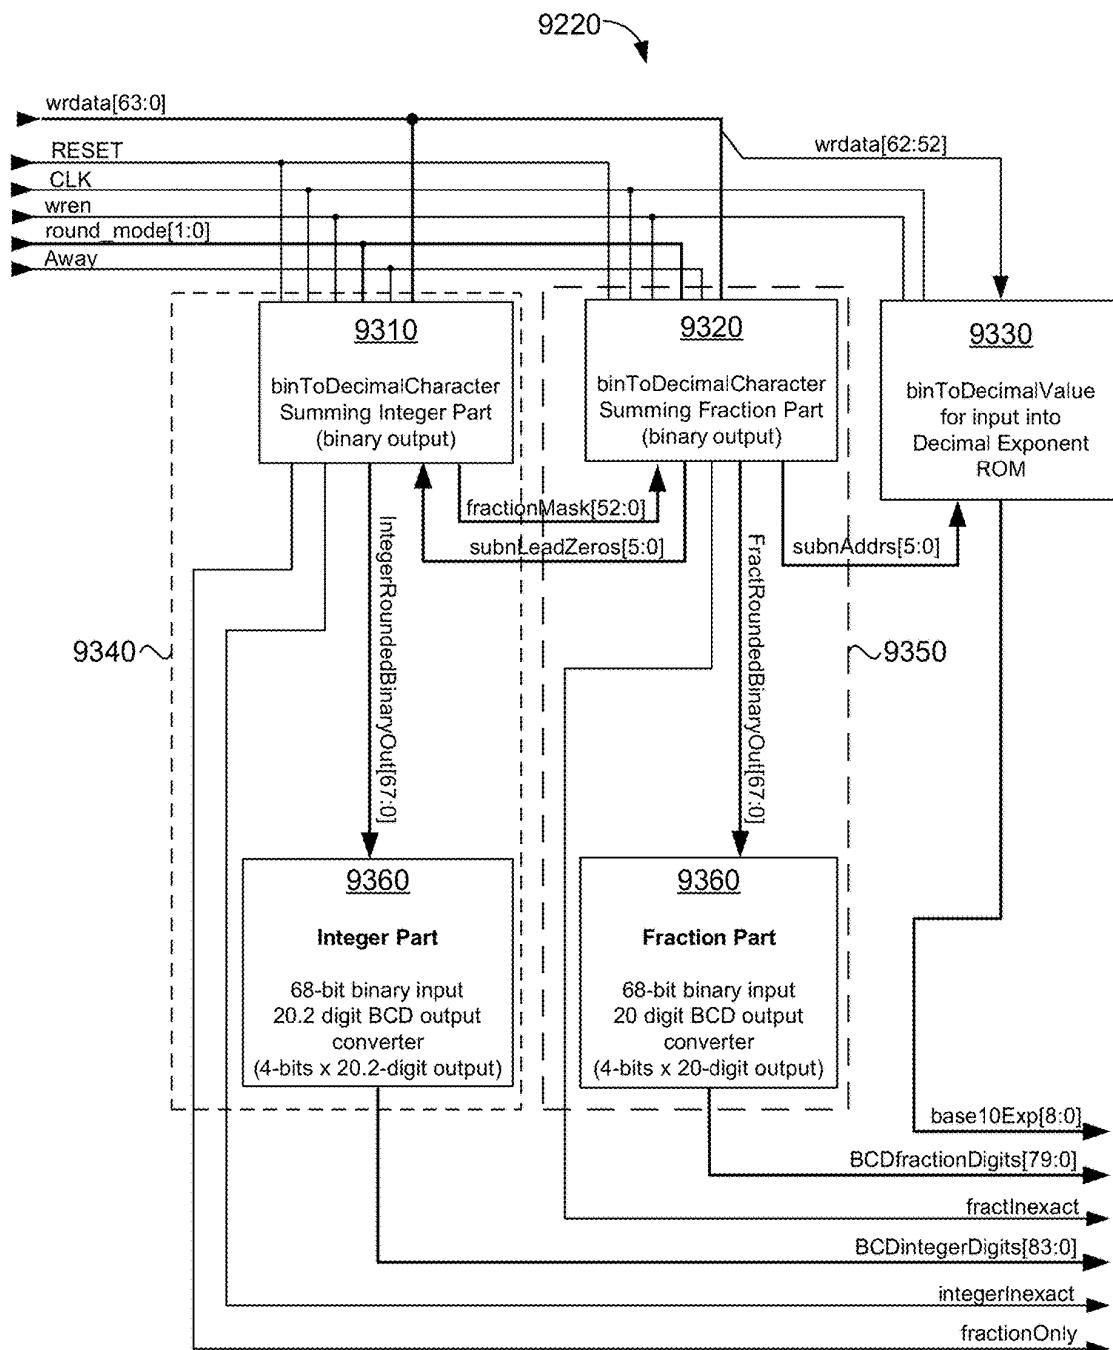


FIG. 31

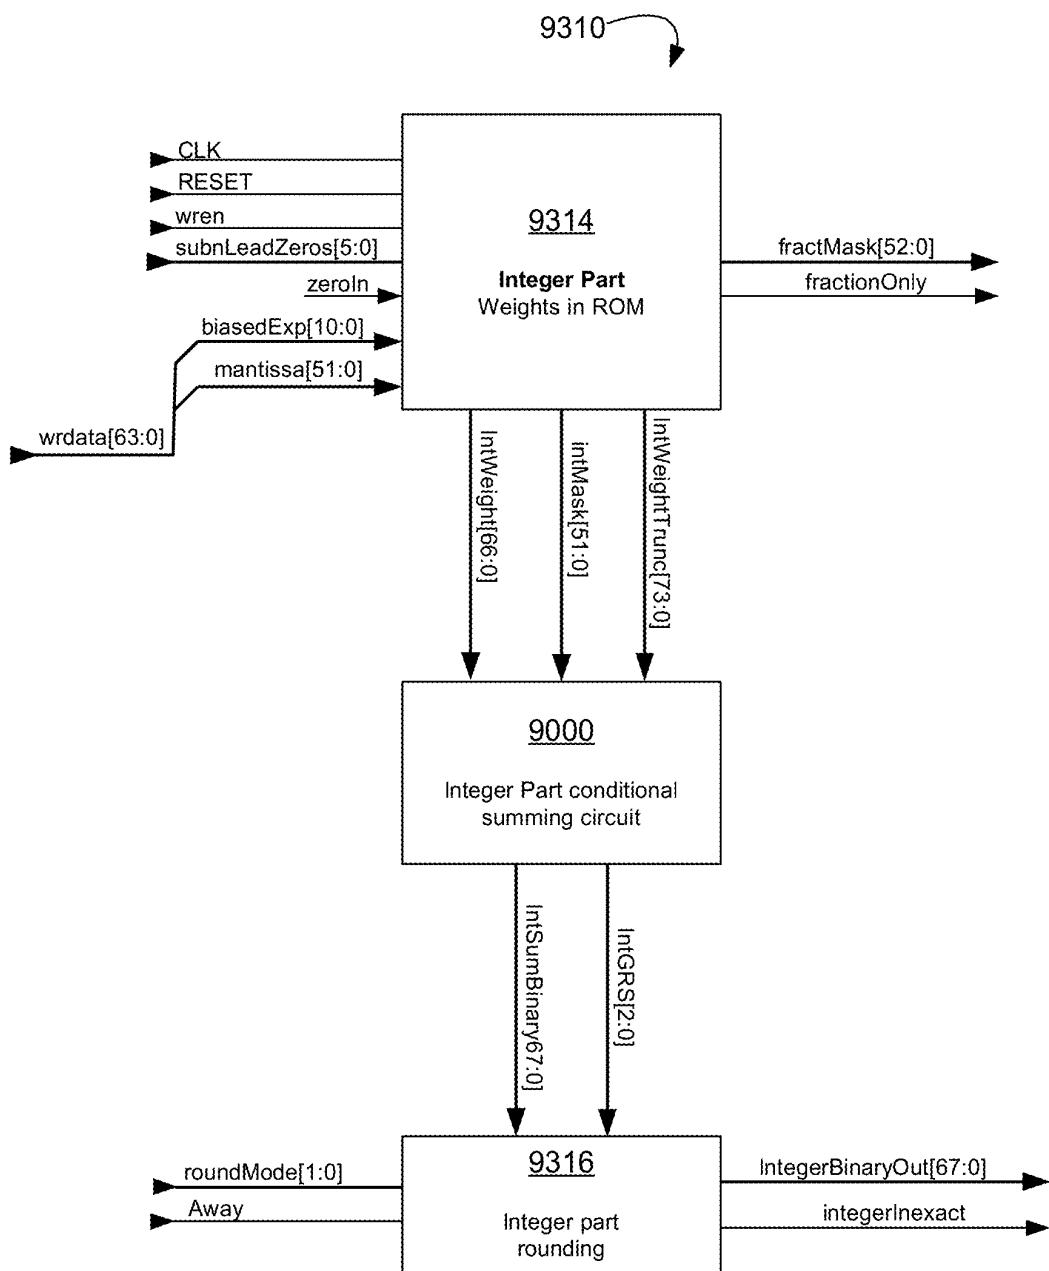


FIG. 32

FIG. 33

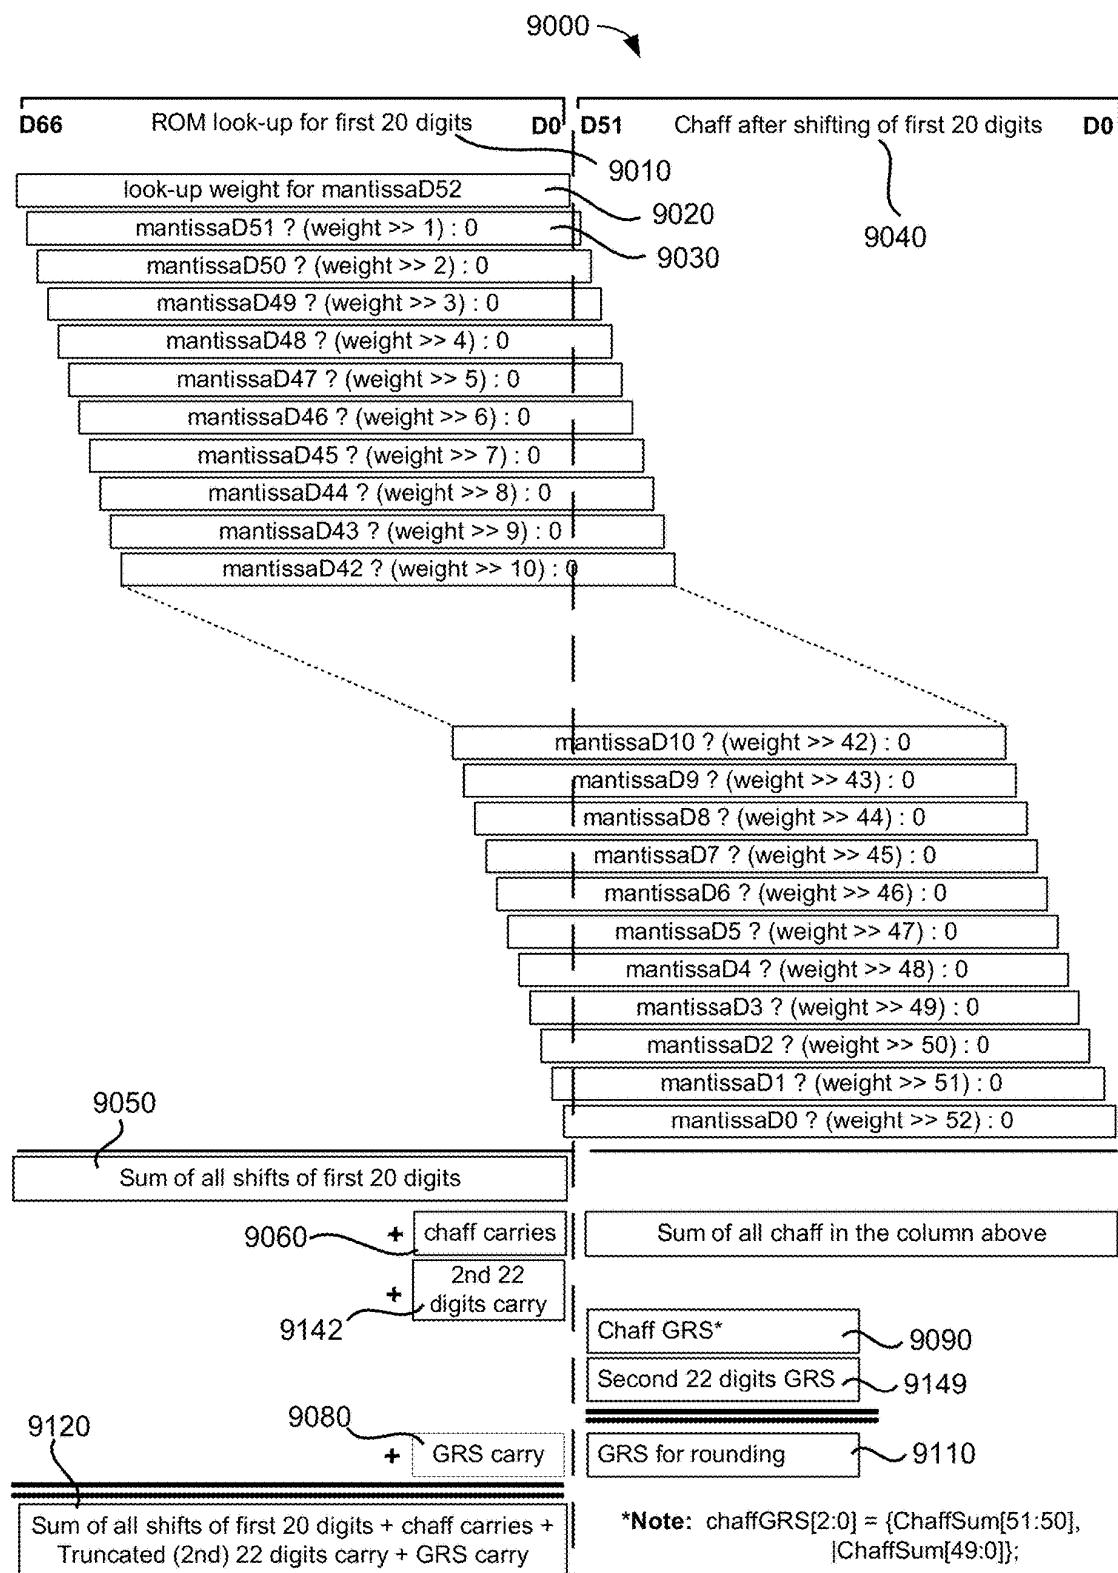
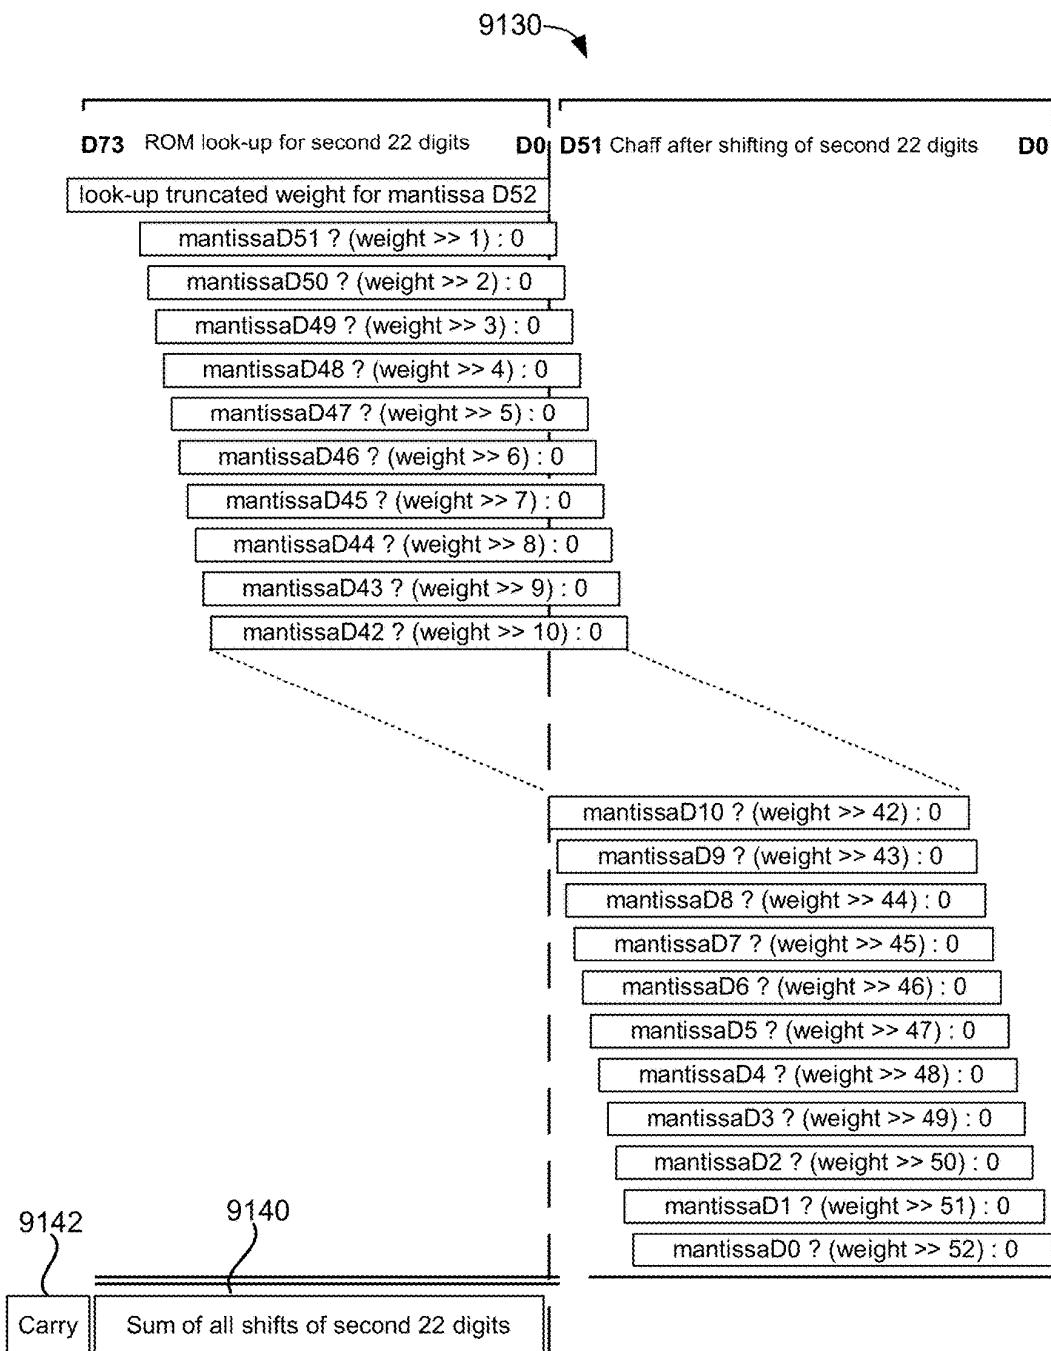


FIG. 34A



9144 ~ Trunc_G = (TruncSum[73:0] >= 74'D50000000000000000000000000000);
 9146 ~ Trunc_R = ((TruncSum[73:0] - 74'D50000000000000000000000000000) >= 74'D05000000000000000000000000000);
 9148 ~ Trunc_S = ((TruncSum[73:0] - 74'D55000000000000000000000000000) >= 74'D00000000000000000000000000001);
 9149 ~ Trunc_GRS = {Trunc_G, Trunc_R, Trunc_S};

FIG. 34B

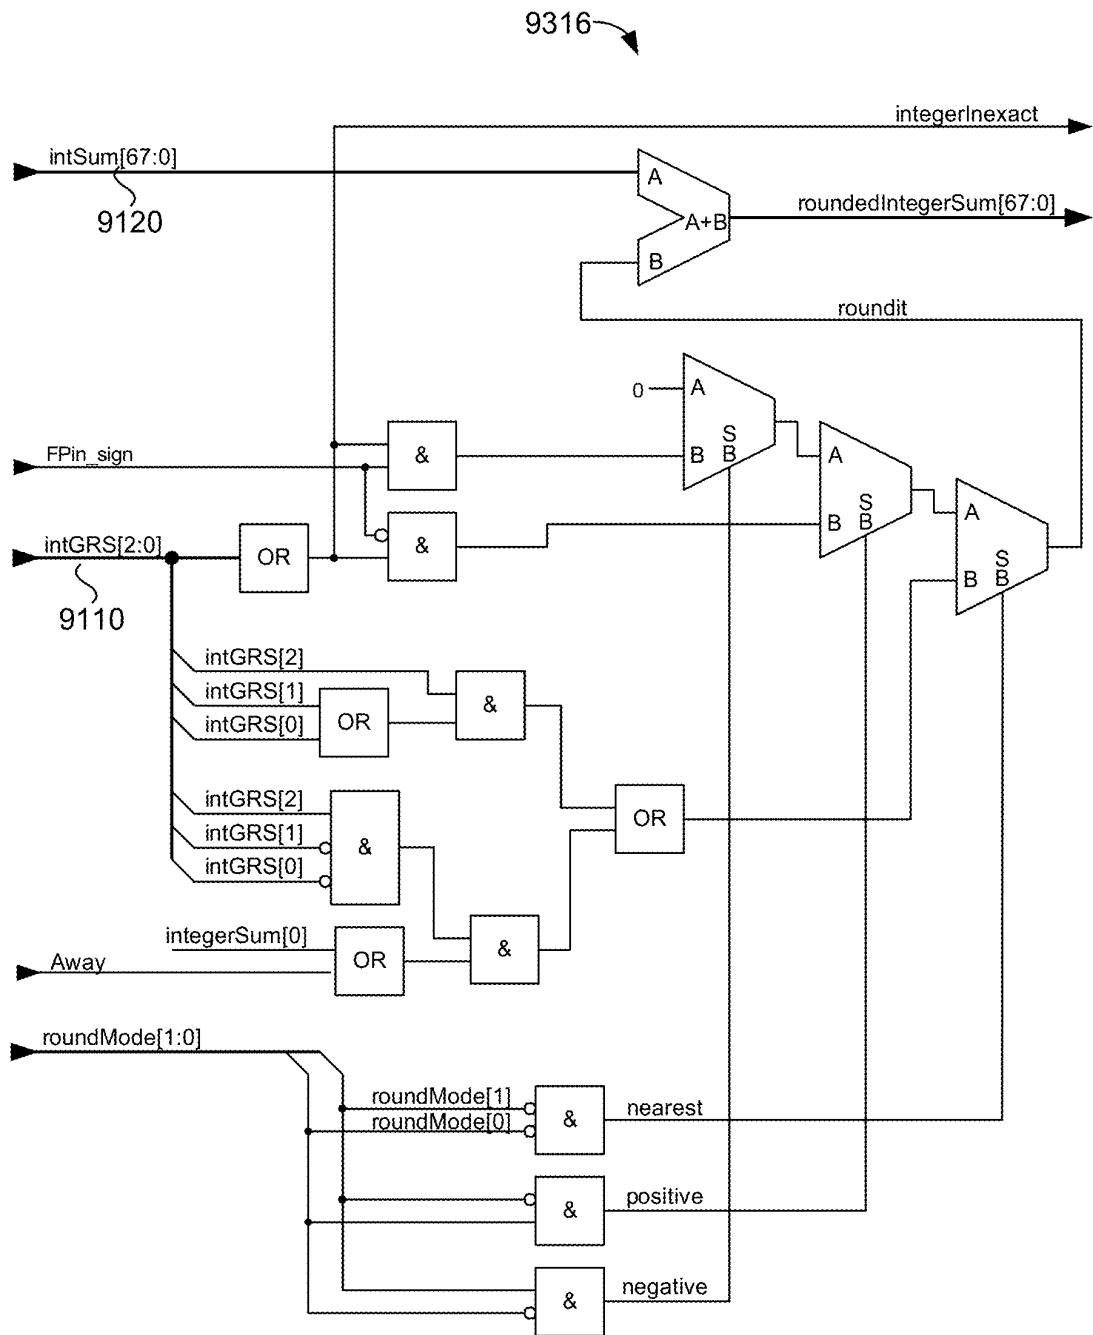


FIG. 35

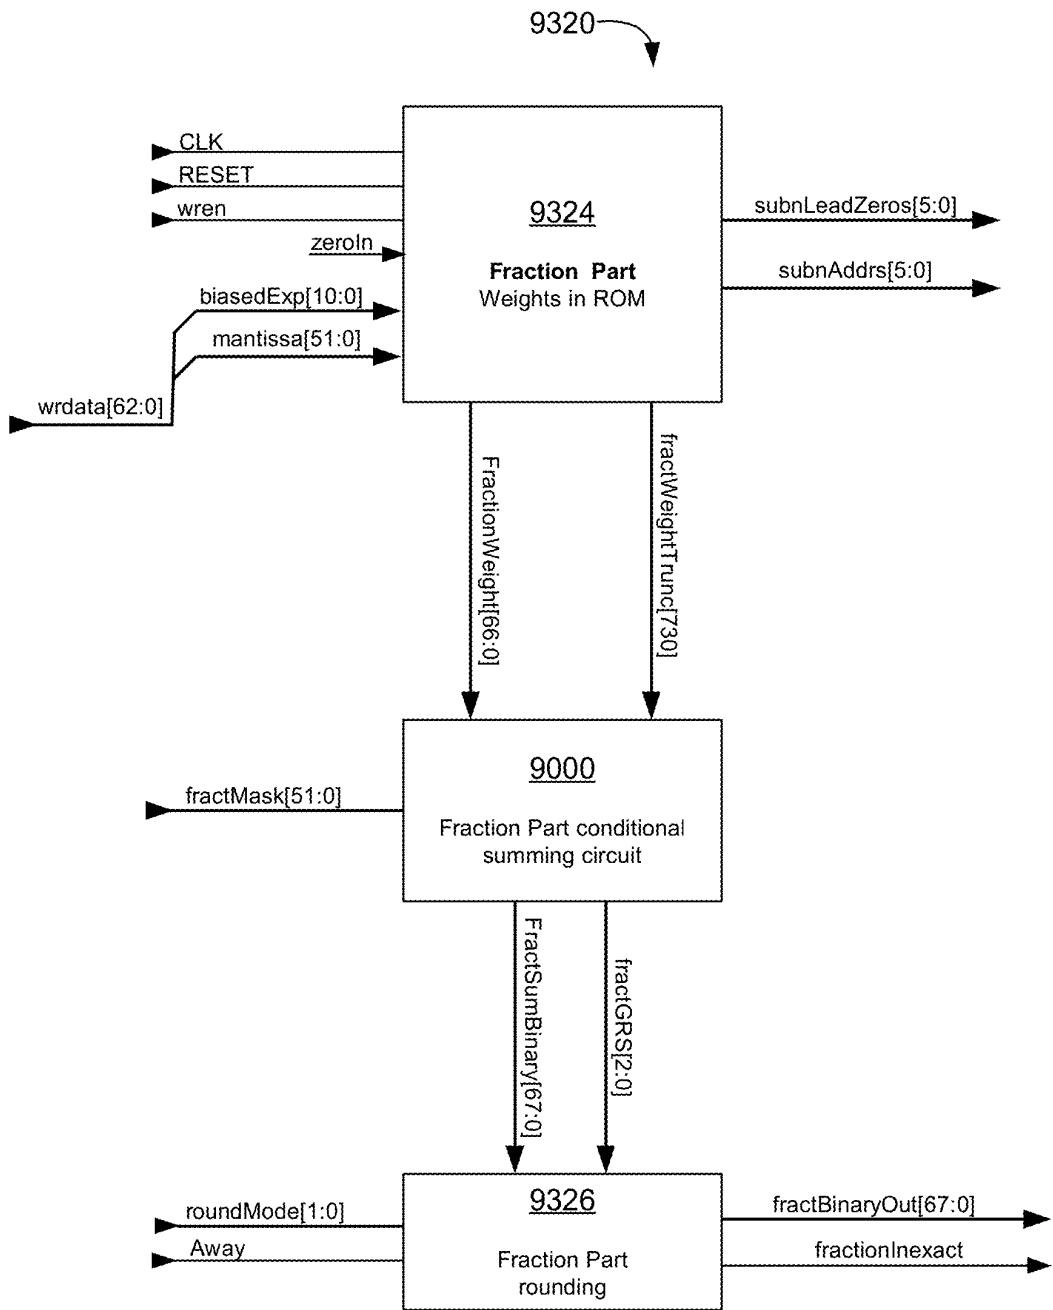


FIG. 36

FIG. 37

9330

```

module baseExpROM (
    CLK,
    RESET,
    rden,
    rdaddrs,
    baseExp,
    subnAddrs
);

input CLK;
input RESET;
input rden;
input {10:0} rdaddrs;
output [8:0] baseExp;
input [5:0] subnAddrs;

reg [8:0] RAMA[2047:0]; // 9 bits x 2048
reg [8:0] RAMB[51:0]; // 9 bits x 52
reg [8:0] baseExp_q;
reg [8:0] subnExp_q;
reg subnormal;

reg [5:0] subnAddrs_del_1,
    subnAddrs_del_2,
    :
    :
    :
    subnAddrs_del_15;

wire [8:0] baseExp;
assign baseExp = subnormal ? subnExp_q : baseExp_q;

always @ (posedge CLK) begin
    if (RESET) baseExp_q <= 0;
    else if (rden) baseExp_q <= RAMA[rdaddrs];
end

always @ (posedge CLK)
    if (RESET) subnormal <= 0;
    else subnormal <= ~rdaddrs;

always @ (posedge CLK)
    if (RESET) subnExp_q <= 0;
    else if (rden && ~rdaddrs) subnExp_q <=
        RAMB[subnAddrs_del_15];
end

always @ (posedge CLK) begin
    if (RESET) begin
        subnAddrs_del_1 <= 0;
        :
        :
        subnAddrs_del_15 <= 0;
    end
    else begin
        subnAddrs_del_1 <= subnAddrs;
        :
        :
        subnAddrs_del_15 <= subnAddrs_del_14;
    end
end
endmodule

```

9336

**Binary To Decimal
Subnormal Exponent
Look-Up ROM
9 Bits x 52- Entry
(RAMB)**

```

initial begin
    RAMB[51] = 308; // -308
    RAMB[50] = 309; // -309
    RAMB[49] = 309; // -309
    RAMB[48] = 309; // -309
    RAMB[47] = 310; // -310
    RAMB[46] = 310; // -310
    RAMB[45] = 310; // -310
    RAMB[44] = 311; // -311
    RAMB[43] = 311; // -311
    RAMB[42] = 311; // -311
    RAMB[41] = 311; // -311
    RAMB[40] = 312; // -312
    RAMB[39] = 312; // -312
    RAMB[38] = 312; // -312
    RAMB[37] = 313; // -313
    RAMB[36] = 313; // -313
    RAMB[35] = 313; // -313
    RAMB[34] = 314; // -314
    RAMB[33] = 314; // -314
    RAMB[32] = 314; // -314
    RAMB[31] = 314; // -314
    RAMB[30] = 315; // -315
    RAMB[29] = 315; // -315
    RAMB[28] = 315; // -315
    RAMB[27] = 316; // -316
    RAMB[26] = 316; // -316
    RAMB[25] = 316; // -316
    RAMB[24] = 317; // -317
    RAMB[23] = 317; // -317
    RAMB[22] = 317; // -317
    RAMB[21] = 317; // -317
    RAMB[20] = 318; // -318
    RAMB[19] = 318; // -318
    RAMB[18] = 318; // -318
    RAMB[17] = 319; // -319
    RAMB[16] = 319; // -319
    RAMB[15] = 319; // -319
    RAMB[14] = 320; // -320
    RAMB[13] = 320; // -320
    RAMB[12] = 320; // -320
    RAMB[11] = 320; // -320
    RAMB[10] = 321; // -321
    RAMB[ 9] = 321; // -321
    RAMB[ 8] = 321; // -321
    RAMB[ 7] = 322; // -322
    RAMB[ 6] = 322; // -322
    RAMB[ 5] = 322; // -322
    RAMB[ 4] = 323; // -323
    RAMB[ 3] = 323; // -323
    RAMB[ 2] = 323; // -323
    RAMB[ 1] = 324; // -324
    RAMB[ 0] = 324; // -324
end

```

9334

**Adjusted Binary To Decimal
Normal Exponent Look-Up ROM
9 Bits x 2049- Entry
(RAMA)**

```

initial begin
    RAMA[2047] = 511; //NaN
    RAMA[2046] = 307 - 19;
    RAMA[2045] = 307 - 19;
    RAMA[2044] = 307 - 19;
    RAMA[2043] = 307 - 19;
    RAMA[2042] = 306 - 19;
    RAMA[2041] = 306 - 19;
    RAMA[2040] = 306 - 19;
    RAMA[2039] = 305 - 19;
    RAMA[2038] = 305 - 19;
    RAMA[2037] = 305 - 19;
    :
    :
    RAMA[1090] = 20 - 19;
    RAMA[1089] = 19 - 19;
    RAMA[1088] = 19 - 19;
    RAMA[1087] = 19 - 19;
    RAMA[1086] = 0; // 18;
    RAMA[1085] = 0; // 18;
    RAMA[1084] = 0; // 18;
    RAMA[1083] = 0; // 18;
    RAMA[1082] = 0; // 17;
    RAMA[1081] = 0; // 17;
    RAMA[1080] = 0; // 17;
    RAMA[1079] = 0; // 16;
    RAMA[1078] = 0; // 16;
    RAMA[1077] = 0; // 16;
    RAMA[1076] = 0; // 15;
    RAMA[1075] = 0; // 15;
    :
    :
    RAMA[1026] = 0; // 0;
    RAMA[1025] = 0; // 0;
    RAMA[1024] = 0; // 0;
    RAMA[1023] = 0; // 0;
    RAMA[1022] = 1; // minus
    RAMA[1021] = 1; // +
    RAMA[1020] = 1; // +
    RAMA[1019] = 2; // v
    RAMA[1018] = 2;
    RAMA[1017] = 2;
    RAMA[1016] = 3;
    RAMA[1015] = 3;
    RAMA[1014] = 3;
    :
    :
    RAMA[ 6] = 307;
    RAMA[ 5] = 307;
    RAMA[ 4] = 307;
    RAMA[ 3] = 308;
    RAMA[ 2] = 308;
    RAMA[ 1] = 308;
    RAMA[ 0] = 308; //subnormal
end

```

FIG. 38

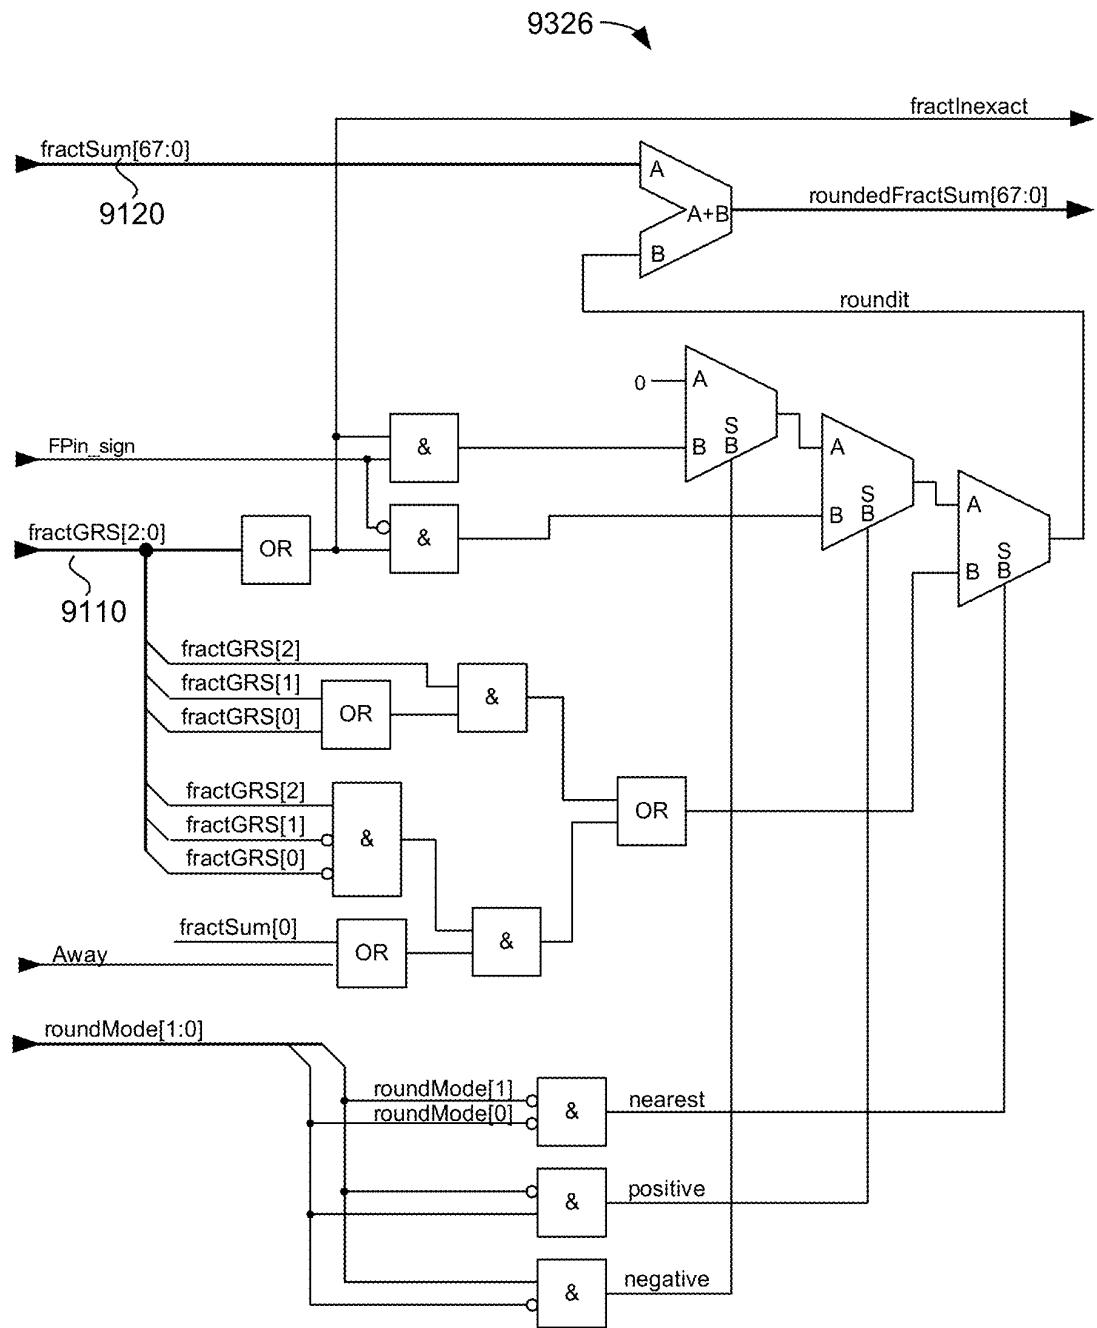


FIG. 39

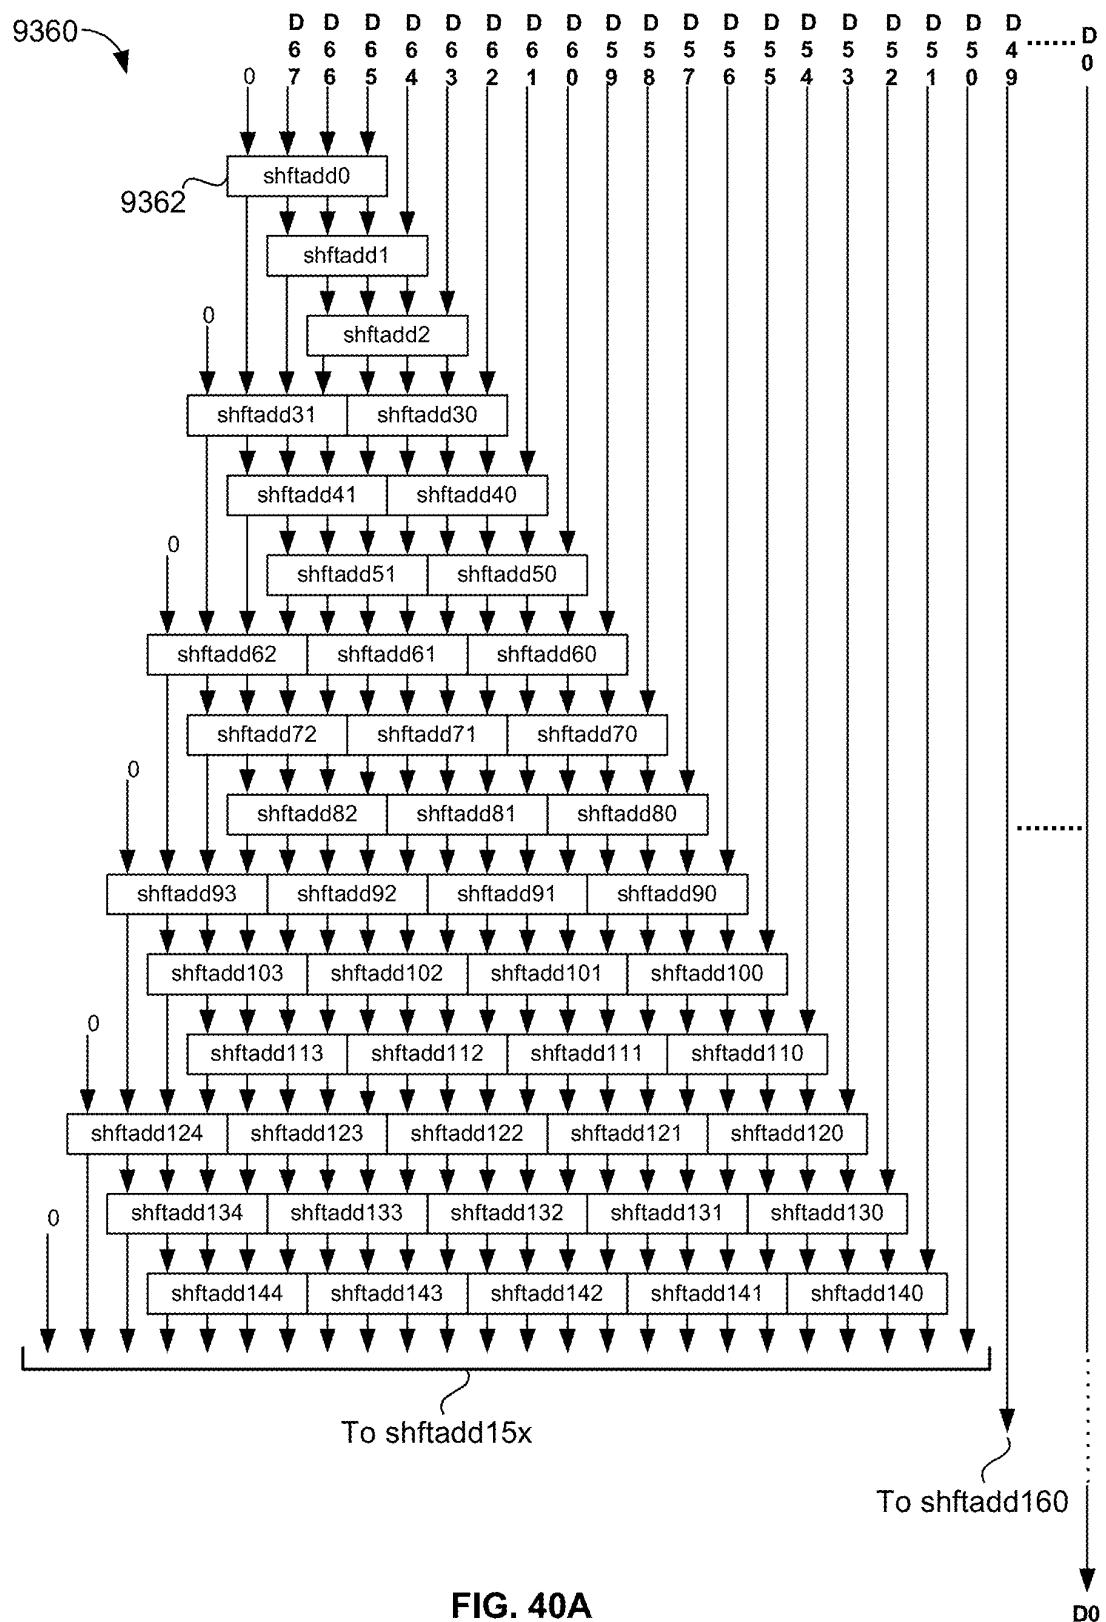


FIG. 40A

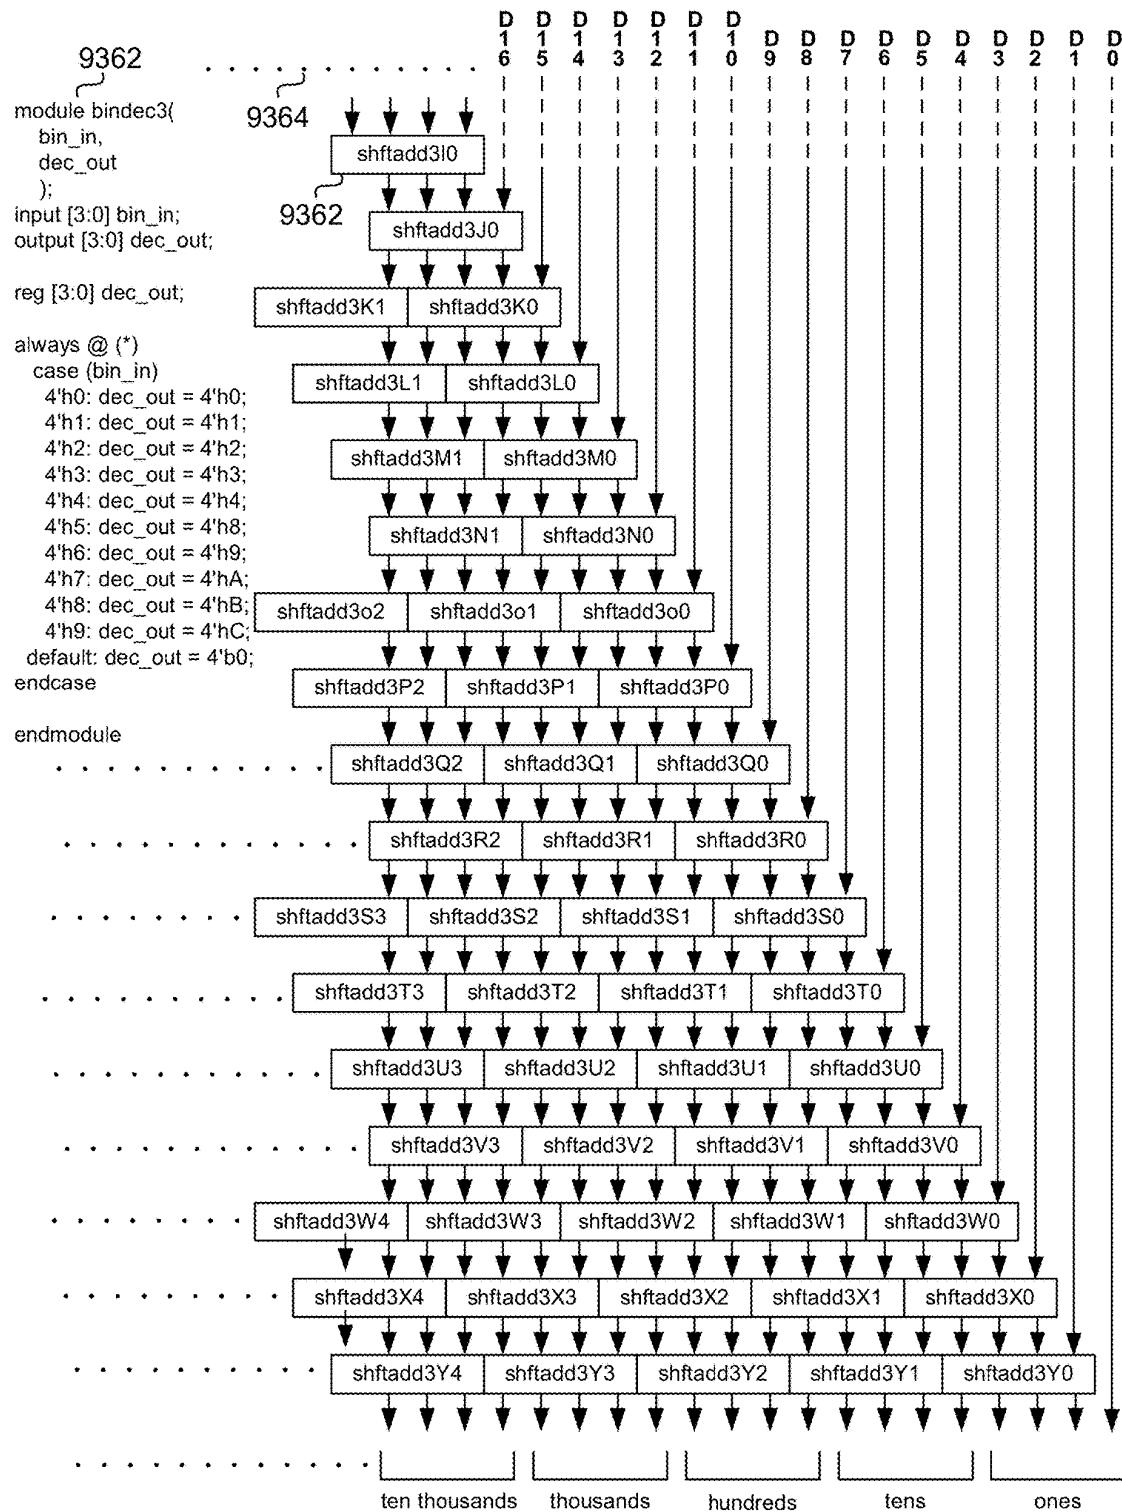


FIG. 40B

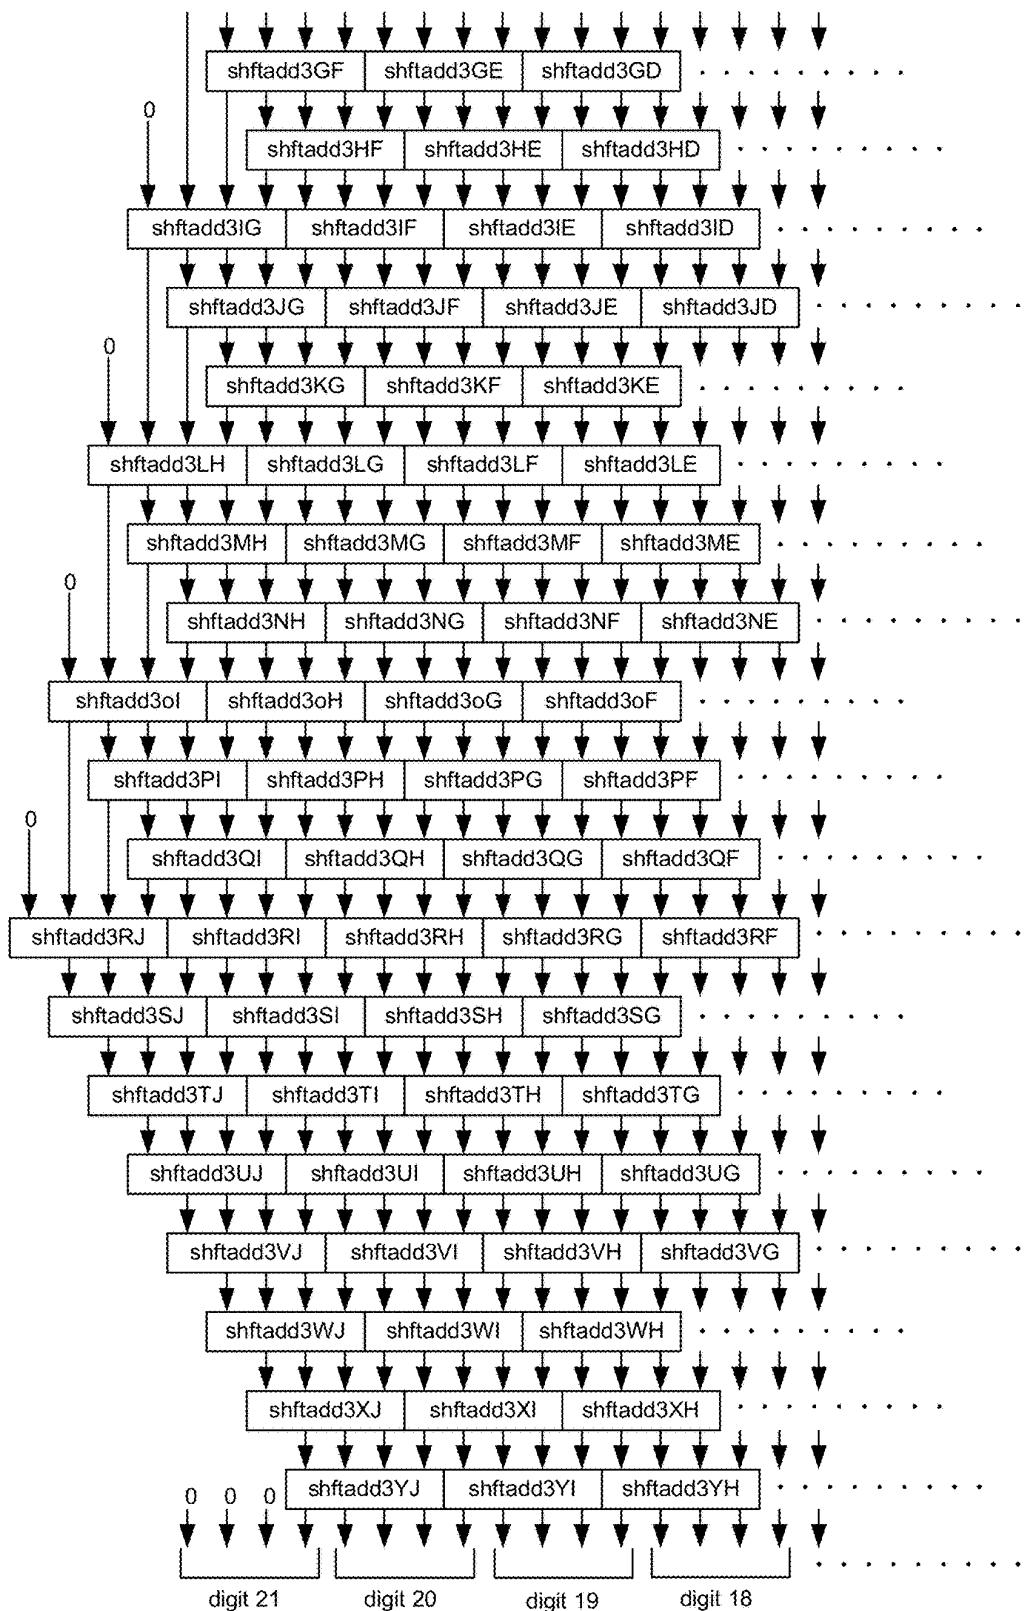


FIG. 40C

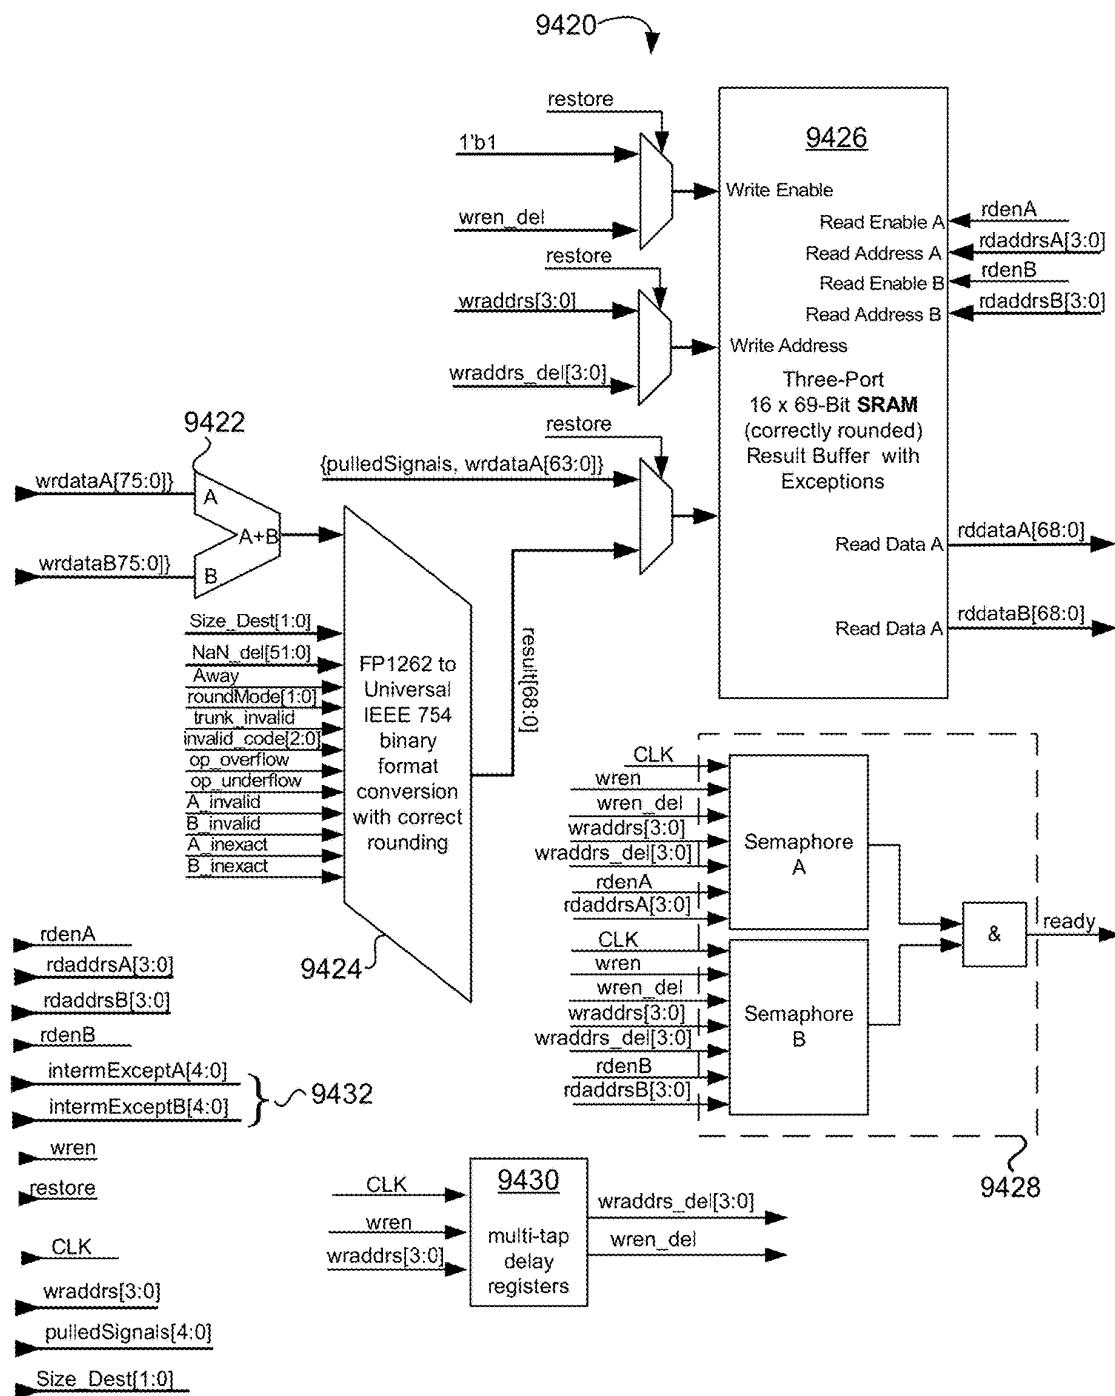


FIG. 41

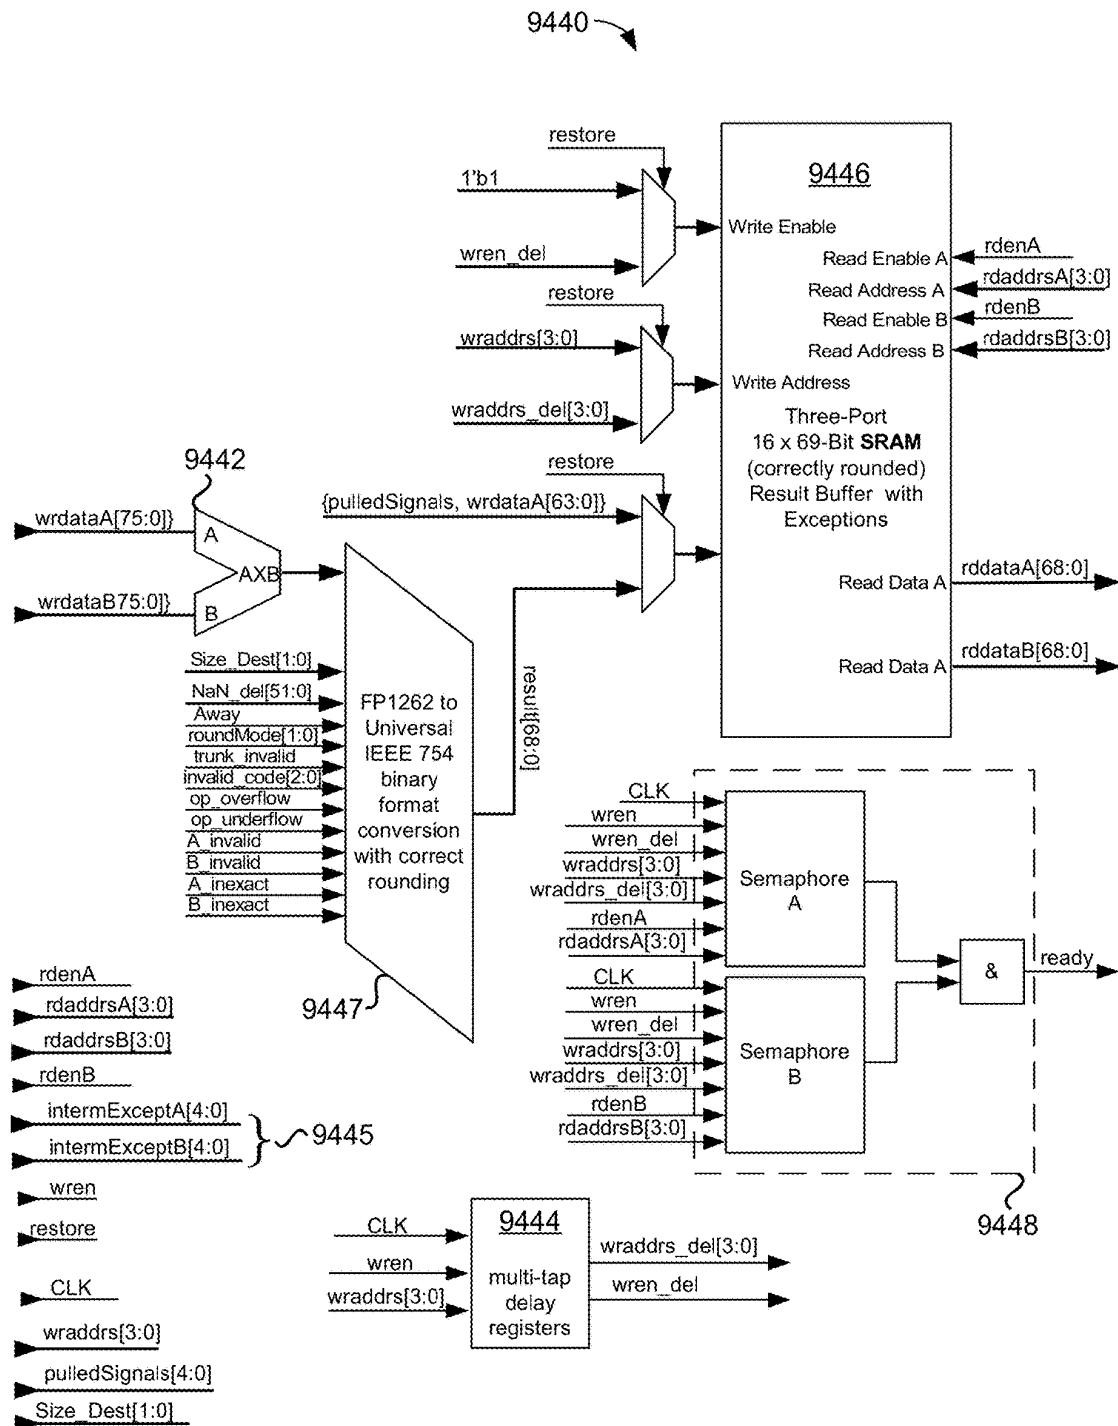


FIG. 42

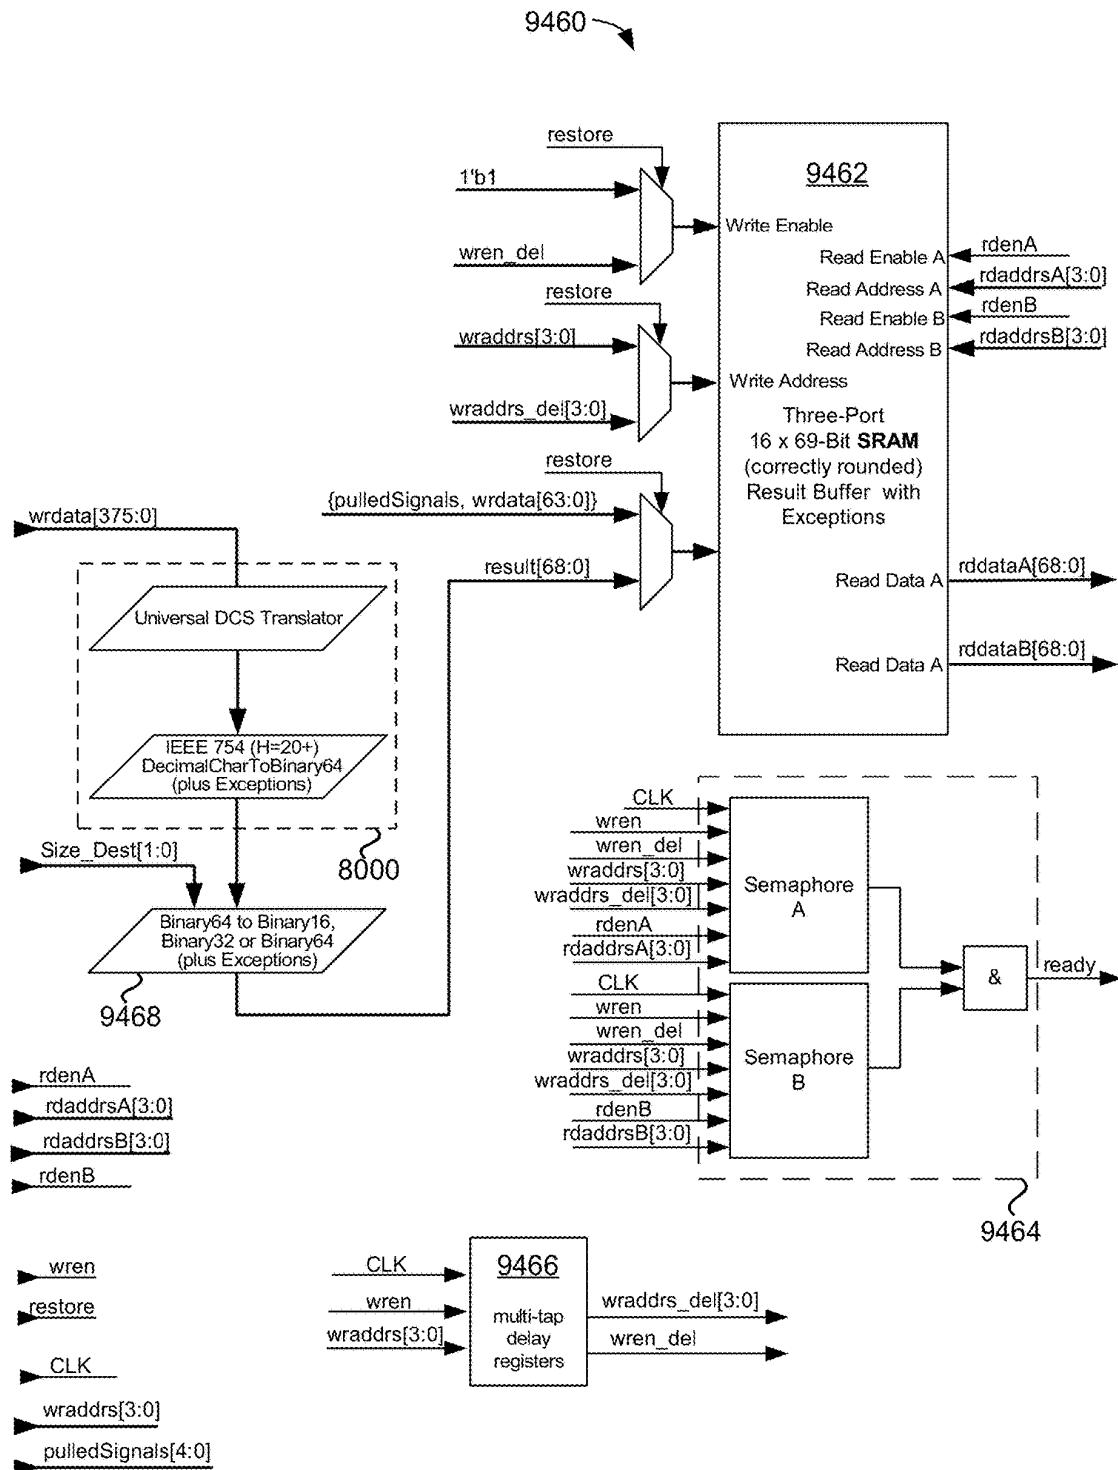


FIG. 43

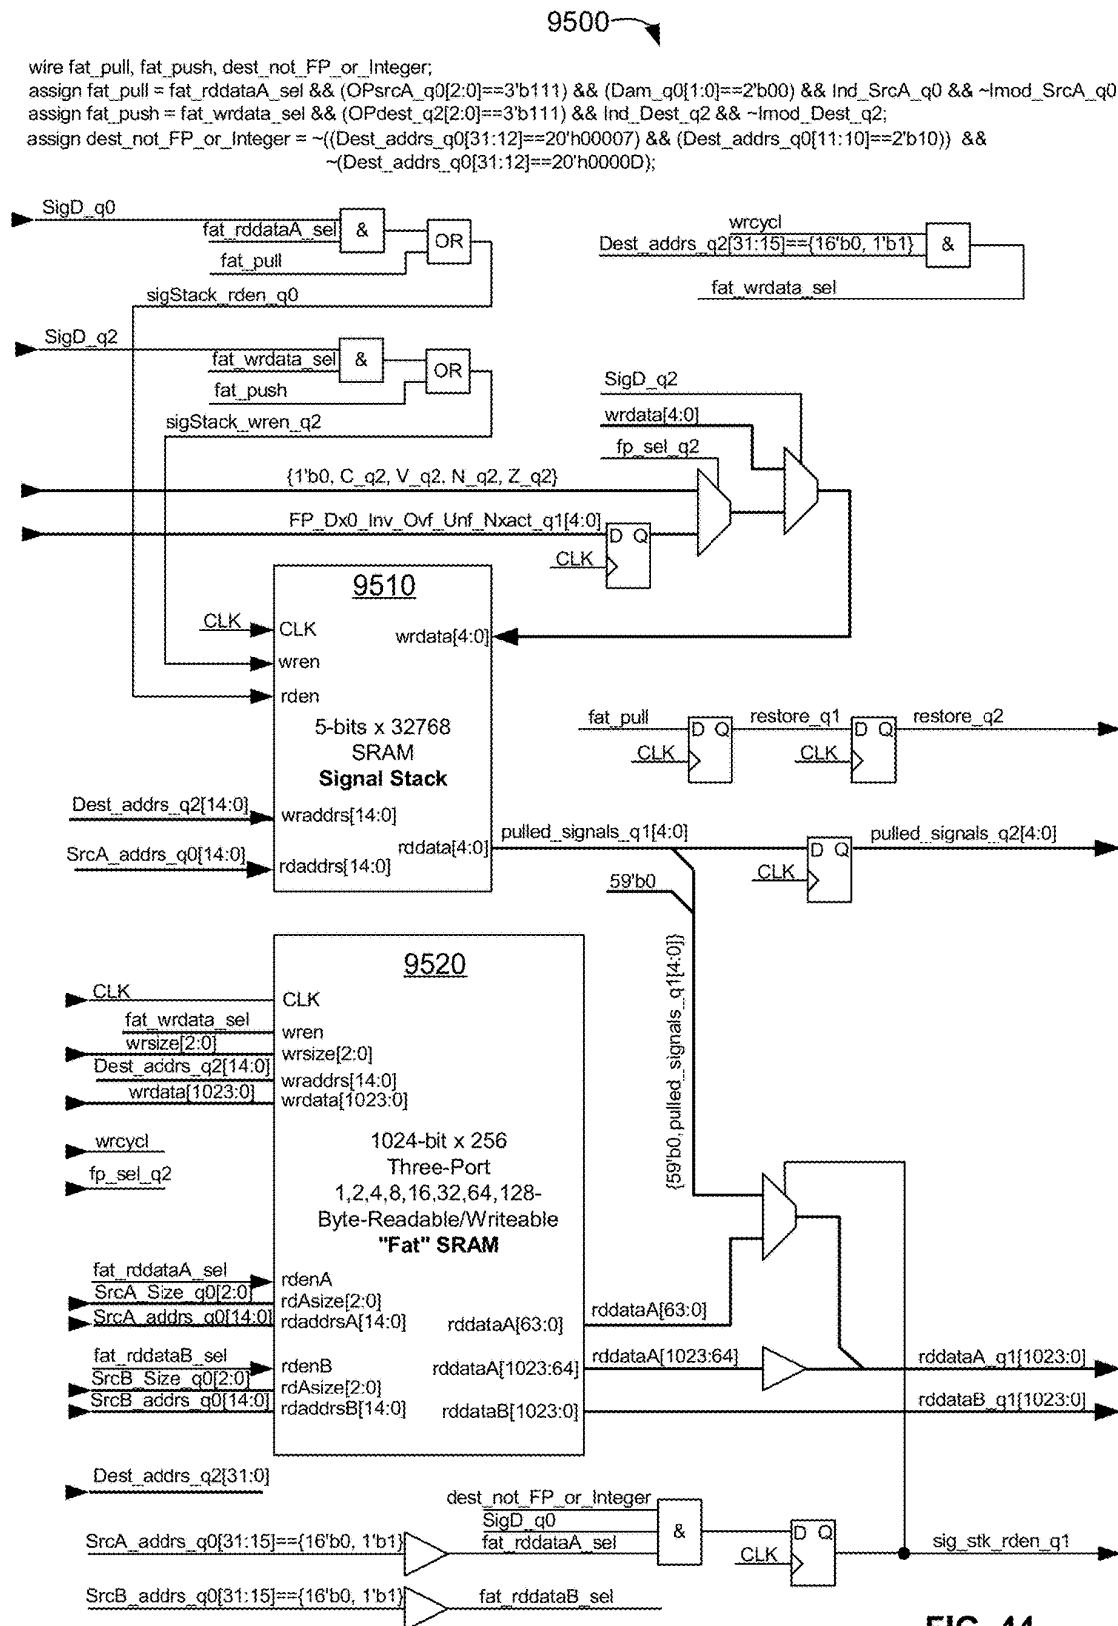


FIG. 44

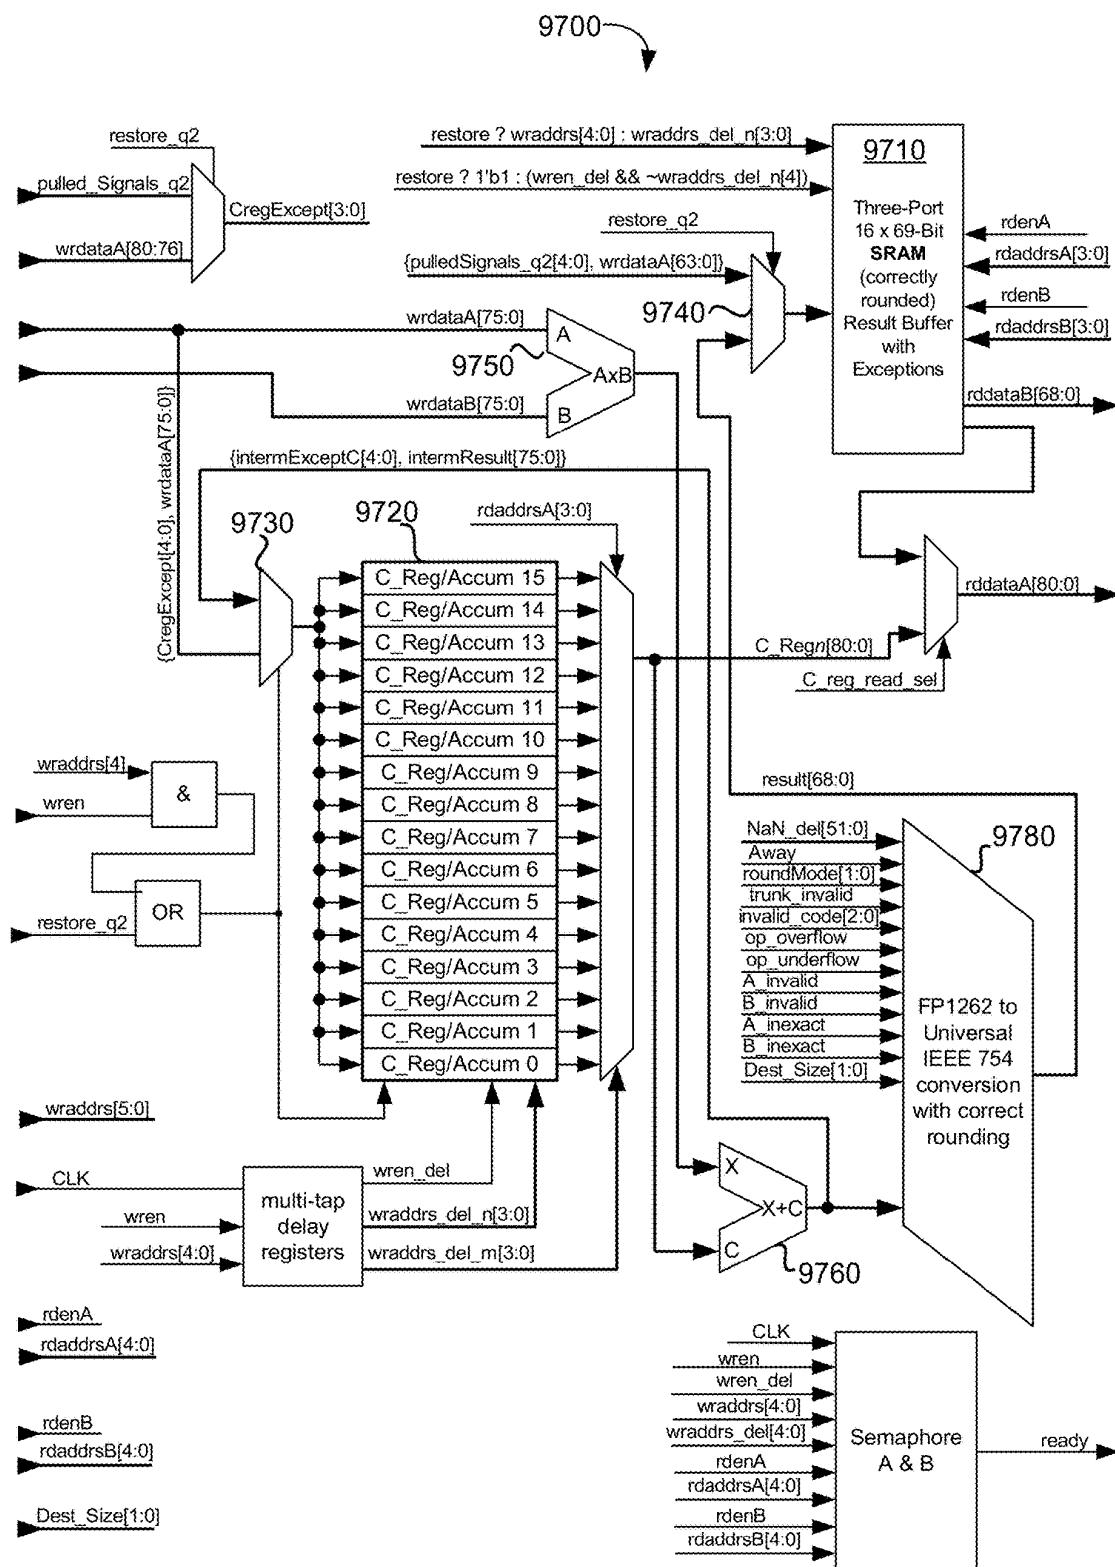
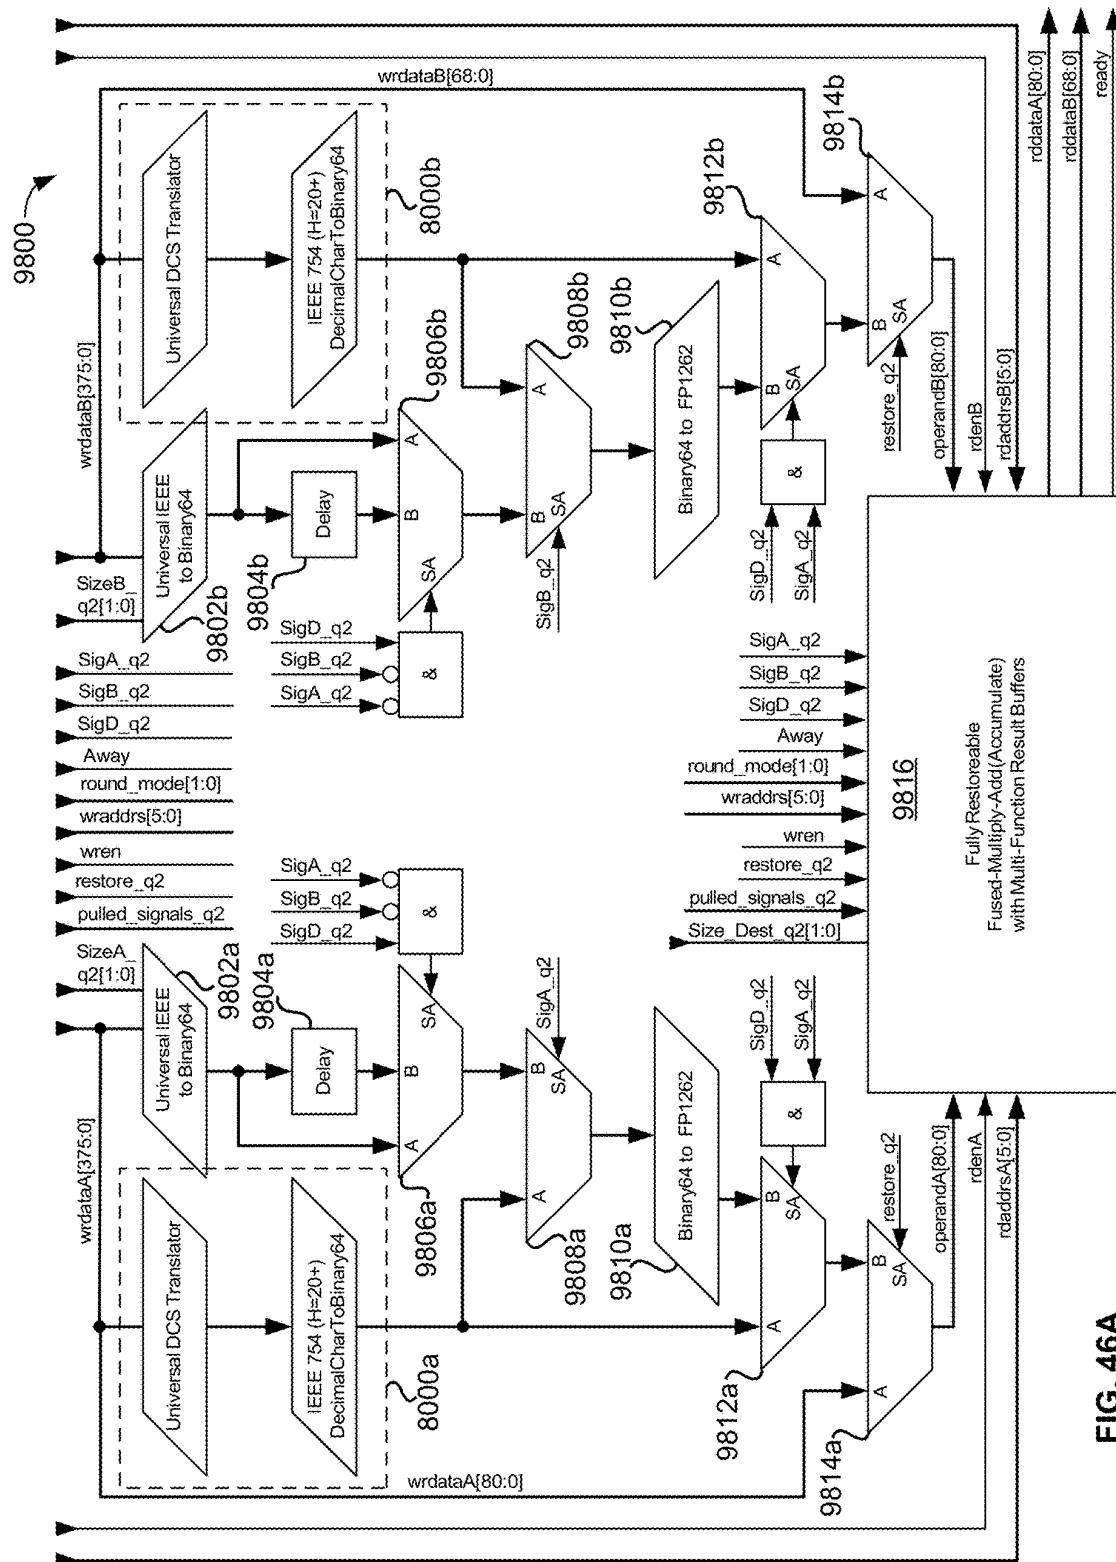


FIG. 45



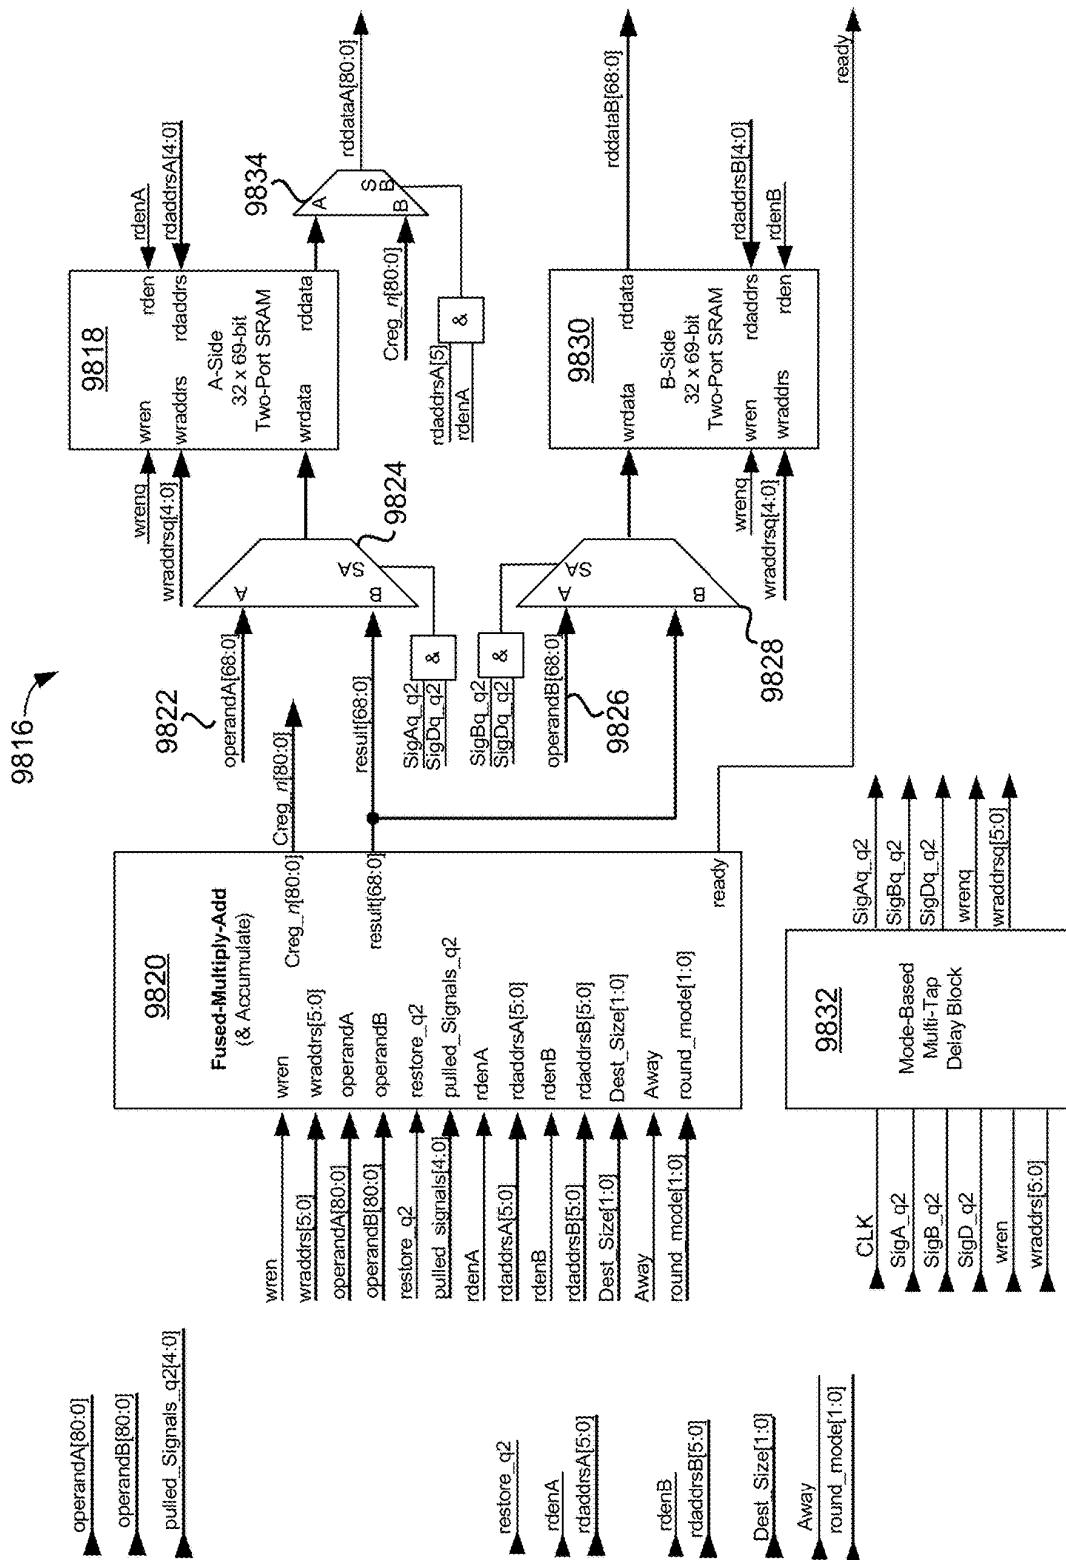


FIG. 46B

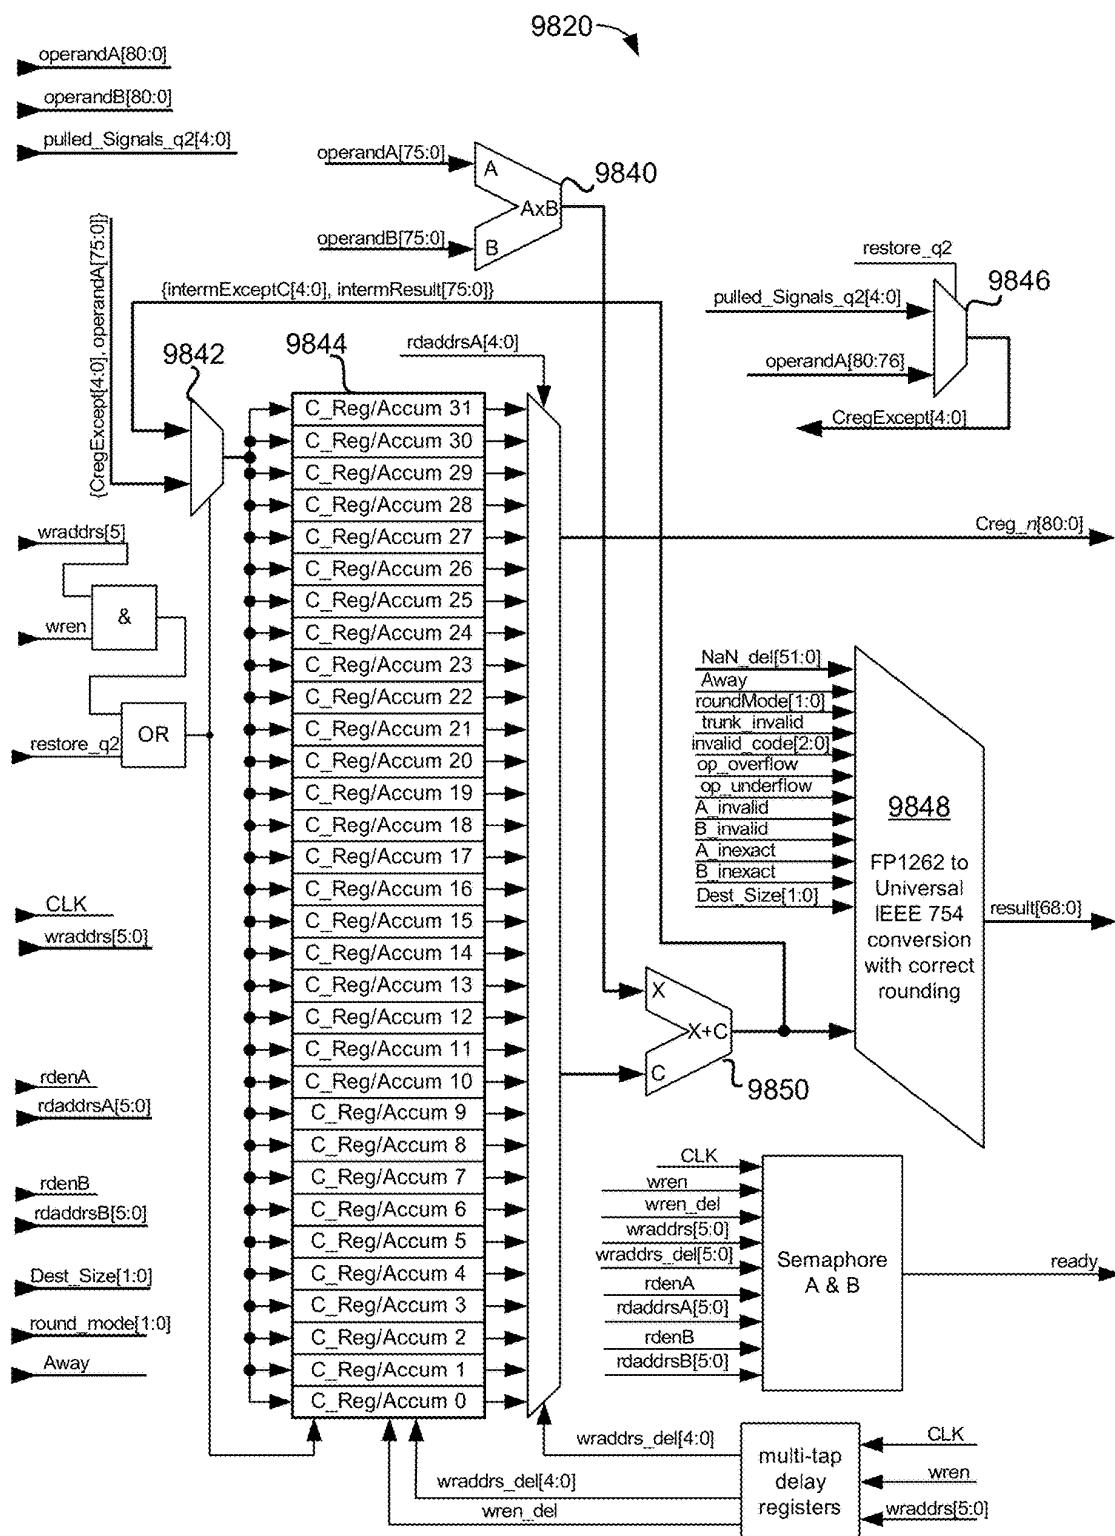


FIG. 46C

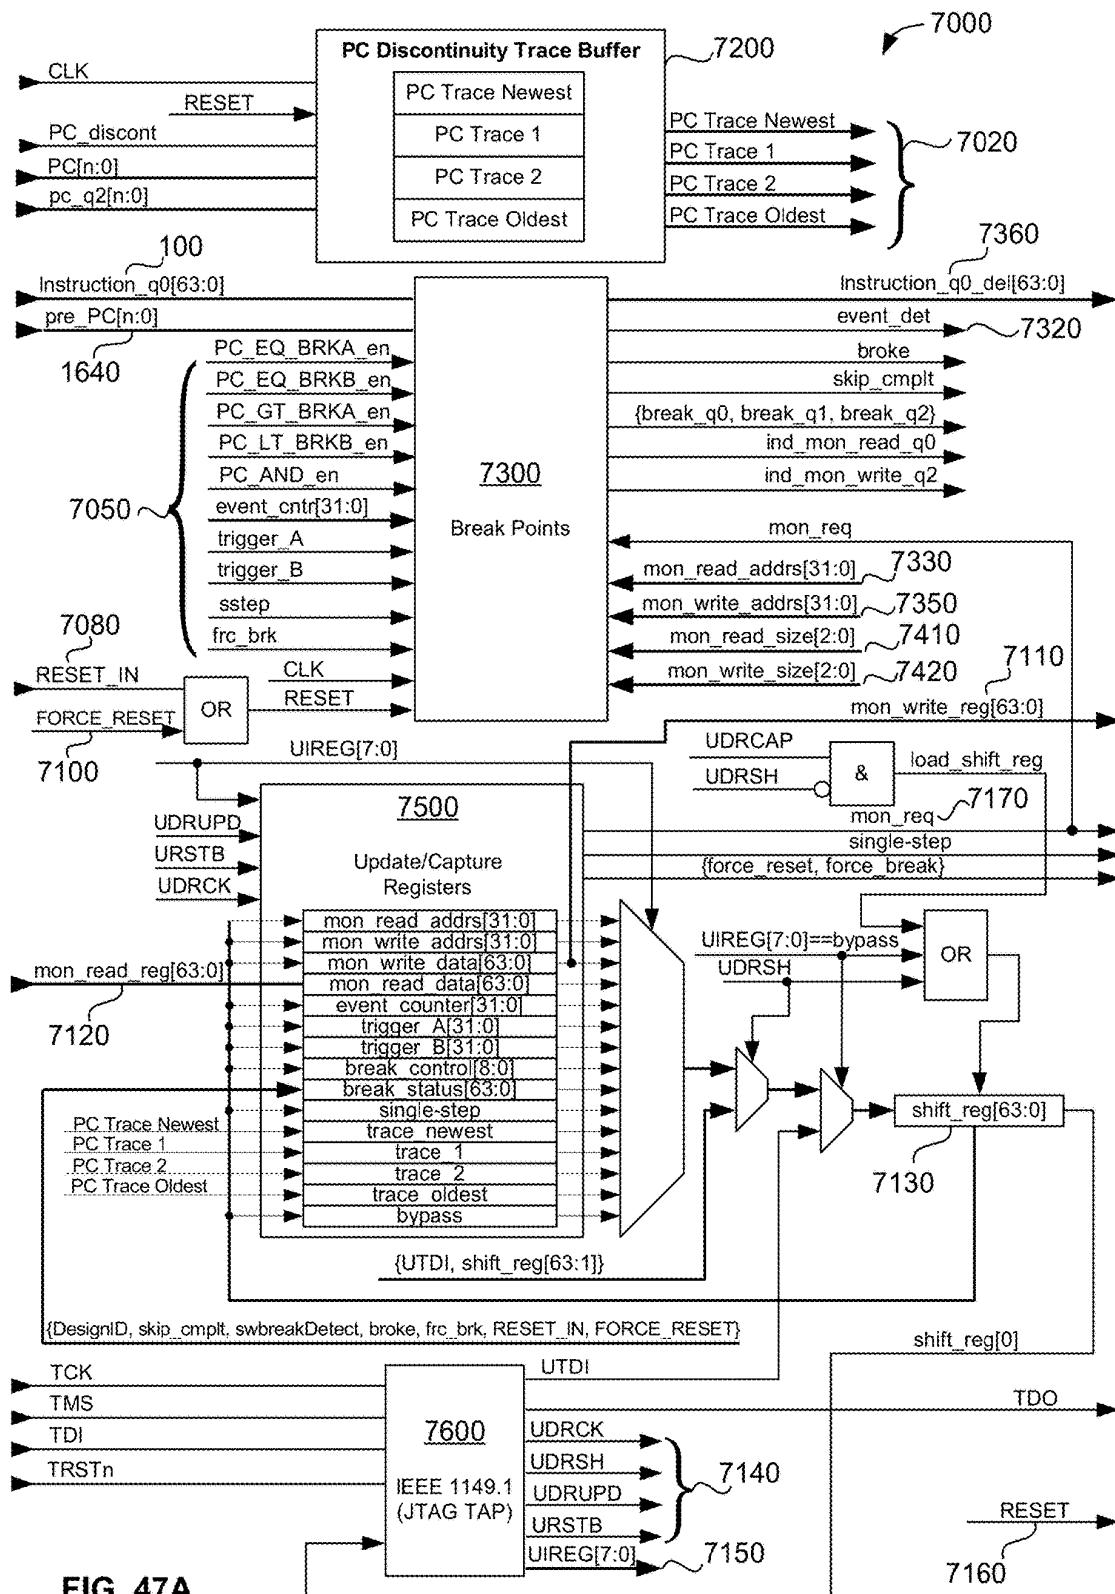


FIG. 47A

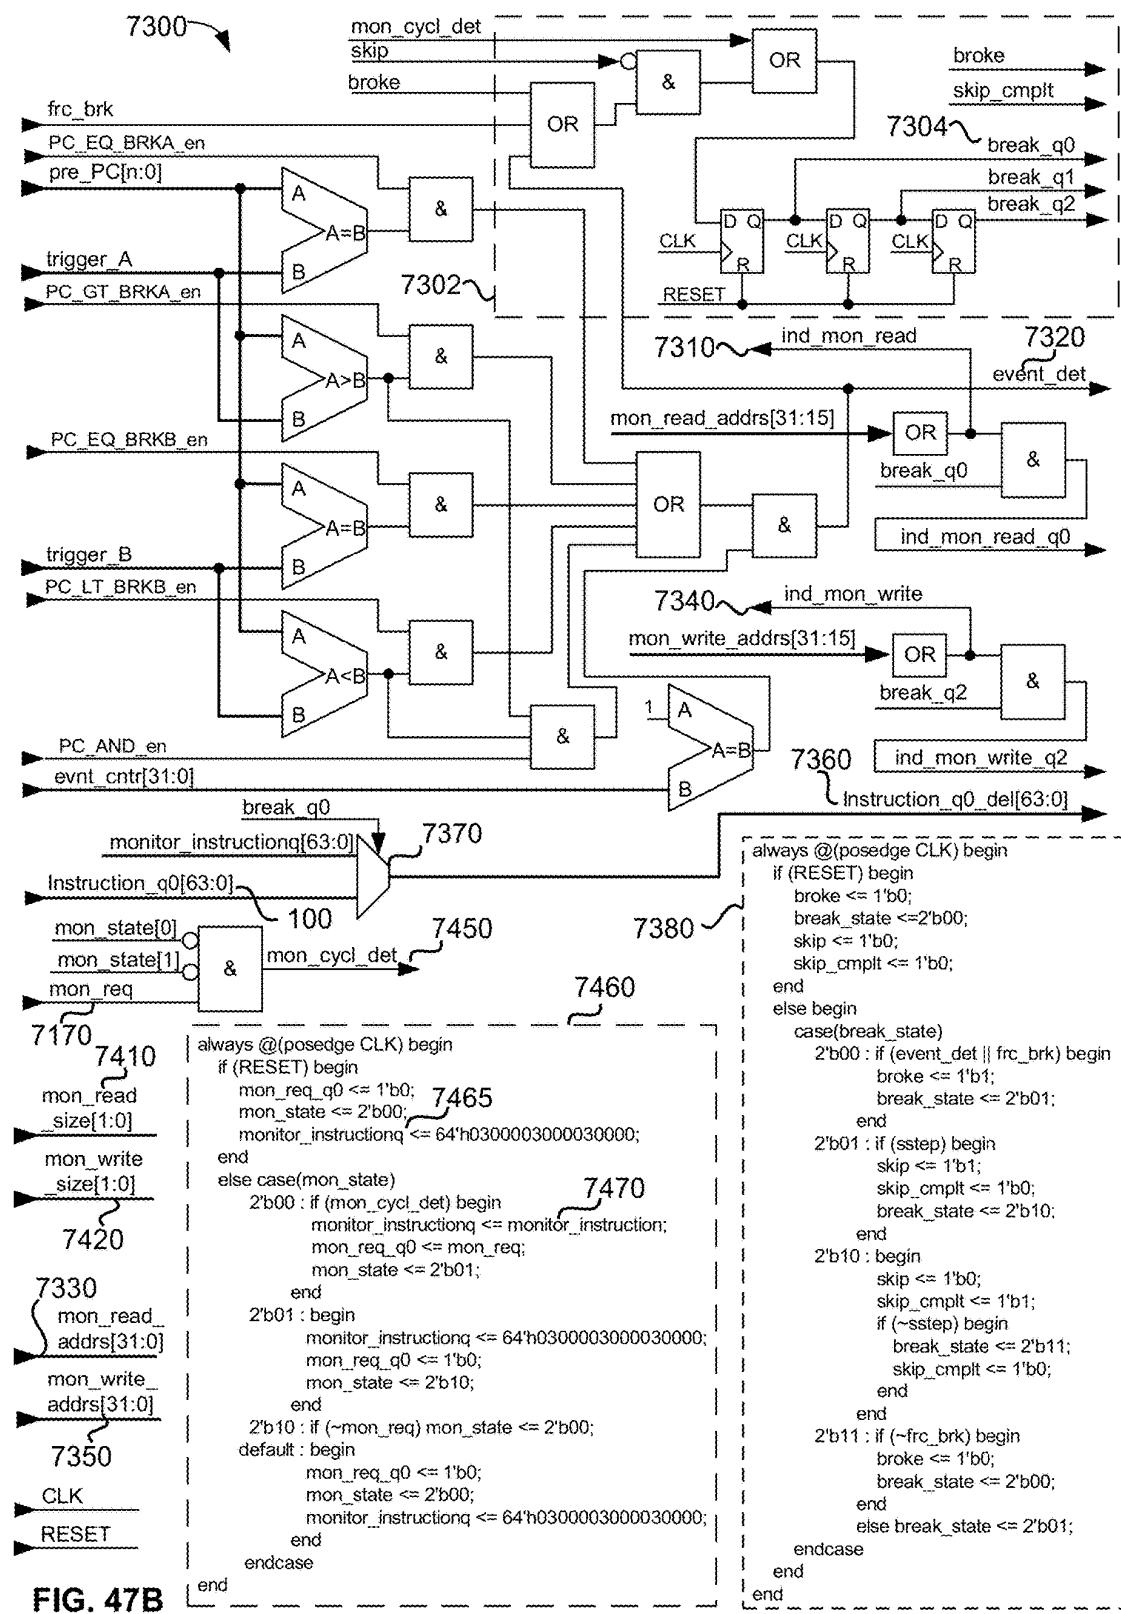


FIG. 47B

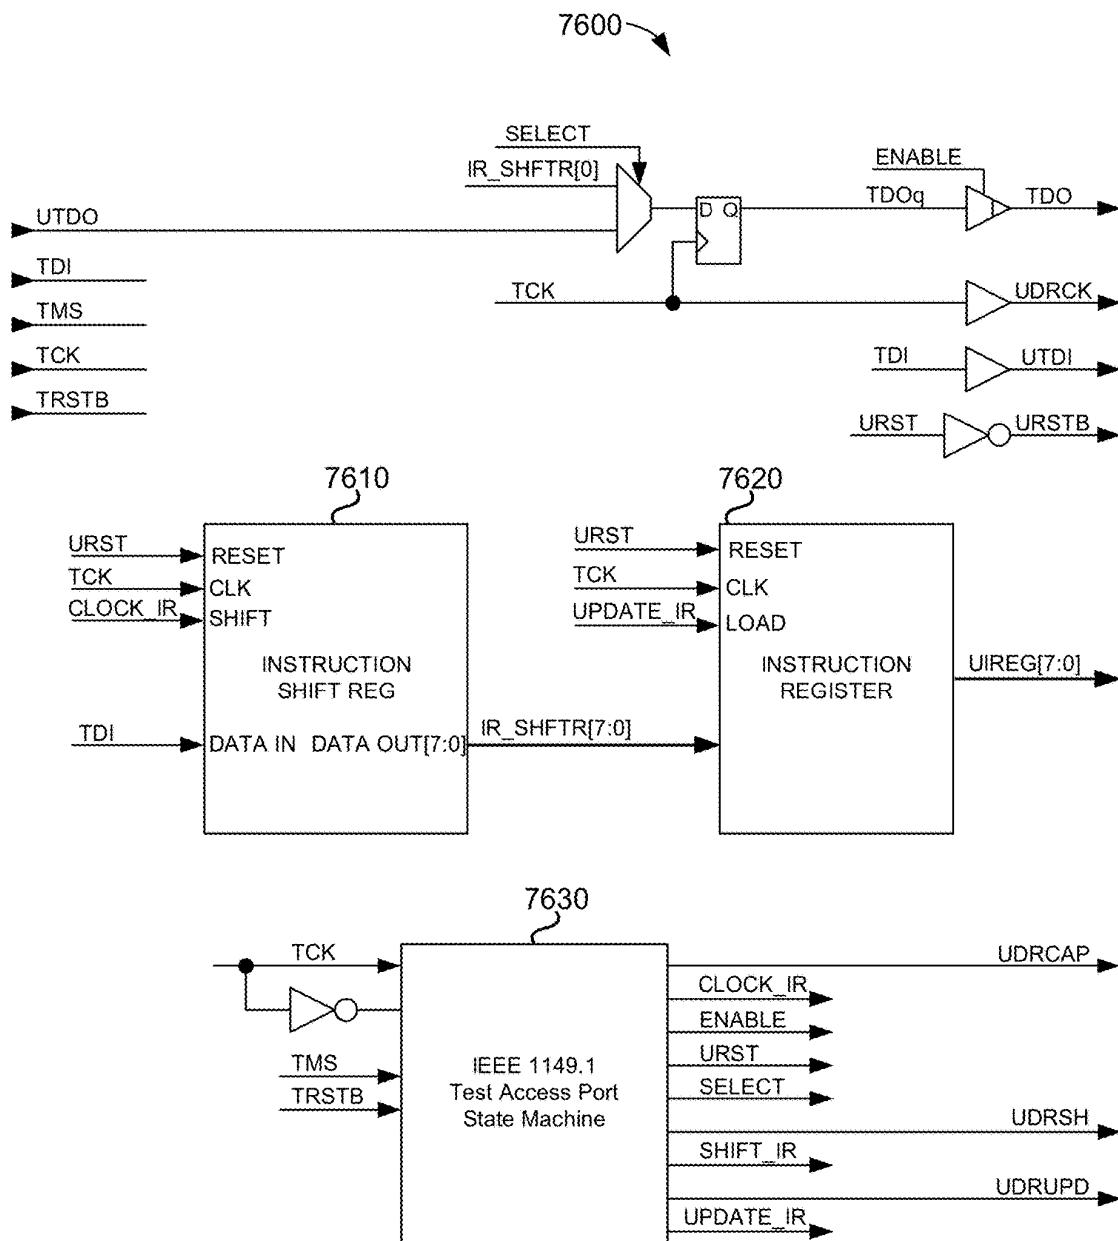


FIG. 47C

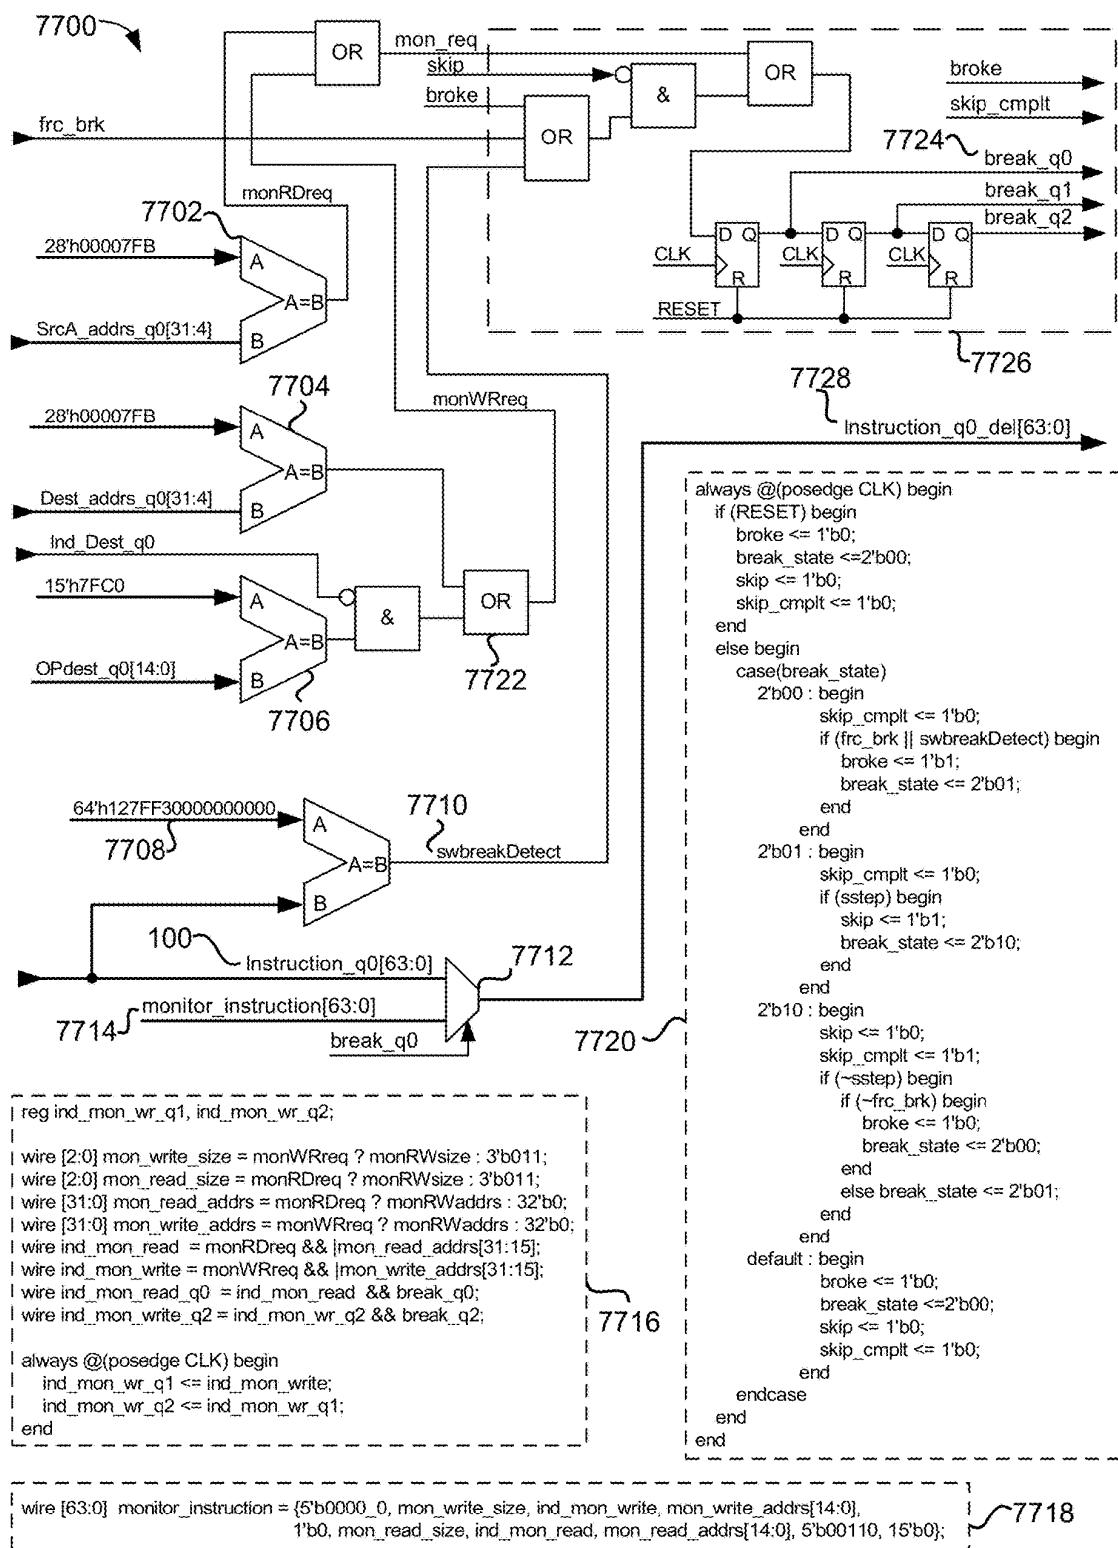


FIG. 48A

```
-----  
| parameter XCU_CNTRL_REG_ADDRS = 15'h7FDF;  
| parameter XCU_STATUS_REG_ADDRS = 15'h7FDE;  
| parameter XCU_PREAMPT_ADDRS = 15'h7FD4;  ↗7740  
| parameter XCU_SSTEP_ADDRS = 15'h7FD3;  
| parameter XCU_FORCE_BREAK_ADDRS = 15'h7FD2;  
| parameter XCU_FORCE_RESET_ADDRS = 15'h7FD1;  
|-----  
| //eXtra Compute Unit (XCU) control register mapped in parent CPU data space  
| always @(posedge CLK)  
|   if (RESET) XCU_CNTRL_REG <= 64'h0000_0000_FFFF_0000; ↗7744  
|   else if ((OPdest_q2==XCU_CNTRL_REG_ADDRS) && wrcycl && ~Ind_Dest_q2 && q2_sel) ↗7746  
|     XCU_CNTRL_REG <= wrsrcAdata;  
|   else if ((OPdest_q2==XCU_FORCE_RESET_ADDRS) && wrcycl && ~Ind_Dest_q2 && q2_sel) ↗7748  
|     XCU_CNTRL_REG[15:0] <= wrsrcAdata[15:0];  
|   else if ((OPdest_q2==XCU_FORCE_BREAK_ADDRS) && wrcycl && ~Ind_Dest_q2 && q2_sel) ↗7750  
|     XCU_CNTRL_REG[31:16] <= wrsrcAdata[15:0];  
|   else if ((OPdest_q2==XCU_SSTEP_ADDRS) && wrcycl && ~Ind_Dest_q2 && q2_sel) ↗7752  
|     XCU_CNTRL_REG[47:32] <= wrsrcAdata[15:0];  
|   else if ((OPdest_q2==XCU_PREAMPT_ADDRS) && wrcycl && ~Ind_Dest_q2 && q2_sel) ↗7754  
|     XCU_CNTRL_REG[63:48] <= wrsrcAdata[15:0];  
|-----  
wire [63:0] XCU_STATUS_REG = XCU_DONE[15:0], XCU_SWBRKDET[15:0], XCU_BROKE[15:0], XCU_SKIPCMPLT[15:0];  
    ↘7756      ↘7758      ↘7760      ↘7762      ↘7764
```

FIG. 48B

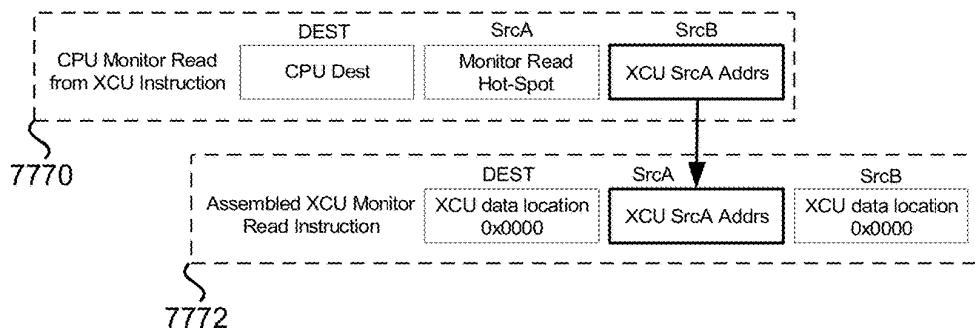


FIG. 48C

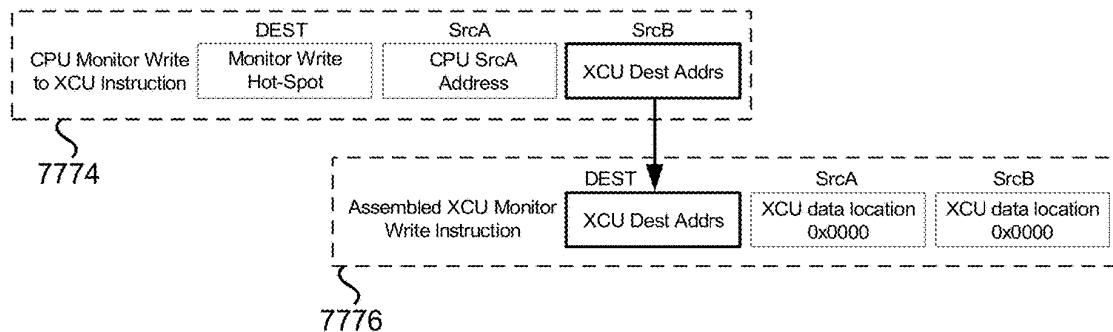


FIG. 48D

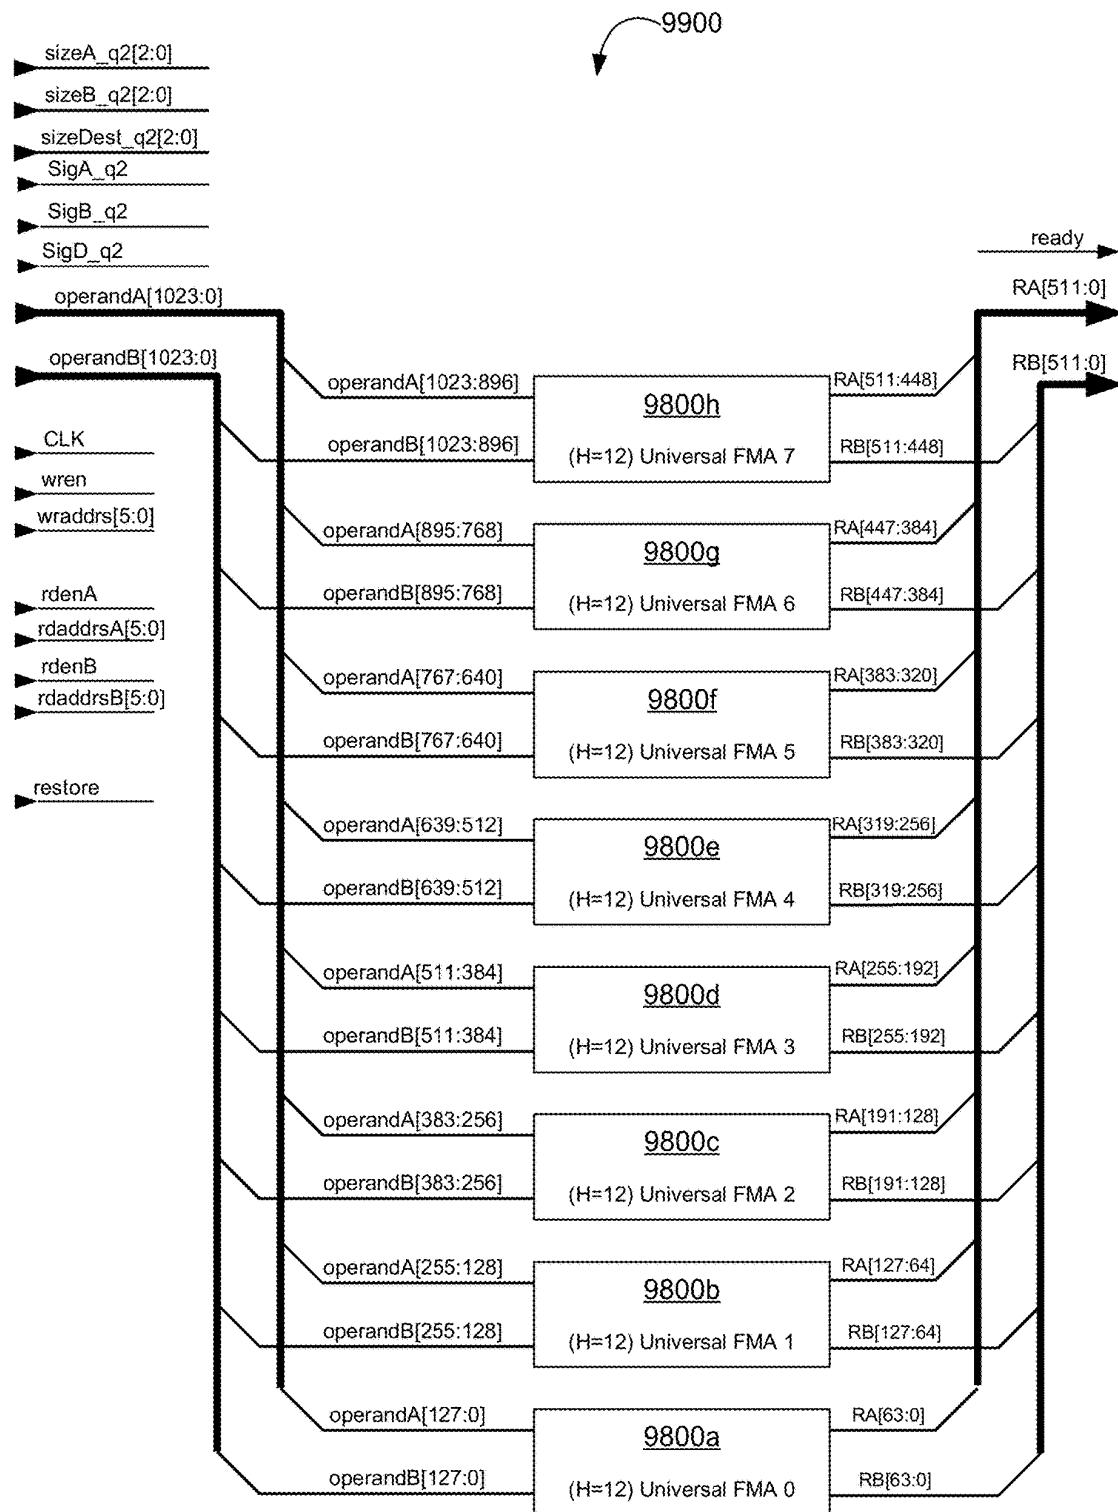


FIG. 49

9950

```

CPU "SYNPL64_TBL"
HOF "bin32"
WDLN 8

; SYNPL 64-Bit CPU Demo 3D Transform Micro-Kernel (SYNPL Assembly language version)

; version 3.05 August 5, 2019
; Author: Jerry D. Harthcock
; Copyright (C) 2019. All rights reserved.

; This program employs from 0 to 16 extra Compute Units (XCUs) to perform a 3D transform (rotate, translate, scale)
; of a .STL 3D object in external memory. If there are 0 XCUs attached to the CPU, then the CPU performs the 3D
; transform solo. Once the CPU determines there is a .stl file to transform, it counts how many XCUs are available,
; then distributes the triangles as evenly as possible among them for processing. Once the CPU determines the XCUs
; have completed the transform, it pulls the transformed triangles out of each XCU's buffer and then pushes them
; out to external memory and automatically changes endian-ness during the push.

; Observe that the .stl format triangles are IEEE 754 binary32 floating-point format, but the hardware
; floating-point operators are binary64 format. The cores do not care because they automatically convert from one
; binary format to the other without explicit instructions to do so.

;private dword storage
bitbucket: EQU 0x00000000 ;this dword location is reserved. Always reads as 0
work_1: EQU 0x00000008
work_2: EQU 0x00000010
work_3: EQU 0x00000018
capture_save: EQU 0x00000020 ;alternate delayed exception capture register 0 save location
capt1_save: EQU 0x00000028 ;alternate delayed exception capture register 1 save location
capt2_save: EQU 0x00000030 ;alternate delayed exception capture register 2 save location
capt3_save: EQU 0x00000038 ;alternate delayed exception capture register 3 save location

;for private storage of parameters for 3D transform
ext_vect_start: EQU 0x000040 ;location in external memory where the first triangle x1 may be found
triangles: EQU 0x000048 ;storage location of number of triangles in this thread's list to process

;dword storage locations for parameters so it will be easy to change to/from double precision
scaleX: EQU 0x00000050 ;scale factor X axis
scaleY: EQU 0x00000058 ;scale factor Y axis
scaleZ: EQU 0x00000060 ;scale factor Z axis
transX: EQU 0x00000068 ;translate amount X axis
transY: EQU 0x00000070 ;translate amount Y axis
transZ: EQU 0x00000078 ;translate amount Z axis

;(32-bit) storage for x1, y1, z1, x2, y2, z2, x3, y3, z3 for assembling half-word pieces from external memory
;x1: EQU 0x00000080
y1: EQU 0x00000084
z1: EQU 0x00000088
x2: EQU 0x0000008C
y2: EQU 0x00000090
z2: EQU 0x00000094
x3: EQU 0x00000098
y3: EQU 0x0000009C

```

FIG. 50A

```

000000A0      =      z3:      EQU      0x000A0
                sin_thetaX: EQU      0x000A4
                cos_thetaX: EQU      0x000A8
                sin_thetaY: EQU      0x000AC
                cos_thetaY: EQU      0x000B0
                sin_thetaZ: EQU      0x000B4
                cos_thetaZ: EQU      0x000B8

000000C0      =      XCUs:      EQU      0x000C0
                result_buf: EQU      0x000C8

000000D0      =      remainder_push: EQU      0x000D0
                remainder_pull: EQU      0x000D8

000000E0      =      PROG_START: EQU      0x80000000
                STL_START: EQU      0x00010000
                buf_START: EQU      0x000000100

000000F0      =      org      0x0
                Constants: DFL      0, load_vects
                prog_len: DFL      0, progend
                ;entry point for this program

00000000 0000000000000000100
00000001 0000000000000000225

00000002      =      ;parameters for this particular 3D transform test run
                    ;transform_3axis_parameters:
                    ;rotate around x axis in integer degrees
rotX:      dfl      0, 29
rotY:      dfl      0, 44
rotZ:      dfl      0, 75
scal_x:    dff      0, 2.0
scal_y:    dff      0, 2.0
scal_z:    dff      0, 2.25
trans_x:   dff      0, 4.75
trans_y:   dff      0, 3.87
trans_z:   dff      0, 2.237
                    ;scale z axis amount real
                    ;translate on x axis amount real
                    ;translate on y axis amount real
                    ;translate on z axis amount real
;RN      siz:dest = (siz:srca, siz:srcb)
                ;size:dest = (siz:srca, siz:srcb)
                ;default interrupt vector locations

000000100
000000100 317FE610000000178
000000101 317FE5100000001AD
000000102 317FE41000000017C
000000103 317FE310000000185
000000104 317FE21000000018E
000000105 317FE110000000197
000000106 317FE0100000001A8

load_vects:      org      0x000000100
                ;load of interrupt vectors for faster interrupt response
                ;these registers are presently not visible to app s/w
                - 2:NMT_VECT = _2:#NMFI_
                - 2:IRQ_VECT = _2:#IRQ_
                - 2:INV_VECT = _2:#INV_
                - 2:DIVX0_VECT = _2:#DIVX0_
                - 2:OVEL_VECT = _2:#OVEL_
                - 2:UNFL_VECT = _2:#UNFL_
                - 2:INEXT_VECT = _2:#INEXT_
                - 2:INEXT_VECT = _2:#INEXT_

```

FIG. 50B

```

000000107 307CDD0020000000020
...   _1::setDVNCZ = _1:#intEnable           ;set interrupt enable bit
000000108 307CD0A000000000010
done:    _1::setDVNCZ = _1:#DoneBit          ;set Done bit
...   _4::TIMER = _4:#600000             ;load time-out timer with sufficient time to process before timeout
000000109 327FFC2000000EA60
...   _4::PCC = (_1:0x00, 0, $)           ;s/w break here (note: $ is current PC)
00000010A 127FF30000000000000
00000010B 1A7FF30000000000000
begin:   _54::PCC = (_1:0x00, 0, threadStart) ;gosub threadStart
00000010C 127FF30000003FFC
...   _4::PCC = (_1:0x00, 0, done)          ;goto done

00000010D 307CDC000000000010
threadStart:
00000010E 02BF727FF2000000
...   _1::clrDVNCZ = _1:#DoneBit           ;clear Done bit to signal host CPU or XCU has started
00000010F 027FF7200C800000
...   _4::AR0 = _4:result_buf            ;load AR0 with pointer to source/destination internal result buffer
...   _4::AR1 = _4:result_buf            ;for XCU X1 of first triangle
000000110 027FF8200C800000
...   _4::ind.0 = _4:@rotX              ;calculate sine of theta X and save
000000111 22785020000200000
...   _4::cosd.0 = _4:@rotX             ;calculate cosine of theta X and save
000000112 2278402000200000
...   _4::ind.1 = _4:@rotY              ;calculate sine of theta Y and save
000000113 2278512000300000
...   _4::cosd.1 = _4:@rotY             ;calculate cosine of theta Y and save
000000114 2278412000300000
...   _4::ind.2 = _4:@rotZ              ;calculate sine of theta Z and save
000000115 2278522000400000
...   _4::cosd.2 = _4:@rotZ             ;calculate cosine of theta Z and save
000000116 2278422000400000
...   _4::ind.1 = _4:@rotX              ;calculate sine of theta X and save
000000117 22000427850000000
...   _4::cosd.1 = _4:@rotX             ;calculate cosine of theta X and save
000000118 22000A82784000000
...   _4::ind.2 = _4:@rotY              ;calculate sine of theta Y and save
000000119 22000AC27851000000
...   _4::cosd.2 = _4:@rotY             ;calculate cosine of theta Y and save
00000011A 22000B02784100000
...   _4::ind.1 = _4:@rotZ              ;calculate sine of theta Z and save
00000011B 22000B42785200000
...   _4::cosd.1 = _4:@rotZ             ;calculate cosine of theta Z and save
00000011C 0200082784200000
...   _4::ind.0 = _4:@rotX              ;calculate sine of theta X and save
00000011D 22000502000500000
...   _4::cosd.0 = _4:@rotX             ;calculate cosine of theta X and save
00000011E 22000582000600000
...   _4::ind.1 = _4:@scal_x            ;scale X factor
00000011F 22000602000700000
...   _4::cosd.1 = _4:@scal_y            ;scale Y factor
000000120 22000682000800000
...   _4::ind.2 = _4:@scal_z            ;scale Z factor
000000121 22000702000900000
...   _4::transX = _4:@trans_X          ;translate X axis amount
000000122 22000782000A00000
...   _4::transY = _4:@trans_Y          ;translate Y axis amount
...   _4::transZ = _4:@trans_Z          ;translate Z axis amount
000000123 017FED1004800000
...   _2::LPCNT0 = _2:triangles        ;load loop counter @ with number of triangles
000000124
loop:    ;scale on X, Y, Z axis
...   _the following routine performs scaling on all three axis first,
...   ;rotate on all three axis second, then translate on all three axis last
;vertex 1
...   _4::FMUL.0 = (_4::*AR0++[4], _4:scaleX)
...   _4::FMUL.1 = (_4::*AR0++[4], _4:scaleY)
...   _4::FMUL.2 = (_4::*AR0++[4], _4:scaleZ)
;vertex 2
000000124 0278202802020050
000000125 0278212802020058
000000126 0278222802020066

```

FIG. 50C

```

000000127 0278232802020050
000000128 0278242802020058
000000129 0278252802020066
00000012A 0278262802020050
00000012B 0278272802020058
00000012C 0278282802020066

        -4:FMUL.3 = (.4; *AR0++[4], -4:scaleX)
        -4:FMUL.4 = (.4; *AR0++[4], -4:scaleY)
        -4:FMUL.5 = (.4; *AR0++[4], -4:scaleZ)
    ;vertex 3
        -4:FMUL.6 = (.4; *AR0++[4], -4:scaleX)
        -4:FMUL.7 = (.4; *AR0++[4], -4:scaleY)
        -4:FMUL.8 = (.4; *AR0++[4], -4:scaleZ)
    ;rotate around X axis
        X1 is now in FMUL_0
        Y1 is now in FMUL_1
        Z1 is now in FMUL_2
        X2 is now in FMUL_3
        Y2 is now in FMUL_4
        Z2 is now in FMUL_5
        X3 is now in FMUL_6
        Y3 is now in FMUL_7
        Z3 is now in FMUL_8

    ;vertex 1
        ; (cos(xrot) * Y1) - (sin(xrot) * Z1)
        -4:FMUL.9 = (.4:FMUL.1, -2:cos_thetaX); FMUL.9 = (cos(xrot) * Y1)
        -4:FMUL.10 = (.4:FMUL.2, -2:sin_thetaX); FMUL.10 = (sin(xrot) * Z1)
        ; (sin(xrot) * Y1) + (cos(xrot) * Z1)
        -4:FMUL.11 = (.4:FMUL.1, -2:cos_thetaX); FMUL.11 = (sin(xrot) * Y1)
        -4:FMUL.12 = (.4:FMUL.2, -2:sin_thetaX); FMUL.12 = (cos(xrot) * Z1)
        -4:FSUB.0 = (.4:FMUL.9, -4:FMUL.10); FSUB.0 = (cos(xrot) * Y1) - (sin(xrot) * Z1)
        -4:FADD.0 = (.4:FMUL.11, -4:FMUL.12); FADD.0 = (sin(xrot) * Y1) + (cos(xrot) * Z1)

    ;vertex 2
        ; (cos(xrot) * Y2) - (sin(xrot) * Z2)
        -4:FMUL.1 = (.4:FMUL.4, -2:cos_thetaX); FMUL.1 = (cos(xrot) * Y2)
        -4:FMUL.2 = (.4:FMUL.5, -2:sin_thetaX); FMUL.2 = (sin(xrot) * Z2)
        ; (sin(xrot) * Y2) + (cos(xrot) * Z2)
        -4:FMUL.13 = (.4:FMUL.4, -2:cos_thetaX); FMUL.13 = (sin(xrot) * Y2)
        -4:FMUL.14 = (.4:FMUL.5, -2:sin_thetaX); FMUL.14 = (cos(xrot) * Z2)
        -4:FSUB.1 = (.4:FMUL.1, -4:FMUL.2); FSUB.1 = (cos(xrot) * Y2) - (sin(xrot) * Z2)
        -4:FADD.1 = (.4:FMUL.13, -4:FMUL.14); FADD.1 = (sin(xrot) * Y2) + (cos(xrot) * Z2)

    ;vertex 3
        ; (cos(xrot) * Y3) - (sin(xrot) * Z3)
        -4:FMUL.9 = (.4:FMUL.7, -2:cos_thetaX); FMUL.9 = (cos(xrot) * Y3)
        -4:FMUL.10 = (.4:FMUL.8, -2:sin_thetaX); FMUL.10 = (sin(xrot) * Z3)
        ; (sin(xrot) * Y3) + (cos(xrot) * Z3)
        -4:FMUL.11 = (.4:FMUL.7, -2:cos_thetaX); FMUL.11 = (sin(xrot) * Y3)
        -4:FMUL.12 = (.4:FMUL.8, -2:sin_thetaX); FMUL.12 = (cos(xrot) * Z3)
        -4:FSUB.2 = (.4:FMUL.9, -4:FMUL.10); FSUB.2 = (cos(xrot) * Y3) - (sin(xrot) * Z3)

```

FIG. 50D

```

0000013E 027802278282782C
    - 4:FADD.2 = (.4:FMUL.11, -.4:FMUL.12) ; FADD.2 = (sin(xrot) * Y3) + (cos(xrot) * Z3)
    ; X1 is now in FMUL.0
    ; Y1 is now in FSUB.0
    ; Z1 is now in FADD.0
    ; X2 is now in FMUL.3
    ; Y2 is now in FSUB.1
    ; Z2 is now in FADD.1
    ; X3 is now in FMUL.6
    ; Y3 is now in FSUB.2
    ; Z3 is now in FADD.2

;rotate around Y axis
;vertex 1
    ; (cos(yrot) * X1) + (sin(yrot) * Z1) ; FMUL.1 = (cos(yrot) * X1)
    - 4:FMUL.1 = (.4:FMUL.0, -2:cos_thetaY) ; FMUL.2 = (sin(yrot) * Z1)
    - 4:FMUL.2 = (.4:FADD.0, 2:sin_thetaY) ; FMUL.3 = (sin(yrot) * Z1)
    ; (cos(yrot) * Z1) - (sin(yrot) * X1) ; FMUL.4 = (cos(xrot) * Z1)
    - 4:FMUL.4 = (.4:FADD.0, -2:cos_thetaY) ; FMUL.5 = (cos(xrot) * Z1)
    - 4:FMUL.5 = (.4:FMUL.0, -2:sin_thetaY) ; FMUL.6 = (sin(xrot) * X1)
    - 4:FMUL.6 = (.4:FADD.0, 2:cos_thetaY) ; FMUL.7 = (cos(xrot) * X1)
    - 4:FMUL.7 = (.4:FMUL.4, -4:FMUL.5) ; FMUL.8 = (sin(yrot) * Z1)
    ; vertex 2
    ; (cos(yrot) * X2) + (sin(yrot) * Z2) ; FMUL.9 = (cos(xrot) * Z2)
    - 4:FMUL.9 = (.4:FMUL.3, -2:cos_thetaY) ; FMUL.10 = (sin(yrot) * X2)
    - 4:FMUL.10 = (.4:FADD.1, 2:sin_thetaY) ; FMUL.11 = (cos(xrot) * X2)
    - 4:FMUL.11 = (.4:FMUL.3, -2:sin_thetaY) ; FMUL.12 = (sin(yrot) * Z2)
    ; vertex 3
    ; (cos(yrot) * X3) + (sin(yrot) * Z3) ; FMUL.13 = (cos(yrot) * X3)
    - 4:FMUL.13 = (.4:FMUL.6, -2:cos_thetaY) ; FMUL.14 = (sin(yrot) * Z3)
    - 4:FMUL.14 = (.4:FADD.2, -2:sin_thetaY) ; FMUL.15 = (sin(yrot) * X3)
    ; (cos(yrot) * Z3) - (sin(yrot) * X3) ; FADD.5 = (cos(yrot) * X3) + (sin(yrot) * Z3)
    - 4:FADD.5 = (.4:FMUL.11, -.4:FMUL.12) ; FSUB.5 = (cos(yrot) * Z3) - (sin(yrot) * X3)
    ; X1 is now in FADD.3
    ; Y1 is now in FSUB.0
    ; Z1 is now in FSUB.3
    ; X2 is now in FADD.4

```

FIG. 50E

```

;
Y2 is now in FSUB_1
;
Z2 is now in FSUB_4
;
X3 is now in FADD_5
;
Y3 is now in FSUB_2
;
Z3 is now in FSUB_5

;rotate around Z axis

;vertex 1
;
Y2 is now in FSUB_1
;
Z2 is now in FSUB_4
;
X3 is now in FADD_5
;
Y3 is now in FSUB_2
;
Z3 is now in FSUB_5

;vertex 2
;
Y2 is now in FSUB_1
;
Z2 is now in FSUB_4
;
X3 is now in FADD_5
;
Y3 is now in FSUB_2
;
Z3 is now in FSUB_5

;vertex 3
;
Y2 is now in FSUB_1
;
Z2 is now in FSUB_4
;
X3 is now in FADD_5
;
Y3 is now in FSUB_2
;
Z3 is now in FSUB_5

;vertex 4
;
Y2 is now in FSUB_1
;
Z2 is now in FSUB_4
;
X3 is now in FADD_5
;
Y3 is now in FSUB_2
;
Z3 is now in FSUB_5

;vertex 5
;
Y2 is now in FSUB_1
;
Z2 is now in FSUB_4
;
X3 is now in FADD_5
;
Y3 is now in FSUB_2
;
Z3 is now in FSUB_5

;vertex 6
;
Y2 is now in FSUB_1
;
Z2 is now in FSUB_4
;
X3 is now in FADD_5
;
Y3 is now in FSUB_2
;
Z3 is now in FSUB_5

;vertex 7
;
Y2 is now in FSUB_1
;
Z2 is now in FSUB_4
;
X3 is now in FADD_5
;
Y3 is now in FSUB_2
;
Z3 is now in FSUB_5

;vertex 8
;
Y2 is now in FSUB_1
;
Z2 is now in FSUB_4
;
X3 is now in FADD_5
;
Y3 is now in FSUB_2
;
Z3 is now in FSUB_5

;vertex 9
;
Y2 is now in FSUB_1
;
Z2 is now in FSUB_4
;
X3 is now in FADD_5
;
Y3 is now in FSUB_2
;
Z3 is now in FSUB_5

;vertex 10
;
Y2 is now in FSUB_1
;
Z2 is now in FSUB_4
;
X3 is now in FADD_5
;
Y3 is now in FSUB_2
;
Z3 is now in FSUB_5

;vertex 11
;
Y2 is now in FSUB_1
;
Z2 is now in FSUB_4
;
X3 is now in FADD_5
;
Y3 is now in FSUB_2
;
Z3 is now in FSUB_5

```

FIG. 50F

```

;now translate on X, Y, Z axis
;vertex 1
    - 4:FADD.0 = (.4:FSUB.6, -4:transX)
    - -4:FADD.1 = (_4:FADD.6, -4:transY)
    - -4:FADD.2 = (_4:FSUB.3, -4:transZ)
;vertex 2
    - 4:FADD.9 = (.4:FSUB.7, -4:transX)
    - -4:FADD.10 = (_4:FADD.7, -4:transY)
    - -4:FADD.11 = (_4:FSUB.4, -4:transZ)
;vertex 3
    - 4:FADD.12 = (.4:FSUB.8, -4:transX)
    - -4:FADD.13 = (_4:FADD.8, -4:transY)
    - -4:FADD.14 = (_4:FSUB.5, -4:transZ)

;copy transformed X1 to alignable memory
    - 4:*AR1++[4] = -4:FADD.0
    - -4:*AR1++[5] = -4:FADD.1
    - -4:*AR1++[6] = -4:FADD.2
    - -4:*AR1++[7] = -4:FADD.9
    - -4:*AR1++[8] = -4:FADD.10
    - -4:*AR1++[9] = -4:FADD.11
    - -4:*AR1++[10] = -4:FADD.12
    - -4:*AR1++[11] = -4:FADD.13
    - -4:*AR1++[12] = -4:FADD.14
;copy transformed Y1 to alignable memory
;copy transformed Z1 to alignable memory
;copy transformed X2 to alignable memory
;copy transformed Y2 to alignable memory
;copy transformed Z2 to alignable memory
;copy transformed X3 to alignable memory
;copy transformed Y3 to alignable memory
;copy transformed Z3 to alignable memory

00000169 027800C2781820068
0000016A 027800D2780020078
0000016B 027800E2781520078
0000016C 02802127800000000
0000016D 02802127801000000
0000016E 02802127802000000
0000016F 02802127803000000
00000170 02802127804000000
00000171 02802127805000000
00000172 02802127806000000
00000173 02802127807000000
00000174 02802127808000000
00000175 127FF417FE043FAF
00000176 0000000000000000
00000177 027FF528047000000

; interrupt/exception trap service routines
NMI_:
    - 4:PCS = (.2:IPCNT0, 16, loop) ;NEXT continue until done
    - ;this NOP needed here to prevent triggering SP increment below
    - 4:PC = _4:*SP++[8]
    - ;return from interrupt

INV_:
    - s4:*SP--[8] = -4:PC_COPY ;save return address from non-maskable interrupt
    - ;time-out timer in this instance)
    - -4:TIMER = -4:#60000 ;put a new value in the timer
    - s4:PC = _4:*SP++[8] ;return from interrupt

INV_:
    - s4:*SP--[8] = -4:PC_COPY ;save return address from maskable invalid operation exception
    - -8:capt0_save = -8:CAPTURE0 ;read out CAPTURE0 register and save it
    - -8:capt1_save = -8:CAPTURE1 ;read out CAPTURE1 register and save it
    - -8:capt2_save = -8:CAPTURE2 ;read out CAPTURE2 register and save it
    - -8:capt3_save = -8:CAPTURE3 ;read out CAPTURE3 register and save it
    - 1:longSig = -1:#invalid ;lower invalid signal
    - 1:razSig = -1:#invalid ;raise invalid flag
    - -4:TIMER = -4:#60000 ;put a new value in the timer
    - s4:PC = _4:*SP++[8] ;return from interrupt

DIV0_:
    - s4:*SP--[8] = -4:PC_COPY ;save return address from maskable divide by 0 exception
    - -8:capt0_save = -8:CAPTURE0 ;read out CAPTURE0 register and save it
    - -8:capt1_save = -8:CAPTURE1 ;read out CAPTURE1 register and save it
    - -8:capt2_save = -8:CAPTURE2 ;read out CAPTURE2 register and save it
    - -8:capt3_save = -8:CAPTURE3 ;read out CAPTURE3 register and save it

```

FIG. 50G

FIG. 50H

```

000001AF 307CD000000000010 - _clrDNCZ = 1:doneBit
000001AF - push_thread = 1:doneBit
000001AF - 2:#(XCU15|XCU14|XCU13|XCU12|XCU11|XCU10|XCU9|XCU8|XCU7|XCU6|XCU5|XCU4|XCU3|XCU2|XCU1|XCU0)
000001B0 317FD11000000FFFF - 2:#(XCU15|XCU14|XCU13|XCU12|XCU11|XCU10|XCU9|XCU8|XCU7|XCU6|XCU5|XCU4|XCU3|XCU2|XCU1|XCU0)
000001B1 317FD21000000FFFF - 2:#(XCU15|XCU14|XCU13|XCU12|XCU11|XCU10|XCU9|XCU8|XCU7|XCU6|XCU5|XCU4|XCU3|XCU2|XCU1|XCU0)
000001B2 317FD110000000000 - 2:#(XCU15|XCU14|XCU13|XCU12|XCU11|XCU10|XCU9|XCU8|XCU7|XCU6|XCU5|XCU4|XCU3|XCU2|XCU1|XCU0)

```

;at this point all XCUs should be in h/w break state doing absolutely nothing

```

00000001B2 227FFFC200001000000
          - _4:AR5 = _4:@prog_len
          - _4:AR0 = _4:#0x8000000000 ;load AR0 with pointer to location of beginning of thread to be pushed
          - _4:AR1 = _4:#0x8000000000 ;into XCU program memories be sure to set MSB of pointer to access program
          - _2:REPEAT = _4:AR5 ;memory indirectly place the thread starting at 0x00000000 in XCU program
          - _8:XCU_MON_Poke_ALL = (8:*AR1++[1], _8/*AR0++[1]);memory (setting MSB of address) forces data to be written to program
          - ;memory instead of data memory push the the 3D transform thread into each
          - ;entire thread is pushed into XCU using this
          - ;instruction sequence.

00000001B7 327FF720000100000
          - _4:AR0 = _4:#$T1_START ;load AR0 with address of external RAM location where raw
          - _4:ext_vect_start = _4:#$T1_START + 96 ;;STL file begins this is the location of the first triangle
          - _4:result_buf = _4:#buf_START ;X1 in external RAM
          - _4:triangles = _4:#AR0[80] ;set destination "signal" bit to signal reverse
          - ;endianness and get number of triangles
          - ;test DONE bit for XCU15
          - ;divide #triangles x 16
          - ;goto push_XCUs

00000001B8 320004020001000050
          - _4:AR0 = _4:#$T1_START ;load AR0 with address of external RAM
          - _4:ext_vect_start = _4:#$T1_START + 96 ;;STL file begins this is the location of the first triangle
          - _4:result_buf = _4:#buf_START ;X1 in external RAM
          - _4:triangles = _4:#AR0[80] ;set destination "signal" bit to signal reverse
          - ;endianness and get number of triangles
          - ;test DONE bit for XCU15

00000001B9 3200042000000000100
          - _4:AR0 = _4:#$T1_START ;load AR0 with address of external RAM
          - _4:ext_vect_start = _4:#$T1_START + 96 ;;STL file begins this is the location of the first triangle
          - _4:result_buf = _4:#buf_START ;X1 in external RAM
          - _4:triangles = _4:#AR0[80] ;set destination "signal" bit to signal reverse
          - ;endianness and get number of triangles
          - ;test DONE bit for XCU15

00000001BA 0A000482280000000000
          - _4:AR0 = _4:#$T1_START ;load AR0 with address of external RAM
          - _4:ext_vect_start = _4:#$T1_START + 96 ;;STL file begins this is the location of the first triangle
          - _4:result_buf = _4:#buf_START ;X1 in external RAM
          - _4:triangles = _4:#AR0[80] ;set destination "signal" bit to signal reverse
          - ;endianness and get number of triangles
          - ;test DONE bit for XCU15

00000001BB 127FF537FDE0004
          - _4:PCC = (_8:XCU_S_R, 63, _8_XCUS)
          - _1:XCUS = _1:#16
          - _4:shift_0 = (_4:triangles, RIGHT, 4) ;divide #triangles x 16
          - _4:PCC = (_1:0x00, 0, push_XCUs) ;goto push_XCUs

00000001BC 30000C000000000010
          - _4:PCC = (_8:XCU_S_R, 55, _4_XCUS) ;test DONE bit for XCU7
          - _1:XCUS = _1:#8
          - _4:shift_0 = (_4:triangles, RIGHT, 3) ;divide #triangles x 8
          - _4:PCC = (_1:0x00, 0, push_XCUs) ;goto push_XCUs

00000001BD 127FE0020048800604
          - _4:PCC = (_8:XCU_S_R, 51, _2_XCUS) ;test DONE bit for XCU3
          - _1:XCUS = _1:#4
          - _4:shift_0 = (_4:triangles, RIGHT, 2) ;divide #triangles x 4
          - _4:PCC = (_1:0x00, 0, push_XCUs) ;goto push_XCUs

00000001BF 127FF537FDE0004
          - _4:PCC = (_8:XCU_S_R, 49, _1_XCU) ;test DONE bit for XCU1
          - _1:XCUS = _1:#2
          - _4:shift_0 = (_4:triangles, RIGHT, 1) ;divide #triangles x 2
          - _4:PCC = (_1:0x00, 0, push_XCUs) ;goto push_XCUs

00000001C0 30000C000000000004
          - _4:PCC = (_8:XCU_S_R, 48, NO_XCUS) ;test DONE bit for XCU0
          - _1:XCUS = _1:#1
          - _4:shift_0 = (_4:triangles, COPY, 1) ;no shift--just copy into shift_0
          - _4:PCC = (_1:0x00, 0, push_XCUs) ;goto push_XCUs

00000001C1 127FF537FDE0004
          - _4:PCC = (_1:#0, solo_process) ;GOTO solo process ;the CPU has to do the
          - ;3D transform solo

00000001C2 127FF53000000000007
          - push_XCUs: ;each XCU PC will be initialized to begin executing here
          - _4:AR3 = _4:#load_vects ;this is the PC address from which all threads
          - _4:AR4 = _4:#begin ;begin processing (ie, exit out of break)
          - _4:work_1 = _4:#{XCU_MON_REQUEST.0} ;get base address of pushXCU operator

```

FIG. 50I

```

0000001D4 027EC0200008000C0          ;add number of XCUs to get most significant address +1
0000001D5 027E80027F600000C0          ;determine any remainder
0000001D6 027E002004827E80          ;sub.0 now contains any remainder
0000001D7 02000027FA00000000          ;copy result of remainder calc into remainder
                                         ;so it can be used later
                                         ;copy result of remainder calc into remainder
                                         ;so it can be used later
                                         ;copy result of remainder calc into remainder
                                         ;so it can be used later
                                         ;preset PCS of all XCUs at once to point to entrypoint of
                                         ;initialization sequence
                                         ;input/result_buf);push the location of the beginning of XCU
0000001D8 020000827EA000000          ;4:remainder_pull = _4:sub.0
                                         ;4:XCU_MON_POKE_ALL = (_4:AR3, _4:PC)
                                         ;4:XCU_MON_POKE_ALL = (_4:result_buf, _4:result_buf);push the location of the beginning of XCU
                                         ;input/result buffer
                                         ;current XCU base address for that XCU
                                         ;address in external RAM of where the first
                                         ;triangle X1 is located
                                         ;for the number of XCUs ...
0000001D9 027FC027FFA27FF5          ;2:LPCNT1 = _2:XCU_S
                                         ;_2:0x0000 = _2:_AR2--[1]
                                         ;_4:AR0 = _4:result_buf
                                         ;_4:add.1 = (_4:shift.0, _4:#0)           ;bump by -1 XCU number
                                         ;load AR0 with pointer to destination result buffer for
                                         ; XCU X1 of first triangle
                                         ;copy calculated triangles/XCUs
                                         ;into add.1 for future use
0000001DA 027FFC028EC8200C8          ;_4:compare = (_4:remainder_push, _2:#0@0)    ;see if there was any remainder from original
                                         ;triangles/XCU calculation of first triangle
                                         ;if no remainder, skip over a push of one more
0000001DB 017FFE100C00000000          ;_4:PCS = (_8:STATUS, Z, no_remainder_push)   ;triangle for the current XCU
                                         ;_4:sub.0 = (_4:remainder_push, _1:#1)        ;decrement any remainder by 1
                                         ;_4:add.1 = (_4:shift.0, _4:#1)            ;add.1 now contains the number of triangles this
                                         ;_4:remainder_push = _4:sub.0                ;particular XCU is to process
                                         ;no_remainder_push:
                                         ;_4:_AR2++[0] = (_2:_add.1, _2:triangles) ;poke the triangle batch size for this XCU into its
                                         ;"triangles" location
                                         ;_2:LPCNT0 = _2:_add.1
                                         ;_2:REPEAT = _2:#17                         ;push 18 half-words into target XCU
                                         ;_54:_AR2++[0] = (_2:_AR1++[2], _2:_AR0++[2]) ;(for a total of 9 32-bit floats per triangle)
                                         ;_4:PCS = (_2:_LPCNT1, 16, push_inner)       ;(AR2 contains the current XCU number)
                                         ;_4:PCS = (_2:_LPCNT0, 16, push_outer)        ;push source pointer by 14 to skip over .STL attribute
                                         ;_4:PCS = (_2:_LPCNT1, 16, push_outer)        ;and NORM fields NEXT_LPCNT0 decrement and
                                         ;NEXT_LPCNT1 decrement number of XCUs in LPCNT0
                                         ;_4:PCS = (_2:_LPCNT1, 16, push_outer)        ;jump if result not zero
                                         ;NEXT_LPCNT1 decrement number of XCUs in LPCNT0
                                         ;and jump if not zero
                                         ;_2:XCU_FORCE_BREAK = _2:#0                  ;clear all h/w breakpoints
                                         ;_2:XCU_SSTEP = _2:#{XCU1|XCU12|XCU13|XCU12|XCU11|XCU10|XCU9|XCU8|XCU7|XCU6|XCU5|XCU4|XCU3|XCU2|XCU1|XCU0}
                                         ;_2:XCU_SSTEP = _2:#0
                                         ;like the CPU before it was interrupted to invoke this process, the XCU's will
                                         ;now encounter a "sw" breakpoint at which point the CPU will change their PC's
                                         ;to threadStart to begin processing

```

FIG. 50J

```

000001F0 127FF337F0E800000          _4:PCC = (_8:XCU_S_R, 32, waitForXCUbreak0) ;wait for XCU_0 to hit s/w breakpoint
000001F1 00000000000000000000          -
000001F2 00000000000000000000          -
000001F3 027FC027FB227FF5          _4:XCU_MON_PKE_ALL = (_4:AR4, _4:PC) ;push "begin" into all XCU PCs simultaneously
000001F4 317FD3100000FFF          _2:XCU_SSSTEP = _2:#[XCU15|XCU14|XCU13|XCU12|XCU11|XCU10|XCU9|XCU7|XCU6|XCU5|XCU4|XCU3|XCU2|XCU1|XCU0]
000001F5 317FD310000000000          _2:XCU_SSSTEP = _2:#0

000001F6 127FF437F0DEC0000          waitForNotDone0: _4:PCC = (_8:XCU_S_R, 48, waitForNotDone0) ;wait for XCU_0 to bring its DONE bit low, indicating processing has started
000001F7 127FF337F0DEC0000          waitForDone0: _4:PCC = (_8:XCU_S_R, 48, waitForDone0) ;wait for XCU_0 to bring its DONE bit high, indicating completion
000001F8 027FF917EC000000          now that XCU0 is done processing its triangles, it's time to start pulling them out and pushing them back into external memory
000001F9 027FF8000000000          _4:AR2 = _2:add.0 ;previously calculated current XCU base address for that XCU
000001FA 017FFEE100C000000          _4:AR1 = _4:xt.vect_start ;address in external RAM of where the first triangle X1 is located
000001FB 0000001BF0A000000          _2:LPCNT1 = _2:XCU5 ;for the number of XCU5
pull_outer:
000001FC 027FF7200C80000          _2:0x00000 = _2:#AR2--[1] ;bump by -1 XCU number
000001FD 027FC127F60200000          _4:AR0 = _4:result_buf ;load AR0 with pointer to destination result buffer for XCU X1
000001FE 127FCF200D810000          _4:compare = (_4:remainder_pull, _2:#0x00) ;copy calculated triangles/XCU into add.1 for future use
000001FF 127FF437F100004          _4:PCC = (_8:STATUS, Z, no_remainder_pull) ;see if there was any remainder from original
00000200 127EA0200D8000001          _4:sub.0 = _4:remainder_pull ;if no remainder, skip over a push of one more triangle for the current XCU
00000201 127EC127F6020001          _4:add.1 = _4:shift.0, _4:#1 ;decrement any remainder by 1
00000202 0200D87EA000000          _4:remainder_pull = _4:sub.0 ;add.1 now contains the number of triangles this particular XCU is to process
no_remainder_pull:
00000203 017FED17EC100000          _2:LPCNT0 = _2:add.1 ;for the number of triangles per XCU
00000204 317FE1000000011          _2:REPEAT = _2:#17 ;pull 18 half-words from target XCU
00000205 0980111800218010          _2:#AR1++[2] = (_2:#AR2++[0], _2:#AR0++[2]) ;(for a total of 9 32-bit floats per triangle)
00000206 0000000000000000          _2:#set destination "signal" bit to
00000207 0000000000000000          _2:#reverse endian-ness just before pull ;reverse endian-ness just before pull
00000208 0100001801100000          _2:0x0000 = _2:#AR1++[14] ;nop
00000209 127FF417FED43FF8          _4:PCS = (_2:LPCNT0, 16, pull_inner) ;jump source pointer by 14 to skip over .STL attribute
0000020A          pull_next_XCU: ;and NORM fields ;NEXT LPCNT0 decrement and jump if result not zero

```

FIG. 50K

```

0000020A 127FF417FE43FFF1          ;NEXT LPCNT0 decrement and jump if result not zero
0000020B 307CDD000000000010          ;set Done bit
0000020C 0A7FF5F8047000000          ;return from interrupt--we are done

0000020D 327FF720001000000
    solo_process:
        4:AR0 = 4:#STL_START          ;load AR0 with address of external RAM location where
                                    ;raw STL file begins
        4:ext_vect_start = 4:#STL_START + 96      ;this is the location of the first triangle X1 in external
                                    ;RAM set destination sign extend bit to signal reverse
        s4:triangles = 4:#AR0[80]          ;endianness and get number of triangles
        4:AR0 = 4:#buf_START           ;Load AR0 with pointer to destination result buffer for
                                    ;XCU X1 of first triangle
        4:AR1 = 4:ext_vect_start      ;address in external RAM of where the
                                    ;first triangle X1 is located
        2:LPCNT0 = 2:triangles        ;FOR LPCNT0 pull triangles in from external memory into
                                    ;internal working memory
        pull_solo:
            2:REPEAT = 2:#17          ;reverse endian-ness just before push by setting
            s2:#AR0++[2] = 2:#AR1++[2]  ;destination sign extend bit (ie, "S2")
            2:0x0000 = 2:#AR1++[14]    ;bump source pointer by 14 to skip over
                                    ;.STL attribute and NORM fields
            4:PCS = (2:LPCNT0, 16, pull_solo)  ;NEXT LPCNT0 decrement and jump if result not zero
            S4:PCC = (1:0x00, 0, threadStart)  ;gosub threadStart compute the transform of
            4:AR1 = 4:ext_vect_start      ;entire 3D object--solo
            4:AR0 = 4:#buf_START         ;FOR LPCNT0 push computed transform result back out
                                    ;to external memory
        push_solo:
            2:LPCNT0 = 2:triangles      ;FOR LPCNT0 push computed transform result back out
            2:REPEAT = 2:#17            ;to external memory
            s2:#AR1++[2] = 2:#AR0++[12]
            2:0x0000 = 2:#AR1++[14]    ;NEXT LPCNT0 decrement and jump if result not zero
            4:PCS = (2:LPCNT0, 16, push_solo)  ;set Done bit
            1:setDNCZ = 1:#DoneBit
            s2:PC = 4:#SP++[8]
        progend:
    end

```

FIG. 50L

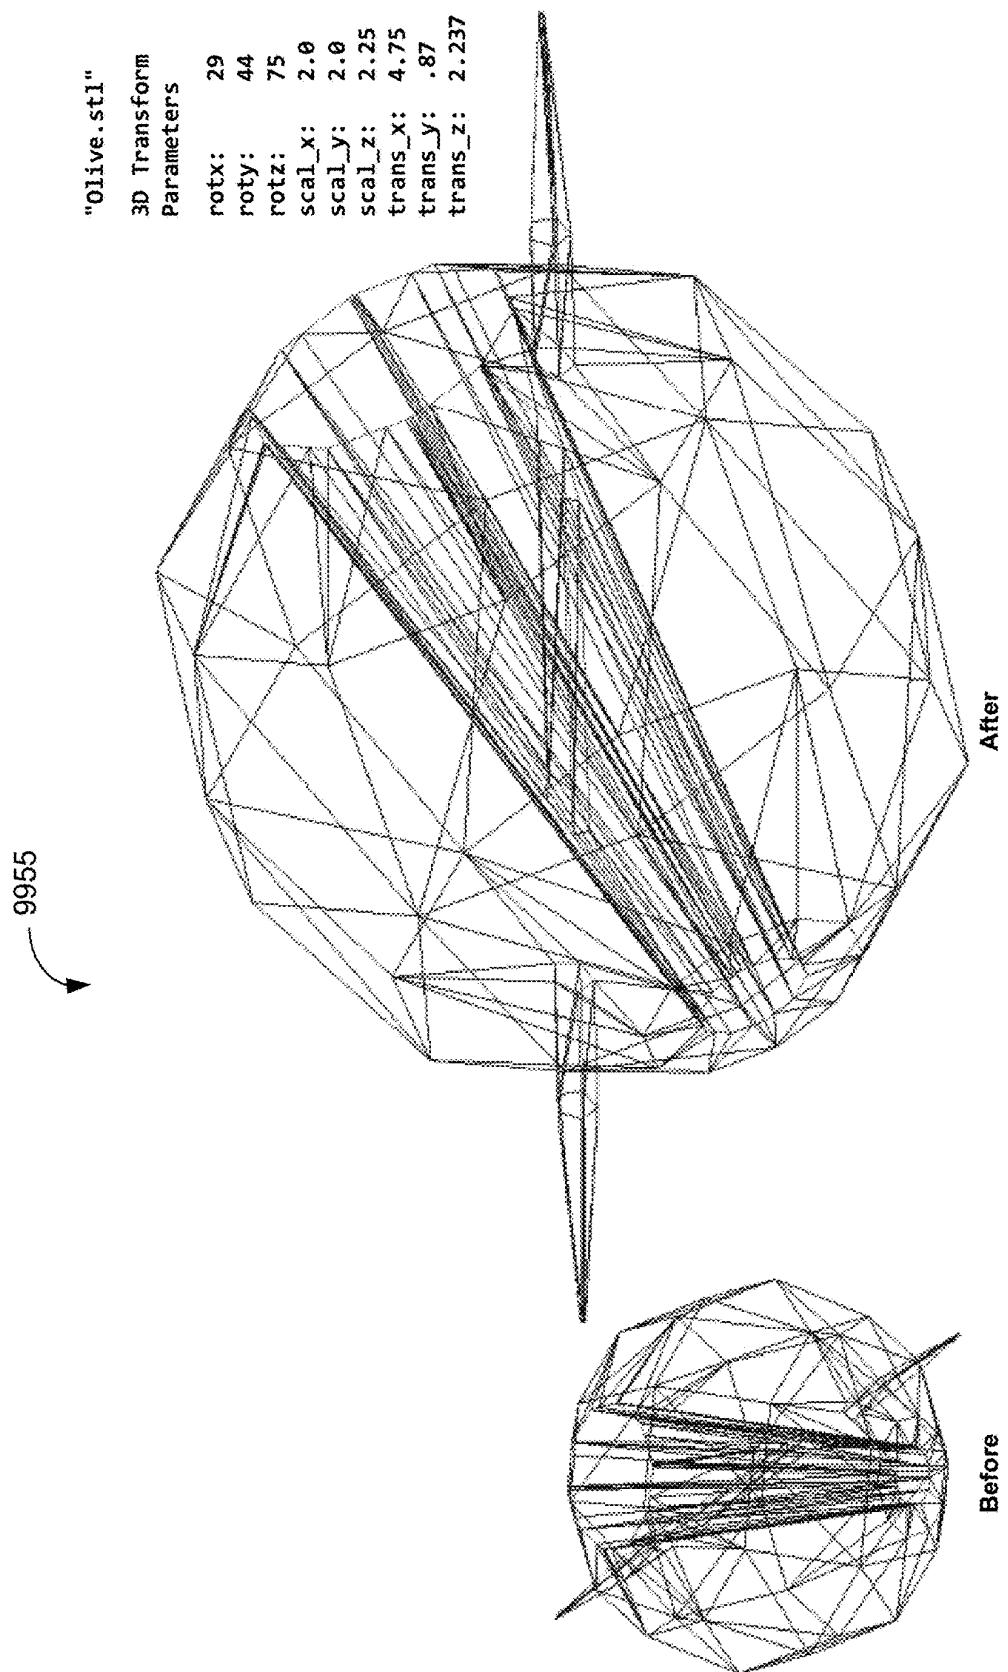


FIG. 51

1

**UNIVERSAL FLOATING-POINT
INSTRUCTION SET ARCHITECTURE FOR
COMPUTING DIRECTLY WITH DECIMAL
CHARACTER SEQUENCES AND BINARY
FORMATS IN ANY COMBINATION**

**CROSS-REFERENCE TO RELATED
APPLICATIONS**

This application claims the priority benefit under 35 U.S.C. § 119(e) of U.S. Provisional Application No. 62/886, 570 filed on Aug. 14, 2019, the disclosure of which is incorporated herein by reference in its entirety.

COPYRIGHT NOTICE

A portion of the disclosure of this patent document contains material which is subject to copyright protection. The copyright owner has no objection to the facsimile reproduction by anyone of the patent document or the patent disclosure, as it appears in the Patent and Trademark Office patent file or records, but otherwise reserves all copyright rights whatsoever.

TECHNICAL FIELD

The present disclosure relates generally to digital processors, and more particularly to a universal floating-point Instruction Set Architecture (ISA) computer.

BACKGROUND

Conventional architectures such as Reduced Instruction Set Architecture (RISC), Complex Instruction Set Computing (CISC), and General-Purpose computing on Graphics Processing Units (GP-GPU) are woefully inadequate for applications involving computations on large floating-point data sets where the raw input is in the form of human-readable decimal character sequences of 9-20 decimal digits or more. Among the reasons they are inadequate is because nine decimal digits is nine bytes, not counting “+” or “-” sign, “.” character, or exponent characters. These characters, together, can add an additional seven bytes to the overall length of the decimal character representation. Stated another way, for a 9-digit representation in formal scientific floating-point notation, up to 16 bytes are required just to store a single representation.

For a 20-byte representation, up to 27 bytes are required. Conventional architectures require that these long character sequences be explicitly converted to IEEE 754-formatted binary representations prior to actual use in computation. Among the first steps in this process is simply loading one or more characters of a given string into a working register, which is typically no wider than 64 bits (8 characters/bytes). Thus, these conventional machines cannot directly compute with these character sequences because they can typically read or write no more than 64 bits (8 characters) at a time. To carry out an IEEE 754-2008 “H=20” double-precision conversion, conventional computers must go through a thousand or more clock cycles of computation, especially for double-precision (binary64), for a single decimal character sequence representation.

Assuming, arguendo, that these machines have special hardware that can do a portion of the conversion process using just three or four specialized instructions after the entire raw character sequence is loaded into a supposed special register, a double-precision binary64 conversion

2

could be completed with a latency of just 30 or so clock cycles. However, there is still the problem of how to hide this latency.

One common method for hiding latency is designing the architecture so that it is of the interleaving, multi-threading type. Basically, an interleaving, multi-threading architecture is one in which the workload is distributed among threads that overlap during execution, in a round-robin fashion, with each thread having its own time-slot for a pre-determined number of clocks, usually just one clock in heavy workload environments. With special pipelined hardware for doing these conversions, all the thread has to do during its first several time-slots is push the entire decimal character representation into this special hardware, such that, once this 10 first time slot has been consumed and all the time slots of all the other threads have been likewise consumed and the first thread's next time slot arrives, all the first thread then has to do is read the converted result out and store it elsewhere, with the other threads doing the same.

15 The problem with employing an interleaving, multi-threading approach to hiding extra-long execution cycles is that even with 8 threads, this is not sufficient to completely hide a 30-or-more-clock conversion pipeline. Eight (8) threads can only hide 8 clocks worth of latency. In this 20 scenario, the overall pipeline of the processor will stall for 22 clocks per conversion while waiting for the current conversion to complete. This strategy works fine for hardware operators with latencies of 8 or fewer clocks, but not for something that has a 30-or-more-clock latency, such as 25 30 would be the case for such hardware IEEE 754-2008 mandated operators as convertToDecimalCharacter, convertFromDecimalCharacter, square Root, and Remainder to name a few.

One might propose designing a processor that has 30 threads so that latencies of all types of operations up to 30 clocks deep can be completely hidden. But this is simply not practical. The main reason is that, while all these conventional threads share the same hardware operators, opcodes, and the like, each thread requires its own program counter 35 (PC), Stack Pointer, working register set, Status Register, and so forth. This creates logic bloat when you start sprinkling this type processor on a chip. Not only that, but such a huge processor would be extremely difficult to debug. Imagine a hydra with 30 heads all joined at the hip as would 40 be the case for a 30-thread interleaving, multi-threading machine. The place where they join together, i.e., where the instruction's opcode decode and execution unit is situated, 45 creates an inherent bottleneck.

Moreover, assuming such a processor will always have a 50 workload heavy enough to keep all threads occupied all the time, more time and resources will be required just to divide up, distribute, and manage the workload among 30 threads (again, the idea of interleaving threads being to hide long execution latencies). Without a workload sufficient to keep 55 all threads occupied all the time, the hardware to implement 30 threads would be a colossal waste of resources that could be better spent elsewhere or not at all.

SUMMARY

The present disclosure provides a new kind of universal floating-point ISA that can “push” and “pull” in a single clock cycle, dual operands including not only long decimal character sequences up to 128 characters in length each, but also “gobs” of data that can be a mix of character sequences, IEEE 754-2008 binary formatted floating-point numbers, integers, and basically any combination thereof, into and out 60 65

of memory-mapped logical, integer and floating-point operators. In certain heavy (“big data”) workload environments, the disclosed ISA can completely hide long latencies without the use of any interleaving, multi-threading methods, or hardware. In fact, the disclosed universal floating-point ISA is a new kind of processor that does NOT employ opcodes at all, but rather is a pure “mover”-style architecture that can push and pull “gobs” of data without the encumbering prior-art “opcode” fetch, decode and execute bottleneck.

Additionally, the new universal floating-point ISA has direct, indirect, immediate, and table-read addressing modes, with the indirect addressing mode having both programmable offset and auto-post-modification of index register capability and a memory-mapped “REPEAT” counter so that, with a single instruction, entire vectors can easily and efficiently be pulled from and pushed into stand-alone operators capable of accepting dual operands of 1 to 128 bytes each—every clock cycle—all without the use of opcodes in the instruction.

Rather than opcodes, the disclosed universal floating-point ISA employs stand-alone, memory-mapped operators, complete with their own pipeline that is completely decoupled from the processor’s primary push-pull pipeline, automatic semaphore, and their own private, directly and indirectly addressable, three-ported (one write-side and two read-side) result buffers. Results, including their exception signals, automatically spill into the private result buffers and are automatically copied into the processor Status Register when pulled from such result buffers. The two read-side ports of the result buffers enable simultaneous pulling-thenpushing, with a single instruction, two operands (operandA and operandB) into the same or different operator in a single clock.

Likewise, the disclosed universal floating-point ISA has at least some directly and indirectly addressable three-port Static Random-Access Memory (SRAM). At least some of this SRAM is able to simultaneously read 1, 2, 4, 8, 16, 32, 64, and 128 bytes at a time on both side A and side B in a single clock.

Furthermore, the disclosed universal floating-point ISA is capable of being easily and efficiently preempted during execution, with all operator result buffers being capable of being preserved and fully restored upon completion of the preempting process, interrupt service routines, and subroutine calls. Because some of the results (together with their exception signals stored in their respective result buffers) can be up to 1029 bits wide, the disclosed ISA employs a new kind of stack that can receive, in a single clock/push, all these bits for context save operations, and when the routine is completed, restore such buffer to its original state with a single pull.

Additionally, the disclosed universal floating-point ISA is capable of scaling its computing power by attaching to the parent CPU, eXtra (child) Compute Units (XCUs) that execute the same instruction set as the parent CPU and offload work the parent CPU pushes into them and pulls from them.

Additionally, the disclosed universal floating-point ISA has a “natural” real-time debug and monitoring capability designed into it that can perform real-time data exchange tasks via a host IEEE 1149.1 (JTAG) and/or host system interface, on-the-fly, without the use of interrupts, opcodes, or direct memory access (DMA) hardware. The real-time monitoring and debug interface is able to easily and effi-

ciently set hardware or software breakpoints, reset, single-step and provide at least a minimum level of real-time trace capability.

For web-based artificial intelligence (AI) applications where data is stored primarily in human-readable form, the disclosed universal floating-point ISA has the ability, in hardware, to compute directly with dual decimal character sequences up to IEEE 754-2008 “H=20” in length, without first having to explicitly perform a conversion-to-binary format process in software before computing with these human-readable floating-point or integer representations. Oftentimes, vectors of data can comprise hundreds, if not thousands, of decimal character sequence entries that conventional processors must explicitly convert beforehand to binary representations before entering the actual computational stream. By way of example, this means that if a vector of dual operands comprising decimal character sequences of at least twenty decimal digits and one hundred entries deep per operand, it could potentially take a conventional processor hundreds of thousands to well over a million clock cycles just to convert these strings to binary representation prior to actually computing with them.

The disclosed universal floating-point ISA performs all the above functions by dispensing with “opcodes” altogether; and employing: special three-port, 1024-bit wide SRAMS; a special dual asymmetric system stack; memory-mapped stand-alone hardware operators with private result buffers with simultaneously readable side-A and side-B read ports; and dual hardware H=20 convertFromDecimalCharacter conversion operators.

With the disclosed universal floating-point ISA and related hardware, the conversion of IEEE 754-2008 H=20 decimal character sequences to IEEE 754 binary format is automatic, if desired, and, for big data scenarios, such conversion can actually be free, in terms of clock cycles and explicit instructions to perform these conversions. For example, with a minimum dual operand vector of just thirty-two entries each of twenty-eight digit decimal sequences being fed into this design’s multi-function, universal Fused-Multiply_Add operator, the net clock requirement is only thirty-two clocks, which includes the automatic conversion of the sixty-four twenty-eight decimal character sequence representations. Meaning the conversion to binary step is free. This is due to the fact that by the time thirty-two H=20 decimal character sequences are pushed into the present invention’s hardware convertFromDecimalCharacter operator using the memory-mapped “REPEAT” operator, the results from the first push are already available for reading.

Stated another way, using the REPEAT operator to push qty. 32 H=20 decimal character sequence operands in the convertFromDecimalCharacter operator and then immediately use the REPEAT operator to pull them out in rapid succession, results in one clock to perform the push and one clock to perform the pull, per conversion. Meaning that the clocks required to do the actual conversion AND the target computation, if any, are “free” and completely hidden. Since the present invention is capable of performing two such conversions every clock cycle, such equates to only 0.5 clocks per conversion of 64 H=20 decimal character sequences, which includes both the push and the pull, due to the fact that the instant invention’s ability to push two H=20 decimal character sequences simultaneously, every clock cycle.

The disclosed universal floating-point ISA includes a “Universal” Fused-Multiply-Add (FMA) operator that is also “multi-mode”. Universal in this context means the

FMA operator can directly accept IEEE-754 half-precision binary16, single-precision binary32, double-precision binary64, decimal character sequences up to 28 decimal digits in length, including decimal character sequences with “token” exponents such as the trailing letters “K”, “M”, “B”, “T” and the character “%”—in any combination—without the CPU or XCU having to first explicitly convert them to binary64. Such automatic conversion is designed into this Universal FMA’s pipeline hardware. Here, “K” is the token for exponent “e+003”, “M” for “e+006”, “B” for “e+009”, “T” for “e+012” and “%” for “e-002”. Moreover, results of each Universal FMA computation are automatically converted into binary16, binary32, or binary64 format as specified by the originating instruction, such conversion logic also built into the this operator’s pipeline hardware.

In the context of the Universal FMA operator, “multi-mode” means the Universal FMA operator can operate as a conventional FMA by bypassing the automatic decimal-character-to-binary conversion circuit, thereby making its pipe much shorter. The Universal FMA operator can also be used in a second mode as a “single” or “dual” decimal-character-to-binary conversion operator—only, by bypassing the Universal FMA operator and simultaneously converting and writing two results into this operator’s result buffers every clock cycle. Such a capability is essential mainly due to the fact each decimal-character-to-binary conversion circuit is rather large and, as such, in many applications it would be inefficient to incorporate a third stand-alone decimal-character-to-binary conversion circuit to support conversions in computations not involving the Universal FMA operator.

As a third mode, the Universal FMA operator can be used as a very powerful Sum-of-Products operator, which is similar to a generic FMA, except the Sum-of-Products includes 32 “fat” accumulators (one for each operator input buffer) on the output of its adder for accumulation of each iteration and to supply this accumulation back into this adder’s “C” input for the next summation. This accumulator is special, because it has an additional 10 lower bits and its contents are never rounded, meaning that the intermediate value actually stored in a given FMA accumulator is never rounded, but such value IS rounded during final conversion to one of the IEEE 754 binary formats specified by the originating instruction immediately before automatic storage into one of the 32 result buffers specified in the originating instruction. Hence the word “Fused” in the IEEE 754-2008 mandated “Fused-Multiply-Add” operation. Here, “Fused” means no rounding during the operation. The extra 10 bits of fraction in the accumulator are there mainly to help absorb underflows until the final iteration. This amount can be increased or decreased by design as the application requires. This sum of products mode is essential for fast, efficient, and precise computation of tensors.

To enable single-clock movement of dual operands (ie, operandA and operandB), whose length can be up to 128 bytes each (for a total of up to 256 bytes), it is now apparent that a special kind of new and novel processor is needed that has embedded in it, specially designed dual-bus infrastructure and memory that can be simultaneously written at 1 to 128 bytes and read at up to 256 bytes per clock cycle, all without the use of opcodes. The disclosed universal floating-point ISA is such a processor.

In one embodiment, the present disclosure provides an opcode-less universal floating-point Instruction Set Architecture (ISA) computer implemented entirely in hardware. The ISA computer includes a program memory and hardware circuitry connected to the program memory. The

hardware circuitry is configured to compute directly with human-readable decimal character sequence floating-point representations without first having to explicitly perform a conversion-to-binary-format process in software before computing with the human-readable decimal character sequence floating-point representations. The hardware circuitry is configured to accept, in any combination, human-readable decimal character sequence floating-point representations or IEEE 754-2008 standard binary arithmetic format floating-point representations, wherein the human-readable decimal character sequence floating-point representations are up to IEEE 754-2008 “H=20” in length. The IEEE 754-2008 standard binary arithmetic format floating-point representations may be IEEE 754-2008 standard binary16 (half-precision), binary32 (single-precision), and binary64 (double-precision) floating-point representations.

In another embodiment, the present disclosure provides a method for converting long decimal character sequences using relatively small dual-half system abbreviated look-up tables and on-the-fly interpolation of binary weights derived from these tables.

In another embodiment, the present disclosure provides a novel asymmetric dual hardware stack for saving in one clock cycle and restoring in one clock cycle, contents of certain memory-mapped hardware operator result buffers, along with their status/exception signals simultaneously.

In yet another embodiment, the present disclosure provides a universal, multi-function Fused-Multiply-Add-Accumulate floating-point operator embedded in the ISA computer for directly computing long series vectors whose data comprise long decimal character sequences and/or binary formats in any combination without first having to explicitly convert them from-to any of these formats beforehand. The multi-function universal FMA operator can be used for fused-multiply-add, sum-of-products, and simultaneous dual decimal-character-to-binary format conversion operations. Because the disclosed universal floating-point ISA can move dual “GOBs” of data of up to 128 bytes each and every clock cycle into memory-mapped operators, such operators can be parallel clusters of operators as in the case of tensors employed for artificial intelligence (AI), deep learning, neural networks, and the like. For example, in the case of binary16 (half-precision) vectors, the disclosed universal floating-point ISA can write to up to quantity 64 sum-of-products operators simultaneously in a single clock cycle (quantity 32 for binary32 and quantity 16 for binary64 formatted numbers). The dual data being written (i.e., operandA and operand B) into such operators can be decimal-character sequences (up to 28 decimal digits in length each), decimal-character sequences with token exponents, half-precision binary16, single-precision binary32, and/or double-precision binary64 formatted representations—in any combination. These formats are automatically converted by novel hardware embedded in such operator(s) pipelines and, once a binary64 result is obtained, again converted to the target binary format immediately prior to automatic storage into one of several operator result buffers specified by the original single “pull”/“push” instruction executed.

Because this novel architecture employs no op-codes, the core processor instruction pipeline knows how to do one thing and one thing only—simultaneous “pull” and “push” single or dual operands—which requires no opcodes.

Further features and benefits of embodiments of the disclosed apparatus will become apparent from the detailed description below.

BRIEF DESCRIPTION OF THE DRAWINGS

In the following section, the invention will be described with reference to exemplary embodiments illustrated in the figures, in which:

FIG. 1 is a diagram of bit fields in an exemplary embodiment of the present disclosure's 64-bit instruction word;

FIG. 2 is a diagram of a special case for the srcB field of the instruction word specifying the number of bits to shift and the shift type for use with the 64-bit SHIFT operator in an exemplary embodiment of the present disclosure;

FIG. 3 is a diagram of a special case for the srcB field of the instruction word specifying the bit position to test and the displacement amount for use with the conditional relative branch operator in an exemplary embodiment of the present disclosure;

FIG. 4 is a diagram illustrating an exemplary embodiment of the present disclosure's default floating-point decimal character sequence input and output format 300 used by the processor's memory-mapped hardware IEEE 754-2008 H=20 convertFromDecimalCharacter, convertToDecimalCharacter, and Universal FMA operators;

FIG. 5 is a diagram illustrating examples of various decimal character sequences, including some with token exponents, their translation to the default format, and their respective IEEE 754 binary64 equivalent representations in an exemplary embodiment of the present disclosure;

FIG. 6 is a block diagram of an exemplary embodiment of the basic architecture of the disclosed universal floating-point ISA, showing the CPU or parent processor;

FIG. 7 is a block diagram of an exemplary embodiment of the basic architecture of an eXtra Processing Unit (XCU) or a child processor;

FIG. 8 is a block diagram illustrating an arrangement of a CPU/parent processor and one to sixteen XCU/child processors in an exemplary embodiment of the present disclosure;

FIG. 9 is a modified Harvard model data memory-map employed by both the CPU and XCU(s) in an exemplary embodiment of the present disclosure;

FIG. 10 is a modified Harvard model program memory-map employed by both the CPU and XCU(s) in an exemplary embodiment of the present disclosure;

FIG. 11 is an exemplary memory-map of the disclosed universal floating-point ISA's memory-mapped programming model register set in an exemplary embodiment of the present disclosure;

FIG. 12A is an exemplary memory-map of the disclosed universal floating-point ISA's hardware-implemented IEEE 754-2008 Mandated Computational floating-point Operators in an exemplary embodiment of the present disclosure;

FIG. 12B is an exemplary memory-map of the disclosed universal floating-point ISA's hardware-implemented IEEE 754-2008 Mandated Comparison Predicates operators in an exemplary embodiment of the present disclosure;

FIG. 12C is an exemplary memory-map of the disclosed universal floating-point ISA's hardware-implemented IEEE 754-2008 Mandated dual-operand, non-computational, non-exceptional operators in an exemplary embodiment of the present disclosure;

FIG. 12D is an exemplary memory-map of the disclosed universal floating-point ISA's hardware-implemented IEEE 754-2008 Mandated single-operand, non-computational, non-exceptional operators in an exemplary embodiment of the present disclosure;

FIG. 12E is an exemplary memory-map of the disclosed universal floating-point ISA's hardware-implemented IEEE

754-2008 miscellaneous mandated operators in an exemplary embodiment of the present disclosure;

FIG. 12F is an exemplary memory-map of the disclosed universal floating-point ISA's hardware-implemented, stand-alone floating-point computational operators not mandated by IEEE 754-2008 in an exemplary embodiment of the present disclosure;

FIG. 13 is an exemplary table showing the disclosed universal floating-point ISA's hardware-implemented native logical and integer arithmetic operators in an exemplary embodiment of the present disclosure;

FIG. 14A is a simplified schematic diagram of an exemplary embodiment of a circuit used to implement the disclosed universal floating-point ISA's memory-mapped hardware Auxiliary Registers (ARn) for indirect addressing;

FIG. 14B is a simplified schematic diagram of an exemplary embodiment of the disclosed universal floating-point ISA's Stack Pointer (SP) for indirect addressing, which is part of the auxiliary register logic block of FIG. 14A;

FIG. 14C is a simplified schematic diagram illustrating how SourceA, SourceB and Destination direct and indirect addresses are generated from the disclosed universal floating-point ISA's instruction in an exemplary embodiment of the present disclosure;

FIG. 15 is a schematic diagram of an exemplary embodiment of the CPU and XCU pre_PC used for accessing the next instruction;

FIG. 16 is a schematic diagram of an exemplary embodiment of the CPU's and XCU's memory-mapped program counter (PC) operator;

FIG. 17 is a schematic diagram of an exemplary embodiment of the CPU's and XCU's memory-mapped PC-COPY register;

FIGS. 18A and 18B are an exemplary table illustrating the bits of the CPU's and XCU's memory-mapped STATUS register/operator and their respective functions in an exemplary embodiment of the present disclosure;

FIG. 18C is a block diagram illustrating an exemplary embodiment of an arrangement, mapping, and implementation of IEEE 754-2008 mandated Comparison Predicates, dual-operand and single-operand non-computational, non-exceptional operators in relation to their respective bits in the CPU's and XCU's memory-mapped STATUS register/operator;

FIG. 18D is a schematic diagram illustrating, along with their respective bit positions, exemplary logic for carrying out in hardware, bit manipulation of the Enable Alternate Immediate exception handling (bits 31-35) as a "group" within the memory-mapped STATUS register/operator in an exemplary embodiment of the present disclosure;

FIG. 18E is a schematic diagram illustrating, along with their respective bit positions, exemplary logic for carrying out in hardware bit manipulation of the Raise No Flag specifiers for the five IEEE 754-2008 exceptions (bits 26-30) as a "group" within the memory-mapped STATUS register/operator in an exemplary embodiment of the present disclosure;

FIG. 18F is a schematic diagram illustrating, along with their respective bit positions, exemplary logic for carrying out in hardware bit manipulation for the IEEE 754-2008 Inexact, Underflow, Overflow, Divide-by-Zero and Invalid flags (only the first three bits, bits 23-25, are shown due to space limitations) as a "group" within the memory-mapped STATUS register/operator in an exemplary embodiment of the present disclosure;

FIG. 18G is a schematic diagram illustrating, along with their respective bit positions, exemplary logic for carrying

out in hardware bit manipulation for the IEEE 754-2008 Inexact, Underflow, Overflow, Divide-by-Zero and Invalid flags (only the last two bits, bits 21 and 22, are shown due to space limitations) as a “group” within the memory-mapped STATUS register/operator in an exemplary embodiment of the present disclosure;

FIG. 18H is a schematic diagram illustrating, along with their respective bit positions, exemplary logic for carrying out in hardware bit manipulation for the IEEE 754-2008 Inexact, Underflow, Overflow, Divide-by-Zero and Invalid “signals” (only the first three bits, bits 18-20, are shown due to space limitations) as a “group” within the memory-mapped STATUS register/operator in an exemplary embodiment of the present disclosure;

FIG. 18I is a schematic diagram illustrating, along with their respective bit positions, exemplary logic for carrying out in hardware bit manipulation for the IEEE 754-2008 Inexact, Underflow, Overflow, Divide-by-Zero and Invalid “signals” (only the last two bits, bits 16 and 17, are shown due to space limitations) as a “group” within the memory-mapped STATUS register/operator in an exemplary embodiment of the present disclosure;

FIG. 18J is a schematic diagram illustrating, along with their respective bit positions, exemplary logic for carrying out in hardware bit manipulation for the CPU and XCU logical and integer arithmetic Carry (“C”), Negative (“N”), Done, Interrupt Enable (“IE”), Zero (“Z”), and Overflow (“O”) flags (only bits 1, 2, 4, and 5 are shown due to space limitations) as a “group” within the memory-mapped STATUS register/operator in an exemplary embodiment of the present disclosure;

FIG. 18K is a schematic diagram illustrating, along with their respective bit positions, exemplary logic for carrying out in hardware integer comparisons in addition to bit manipulation for the CPU and XCU logical and integer arithmetic Zero (“Z”) and Overflow (“O”) flags (bits 0 and 3) as a “group” within the memory-mapped STATUS register/operator in an exemplary embodiment of the present disclosure;

FIGS. 18L and 18M are a schematic diagram illustrating, along with their respective bit positions, exemplary logic for carrying out in hardware bit manipulation for the IEEE 754-2008 “recommended” substitutions for abrupt underflow, substitute X, substitute xor(X), inexact, underflow, overflow, substitute overflow, divide-by-zero, and invalid exceptions (bits 56-58) as a “group” within the memory-mapped STATUS register/operator in an exemplary embodiment of the present disclosure;

FIG. 18N is a schematic diagram illustrating, along with their respective bit positions in the STATUS register/operator, exemplary logic for carrying out in hardware bit manipulation as a group, dynamic rounding mode attributes mandated by IEEE 754-2008, namely, the encoded Rounding Mode bits 1 and 0, Away bit, Enable dynamic rounding mode, and default override bit (bits 51-55) in an exemplary embodiment of the present disclosure;

FIG. 18O is a schematic diagram illustrating an exemplary embodiment of memory-mapped logic for carrying out in hardware the IEEE 754-2008 mandated testing for, as a group, an exception flag raised condition using a “testSavedFlags” or “testFlags” memory decode, as well as restoring this status bit using the “loadStatusReg” memory decode logic for these memory-mapped operators;

FIG. 19 is a schematic diagram illustrating an exemplary embodiment of the disclosed universal floating-point ISA’s memory-mapped hardware REPEAT counter circuit;

FIG. 20 is a schematic diagram illustrating an exemplary embodiment of the disclosed universal floating-point ISA’s memory-mapped hardware loop-counter operators;

FIG. 21 is a schematic diagram illustrating an exemplary embodiment of the disclosed universal floating-point ISA’s optional floating-point exception capture module;

FIG. 22A is a block diagram illustrating an exemplary embodiment of the disclosed universal floating-point ISA’s hardware implementation of IEEE 754-2008 mandated computational operator module showing dual operand inputs, their exception signals, dual result outputs and ready semaphore output;

FIG. 22B illustrates exemplary memory-mapped hardware implementations of IEEE 754-2008 convertToDecimalCharacter, Addition, Fused-Multiply-Add 9440, and convertFromDecimalCharacter operator module inputs and outputs in an exemplary embodiment of the present disclosure;

FIG. 23 is a block diagram illustrating an exemplary embodiment of the disclosed universal floating-point ISA’s logical and integer arithmetic operator module illustrating dual operand inputs, their signals, dual result outputs and ready semaphore output;

FIG. 24 is a block diagram illustrating an exemplary embodiment of the disclosed universal floating-point ISA’s hardware implementation of a double-precision IEEE 754-2008 H=20+ convertFromDecimalCharacter operator illustrating a virtually identical dual half-system approach;

FIG. 25 is a schematic diagram illustrating an exemplary embodiment of a circuit employed by the integer part quantizer/encoder of FIG. 24 to compute/encode the integer part intermediate mantissa;

FIG. 26 is a schematic diagram of an exemplary embodiment of a circuit employed by the fraction part quantizer/encoder of FIG. 24 to compute/encode the fraction part intermediate mantissa;

FIG. 27 is a schematic diagram illustrating an exemplary embodiment of the disclosed universal floating-point ISA’s convertFromDecimalCharacter operator’s look-up ROMs for the integer part, showing the interpolation method for determining the weights and binary exponents derived from a decimal exponent input obtained from the original decimal character sequence input;

FIG. 28A is a schematic diagram illustrating an exemplary embodiment of the disclosed universal floating-point ISA’s convertFromDecimalCharacter operator’s look-up ROMs for the fraction part, illustrating the interpolation method for determining the weights and binary exponents derived from a decimal exponent input obtained from the original decimal character sequence input;

FIG. 28B is a schematic diagram illustrating an exemplary embodiment of the disclosed universal floating-point ISA’s convertFromDecimalCharacter operator’s look-up ROMs for the fraction part subnormal exponent inputs, illustrating the interpolation method for determining the weights and binary exponents derived from a decimal exponent input obtained from the original decimal character sequence input;

FIG. 29 is a block diagram illustrating an exemplary embodiment of the disclosed universal floating-point ISA’s hardware implementation of a stand-alone, fully pipelined, memory-mapped H=20+ convertToDecimalCharacter component/module mandated by IEEE 754-2008;

FIG. 30 is a block diagram illustrating an exemplary embodiment of the binary64 to H=20 decimal character sequence converter of FIG. 29;

FIG. 31 is a block diagram illustrating an exemplary embodiment of the binary-to-decimal-character conversion engine of FIG. 30 illustrating virtually identical dual half-systems, one for the integer part and one for the fraction part;

FIG. 32 is a block diagram illustrating an exemplary embodiment of the integer part binary-to-decimal-character summing circuit of FIG. 31, including integer part weights look-up ROM block, conditional summing circuit, and rounding circuit;

FIG. 33 is a partial detail of an exemplary embodiment of the contents of the integer part binary-to-decimal-character mantissa D52 weight look-up ROM;

FIG. 34A is a diagram illustrating an exemplary embodiment of the method/algorithm used for computing both the integer part intermediate value and fraction part intermediate value that are submitted to their respective BCD converter circuits, including the method for obtaining a Guard, Round, and Sticky bit for each part, of the disclosed universal floating-point ISA's double-precision IEEE 754-2008 H=20+ convertToDecimalCharacter operator;

FIG. 34B is a diagram illustrating an exemplary embodiment of the method/algorithm used for computing the sum of the truncated part (i.e., second 22 digits) used in the computation of the respective integer part intermediate value and fraction part intermediate value of the disclosed universal floating-point ISA's convertToDecimalCharacter operator, including a method for deriving a truncated part GRS used in the final sum;

FIG. 35 is a schematic diagram illustrating an exemplary embodiment of an integer part rounding circuit that correctly rounds the integer part intermediate result prior to submission of the intermediate result to the BCD conversion circuit;

FIG. 36 is a block diagram illustrating an exemplary embodiment of the fraction part binary-to-decimal-character summing circuit of the fraction-part half-system, comprising a fraction part weights look-up ROM block, conditional summing circuit, and rounding circuit;

FIG. 37 is a partial detail illustrating an exemplary embodiment of the convertToDecimalCharacter fraction part ROM weight look-up contents showing the first 20 digits and the second 22 digits (truncated part of the weight), along with the actual Verilog RTL source code employed to obtain a mantissa mask used during the hardware computation;

FIG. 38 is a partial detail illustrating an exemplary embodiment of the look-up ROM block 9330 and actual Verilog RTL source code used by the convertToDecimalCharacter operator for converting the adjusted binary exponent input to an adjusted decimal exponent for both normal and subnormal numbers;

FIG. 39 is a schematic diagram illustrating an exemplary embodiment of the rounding circuit used for rounding the fraction part intermediate result prior to submission of the intermediate result to the BCD conversion circuit used in the disclosed universal floating-point ISA's convertToDecimalCharacter hardware operator;

FIGS. 40A, 40B, and 40C are block diagrams that together show, respectively, the upper left-most, lower right-most, and lower left-most sections of the fully pipelined binary-to-binary-coded-decimal (BCD) conversion block used by the integer part half-system and the fraction part half-system to convert their respective rounded 68-bit binary outputs to BCD.

FIG. 41 is a block diagram illustrating an exemplary embodiment of a memory-mapped, fully restorable, hardware-implemented, double-precision floating-point “addi-

tion” operator module, including 16-entry by 69-bit SRAM result buffer, restore capability, and ready semaphore;

FIG. 42 is a block diagram illustrating an exemplary embodiment of a memory-mapped, stand-alone, fully restorable, hardware-implemented, double-precision floating-point “multiplication” operator module including 16-entry by 69-bit SRAM result buffer, restore capability, and ready semaphore;

FIG. 43 is a block diagram of an exemplary embodiment of a memory-mapped, fully restorable, stand-alone, double-precision (H=20), hardware-implemented double-precision floating-point “H=20” convertFromDecimalCharacter operator module, including 32-entry by 69-bit SRAM result buffer, restore capability, and ready semaphore;

FIG. 44 is a block diagram of an exemplary embodiment of a multi-function dual asymmetric “fat” stack and “fat” SRAM block used for operator context save-restore operations and other general-purpose functions;

FIG. 45 is a block diagram of an exemplary embodiment of a memory-mapped, fully restorable, stand-alone, double-precision floating-point “fusedMultiplyAdd” (FMA) operator module, which is designed to also operate as a sum-of-products operator;

FIG. 46A is a top-level block diagram of an exemplary embodiment of the present disclosure's memory-mapped, stand-alone, fully restorable, multi-function, Universal Fused-Multiply-Add (FMA) (and accumulate) operator module, including dual convertFromDecimalCharacter converters on the input;

FIG. 46B is a block diagram illustrating an exemplary embodiment of the Universal FMA (and accumulate) circuit and “split” SRAM block for storage of either dual convertFromDecimalCharacter results, FMA or sum-of-products results, readable on side A and side B;

FIG. 46C is a block diagram illustrating an exemplary embodiment of the FMA circuit employed by the multi-function universal FMA operator module to perform both pure FMA and sum-of-products computations;

FIG. 47A is a block diagram of an exemplary embodiment of an optional hardware JTAG-accessible, breakpoint, trace buffer, and real-time monitor/debug module that enables on-the-fly, real-time-data-exchange operations between the parent CPU and up to quantity (16) child XCs attached to it within the same device;

FIG. 47B is a block diagram illustrating an exemplary embodiment of a breakpoint, single-step and real-time monitor/debug module;

FIG. 47C (Prior Art) is a block diagram illustrating a conventional industry standard IEEE 1149.1 (JTAG) state machine Test Access Port (TAP);

FIG. 48A is a simplified schematic and pertinent Verilog RTL source code describing behavior of the XCU breakpoint module in an exemplary embodiment of the present disclosure;

FIG. 48B illustrates exemplary snippets of Verilog RTL showing memory mapping and behavioral description of the parent CPU's XCU hardware breakpoint control and status registers in an exemplary embodiment of the present disclosure;

FIG. 48C is a diagram illustrating, in an exemplary embodiment, fusing of the parent CPU monitor read instruction to the child XCU monitor read instruction assembled by the parent CPU;

FIG. 48D is a diagram illustrating, in an exemplary embodiment, fusing of the parent CPU monitor write instruction to the child XCU monitor write instruction assembled by the parent CPU;

FIG. 49 is a block diagram illustrating an exemplary embodiment of a double-quad, single-precision (H=12) Universal FMA operator that can accept in a single push, quantity (8) 16-character decimal character sequence or binary32 format numbers as operandA and quantity (8) 16-character decimal character sequence or binary32 format numbers as operandB, outputting as quantity (8) binary32 format numbers (including corresponding exceptions) for a total 512 bits, or quantity (8) binary32 format numbers only, for a total of 256 bits for each pull;

FIGS. 50A through 50L are a post-assembly listing of an example program written in the instant ISA assembly language that employs up to quantity (16) child XCUs to perform a 3D transform (rotate, scale, and translate) on all three axes, of a 3D object in .STL file format and to write the results of such transform back out to external program memory.format and to write the results of such transform back out to external program memory; and

FIG. 51 is an actual wire-frame “Before” and “After” rendering of a simple cocktail “olive” 3D model in .STL file format performed by from 1 to 16 child XCUs or solo parent CPU using the scale, rotate, and translate parameters shown for each axis.

DETAILED DESCRIPTION

The present disclosure will now be described more fully hereinafter with reference to the accompanying drawings, in which preferred embodiments of the disclosure are shown. In the below, for purposes of explanation and not limitation, specific details are set forth in order to provide a thorough understanding of the present disclosure. It will be apparent to one skilled in the art that the present disclosure may be practiced in other embodiments that depart from these specific details.

The IEEE 754-2008 Standard for Floating-Point Arithmetic is hereby incorporated by reference herein. Likewise, the IEEE 1149.1-1990 Joint Test Action Group (JTAG) standard is hereby incorporated by reference herein.

The disclosed embodiment of the Universal Floating-Point ISA processor employs a doubly modified Harvard architecture memory model, meaning that it comprises separate program and data memory buses, but with the added ability to access program memory as if it were data memory, doubly because the present architecture can push and pull two operands simultaneously. It has immediate, direct and indirect addressing modes, as well as table read from program memory using direct or indirect addressing mode. The modified Harvard model can be implemented as either a single-thread processor or a interleaving, multi-threading processor, wherein the threads share the same operators, but have their own processor registers (such as PC, SP, Auxiliary Registers, Status Register, etc.) and can share some or all the system memory. This ISA can be easily adapted as a Von Neumann memory model as well.

Generally speaking, the source and/or destination operator(s) ultimately determine the meaning of certain bit fields of the instruction.

FIG. 1 is a diagram of bit fields 110 in an exemplary embodiment of the present disclosure’s 64-bit instruction word. The bit fields are broken down as follows:

Bits 63 and 62 make up the RM[1:0] “static” directed rounding mode specifier 150 that can be used by floating-point hardware operators to effectuate a correctly rounded floating-point result according to this specifier. This round-

ing mode specifier is lower priority to the “dynamic” directed rounding mode attribute specifier in the processor’s STATUS register, if enabled.

This “opcode-less” ISA has only four mnemonics: “_”, “P”, “N” and “Z”, which may be used to specify the static rounding direction of a given floating-point instruction. Each assembly line begins with one of these four mnemonics, usually after a tab or unique label identifying the source line. Each of these four single-character mnemonics, when encountered, signal the assembler to insert a “00”, “01”, “10”, or “11” in the RM[1:0] bit positions 150 of the instant instruction as follows:

Mnemonic/2-Bit Value

“_”=00=round using current “default” rounding direction, i.e., “nearest”. Note: the “default” rounding direction can be changed to “away” by setting bit 54 in the STATUS register.
“P”=01=round towards positive infinity
“N”=10=round towards negative infinity
“Z”=11=round towards zero

Note: the two RM bits 150 of the instruction can be overridden by setting bit 53 (enable RM attributes) of the STATUS register. If set, rounding mode attribute bits 51 and 52 of the STATUS register determine rounding direction using the same 2-bit code definitions above.

It should also be noted that IEEE 754-2008 does not, per se, anticipate or comprehend that rounding direction be embedded or signaled in the instruction, but rather as an attribute implemented in software. Consequently, some implementations of the instant ISA may not have use of the RM bits for signaling rounding direction within the instruction itself. In such cases, it may be more desirable to employ these two bits for some other signaling purpose, depending on the needs of the operator and the implementer.

DAM[1:0] Data Access Mode 160 specifies from which memory the operand read(s) are to take place for the instant instruction. The assembler/compiler determines what these two bits should be according to the addressing modes specified in the instant assembly line source statement. Their 2-bit encoded meanings are:

00=both operand A and operand B are read from data memory using either direct or indirect addressing modes.

01=operand A is either direct or indirect and operand B is immediate (i.e., immediately available within lower 16 bits of the srcB field 140 of the instant instruction). A “#” character 230 immediately preceding operandB in the source line signals the assembler that operandB is “immediate” and to insert a “01” into DAM[1:0] bit positions 160 of the instant instruction during assembly.

10=operandA is a table-read from program memory using direct (table read from program memory) addressing mode and operand B is either direct or indirect and NEVER immediate. An “@” character 220 immediately preceding operandA in the source line signals the assembler that operandA resides in program memory and to therefore insert a “10” into DAM[1:0] bit positions 160 of the instant instruction during assembly.

11=32-bit immediate (i.e., immediately available within instant instruction). If there is only one operand, i.e., operandA all by itself on the assembly line and it is immediately preceded by a “#” character, this signals the assembler that the sole operandA is a 32-bit immediate value within the instant instruction and to insert a “11” into DAM[1:0] bit positions 160 of the instant instruction during assembly. Note: there are enough unused bits in this particular mode to actually make this a 40-bit immediate value, but the present assembler does not support 40-bit immediate values.

15

DEST **120**, srcA **130** and srcB **140** Bit Fields

The instant ISA is designed to effectuate in a single clock simultaneous movement of single or dual operands, operandA and operandB, to a destination specified in DEST **120** using the addressing mode, size and signal specified in their respective srcA **130** and srcB **140** bit fields. The source and destination addresses can be processor hardware operator result buffers, data memory, both on-chip and off-chip, program memory, basically anything within the processor's address map, which also includes the entire memory map of any child XCU's attached to the parent CPU.

Contained within each of the DEST **120**, srcA **130** and srcB **140** bit fields are size specifier **180**, direct/indirect addressing mode specifier **190** and, if indirect addressing mode is specified in **190**, which type of indirect addressing mode is to be employed for the instant access as specified by **200**.

SIGA, SIGB and SIGD **170** are general-purpose "signals" whose meaning depends on context, such as implied by source and/or destination operator, if any. For example if the source(s) are SRAM locations and the destination is an integer arithmetic operator, then if the source operand's SIGA/B is set, this would usually signal that the sign of the operand is to be automatically sign-extended in hardware to 64 bits before being pushed into that operator. If the destination SIGD is set, this can be used to signal the operator that it is to employ signed arithmetic, as opposed to unsigned arithmetic, in that particular operation.

For floating-point operations such as the instant invention's Universal Fused-Multiply-Add operator, SIGA and SIGB, when set, are used to signal that the operands are human-readable decimal-character sequences as opposed to IEEE 754 binary format numbers, which is default. SIGD for this operator, when set along with SIGA and SIGB also being set, signals the operator to bypass FMA function and simply convert the two decimal-character sequences to the binary format specified by the destination size bits and automatically store them simultaneously in the result buffer location specified in the DEST **120**. If SIGA and SIGB are clear along with SIGD being set, this signals the Universal FMA to bypass the delay circuit used for coherency when both operands are binary format, which has the effect of substantially shortening the Universal FMA pipeline by about 23 clocks, as that particular operation does not involve a convertFromDecimalChar step. When SIGD is clear (0), this signals that both operands propagate thru the entire length of the pipeline to maintain coherency, regardless of the state of SIGA or SIGB.

SIGA, SIGB and SIGD Summarized

SIGA Signal for operandA field meaning depends on context—either "s" or "_" must immediately precede Size field and can be mixed and matched with the other operand's and/or destination SIG signals.

can mean: For integer or logical operators, "s"=signed (sign-extended);
"_"=default=0=unsigned (zero-extended), depending on target operator.

can also mean: 1=text (decimal character sequence);
0=binary format (mainly used with direct character sequence computations/operators).
Implementer can make it mean anything he/she wants it to mean, depending on the target operator.

SIGB Signal for operandB field meaning depends on context—either "s" or "_" must immediately precede Size field can be mixed and matched.

can mean: For integer or logical operators, "s"=signed (sign-extended);

5

16

"_"=default=0=unsigned (zero-extended), depending on target operator.

can also mean: 1=text (decimal character sequence);
0=binary format (mainly used with direct character sequence computations/operators).

Implementer can make it mean anything he/she wants it to mean, depending on the target operator.

SIGD Signal for DEST field meaning depends on context—either "s" or "_" must immediately precede Size field can be mixed and matched.

can mean: "s"=1=signed (sign-extended) results;
"_"=0=unsigned (zero-extended), depending on target operator

can mean: 1=eXchange ends (ie, flip endian-ness) on and during read of operand A, depending on context. For example, reading/writing from/to external memory versus on-chip memory in cases where external data is little endian.

can also mean: 1=store results as text (decimal character sequence) 0=binary format

can also mean: restore selected operator with value and flags being written.

Implementer can make it mean anything he/she wants it to mean, depending on the target operator.

SIZ[2:0] **180** Size in bytes, of source/destination (shown with SIGn bit cleared, i.e., immediately preceded with "_" character, which means "default").

"_1:"=0 000=1 byte
"2:"=0 001=2 bytes (half-word)
"4:"=0 010=4 bytes (word)
"8:"=0 011=8 bytes (double-word)
"16:"=0 100=16 bytes (gob)—i.e, vector, structure and/or mix of types up to this byte count
"32:"=0 101=32 bytes (gob)
"64:"=0 110=64 bytes (gob)
"128:"=0 111=128 bytes (gob)

SIZ[2:0] Size in bytes, of source/destination (shown with SIGn bit set, i.e., immediately preceded with "s" character).

Preceding the Size **180** specifier with a "s" character, signals the assembler to set the SIG bit for the respective DEST, srcA or srcB field.

"s1:"=1 000=1 byte
"s2:"=1 001=2 bytes (half-word)
"s4:"=1 010=4 bytes (word)
"s8:"=1 011=8 bytes (double-word)
"s16:"=1 100=16 bytes (gob)—i.e, vector, structure and/or mix of types up to this byte count
"s32:"=1 101=32 bytes (gob)
"s64:"=1 110=64 bytes (gob)
"s128:"=1 111=128 bytes (gob)

INDirect Addressing IND **190** Bit
The DEST **120**, srcA **130**, srcB **140** fields each have a IND **190** bit, which, when set, employs the contents of the specified Auxiliary Register as a pointer to either data or program memory for read operations if srcA or srcB field(s), or write operation if in the DEST field. An "*" character immediately preceding one of the specified AR0 thru AR6 Auxiliary Registers ("ARn") specifier **250** or the Stack Pointer (SP) specifier **250**, signals the assembler to set IND bit **190** in the respective field in the instruction.

IND 1=indirect addressing mode for that field. IND 0=direct addressing mode for that field. If IND=0, then the 15 bits specified in **210** of the respective DEST **120**, srcA **130** and/or srcB **140** field(s) is employed as the lower 15 bits of the direct address, with the higher order bits all=0. Direct addressing mode enables "direct" access to the specified

zero-page location without first having to load an Auxiliary Register with a pointer to the location to be accessed.

While the direct addressing mode does have its advantages and uses, it also has its disadvantages. Among them, absent some kind of paging mechanism, the direct addressing mode can only reach the first 32 k locations in the memory map. Another disadvantage is that since the direct address comes from the instruction in program memory, this address cannot be automatically post-modified or offset, unlike Auxiliary Registers (ARn) and Stack Pointer (SP) employed by the indirect addressing mode.

IMOD 200 Bit

There are two indirect addressing modes that may be employed for reads and writes: Indirect with +/-auto-post-modification and Indirect with +/-offset but with no auto-post-modification.

IMOD is only used with IND=1, meaning it is only used with indirect addressing mode for a given field.

IMOD=1 means: use specified ARn contents + (plus) or - (minus) signed AMOUNT field 240 for the effective address for accessing operandA, operandB or DEST. ARn remains unmodified. With IMOD=1, the range of AMOUNT is +1023 to -1024.

IMOD=0 means: use specified ARn contents as pointer for the instant read cycle of operandA or operandB (if specified) or write cycle for DEST. Then automatically post-modify the contents of ARn by adding or subtracting UNSIGNED AMOUNT field to/from it. With IMOD=0, the range of AMOUNT is 0 to 1023 for positive amounts and 0 to 1024 for negative amounts. Note: the programmer enters an unsigned amount in the source line, but the assembler automatically converts this amount to "signed" during assembly.

Direct Addressing Mode General Rules

On the same assembly line, direct addressing mode can be mixed and matched with:

- indirect addressing mode (any combination)
- immediate addressing mode for srcB in case of dual operands
- immediate addressing mode for srcA in case of single operand
- table-read from program memory for srcA in case of dual operands
- table-read from program memory for srcA in case of single operand

Indirect Addressing Mode General Rules

On the same assembly line, indirect addressing mode can be mixed and matched with:

- indirect addressing mode (any combination)
- immediate addressing mode for srcB in case of dual operands
- immediate addressing mode for srcA in case of single operand
- table-read from program memory for srcA in case of dual operands
- table-read from program memory for srcA in case of single operand

#Immediate Addressing Mode General Rules

On the same line, immediate addressing mode can be mixed with:

- direct addressing mode as srcA and immediate as srcB
- indirect addressing mode as srcA and immediate as srcB if srcA is signaled as #immediate and appears as the sole source, then it is 32-bit immediate. Note: 32 bits is a limitation of the present assembler, in that the instruction word can accommodate up to 40 bits in this field.

#immediate addressing mode may not be used with @table-read addressing mode on the same line

@Table-Read Addressing Mode General Rules

On the same line, @table-read addressing mode can be mixed with:

direct or indirect addressing mode as srcB

@table-read addressing mode may not be used with the #immediate addressing mode on the same line

Below are examples of actual usage conforming to the following assembly line format for a single operand, i.e., srcB is absent from the assembly line. Refer to FIGS. 9, 10, 11, 12A thru 12D and FIG. 13 for mapping of SRAM, programming model registers, floating-point, integer and logical operators and their respective labels and assignments used by the assembler:

Label:	RM	SigSizeD:DEST = SigSizeA:srcA
demoA:	_	_8:work.0 = s1:work.1
	_	_2:LPCNT1 = _2:#50
	_	_2:*AR1++[2] = _2:work.3
	_	_8:work.4 = _8:@progTable1
p		_8:cnvFDCS.0 = _32:*AR2[+0]

In the first example given above, "demoA" followed by a semicolon is an optional label. The first "_" character is the mnemonic that specifies in this instance the "default" rounding mode, but that since the destination is not a floating-point operator, means "don't care." The next "_" is the signal 170 for the destination 120 and the "8" character followed by a colon specifies the size 180 in bytes of the value to be stored in direct memory location "work.0", the destination 180 address. The "=" character means "pulled from". The "s" character in front of the "1" character is the signal 170 for the operandA source, signaling that the value being pulled from direct location "work.1" is to be automatically sign-extended to 64 bits while in-route to the destination. The "1" character to the left of the semicolon specifies the size 180 in bytes of operandA being pulled from "work.1", the srcA 130 address.

In the second example shown above, "LPCNT1" is the label specifying the direct address of one of the processor's memory-mapped hardware loop counters as the destination 120 address. In this example, LPCNT is being loaded with an immediate value of 50 decimal. The "#" character 230 appearing to the right of the colon in the srcA 130 field signals immediate addressing mode. Note that in this instance, because a single operand immediate is always 32 bits and only two bytes (16 bits) are specified, the two-byte value of "50" is automatically zero-extended by the assembler to 32 bits before being inserted into the assembled instruction word.

In the third example, "*AR1++[2]" specifies indirect addressing mode for the destination 120. In this instance, it is specifying the use of the contents of Auxiliary Register AR1 as an indirect pointer to where the two-byte value pulled from "work.3" is to be pushed. The "*AR1" followed by "++" signals the assembler that this is post-modified indirect addressing mode and to insert a "0" into the destination 120 IMOD 200 bit position in the assembled instruction word. The "*" character 250 in front of the "AR" 250 characters signals the assembler that the addressing mode for the destination 120 field is indirect and to insert a "1" into the IND bit position 190 for the destination 120 field. The "[2]" characters to the right of the "++" characters specify the amount by which AR1 is to be automatically post-

incremented. If “`--`” had appeared in place of the “`++`”, this would have signaled post-decrement.

In the fourth example, the “`@`” character in front of the label “`progTable1`” signals the assembler that `progTable1` is a direct address for `srcA` field **130** accessing program memory for data and to insert a “`10`” into Data Access Mode bits `DAM[1:0] 160` in the instruction word during assembly. In this instance, 8 bytes will be pulled from program memory location `progTable1` and all 8 bytes will be pushed into direct data memory location `work.4`, when executed.

In the fifth and last example above, a 32-byte decimal character string will be pulled from the data memory location specified by the contents of `AR2` (with offset of `+0`) and pushed into the `convertFromDecimalCharacter` operator input buffer number 0. The unrounded intermediate result of the conversion will be converted to IEEE 754-2008 binary64 (double-precision) format and rounded towards positive infinity before being stored in this operator’s result buffer number 0. If the “`p`” mnemonic had been a “`_`” instead, the intermediate result would have been rounded to nearest even, which is default rounding mode. If the size **180** of the destination **120** field had been a “`4`” instead of an “`8`”, the intermediate result would have been automatically converted to binary32 (single-precision) format instead binary64 format. In this example, the “[+0]” to the right of `*AR` signals the assembler that `srcA 130` is indirect with offset addressing mode and to insert a “`1`” in `srcA` field IMOD bit position **200** of `srcA` field **130** during assembly.

Below are a few examples of actual usage conforming to the following assembly line format for both single and dual operands:

Label:	RM	SigSizeD:DEST = (SigSizeA:srcA, SigSizeB:srcB)
demoB:	n	<code>_8:fmul.0 = (_8:vect+32, _8:scale.1)</code>
demoC:	-	<code>_4:AR0 = _4:#fadd.0</code>
	-	<code>_4:AR3 = _4:#vectStart</code>
	-	<code>_4:AR2 = _4:#vectStart + 8</code>
	-	<code>_2:REPEAT = _2:#15</code>
	-	<code>_8:*AR0++[1] = (_8:*AR3++[16], _8:*AR2++[16])</code>
demoD:	-	<code>_4:AR3 = _4:#vectStart</code>
	-	<code>_4:AR2 = _4:#vectStart + 32</code>
	-	<code>_2:LPCNT0 = _2:#16</code>
loopD:	-	<code>_4:AR0 = _4:#ufma.0</code>
	-	<code>_2:REPEAT = _2:#31</code>
p	-	<code>_8:*AR0++[1] = (s32:*AR3++[64], s32:*AR2++[64])</code>
	-	<code>_4:PCS = (_2:LPCNT0, 16, loopD)</code>

In the line above with the label “`demoB:`” is a double-precision floating-point multiply operation that has its result rounded towards negative infinity. The syntax for operations involving two or more operands expects the operand list to be enclosed in parenthesis. In this instance, the destination, `srcA` and `srcB` are direct addressing mode. `Fmul.0` is the first input buffer to the IEEE 754-2008 floating-point “multiplication” hardware operator. The final result is stored in the operator’s first result buffer location, which is the same location as the input.

The sequence of instructions beginning with the label “`demoC`” illustrates how to use the processor’s memory-mapped hardware REPEAT counter. In this example, the destination, `srcA` and `srcB` employ the indirect addressing mode. A vector comprises 16 entries of dual operands, `operandA` and `operandB` that are binary64 (double-precision) floating-point representations. The destination is the first input buffer to the processor’s IEEE 754-2008 floating-point “addition” operator, `fadd.0`. “`REPEAT n`” means: fetch and execute the following instruction once, then “repeat” it

n additional times after that. Thus, in this instance, the floating-point addition operation is executed a total of 16 times. Because all addressing is indirect, `AR0` must be loaded with the physical address of the first input buffer to the floating-point addition operator. `AR3` must be loaded with the physical address of the first element of the vector and `AR2` must be loaded with the physical address of the second element of the vector, which are `operandA` and `operandB`, respectively. The instruction that immediately follows the REPEAT instruction will execute exactly 16 times. When complete, this floating-point addition’s sixteen result buffers `fadd.0` thru `fadd.15` will contain results for all 16 operations, with all being immediately available for use in subsequent operations by the time they are pulled out using the REPEAT instruction, meaning an apparent zero latency.

The sequence of instructions beginning with the label “`demoD`” above is very similar to “`demoC`”, except it involves computing directly with raw, human-readable decimal character sequence pairs stored in qty. 32, 16-entry vectors of dual decimal character sequences, `operandA` and `operandB`. The target operator is the present invention’s hardware Universal FMA operator employed as a double-precision sum of products operator. It is assumed all of the Universal FMA Creg/accumulators have already been initialized to 0 by the time the sequence starting at “`demoD`” begins. Auxiliary Registers `AR0`, `AR3` and `AR2` are initialized with their respective pointers into the vectors. Because each of the 32 vectors of dual decimal character sequence operands are 16 deep, the processor’s memory-mapped hardware loop counter `LPCNT0` is initialized with the number of iterations, which is 16. The REPEAT counter is initialized with the number of vectors -1, which is 32-1, which is 31.

Because both `operandA` and `operandB` are human-readable decimal character sequences in this instance, such is signaled in the source line by use of the “`s`” character in each operand’s “signal” bit **170**, SIGA and SIGB. If the “signal” character for `srcB` had been a “`_`” instead of an “`s`” character and its size=8 instead of 32, then `operandB` would be treated as a IEEE 754 binary64 format representation, causing a bypass of the `convertFromDecimalCharacter` portion of the pipeline and instead routed thru a delay of equivalent length for coherency.

“`PCS`” in the last line is a conditional load shadow location of the processor’s memory-mapped hardware Program Counter operator. It will accept a push of a new value only if the bit specified in the source line is “set”, i.e., “=1” when pushed. In this instance, the source line is specifying that the PC only be loaded with a relative displacement if bit **16** of the processor’s memory-mapped hardware loop counter `LPCNT0` is “set”, i.e., “=1”. If not already zero, the hardware loop counter will automatically decrement by 1 if tested or pulled, but only if the destination of that instruction is the PCS address. In this instance, Bit **16** is the hardware loop counter’s “Not Zero” signal. If the pre-decremented value of the contents of the hardware loop counter is not zero, then the “Not Zero” signal will be “1”. “`PCS`” is only used to test for “1” to take a branch. To test for “0” to take a branch, use “`PCC`”, which resides at a different shadow location.

Once all the vectors have been pushed into the Universal FMA roughly 592 clocks later, the processor can immediately begin pulling out all 32 Universal FMA result buffers for use in subsequent operations. When results are pulled out in the same order they were pushed in, there is an apparent zero latency, which includes the IEEE 754-2008 H=20

conversion from decimal character sequence AND the actual computations. In this instance, results are rounded only once, towards positive infinity as specified by the “p” mnemonic.

Special Cases for srcB

Referring to FIG. 2 and FIG. 3, at least two special cases for srcB show how srcB bit fields can be further broken down for easier assembly by the assembler and easier readability by the programmer.

SHIFT Operator Example

FIG. 2 is a diagram of a special case for the srcB field of the instruction word specifying the number of bits to shift and the shift type for use with the 64-bit SHIFT operator in an exemplary embodiment of the present disclosure. When the assembler detects a literal shift type such as “LEFT”, “LSL”, “ROL”, etc., 270 in the srcB field of the source assembly line, it parses the argument and forms a srcB bit pattern for the srcB 140 portion of the instruction word. “Bits” 260 specifies the number of bit positions to shift by, and “Shift Type” 270 specifies the type of shift that is to take place. In addition to filling in the “Bits” 260 and “Shift Type” 270 subfields of srcB 140, the assembler also inserts a “01” into the DAM[1:0] 160 bits to indicate that srcB 140 is to be treated as #immediate. Note that the result is a hardwired shift operation because it employs an immediate value for srcB. For computed shift operations, that is, ones that have a computed shift amount and direction, the user can synthesize a srcB in data RAM with software and use it like any other operator.

Example SHIFT Operation Coding:

demoE:	—	_8:SHIFT.0 = (_8:work.5, LEFT, 12)
	—	_8:SHIFT.1 = (_8:AND.2, ASR, 3)
	—	_8:SHIFT.2 = (_8:XOR.7, COPY, 1)

In the first example starting with the line labeled “demoE”, the 64-bit contents of location “work.5” are pushed into the first input buffer, SHIFT.0, of the processor’s memory-mapped hardware SHIFT operator with instructions to shift that value LEFT, 12 places, with the 12 LSBs being replaced with 0 s. The carry (“C”) and overflow (“O”) signals are not affected. The zero (“Z”) signal is set if the result of the shift is zero and cleared if not zero. The negative (“N”) signal is set if the result is negative and cleared if positive. Results, along with the C, V, Z, and N signals, are stored in result buffer SHIFT.0.

In the second example, the 64-bit contents of result buffer “AND.2” are pushed into the second input buffer, SHIFT.1, of the processor’s memory-mapped hardware SHIFT operator with instructions to perform an arithmetic-shift-right 3 places, with the 3 MSBs being replaced with a copy of the original MSB. The carry and overflow signals are not affected. The zero signal is set if the result of the shift is zero and cleared if not zero. The negative signal is set if the result is negative and cleared if positive. Results, along with the C, V, Z, and N signals, are stored in result buffer, SHIFT.1.

In the third example, the 64-bit contents of result buffer “XOR.7” are pushed into the third input buffer, SHIFT.2, of the processor’s memory-mapped hardware SHIFT operator with instructions to do nothing but copy it into result buffer SHIFT.2. The carry and overflow signals are not affected. The zero signal is set if the value copied is zero and cleared if not zero. The negative signal is set if it is negative and cleared if positive. Results, along with the C, V, Z and N signals, are stored in result buffer, SHIFT.2.

Program Counter (PC) Example

FIG. 3 is a diagram of a second special case for the srcB field of the instruction word specifying the bit position to test and the displacement amount for use with the conditional relative branch operator in an exemplary embodiment of the present disclosure. The second special case for srcB formatting involves the processor’s memory-mapped hardware Program Counter (PC). Referring to FIG. 11, the PC has a base address labeled “PC” as well as three shadow addresses labeled “PCC”, “PCS” and “PCR”.

Pushing a value into “PC” direct address will load the PC unconditionally with that absolute, unsigned value and the processor’s PC will begin incrementing from there, effectuating an unconditional JUMP.

Pushing a value into “PCR” direct address will load the PC unconditionally with the result obtained from adding the signed value being pushed to the current PC value, effectuating a unconditional long relative BRANCH.

Pushing a value into “PCC” direct address will load the PC if and only if the condition specified on the assembly line is true. In the case of the PCC address, the condition is true, if and only if, the specified bit position 280 of the memory-mapped location specified in srcA 130 is clear, i.e., “0”. If the condition is true, then and only then will the PC be loaded with the result obtained from adding to the current PC value the signed value 290 being pushed, effectuating a conditional bit-test-and-branch-if-CLEAR operation. Otherwise, if not true, the push will be ignored and the PC will continue incrementing.

Pushing a value into “PCS” direct address will load the PC if and only if the condition specified on the assembly line is true. In the case of the PCS address, the condition is true, if and only if the specified bit position 280 of the memory-mapped location specified in srcA 130 is set, i.e., “1”. If the condition is true, then and only then will the PC be loaded with the result obtained from adding to the current PC value to the signed value 290 being pushed, effectuating a conditional bit-test-and-branch-if-SET operation. Otherwise, if not true, the push will be ignored and the PC will continue incrementing.

Example Program Counter (PC) Operation Coding

restart:	—	_4:PC = _4:#START
ifEqual:	—	_2:COMPARE = (_2:ADD.6, _2:#0x3456)
	—	_4:PCS = (_8:STATUS, Z, waitForXCUBreak0)
	—	_4:PCR = _4:#continue
waitForXCUBreak0:	—	_4:PCC = (_8:XCU_S_R, 32, waitForXCUBreak0)
	—	
	—	

In the first example above beginning with the label “restart”, the PC is unconditionally loaded with the immediate absolute program address “START”, which, hypothetically might be the beginning of a program. When executed, the pre-PC begins fetching instructions from that address.

The second example is a hypothetical test for equality between two 16-bit integers using the processor’s memory-mapped hardware integer compare operator. In this instance, a 16-bit value residing in the processor’s memory-mapped integer ADD operator result buffer 6 is compared with the 16-bit immediate integer 0x3456. On the line after the comparison, a test of the processor’s 64-bit STATUS register is performed to determine whether the zero (“Z”) flag is set, i.e., =“1”. If set, the signed offset, “waitForXCUBreak0”, is added to the PC value at the location of that instruction,

effectuating a relative branch. If not set, the two 16-bit values are not equal, in which case, the push into the PC is ignored and the PC just keeps on incrementing.

The next line is a simple unconditional relative long branch using the processor's memory-mapped PCR operator. In this instance, the PC value at that instruction's program address is added with the immediate signed relative offset, "continue", effectuating an unconditional, relative long branch.

The last example is of the processor's PCC operator. In this instance, had the previous COMPARE operation resulted in the "Z" flag being set, i.e., equality, then program execution would have entered here. Assuming that it did enter here and assuming, hypothetically that the parent CPU had previously set a breakpoint in its child XCU.0, the parent program will just sit here and wait until bit 32 of the parent processor's XCU Status/Control register is set, indicating the breakpoint had been reached by XCU.0. If the test tests true, i.e., the breakpoint has been reached, the parent processor PC will fall out the bottom and continue executing the remainder of the program or thread.

Note: a source line containing only a "_" mnemonic and nothing else that follows, with or without label, is a "No Operation" (NOP). In such case, the instruction generated is a 64-bit 0x0000000000000000. In this case, when executed, operandA and operandB are simultaneously pulled from data RAM location 0x00000000 and immediately pushed right back into the same location. In this architecture, the first eight byte locations in data RAM are always pulled as 0, no matter what values had previously been pushed there. No flags or signals are affected from pulling or pushing any data RAM location. Consequently, the net effect of a "_" all by itself is no operation. The PC does, however, continue to increment.

Default Decimal Character Sequence Format

FIG. 4 is a diagram illustrating an exemplary embodiment of the present disclosure's default floating-point decimal character sequence input and output format 300 used by the processor's memory-mapped hardware IEEE 754-2008 H=20 convertFromDecimalCharacter, convertToDecimalCharacter, and Universal FMA operators. The stand-alone version of the disclosed universal floating-point ISA's convertFromDecimalCharacter operator, if present, only accepts this default format as input.

The instant invention's convertToDecimalCharacter operator produces this default format as output. When pulled from one of its result buffers, the representation's most significant bytes that fall outside the 47 characters shown will automatically be filled with null characters to 64 bytes during the pull.

The Universal FMA operator directly accepts as input for operandA and operandB not only the default format 300, but also the formats 400 shown in FIG. 5, in addition to IEEE 754-2008 binary64, binary32, and binary16 formats, in any combination.

For input as a default format into either the stand-alone convertFromDecimalCharacter operator, if present, or the Universal FMA, the 47th character position of decimal character sequence 370 must be occupied by either a "+" character or a "-" character. Character position numbers 380 are shown for reference. The integer part 350 of the floating-point representation comprises 21 decimal characters. The 46th character position can only be occupied by either a 1 or a 0. The fraction part 340 comprises 20 decimal digits representing the fraction part of the floating-point representation. Character position 5 must be occupied by an "e" character 330. Character position 4 must be occupied by

either a "-" or "+" character 320. Character positions 1 thru 3 must be occupied by a decimal character sequence representing the decimal exponent 310, relative to the decimal digit occupying character position 6. In the example decimal character sequence 370 shown, the actual decimal character floating point representation is: "9845913.87323175849002839455", which is 27 decimal digits long, 7 decimal digits greater in length than the minimum "H=20" mandated by IEEE 754-2008 for binary64 results. Since the decimal place is referenced relative to the last decimal digit at character position 6 and the exponent is "-020", the actual decimal place is determined by counting 20 digits to the left.

For inputting infinities and NaNs, both signaling and quiet, in the default decimal sequence format, refer to 300, 460, 470, 480, 490, and 510 of FIG. 5.

Universal FMA Input Formats

FIG. 5 is a diagram illustrating examples of various decimal character sequences, including some with token exponents, their translation to the default format, and their respective IEEE 754 binary64 equivalent representations in an exemplary embodiment of the present disclosure. The present disclosure's Universal Fused-Multiply-Add (accumulate) operator can accept as input for operandA and operandB IEEE 754-2008 "H=20" decimal-character sequences up to 28 decimal digits in length using the default decimal-character sequence format 300, as well as IEEE 754-2008 binary64, binary32, and binary16 formats. In addition to the default decimal-character sequence format 300, the Universal FMA can accept the character sequence formats 400 shown in FIG. 5. Among them are human-readable decimal character sequences comprising "token" exponents. Token exponents are often found in online financial and statistical reports. Token exponents presently supported by the present invention's Universal FMA include the capital letters "M", "T", "B", "K" and the character "%", representing exponents "e+006", "e+012", "e+009", "e+003" and "e-002" respectively. The Universal FMA can also accept decimal character sequences in scientific notation, integers, fractions, and other non-exponentiated character representations up to twenty-eight decimal digits in length as shown in character sequence formats 400. All of the above representations can be pushed directly into the Universal FMA in any combination.

Reference 410 is a representation for 29 million; 420 is a representation for 73.75 trillion; 430 is a representation for 150 billion; 440 is a representation for 35 thousand; 450 is 50 a representation for 0.5593; 460 is a representation for negative infinity; 470 is a representation for positive infinity; 480 is a representation for a positive, quiet NaN with payload; 490 is a positive signaling NaN with payload; and 510 is an example payload in hexadecimal characters. Payloads for all NaNs must be formatted as shown, including case, prior to submission into the operator.

Other token exponents that could just as easily be implemented on the front-end or even on the back-end of the Universal FMA conversion (if a convertToDecimalCharacter circuit is incorporated into its pipeline) include the SI units (International System of Units). For example, yotta "Y" for "e+024", zetta "Z" for "e+021", exa "E" for "e+018", peta "P" for "e+015", tera "T" for "e+012" (same as "Trillion" above), giga "G" for "e+009", mega "M" for "e+006", kilo "k" for "e+003", hecto "h" for "e+002", deka "da" for "e+001", deci "d" for "e-001", centi "c" for "e-002", milli "m" for "e-003", micro "u" for "e-006", nano

“n” for “e-009”, pico “p” for “e-012”, femto “f” for “e-015”, atto “a” for “e-018”, zepto “z” for “e-021” and yacto “y” for “e-024”.

As can be seen, token exponents are a simple and useful way for reducing the number of characters needed to represent a human readable value used as either input or output in computation. Precisely which tokens to employ/support would most likely depend on target application and implementer objectives. This is important because most prior art machines can typically move data as either 1, 2, 4, 8 bytes per clock—best case. Using tokens allows for easier movement of more precise numbers in fewer clocks because more significant digits can be used that would otherwise be occupied by the “e+xxx” portion of the character sequence representation for a given 8, 16 or 32-byte character sequence representation.

The front-end to the Universal FMA detects if the operand(s) being pushed into it are default character sequence format, or an IEEE 754-2008 binary format. If not already in the default character sequence format, it will automatically translate such character sequence(s) 400 into the default format 300 prior to entering the Universal FMA pipeline. By way of example, Binary64 Result 390 shows the IEEE 754-2008 binary64 representation after conversion. NaN input payload 510 and NaN result payload 520 show that, in the instant implementation, NaN payloads are preserved and propagated during conversion.

CPU (Parent) Block Diagram

FIG. 6 is a block diagram of an exemplary embodiment of the basic architecture 600 of the disclosed universal floating-point ISA, showing the CPU or parent processor. The basic architecture 600 is fully pipelined, meaning that instruction fetch, decode, and execution cycles overlap as do its program and data memory accesses. A pre-PC 610 (as opposed to an ordinary PC) is necessary for generating the physical instruction fetch address because program memory in this implementation comprises synchronous RAM (SRAM) blocks which register the read address on the rising edge of the processor’s clock. Thus, the current instruction address must be presented early so it can be registered on the rising edge of the clock. The PC, unlike the pre-PC, registers the output of the pre-PC on the rising edge of the clock and thus its contents will generally match that of the program SRAM A-side internal address register.

In the instant implementation, the program resides in either synchronous ROM (SROM) or RAM (SRAM) that is three-port 620, having one write side and two read sides. The second read-side port is to accommodate @table-read operations, useful for reading constants, pointers, etc., in tabular form, from program memory. If an application never uses table-read operations, then program memory can be implemented in a purely two-port configuration, thereby reducing the amount of SRAM needed for program storage by 50%. Alternatively, if the target application requires only a relatively few locations for storage of constants in program memory, then program memory configuration can be a hybrid of three and two-port memory, with the first few thousand words implemented in three-port SRAMs and the rest implemented in two-port memory.

Ordinarily, child XCUs should be implemented in SRAM so that the parent CPU can push copies of threads into child XCU program memory for use in performing tasks. The parent CPU program memory will almost always comprise a small micro-kernel in SROM and the rest in SRAM.

In addition to @table-read (read-only) operations directly from the first 32 k locations in program memory, in the present implementation, both parent CPU and child XCU

entire program memory can be accessed “indirectly” for both read and write operations by simply setting the most significant bit (i.e., bit 31) of the Auxiliary Register being used as the pointer for such access. Setting this MSB maps program memory into data memory space, such that 0x80000000 of data memory space maps to 0x00000000 of program memory space, for both read and write operations from/to program memory. Caution should be exercised when performing write operations to program memory during run-time, in that, presently, there are no safeguards, such as lock, supervisor modes, etc., have been implemented to protect a thread from corrupting its own program memory, due to a spurious write or whatever.

Once the contents of the pre-PC has been registered into program memory 620 on the rising edge of the clock, the instruction 100 appears on the A-side output and is presented to the Auxiliary Register and Stack Pointer register bank 630. @Table-reads, if any, appear on the B-side output. All bits in the instruction output of program memory 100 are delayed one or more clocks by Pipeline Control block 700 for use by various operators. Some operators 1000, registers 1100, SRAM 900 and 1200, Cascaded Instructions 1300 and IEEE 1149.1-1990 JTAG Debug Port 1400 may require some or all bits of instruction 100. Cascaded instructions are monitor-read or monitor-write instructions issued by a parent CPU to a child XCU. Thus, the two terms are synonymous when such transactions involve a parent CPU and a child XCU. Note that it is also possible for a child XCU to also become a parent using the same approach and logic, but this may not be the best strategy, in that a better strategy might be employing two or more parent CPUs, each with their own child XCUs attached to them.

During Stage q0, instruction 100 is immediately presented to the Auxiliary Register and Stack Pointer bank 630. If any or all the IND 190 bits are set, then the Auxiliary Register and Stack Pointer register bank 630 logic substitutes the contents of the specified Auxiliary Register(s) for the DEST 120, srcA 130 and/or srcB 140 direct addresses. Otherwise, if IND 190 is not set for a given address field in 100, then the 15 bits for the respective “direct” address field in 120, 130 and/or 140 in instruction 100 is/are, after having been automatically zero-extended to 32 bits, used as the address for the respective push-pull operation of the instruction.

All processor programming model registers, operators 45 and their result buffers, data memory, control and status register and JTAG real-time debug functions are mapped into the processor’s data memory space 640. The beauty and elegance of having all registers, etc. mapped into data memory space, provided they are readable, is that they are all visible to the JTAG 1400 hardware without the need of special opcodes. Program memory, registers, operator input buffers, operator result buffers, etc., are accessed just as easily as accessing data memory, all without the use of any specialized opcodes.

Ideally, data RAM, operator result buffers, processor registers, etc., are implemented with two read-side ports, one for accessing operandA and the other for accessing operandB—simultaneously. Mux 650 selects operandA from the A-side outputs and is registered into register 660 during the rising edge of the Stage q1 clock and automatically sign-extended to 64 bits by SextA 670 if srcA SigA 170 (delayed) is set, otherwise zero-extended.

Likewise, mux 740 selects operandB from the B-side outputs and is registered into register 730 during the rising edge of the Stage q1 clock and automatically sign-extended to 64 bits by SextB 720 if srcB SigB 170 (delayed) is set, otherwise zero-extended.

The CPU can operate stand-alone or as a parent to one or more (up to sixteen) child XCUs connected to it by way of busses **680**, **690**, **710**, **770**, **750** and **760**. Bus **680** comprises an encoded four-bit write select, encoded four-bit read select bus, and a “push-all” signal.

For a cascaded push of data from the parent to the memory-mapped child XCU, the four-bit write select bus portion of bus **680** specifies into which XCU data will be pushed. If data is pushed through the magic “push-all” memory-mapped window of the parent, the data will be pushed into all connected child XCUs simultaneously, in which case the “push-all” signal of bus **680** will be active. The “push-all” feature is handy for copying entire threads from the parent CPU program memory and pushing them into all the child XCUs program RAM in cases where the workload is divided among all XCUs executing the same thread. “Push-all” can also be used for pushing XCU data memory parameters that will be used for a given task. It can be used to preset all XCUs PCs to a known start location or even set a software breakpoint in all XCUs simultaneously.

The four-bit read bus portion of bus **680** is used for selecting which of the connected child XCUs is to be pulled. The parent CPU can pull data from anywhere with the target XCU’s memory map. Since the parent CPU essentially treats any target XCU’s memory-map as its own, it can employ its REPEAT operator to push or pull entire blocks of memory to/from the selected XCU, or in the case of program memory push, entire thread using the REPEAT instruction followed by the instruction that does the push.

There is a cascaded write data bus **690** to the XCUs. Presently bus **690** is only 64 bits wide, but can be enlarged to 1024 or more bits if application requires it.

A cascaded instruction bus **710** runs from the parent CPU to the child XCUs and is essentially “fused” to any attached child XCUs’ instruction bus. Anytime the parent CPU wants to perform a real-time data exchange operation between a given XCU, the CPU forces the XCU to execute the cascade instruction instead of its own, all without the use of interrupts, opcodes, or DMA. This way, the parent CPU has complete visibility and writability into the target XCU’s memory map.

The parent CPU’s memory-mapped XCU control register generates output **770**. With this register, the CPU can individually or altogether force a breakpoint, reset, single-step, or preempt the selected XCU(s).

There is a response data bus **750** from attached XCUs. Bus **750** may be the output of a mux that outputs one of up to sixteen XCU read buses, i.e., the read bus of the selected XCU. Alternatively, since the CPU is capable of reading 1024 bits onto its operandA data bus, the bus **750** may be configured to pull up to sixteen XCUs simultaneously with each XCU providing 64 bits, for a total 1024 bits.

Bus **760** includes DONE, Software_Break_Detect, XCU_BROKE, and SKIP_CMPLT status signals from each of up to sixteen XCUs connected to the parent CPU. DONE[15:0] are the XCU done bits (one for each XCU), which a given XCU uses to signal that it is done performing a task. Software_Break_Detect[15:0] indicate which, if any, XCUs have encountered a software breakpoint. XCU_BROKE[15:0] indicate which, if any, XCUs have hit a hardware breakpoint. SKIP_CMPLT[15:0] indicate which, if any, XCUs have completed a previously issued single-step command.

FIG. 7 is a block diagram of an exemplary embodiment of the basic architecture of a child XCU **800** and shows that the child XCU is virtually identical to the parent CPU except that cascaded instruction and related buses **750**, **760**, **680**, **690**, **710**, and **770** are turned around. The illustrated XCU

has no IEEE 1149.1 (JTAG) debug port, although one could easily be added and included in the same scan chain as the parent if desired. By way of a tiny micro-kernel running on the parent CPU, the CPU can operate as an intermediary between the external debug environment/hardware and any of the child XCUs attached to it, since the CPU has complete real-time visibility into any of them. This saves on the logic needed to implement in the XCUs logic a JTAG-accessible debug port like the CPU has.

It should be further understood that, like the CPU, the child XCUs can be implemented with a minimum suite of operators tailored for a specific application. For example one XCU can be configured with memory-mapped operators that are tailored for DSP applications. Another, or several, XCUs can be configured as Tensor Processing Units (TPUs) or GP-GPUs. Others can be configured as a general-purpose microcontrollers with few or no hardware floating-point operators. The major point here is, because the CPU and XCUs do not employ opcodes in the instruction set, such can be easily customized to perform virtually any task for virtually any application, with all using the same software development and debug tools. Stated another way, no new C compiler, assembler, debugger, development environment tools need to be developed just because this ISA has been targeted as a DSP, TPU, microcontroller, or floating-point processor, because, fundamentally, the core is just a push-pull shell that specific memory-mapped hardware operators, each with their own independent pipeline(s) are plugged into, the instruction set of which is optimized to do this push-pull function very efficiently and effectively.

FIG. 8 is a block diagram illustrating an arrangement of a CPU/parent processor **600** and one to sixteen XCU/child processors **800** in an exemplary embodiment of the present disclosure. FIG. 8 shows how the parent CPU **600** and child XCUs **800** are arranged when configured as a multi-core processor, including communication/control/status/response buses.

FIG. 9 is a modified Harvard model data memory-map **640** employed by both the CPU and XCU(s) in an exemplary embodiment of the present disclosure. The first 16 k bytes of data memory space is occupied by a three-port SRAM **1200** and is both directly and indirectly accessible on 1, 2, 4, and 8-byte boundaries. The contents of physical addresses 0x00000000 thru 0x00000007 in data memory space are always pulled as 0x00000000 no matter what has been previously written to them. Among the reasons for this is the case where the current instruction only has operandA, the B-side of data RAM will always read as 0x00000000 because when there is no operandB specified in the instruction, its bit field is filled with zeros, which maps as a direct access of byte location 0x00000000, which is zero-extended to 64 bits. Another use of this feature is it is a handy way of forcing a 0x00000000 on the A-side of the bus using a SrcA address of “0”. This, in combination with the logical OR operator, can be used to transfer the contents of the B-side of the bus to the A-side by pushing location “0” and the B-side into the OR operator and reading the result out on the A-side.

Implemented memory-mapped hardware floating-point operators **3000** reside in directly-addressable data memory starting at data memory location 0x00007CFF and growing downward to location 0x00004000. For a complete map of IEEE 754-2008 mandated floating-point operators and their respective assembler labels, refer to FIG. 12A thru FIG. 12E. These floating-point operators will produce exception “signals” that are automatically stored in the specified result buffer simultaneously with the result of the computation as

the five most significant bits. Thus, for example, a binary64 result will be stored in the specified result buffer location as 69 bits because the five most significant bits are exception signals that may or not be set. These signals will not affect the processor's corresponding STATUS Register bits unless and until such floating-point result is actually pulled from the specified result buffer. The floating-point operator status bits within the STATUS Register abide in the IEEE 754-2008 mandated rules for signaling and flagging exceptions, depending on whether the processor at the time of the pull is operating in accordance with default exception handling or delayed exception handling for floating-point operations. This aspect is discussed in the STATUS Register section below.

Note that more floating-point operators can be included by growing the map downward from 0x00004000. Also note that the present invention employs IEEE 754-2008 double-precision hardware operators and each input buffer and corresponding result buffer output are mapped on single-location boundaries. If desired, double-precision floating-point operators can be substituted with single-precision operators. In such case, the mapping is still the same because the inputs are still mapped on single-location boundaries. For example, the floating-point addition operator presently has sixteen result buffers occupying sixteen consecutive locations starting from data memory address 0x00007800 and continuing upward to 0x0000780F. Such arrangement and strategy permits use of the REPEAT operator in combination with the indirect addressing auto-post-modify mode for performing up to sixteen floating-point addition operations in rapid succession, as well as for immediately turning around pulling the results out in rapid succession.

It should be understood that it is possible to shadow-map a double-precision, single-precision and half-precision floating-point operator at the same location without conflict. The method for doing that is to use the Size **180** field of the instruction to select to which operator operandA and operandB is to be pushed or from which operator the result will be pulled. The advantage of shadowing hardware operators like this is that the lower-precision versions of the same operator typically not only clock faster but also use substantially less logic. Thus, for some implementations, it may be advantageous in certain scenarios to employ the lower precision version of a given hardware operator shadowed behind the larger and higher precision operator at the same base location.

Integer arithmetic and logical operators **4000** are located as shown in the processor's data memory map. For a list of these operators, their mapped locations and corresponding assembler labels, refer to the table in FIG. 13. Integer arithmetic and logical operators **4000** residing in the processor's data memory map are the only ones that might affect the Z, C, N, and V flags within the processor's STATUS Register, but only when pulled from such operator's result buffer(s). The shaded area just below the integer arithmetic and logical operators **4000** shows the directly addressable memory locations that can be used for adding more integer and/or logical operators.

An exemplary range of Memory-mapped hardware operator locations for logical, integer arithmetic and floating-point operators are in the data memory range **1000**. These operators can automatically affect certain bits within the processor's STATUS Register when a result is pulled from such operator's result buffer. Except for explicit operations on the STATUS Register itself, or any of its "shadow" operator locations, presently, pulling from any other location within the processor's memory map will have no effect on the bits

within the STATUS Register. For instance, simply pulling from a data RAM location and pushing it to another or same data RAM location will have no effect on the STATUS Register.

5 The instant architecture employs magic memory-mapped "hot-spots" within block **1300** to effectuate real-time data exchange operations between the parent CPU and any child XCUs attached to it. Operation of these hot-spots is discussed in the XCU management section below.

10 Optional floating-point exception capture registers/buffers **990** (if present) may also be mapped in data memory as shown. If present, these registers can be configured to capture/intercept certain diagnostic information about a given computation when the result of that computation is 15 pulled from its result buffer and is used when delayed exception handling is enabled by setting certain bits in the processor's STATUS Register. Use of this feature is discussed in the exception capture section below.

The processor's programming model registers such as PC, 20 PC_COPY, STATUS Register, Auxiliary Registers AR0 thru AR6, Stack Pointer, etc. are mapped in data memory at address range **2000**. These registers are directly-addressable. Note that location 0x00007FFF is the highest (last) directly-addressable location in the processor's memory map. Refer 25 to FIG. 11 for a complete list of all the processor's programming model registers, their corresponding locations, assembler labels and state after reset **2000**. Note that most processor programming model registers in **2000** can only be pushed using direct addressing mode, even though the 30 general rule with this architecture is that anything that is directly addressable is also indirectly addressable, i.e., resides BELOW 0x00008000 in the processor's memory map. Objects residing ABOVE 0x00007FFF can only be accessed indirectly using Auxiliary Registers(s) or Stack Pointer.

35 Address range **9500** shows a dual asymmetric SRAM arrangement and their mapping in indirectly addressable memory space. This memory is special because it can be used as general-purpose data RAM for storage of up to 40 1024-bit wide data. It can be pushed and pulled in sizes of 1, 2, 4, 8, 16, 32, 64 or 128 bytes. It has one write-side, a read-side-A and a read-side-B so that up to two 128-byte gobs can be pulled from it simultaneously.

Address range **900** (see FIGS. 6 and 7) may include a 32 45 k-byte by 1024-bit, three-port SRAM **9520** as explained above, as well as a two-port 32 k by 5-bit SRAM **9510** mapped to the same location range as SRAM **9520**. This 5-bit SRAM **9510**, together with fat SRAM **9520** working in concert, can be employed as a special kind of stack for saving and restoring the contents of any operator's result buffers. For instance, there is logic that detects, in a given instruction, when the Stack Pointer (SP) is employed as the destination pointer pointing to **9500** and the source address, whether direct or indirect, is a memory-mapped hardware 50 operator result buffer location. When this condition exists, the logic determines automatically that this is a result buffer push onto the stack. When this happens, the result pulled from the specified result buffer is pushed onto the stack and its corresponding exception/condition signals are pushed 55 onto the 5-bit appendage SRAM at its corresponding location. The SP contents is then automatically decremented by the amount specified in the originating instruction.

Conversely, if this special logic detects in a given instruction that the SP is being used as a source pointer pointing to **9500** and the destination is a memory-mapped hardware logic, integer or floating-point input buffer location, such 60 logic will determine that the operation is a stack pull,

effectuating an operator result buffer restore operation that restores not only the original result but also the five exception/condition signals simultaneously.

Note that the 5-bit SRAM **9510** only accommodates five bits at a time, while the fat SRAM **9520** accommodates sizes of 1 byte to 128 bytes. SRAMs **9510** and **9520** share the same address buses. This means that, for example, if there are two consecutive pushes of 128 bytes each and the SP is decremented by 128 bytes each time, a total of 256 consecutive byte locations are written into **9520** while only a total of 10 bits are written into **9510**, five bits at the initial location specified by the SP and five more bits specified by the SP now decremented by 128. Meaning, the contents of **9510** of the 127 locations between the first push and the second push remain unchanged.

It should also be noted that additional logic has been added to **9500** to make the contents of **9510** visible and changeable by the user and debuggers. This is discussed in the asymmetric stack section below.

Block **950** is an optional data SRAM block in indirectly-addressable data memory space. It can be implemented as two-port or three-port SRAM, as conventional 1, 2, 4, and 8-byte sizes or as fat 1, 2, 4, 8, 16, 32, 64, and 128-byte sizes. Alternatively, depending on the application, it can be omitted altogether.

Block **940** is optional. It may be implemented or omitted. It may be two-port or three-port, deeper, shallower, or wider, depending on the target application and desires/objectives of the implementer.

Block **930** may be mapped as an indirectly-addressable external memory access window, providing access to a gigabyte of external indirectly addressable memory. If required, block **930** can be made to grow downward towards block **940**, thereby providing a two-gigabyte external access window. Alternatively, without too much problem, block **930** external access window logic may be modified to grow the access window towards block **920**, thereby providing a four-gigabyte external access window if need be.

With the disclosed universal floating-point ISA being implemented as a modified Harvard memory model, it is desirable to be able to access tables and other data, including threads, stored in physical program memory **620** as if it were actually residing in data memory space **640** using the indirect address mode. Program Memory Access Window **920** in data memory space can be used by the processor to indirectly access the entire physical program memory space as if it were actually data memory. When data memory is accessed in the indirect range between 0x80000000 and 0xFFFFFFFF inclusive, memory block selection logic within the disclosed ISA will access program memory space instead of data memory space, such that, for example, pulling from indirect data address 0x80000000 will yield the 64-bit contents of program memory location 0x00000000. Each 64-bit value in program SRAM occupies one location. Thus, if attempting to perform a block transfer from program memory to data memory using the REPEAT instruction, the Auxiliary Register used as the source pointer must be incremented/decremented by 1, while the Auxiliary Register used as the destination pointer into data memory must be incremented/decremented by 8. This is because data SRAM is byte-addressable while program memory can only be accessed 8-bytes (64 bits) at a time. Stated another way, instructions in program memory are 64 bits and each occupies exactly one location. When pushed into data SRAM, the same program instruction or table data occupies exactly eight byte locations.

FIG. 10 is a modified Harvard model program memory-map employed by both the CPU and XCU(s) in an exemplary embodiment of the present disclosure. Program memory space **620** may be two-ported or three-ported. All 5 instructions are 64 bits wide and are fetched in a single clock cycle. Thus program memory must be 64 bits wide as shown. The depth of program memory can, in theory, be 2^{64} locations. However, for practical reasons when the design is implemented in a Field-Programmable Gate Array (FPGA), program memory is limited to 2^{20} 64-bit words 10 **1060**. @Table-read operations from program memory are pulled on 64-bit (8-byte) boundaries only, one 8-byte word per location. Meaning, the PC increments by 1, not 8. Likewise, for @Table-read operations, if program memory 15 is accessed using indirect addressing mode and the REPEAT operator for block transfers, the instruction used for such accesses Auxiliary Register must be specified to increment (or decrement) by 1. If the destination of such block transfer 20 is data RAM, the Auxiliary Register used as a pointer to such data RAM must be specified in the instruction to increment by 8 and must start on a boundary evenly divisible by 8. This is because data RAM is accessible on byte boundaries, unlike program memory, which is only accessible on 8-byte 25 boundaries.

FIG. 11 is an exemplary memory-map **2000** of the disclosed universal floating-point ISA's memory-mapped programming model register set in an exemplary embodiment of the present disclosure. The memory-map shows the 30 memory-mapped processor programming model registers, their corresponding location in data memory space, their assigned assembler labels, reset state, and modes of access. Most of these registers may only be pushed using the direct 35 addressing mode, even though, ordinarily, anything in this architecture that is directly addressable is also indirectly addressable, but not the other way around.

FIG. 12A is an exemplary memory-map **3000** of the disclosed universal floating-point ISA's hardware-implemented IEEE 754-2008 Mandated Computational floating-point Operators in an exemplary embodiment of the present disclosure. For each of these hardware-implemented operators, a set of at least sixteen input/result buffer locations and the address of each operator's first input/result buffer location are provided in data memory space. Also shown are the 40 assigned assembler labels unique to each input/result buffer. For operators that require more than sixteen clocks to complete, the number of input/result buffers is thirty-two. Also shown are the input sizes each operator will accept and the result sizes they can produce.

FIG. 12B is an exemplary memory-map of the disclosed universal floating-point ISA's hardware-implemented IEEE 754-2008 Mandated Comparison Predicates operators **3020**, in an exemplary embodiment of the present disclosure. Unlike the operators in memory-map **3000**, the hardware-implemented comparison predicates **3020** have no result buffer. Instead, operandA and operandB are pushed into the direct address in FIG. 12B corresponding to the comparison to be made using the corresponding assembler label or explicit direct address. If the comparison is TRUE, the 50 “IEEE compare True” bit **5030** (i.e., bit-**50** of the processor's STATUS Register **5000** of FIG. 18A), is set automatically. If the comparison is FALSE, then bit-**50** is cleared as a result of the comparison. FIG. 12B also shows that the IEEE 754-2008 binary formats of operandA and operandB can be mixed and matched. This is because the logic behind the floating-point comparison hardware automatically converts binary16 and binary32 formatted numbers to binary64 55 60 65

before the comparison is made. Conversions from binary16 or binary32 to binary64 format, if any, are always exact.

FIG. 12C is an exemplary memory-map of the disclosed universal floating-point ISA's hardware-implemented IEEE 754-2008 Mandated dual-operand, non-computational, non-exceptional operators **3030** in an exemplary embodiment of the present disclosure. The non-exceptional operators work in a manner similar to the comparison predicates **3020** of FIG. 12B. The main difference is that the two non-exceptional operators **3030** produce no exceptions. When operandA and operandB are pushed into the "totalOrder" operator input located at the direct address "tOrd", i.e. 0x00007CE9, and the result of the total order test/comparison is TRUE, then the "total Order True" bit **5070** (i.e., bit-**46**) of the processor's STATUS Register (see FIG. 18A) is automatically set. If the result of total order test is FALSE, then bit-**46** is automatically cleared.

The totalOrderMag hardware operator of **3030** works in the same manner as the totalOrder operator, except the operation tests/compares the absolute values (magnitudes) of operandA and operandB when pushed into the totalOrderMag operator input, tOrdM at direct data address 0x00007CE8. If the result of the totalOrderMag test is TRUE, then the total Order Mag True bit **5060** (bit-**47**) of the processor's STATUS Register (see FIG. 18A) is set; otherwise it is cleared.

FIG. 12D is an exemplary memory-map of the disclosed universal floating-point ISA's hardware-implemented IEEE 754-2008 Mandated single-operand, non-computational, non-exceptional operators **3040** in an exemplary embodiment of the present disclosure. Each of these hardware-implemented operators has its own input buffer into which operandA is pushed. For example, to test whether an operand is canonical, push the operand into the "isCanonical" input at direct address "isCanonical" (i.e., location 0x00007CD8). If canonical, the IEEE "is True" bit **5040** (bit-**49**) in the processor's STATUS Register (see FIG. 18A) is automatically set; otherwise it is cleared.

The last non-computational operator listed in **3040**, "Class" works differently than the "is" tests listed above it, in that pushing an operand into the Class input at direct address, "clas" (location 0x00007FCE) sets one of the mutually exclusive Class bits **5080** in the processor's STATUS Register (see FIG. 18). For example, if the operand pushed into "clas" is a negative subnormal number, then the negative subNormal bit (bit-**40**) of the STATUS Register is set; otherwise it is cleared. Only one of these 5080 bits is set and all the others are cleared after each push. In addition, after each push into the Class operator, the class is encoded into a four-bit value corresponding to the class indicated in the STATUS Register when executed. This four-bit class code can be read at direct location 0x00007FD6. The Class encodings for this four-bit code have the following meanings:

Code	Meaning
0x1	signaling NaN
0x2	quiet NaN
0x3	negative Infinity
0x4	negative Normal
0x5	negative Subnormal
0x6	negative Zero
0x7	positive Zero
0x8	positive Subnormal

-continued

Code	Meaning
0x9	positive Normal
0xA	positive Infinity

Thus, after pushing an operand into the Class operator, there are two ways to evaluate its resulting class. The first method is to simply test one of the ten 5080 bits of the STATUS Register corresponding to a specific class. The other method is to compare the value that was automatically stored at the class pull location 0x00007FD6 with one of the four-bit class codes listed in the table above. In other words, the four-bit code indicates the operand's class according to the Class code table above.

FIG. 12E is an exemplary memory-map of the disclosed universal floating-point ISA's hardware-implemented IEEE 754-2008 miscellaneous mandated operators **3050** in an exemplary embodiment of the present disclosure. Each of the hardware-implemented operators **3050** have a corresponding memory-mapped direct push location. See FIG. 12E far right column for examples of how to use the miscellaneous mandated operators **3050**.

FIG. 12F is an exemplary memory-map of the disclosed universal floating-point ISA's hardware-implemented, stand-alone floating-point computational operators not mandated by IEEE 754-2008 in an exemplary embodiment of the present disclosure. The table of FIG. 12F gives a list of the implemented, but not mandated, computational operators **3060** along with their respective assembler labels and first and last direct input buffer addresses.

FIG. 13 is an exemplary table showing the disclosed universal floating-point ISA's hardware-implemented native logical and integer arithmetic operators **4000** in an exemplary embodiment of the present disclosure. The table further illustrates corresponding assembler labels, and direct addresses of the first input/result buffer location. Corresponding input and result sizes are also shown. The number of bytes shown for the result size is always zero-extended to 64-bits with the 4-bit condition code concatenated as the upper Most Significant Bits (MSBs), creating a 68-bit pull. When pulled, the 4-bit condition code resulting from the operation is automatically registered into the corresponding bits of the processor's STATUS Register. For instance, the condition signals that are pulled include integer/logical Z, C, N, and V, which are the integer/logical zero, carry, negative, and overflow signals, respectively. These condition code "signals" are automatically registered as condition code "flags" in the STATUS Register as a result of the pull.

FIG. 14A is a simplified schematic diagram of an exemplary embodiment of a circuit **630** used to implement the disclosed universal floating-point ISA's memory-mapped hardware Auxiliary Registers (ARn) for indirect addressing. The contents of AR0 thru AR6 are primarily employed as indirect pointers for the ISA's indirect addressing mode. They can also be used to hold calculated amounts to be loaded into the processor's REPEAT counter in computed REPEAT scenarios. Stated another way, the REPEAT counter can only be loaded with an #immediate value in the instruction or directly with an amount being held in a ARn. #Immediate values in the instruction are not variable, while the contents in a given ARn are variable and thus can be computed amounts.

Example A

_2: REPEAT=_2:#23
_8.*AR0++[8]=_8.*AR1++[8]

In Example A above, REPEAT is loaded with the immediate value 23. Thus the instruction that follows it will be executed exactly 24 times. But when a routine requires that the repeat amount be variable, i.e., not fixed like it is in Example A, such a requirement can be satisfied by computing an amount and storing that amount in an ARn **6090** and then push that value into the REPEAT counter.

Example B

```
_2: ADD.3=(_2: work.7, _1:#4)
_2: AR6=_2: ADD.3
_2: REPEAT=_2: AR6
_8:*AR0++[8]=_8:*AR1++[8]
```

In Example B above, the REPEAT counter is loaded with a value previously computed and stored in AR6. It should be noted that pushing a value into the REPEAT counter behaves differently than pushing a value into data RAM, in that the value being pushed into the REPEAT counter is registered during the rising edge of Stage q1 rather than Stage q2. This is so that REPEAT counter can be immediately used by the following instruction. Otherwise, due to the instruction pipeline, the REPEAT counter will not have been actually loaded by the time the instruction that follows needs it. Thus, as shown in Example B above, REPEAT is “immediately” loaded with the contents of AR6 “as if” such contents were an #immediate value. As can be seen, ARns have uses beyond just indirect indexing into memory.

The circuit **630** shows that, like the REPEAT counter, #immediate loads of a ARn occur a clock early, i.e., during the rising edge of Stage q1. This is so that their contents can be immediately used as a pointer by the following instruction. Otherwise, another instruction would have to be inserted between the ARn load and the instruction that uses such ARn contents **6040** as a pointer for SrcA address **6010**, SrcB address **6020**, or Destination address **6030**.

The circuit **630** further shows that the contents of a ARn can be automatically post-incremented or post-decremented by an amount specified in an instruction using the adder **6080** dedicated to that ARn.

For employment as pointers into indirect memory space, the circuit **630** shows that ARns **6090** provide two modes of carrying this out. The first mode is automatic post-modification mode as previously described using **6080**. The second mode is variable offset with no post-modification mode. To implement this mode, hardware adders **6050**, **6060** and **6070** are situated between ARn **6090** to provide the computed offset for SrcA address **6010**, SrcB address **6020**, or Destination address **6030**. The IMOD 200 bit in DEST **120**, srcA **130** and srcB **140** (see FIG. 1) determines which indirect addressing mode for each will be employed, but only if IND **190** bit for that is also set.

FIG. 14B is a simplified schematic diagram of an exemplary embodiment of the disclosed universal floating-point ISA’s Stack Pointer (SP) for indirect addressing, which is part of the auxiliary register logic block **630** of FIG. 14A. The circuit is very similar in operation to the ARn circuit **630** in the following ways. It has a register **6170** to hold the value of the SP **6130**. It has an adder dedicated for automatic post-modification operations by a signed amount specified in the instruction. It has adders **6140**, **6150**, and **6155** situated between **6170** and SrcA_SP **6100**, SrcB_SP **6110** and Dest_SP **6120** dedicated to computing an offset address for each of them. The SP differs slightly from ARns in that **6100** and **6110** are always a result of an offset amount, while **6120** can be either offset mode or the automatically post-modified amount (i.e., with no offset). Stated another way,

like an ARn, the SP is employed as an indirect pointer for use as either a srcA address, srcB address, or Destination address (or any combination thereof).

The disclosed ISA allows for adopting the convention that a stack “push” pushes downward to a lesser address and that a stack “pop” pops upward to a greater address. To effectuate a stack “pop”, the SP direct address must appear in a source operand field of the instruction. To effectuate a stack “push”, the SP direct address must appear in the destination field of the instruction and must be automatically post-decremented by an amount equal to the size of the word just pushed so as to point to the next available location downward in the stack.

However, to effectuate a stack “pop”, the indirect address must be “pre”-incremented by an amount equal to the size of the previous stack “push”. This is where the SP and ARns differ and FIG. 14B shows this difference.

Example C

OVFL_:	<pre>s4:*SP-[8] = _4:PC_COPY _8:capt0_save = _8:CAPTURE0 _8:capt1_save = _8:CAPTURE1 _8:capt2_save = _8:CAPTURE2 _8:capt3_save = _8:CAPTURE3 _1 :lowSig = _1:#overflow _1 :razFlg = _1:#overflow _4:TIMER = _4:#60000 s4:PC = _4:*SP++[8]</pre>
--------	-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------

Example C above is a hypothetical exception interrupt service routine for floating-point overflow using alternate immediate exception handling. In the first line, the PC is pushed onto the stack using SP as a pointer and automatically post-decremented by 8. Note that a size of 4 could also have been used, but a size of 8 was chosen here to maintain size coherence with possible binary64 floating-point values that might also be pushed at some later point in the service routine. The last line in the routine shows SP effectuating a stack “pop” which restores the PC with program address it would have fetched from but for the interrupt, effectuating a return from interrupt operation. Note the signal “s”-bit for the destination is set for the first and last instruction of the above service routine. This is simply to show how the signal-bit of a given destination, operandA or operandB field can be used to “signal” something has happened or make a distinction of some kind. In this instance, it merely signals the processor has entered (first occurrence) and exited (second occurrence) a interrupt service routine, distinguishing it from a mere subroutine, for example.

FIG. 14C is a simplified schematic diagram illustrating how SourceA, SourceB, and Destination direct and indirect addresses are generated from the disclosed universal floating-point ISA’s instruction in an exemplary embodiment of the present disclosure. In particular, FIG. 14C shows how addresses corresponding to instruction **100** fields DEST, **120** srcA **130** and srcB **140** are selected. In the schematic, “ind_SrcA_q0” is IND **190** of srcA **130** field, OpsrcA_q0 [14:0] is **210** of srcA **130** field and OpsrcA_q0[2:0] is **250** of srcA field. If Ind_SrcA_q0 (IND **190**) is a “1” then SrcA_addrs_q0 **6200** is chosen from among the contents of one of ARn or the Stack Pointer SP, depending on the value of OpsrcA_q0[2:0] **250** in the instruction. Otherwise, if ind_SrcA_q0 is a “0”, then the zero-extended direct address OpsrcA_q0[15:0] is used as the srcA q0 address. The logic for determining the addresses SrcB_addrs_q0 **6210** srcB q0 address, Dest_addrs_q2 **6220** destination q2 address, and

Dest_addrs_q0 **6230** destination q0 address is the same as for SrcA_addrs_q0, except they each use their own respective IND **190** bits in their respective instruction fields.

Note that pulls from memory occur before pushes. Thus, srcA and srcB addresses must be produced during Stage q0. Pushes occur during the rising edge of Stage q2, thus the destination address must be delayed by two clocks before use as a push address. There are a couple exceptions to this general rule. The first exception is the case where a parent CPU needs to perform a real-time data exchange between it and a XCU connected to it. In such case, the destination address for the real-time data exchange must be pushed simultaneously with the srcA address into the XCU data exchange “hot-spot” in the parent processor’s memory map, in that they are both used by the target XCU’s cascaded instruction logic to synthesize a new instruction on-the-fly that is substituted and executed instead of the one that was just fetched from the XCU’s own program memory during that cycle.

The other exception is for use by the processor’s logic for selecting between fat stack or 5-bit signal stack portions of the dual asymmetric stack during a read of either. For more information on this, refer to “fat stack” block **9500** in FIG. **44**.

For real-time-data-exchange operations between a debug interface such as JTAG and the CPU or real-time-data-exchange operations between a parent CPU and child XCUs, an extra level of address selection is included on the SrcA_addrs_q0 output **6200** and on the Dest_addrs_q2 output **6220** by way an additional mux. Thus, during q0 of a real-time-data-exchange operation, if the monitor read address **7330** is greater than 0x00007FFF, ind_mon_read **7310** will be driven high by logic in the hardware breakpoint module **7300** of debug module **7000** of FIG. **47A**, indicating a indirect monitor read cycle, in which case the monitor read address **7330** is driven onto srcA_addrs_q0 output **6200**. Likewise, if the monitor write address **7350** is greater than 0x00007FFF, ind_mon_write **7340** will be driven high by logic in the hardware breakpoint module **7300** of debug module **7000**, indicating a indirect monitor write cycle, in which case the monitor write address **7350** is driven onto Dest_addrs_q2 output **6220**.

FIG. **15** is a schematic diagram of an exemplary embodiment of the CPU and XCU pre_PC used for accessing the next instruction. FIG. **15** shows the logic **610** for implementing the pre_PC **1640**, which requires the contents of the processor’s program counter (PC) **1620** as input, along with several other inputs. As mentioned previously, a “pre”_PC is required because the SRAM blocks in modern FPGAs are synchronous, thus the read address into program memory must be registered internally a clock cycle ahead of the time the contents are actually needed. It can be seen from the pre_PC **1640**, that in the disclosed embodiment, the pre_PC and, consequently the PC, has an address reach of one megawords (eight megabytes) and, consequently, a 20-bit program counter. This, of course, can be increased to as much as 64-bits if desired, but such would provide excess capability for implementations fully embedded in an FPGA. The pre_PC is not directly readable by software.

FIG. **16** is a schematic diagram **615** of an exemplary embodiment of the CPU’s and XCU’s memory-mapped program counter (PC) operator **1620**, which is directly readable and writable by software. As shown, the PC can be loaded from a number of sources, most of them driven by prioritized events that do not include a direct, unconditional write operation. For example, “Id_vector” when an interrupt is being acknowledged is, apart from an active RESET, the

highest priority. If “Id_vector” is active, PC will automatically be loaded with the presented vector[19:0].

Next in priority is “rewind_PC”. If rewind_PC is active and no breakpoint is active, then the PC **1620** is automatically loaded with the previous PC value. Rewind_PC is usually the result of a write-side collision or a not-ready condition.

Next in priority is a “bitmatch” **1610** condition as a result of bit-test logic that can only go active “true” if and only if the specified bit of the data being pulled corresponds to the PC conditional load destination address of the instruction. For example, a bit test is “true” if the bit being tested is a “1” AND the destination address is PCS (load-Program-Counter-if-Set address). A bit test is also “true” if the bit being tested is a “0” AND the destination address is PCC (load-Program-Counter-if-Clear address). If the test is “true”, the PC is loaded with the sum of the signed relative displacement specified in the instruction and the PC value at the time of the fetch delayed by two clocks (i.e., pc_q2). Note, in terms of direct addresses, “BTBC_” and “BTBS_” are synonymous with “PCC” and “PCS”, respectively. “JMPA_” and “BRAL_” are synonymous with “PC” and “PCR” respectively.

Next in priority is direct, unconditional write of the PC **1620** in software employed as a long absolute jump.

Last in priority, when not in an active break state or an active REPEAT cycle, the PC is loaded with the contents of the pre_PC **1640** of FIG. **15**. If the REPEAT counter is not zero or a break condition is active, the PC remains unchanged.

Anytime the PC **1620** is loaded with a discontinuous value (i.e., a value not exactly equal to PC+1), “discont_out” **1600** goes active “1” to indicate a PC discontinuity. The discont_out **1600** is used by the debug logic block’s PC discontinuity trace buffer, if present. The discont_out **1600** is also used by the logic needed to kill/disable the processor’s write signal that would have otherwise occurred two clocks later in the processor’s push-pull pipeline, effectuating a pipeline flush. This is necessary because, by the time a branch actually executes, there are still two previously-fetched instructions in the pipeline that are about to execute with a corresponding active write. “Write-disable” **1630** is necessary to kill the corresponding writes of the instruction pipeline when a discontinuity occurs.

Further note that, like the pre_PC **1640**, the PC **1620** is preset to 0x00100 during reset. This allows the first 100 hex locations (i.e., first 2048 bytes) in program memory to be used for storage of constants, if needed.

FIG. **17** is a schematic diagram of an exemplary embodiment of the CPU’s and XCU’s memory-mapped PC-COPY register **618**. On RESET, PC_COPY **1650** is preset to 0x00100, the same as PC **1640** and pre_PC **1620**. The main purpose of PC_COPY is to provide a return address anytime an interrupt or subroutine call is made. Thus, upon entry into a subroutine or service routine, PC_COPY (and not the PC, per se) must be pushed onto the system stack to preserve the return address. Once the subroutine or service routine is complete, then a return can be effectuated by popping the top of the stack into the PC, as the last step. Thus PC_COPY **1650** only changes when an event occurs that causes a PC discontinuity.

PC_COPY **1650** is readable and is located at location 0x00007FF2 in the processor’s data memory among the other programming model registers **2000** of FIG. **11**.

FIGS. **18A** and **18B** are an exemplary table illustrating the bits of the CPU’s and XCU’s memory-mapped STATUS register/operator and their respective functions in an exem-

plary embodiment of the present disclosure. FIG. 18A shows bits **63** thru **16**, and FIG. 18B shows bits **15** thru **0**. Taken together, FIG. 18A and FIG. 18B show all the bits and their functional names comprising the processor's memory-mapped, 64-bit STATUS Register **5000**. Most of these bits enable carrying out in hardware all the default and alternate modes mandated or recommended in the IEEE 754-2008 floating-point specification.

As shown, most of the bits are grouped, and these grouped bits have a common bit-set and bit-clear address, such that when a “1” for the corresponding bit position is written to the group's bit-set address, that bit position is set. If a “1” for that bit position is written to that group's bit-clear address, that bit is cleared. In this manner, any combination of bits within a group can be easily and efficiently set or cleared with a single push. For example, bits **0** thru **5** (FIG. 18B) are the “processor integer arithmetic flags (condition code), status and control bits” group **5190**, and bits **16** thru **20** (FIG. 18A) are the IEEE 754-2008 exception signals group **5120**, and so on.

The main purpose for grouping the bits in this manner is that they can easily be manipulated in hardware, as a group. The bits **0** thru **5** comprise “processor integer arithmetic flags (condition code), status and control bits” that really have nothing to do with floating-point operations, per se, but rather relate to integer and logical operations. Like the other groups, dedicated logic within certain memory-mapped locations that shadow the STATUS Register has been added to make manipulating them together as a group a simple operation. Without this special logic in hardware, manipulating these bits as a group is a cumbersome process.

The following example code shows a conventional, yet cumbersome, way to clear the C flag and set the Z flag using the logical AND and OR operators:

```
demoG:   _8:work.9 = _8:STATUS
          _1:AND.0 = (_1:work.9, _1:#0x3D)
          _1:OR.0 = (_1:AND.0, _1:#0x01)
          _1:work.9 = _1:OR.0
          _8:STATUS = _8:work.9
```

Another conventional method for setting and clearing individual bits is to use the processor's memory-mapped BSET and BCLR operators, which can only manipulate one bit at a time. For example:

```
demoH:   _8:BCLR.0 = (_8:STATUS, _1:#C)
          _8:BSET.0 = (_8:BCLR.0, _1:#Z)
          _8:STATUS = _8:BSET.0
```

While cumbersome, both demoG and demoH work fine. The main problems with using them is that, due to pipeline issues regarding intermediate results not being immediately ready for use in subsequent steps, a pipeline stall will be incurred after each step in the process. Not only that, but demoH is slightly more efficient than demoG and eventually becomes less efficient than demoG when more than one bit needs to be set or cleared at a time because the BSET and BCLR operators presently only operate on one bit at a time.

To solve this problem and make manipulation of bit groups within the present invention's STATUS Register easier and more efficient to carry out, special logic has been added to it, making the memory-mapped STATUS Register an operator in its own right. The following code demonstrates bit group manipulation within the STATUS Register itself:

demoJ:	—	$-1:0x7CDD = _1:#0x05$
demoK:	—	$_1:0x7CDC = _1:#0x32$

In the “demoJ” example above, the Z flag and the N flag of the STATUS Register are simultaneously set using a single instruction. Other bits that may be already set, remain set, i.e., undisturbed. Referring to FIG. 18B for those bits, notice that there is a column in the table corresponding to a “setBit” address and a “clearBit” address. Thus, in demoJ, the single-byte location **0x7CDD** is the destination for the “set-bit(s)” operation within the STATUS Register itself. The #immediate value of **0x05** corresponds to the bit respective bit positions of the **5190** bit group. For example, **000101** binary corresponds to the relative bit positions of the Z flag and the N flag within the **5190** bit group. If setting all the bits within the group is desired, then **111111** binary (or **0x3F** hex) would be used.

Likewise, demoK shows how to clear desired bits within the **5190** group. The method is identical to setting bits as shown in demoJ, except the pattern of bits specifying which bits to clear are pushed into location **0x7CDC** (the “clearBit” address) instead of **0x7CDD**. Bits unspecified in the value being pushed remain undisturbed.

Some bits within **5000** are read-only, meaning they cannot be directly written to. For example, bit **6**, when set, means **A>B** as shown at **5180**. Bit **7**, when set, means **A<=B** as shown at **5170**. Bit **8**, when set, means **A>=B** as shown at **5160**. A test of bits **6, 7, 8** should be tested only after a COMPARE (integer compare) operation. For example:

demoL:	—	$_8:COMPARE = (_8:work.10, _8:workx.11)$
	—	$_4:PCS = (_8:STATUS, 6, skip1)$
	—	$_8:ADD.7 = (_8:ADD.7, _1:#0x01)$
skip1:	—	$_8:work.12 = _8:ADD.7$

Bit **9**, the “exception Source” bit **5150** indicates which result bus (i.e., read-side A or read-side B), caused a floating-point exception, if any. A “1” in bit **9** may indicate that a pull from the B-side of the floating-point operator block caused the instant exception. A “0” may indicate that the exception, if any, was caused by a pull from the floating-point operator block A-side.

Bit **10** is a read-only bit that reflects the state of the interrupt request IRQ **5140** input of the processor. A “1” may indicate that an interrupt is being requested. If bit **5**, the interrupt enable bit in group **5190**, is “1” at the time of the request, and that interrupt or non-maskable interrupt (NMI) is not already in service, the PC is automatically loaded with the address of that interrupt's service routine. If bit **5** is “0”, interrupts are disabled and any request on the IRQ bit **5140** is ignored. Note that in the disclosed embodiment, interrupt vectors must be loaded into their respective registers during power-up initialization. The vector register locations are given in table **2000** of FIG. 11.

Bits **11** thru **15** are spare bits **5130** that are presently unused.

Bits **16** thru **20** are the IEEE 754-2008 exception “signal” bits **5120** implemented in hardware. Anytime a result is pulled from a floating-point result buffer, if an exception occurred during its computation, under alternate immediate exception handling, the corresponding exception will be signaled automatically in bits **5120** as a result of the pull.

Under default exception handling, if an overflow condition is indicated as a result of a pull from a floating-point operator, only the inexact signal (bit **20** of group **5120**) and

41

overflow Flag (bit **23** of group **5110**) will be automatically set, in that overflows are always inexact. If an underflow is indicated as a result of a pull from a floating-point operator under default exception handling, only the inexact signal in bit **20** is automatically set, but only if the underflowed result is also inexact. Under default exception handling for underflow, if the underflowed result is exact, underflow is not flagged and inexact is not signaled.

Bits **21** thru **25** are the IEEE 754-2008 exception “flags” **5110** implemented in hardware. Under default exception handling, anytime a result is pulled from a floating-point result buffer, if an exception occurred during its computation, the exception will be flagged (set=“1”) automatically in bits **5110**, provided that its corresponding “razNoFlag” bit in group **5100** is not enabled.

Under alternate immediate exception handling, it is up to the corresponding exception interrupt service routine to explicitly set the corresponding exception flag in group **5110** and explicitly clear the corresponding exception signal in group **5120** that was automatically raised, causing the floating-point exception interrupt.

Bits **26** thru **30** are the IEEE 754-2008 “raise no flag” bits **5100** implemented in hardware. Under the standard’s default exception handling, if a corresponding “razNoFlag” bit is enabled, any result pulled from a floating-point result buffer that experienced an exception during its computation, such exception will be ignored by the **5110** logic. Note that upon reset of the processor, bit **30** of group **5100** is preset to “1”. This is because, under default exception handling for inexact, the inexact “flag” is not raised. Under alternate immediate exception handling for inexact, it is up to the implementor’s inexact exception interrupt service routine to explicitly raise the inexact flag, if signaled in group **5120**, and such does not necessarily depend on the state of bit **30** in group **5100**.

It should be noted that bit groups **5120**, **5110**, **5100**, and **5090** can be manipulated as a group in the same manner as bit group **5190** described previously.

Bits **36** thru **45** are the mutually exclusive “one-of” “class” bits **5080** that, together, indicate the IEEE 754-2008 class of the number pushed into the Class hardware operator “clas” at direct location **0x7FCE** (see FIG. 12D). Stated another way, after executing the Class operation, only one of these ten bits will be a “1”, while all the others will be “0”.

Bit **46** stores the result **5070** of this processor’s IEEE 754-2008 hardware total Order operator (see FIG. 12C).

Bit **47** stores the result **5060** of this processor’s IEEE 754-2008 hardware total Order Mag operator (see FIG. 12C).

A “1” in **5050** (bit **48**) indicates whether any of the **5110** flags are presently raised as a result of either a “testSavedFlags” or “testFlags” push. A “0” in **5050** indicates no flags are presently raised. Table 3050 of FIG. 12E gives the direct address of these two hardware operators, as well as examples of how to use them.

Bit **49** stores the result **5040** of this processor’s IEEE 754-2008 hardware “Is” operators in **3040** of FIG. 12D. For example, to test a number to see if it is a normal number, push it into the “isNormal” operator located at **0x00007CD1**. If it is in fact a normal number, when the operation is completed, the “isTrue” bit (bit **49**) will be a “1”, otherwise it will be cleared to “0”.

Bit **50** stores the result **5030** of this processor’s IEEE 754-2008 comparison predicates operators shown in **3020** of FIG. 12B. If the comparison is true, then **5030** is set; otherwise it is cleared. For example, to perform an IEEE 754-2008 “compareSignalingGreaterEqual” operation, push

42

the two operands into cmpSGE (location **0x00007CF9**). Then simply perform a bit-test and branch operation as follows:

5	demoM:	—	_8:cmpSGE = (_8:bin64A, _8:bin64B)
		—	_4:PCS = (_8:STATUS, 50, true1)
		—	_8:FSUB.0 = (_8:bin64A, _8:bin64B)
	true1:	—	_8:FMUL.0 = (_8:FSUB.0, _8:FSUB.0)

10 Bits **51** thru **54** are the disclosed processor’s hardware implementation of the IEEE 754-2008 directed rounding mode attributes. These bits **5020** can be configured with a single write to direct address **0x00007FD8**. With bit **53** (the enable RM attribute bit) set to “1” the rounding mode specified in bits **52** and **51** overrides the static rounding mode specified in bits **63** and **62** in the instant instruction. Below is a table showing the four rounding modes according to bits **52** and **51** (assuming bit **53** is set):

Code	Rounding Mode Attribute
00	Use “default” rounding mode
01	Round towards positive infinity
10	Round towards negative infinity
11	Round towards zero (i.e., truncate)

30 Bit **54** (“Away”) determines the default rounding mode for both the dynamic round mode attributes specified by bits **51** and **52** of the STATUS Register when the enable RM attribute bit (bit **53**) is set (“1”) and the static rounding mode (i.e., the rounding mode specified in the instant instruction) when bit **53** is clear (“0”). If Away is “0”, then the default rounding mode is round to nearest even. If Away is “1”, then 35 the default rounding mode is round away from zero.

35 Bit **55** is the “override RM bits” control bit. It is intended for use as an aid in debugging. When set, it will force the rounding direction to nearest even, no matter what modes are specified in group **5020**.

40 Bits **56** thru **63** are for use as the substitute enable bits **5010** recommended in IEEE 754-2008 for resuming alternate exception handling. During an exception service routine under resuming alternate exception handling, these bits can be tested for determining whether a result should be substituted, depending on the exception in service. Like bit groups **5190**, **5120**, **5110**, **5100**, and **5090**, the bits in **5010** can be manipulated in like manner as a group using a single instruction.

45 FIG. 18C is a block diagram illustrating an exemplary embodiment of an exemplary arrangement, mapping, and implementation of IEEE 754-2008 mandated Comparison Predicates, dual-operand and single-operand non-computational, non-exceptional operators in relation to their respective bits in the CPU’s and XCU’s memory-mapped STATUS register/operator **5000** (FIGS. 18A-18B).

50 FIG. 18D thru FIG. 18O are exemplary schematic diagrams of the present invention’s memory-mapped multi-function STATUS Register hardware operator. The STATUS Register is, in itself, an operator because all the logic necessary to carry out not only the IEEE 754-2008 comparison predicates, dual and single-operand non-computational non-exceptional operations, integer compare, and efficient manipulation of bit groups is all part of the STATUS Register operator’s logic, in that each operator has a unique memory-mapped input address that shadows the STATUS Register’s primary programming model direct address at location **0x00007FF1** in data memory. For efficient context

save/restore operations, this primary programming model direct address can be used to save/restore the entire 64-bit contents of the STATUS Register with a single pull/push.

FIG. 18D is a schematic diagram illustrating, along with their respective bit positions, exemplary logic for carrying out in hardware, bit manipulation of the Enable Alternate Immediate exception handling (bits 31-35) as a “group” within the memory-mapped STATUS register/operator in an exemplary embodiment of the present disclosure.

FIG. 18E is a schematic diagram illustrating, along with their respective bit positions, exemplary logic for carrying out in hardware bit manipulation of the Raise No Flag specifiers for the five IEEE 754-2008 exceptions (bits 26-30) as a “group” within the memory-mapped STATUS register/operator in an exemplary embodiment of the present disclosure.

FIG. 18F is a schematic diagram illustrating, along with their respective bit positions, exemplary logic for carrying out in hardware bit manipulation for the IEEE 754-2008 Inexact, Underflow, Overflow, Divide-by-Zero and Invalid flags (only the first three bits, bits 23-25, are shown due to space limitations) as a “group” within the memory-mapped STATUS register/operator in an exemplary embodiment of the present disclosure.

FIG. 18G is a schematic diagram illustrating, along with their respective bit positions, exemplary logic for carrying out in hardware bit manipulation for the IEEE 754-2008 Inexact, Underflow, Overflow, Divide-by-Zero and Invalid flags (only the last two bits, bits 21 and 22, are shown due to space limitations) as a “group” within the memory-mapped STATUS register/operator in an exemplary embodiment of the present disclosure.

FIG. 18H is a schematic diagram illustrating, along with their respective bit positions, exemplary logic for carrying out in hardware bit manipulation for the IEEE 754-2008 Inexact, Underflow, Overflow, Divide-by-Zero and Invalid “signals” (only the first three bits, bits 18-20, are shown due to space limitations) as a “group” within the memory-mapped STATUS register/operator in an exemplary embodiment of the present disclosure.

FIG. 18I is a schematic diagram illustrating, along with their respective bit positions, exemplary logic for carrying out in hardware bit manipulation for the IEEE 754-2008 Inexact, Underflow, Overflow, Divide-by-Zero and Invalid “signals” (only the last two bits, bits 16 and 17, are shown due to space limitations) as a “group” within the memory-mapped STATUS register/operator in an exemplary embodiment of the present disclosure.

FIG. 18J is a schematic diagram illustrating, along with their respective bit positions, exemplary logic for carrying out in hardware bit manipulation for the CPU and XCU logical and integer arithmetic Carry (“C”), Negative (“N”), Done, Interrupt Enable (“IE”), Zero (“Z”), and Overflow (“0”) flags (only bits 1, 2, 4, and 5 are shown due to space limitations) as a “group” within the memory-mapped STATUS register/operator in an exemplary embodiment of the present disclosure.

FIG. 18K is a schematic diagram illustrating, along with their respective bit positions, exemplary logic for carrying out in hardware integer comparisons in addition to bit manipulation for the CPU and XCU logical and integer arithmetic Zero (“Z”) and Overflow (“0”) flags (bits 0 and 3) as a “group” within the memory-mapped STATUS register/operator in an exemplary embodiment of the present disclosure. Note that the Z and V flags can be bit-manipulated with N, Done, IE, and C as a group, but the Z and V flags have

additional logic to support integer comparisons within the memory-mapped STATUS register/operator.

FIG. 18L is a schematic diagram illustrating, along with their respective bit positions, exemplary logic for carrying out in hardware bit manipulation for the IEEE 754-2008 “recommended” substitutions for abrupt underflow, substitute X, substitute xor(X), inexact, underflow, overflow, divide-by-zero and invalid exceptions as a “group” within the memory-mapped STATUS register/operator in an exemplary embodiment of the present disclosure (only the first five bits, bits 59-63 are shown due to space limitations).

FIG. 18M is a continuation of FIG. 18L above illustrating schematics for the substitute overflow, divide-by-zero and invalid exceptions (bits 56-58).

FIG. 18N is a schematic diagram illustrating, along with their respective bit positions in the STATUS register/operator, exemplary logic for carrying out in hardware bit manipulation as a group, dynamic rounding mode attributes mandated by IEEE 754-2008, namely, the encoded Rounding Mode bits 1 and 0, Away bit, Enable dynamic rounding mode, and default override bit (bits 51-55) in an exemplary embodiment of the present disclosure.

FIG. 18O is a schematic diagram illustrating an exemplary embodiment of memory-mapped logic for carrying out in hardware the IEEE 754-2008 mandated testing for, as a group, an exception flag raised condition using a “testSavedFlags” 5320 or “testFlags” 5330 memory decode, as well as restoring this status bit using the “loadStatusReg” 5310 memory decode logic for these memory-mapped operators that either set or clear the single-bit register 5370, “aFlagRaised”, which is Status Register bit [48] 5380. As can be from this schematic, whenever the CPU writes directly to the Status Register to restore all restorable bits in it, 5370 is loaded with wrdata[48] 5360 of the write data bus.

To perform an IEEE 754-2008 “testFlags” operation, simply push a 5-bit value, whose “1” bits correspond in position to the exception flag to be tested, into the “testFlg” operator “tstFlg” location 0x7CE3 in data memory (FIG. 12E for examples on use). As can be seen from FIG. 18O, OR gate 5340 will be driven high if any combination of one or more Status Register exception flags is high AND write-data bit corresponding to such flag is also high. Thus, if anything is pushed into direct byte location “tstFlg” 0x7CE3, signal “testFlags” 5330 will go high and the output of OR gate 5340 will be registered into the “aFlagRaised” register 5370.

Similarly, “testSavedFlags” 5320 compares respective bit positions of exception flags that were previously saved and later read back to do the bit test as a group. In this instance, when the previously saved exception bits are read back, they will be on the five LSBs of the wrdataA (operandA) bus and are compared with the corresponding bits on the wrdataB (operandB) bus. If any operandA bit and corresponding operandB bit matches are both “1”, this condition is registered in 5370. Thus, the instant circuit provides an efficient means in hardware to test saved exception flags as a group using a single push into location tstSavFlg operator direct data location 0x7FCE. Refer to FIG. 12E for an example on how to use this operator.

FIG. 19 is a schematic diagram illustrating an exemplary embodiment of the disclosed processor’s memory-mapped hardware REPEAT counter circuit 1900, which is mapped at location 0x000007FEF in directly addressable data memory space. Once loaded, it begins to decrement by one every clock cycle until it reaches zero. During the time it is not zero, the pre_PC and the PC remain frozen, causing the current instruction to execute a total of REPEAT+1 times.

As shown, there are actually two repeat counters, not just one. This is because an instruction being repeated using the repeat counter is non-atomic with respect to enabled interrupts. This means that if an enabled interrupt is asserted and acknowledged, the “REPEAT_a” counter **1910** suspends down-counting until the interrupt service routine is completed, at which time it automatically resumes down-counting until it reaches zero.

There are two virtually identical REPEAT counters: “REPEAT_a” **1910** and “REPEAT_int”**1940**. “REPEAT_a” is automatically selected for use “outside” of interrupt service routines. “REPEAT_int” is automatically selected for use “inside” interrupt service routines. This strategy allows instructions being repeated to be interrupted. If an enabled interrupt is acknowledged, counting of REPEAT_a is temporarily suspended automatically and its contents remains undisturbed during service of the interrupt.

While inside an interrupt service routine, REPEAT_int may be loaded and used in the same way as using REPEAT outside an interrupt service routine. Since the described embodiment does not permit nested interrupts, only one REPEAT_int counter is required. If the implementer incorporates an interrupt controller that permits nested interrupts, then a REPEAT_int counter for use by each interrupt should be included. Alternatively, a state-machine can be employed to suspend decrementing REPEAT upon acknowledgment of an interrupt, allowing time for the current contents of REPEAT to be preserved by pushing it onto the stack. This process is tricky and the process of preserving and restoring REPEAT cumbersome, which is among the reasons the strategy of having two separate REPEAT counters that shadow each other was chosen for the described embodiment.

Encountering a breakpoint will suspend decrementing of the REPEAT operator until the processor exits the breakpoint. The current repeat count **1920** may be pulled during an active breakpoint. If the breakpoint occurs during an interrupt service routine, then the current contents of REPEAT_int **1940** is read as the current repeat value **1920**. If the breakpoint occurs outside an interrupt service routine, then the current contents of REPEAT_a **1910** is read as the current repeat value **1920**.

A parent CPU can read the contents of a child XCU REPEAT counter. The JTAG real-time debug module, if present, can read the contents of the CPU REPEAT counter, anytime, on-the-fly or during a breakpoint.

From FIG. 19, it can be seen that the REPEAT counter can only be loaded using #immediate addressing mode or using the contents of one of the Auxiliary Registers (ARn). The reason for this is because the immediate value is immediately available within the instant instruction and does not need to be read out of RAM. As such, the REPEAT counter is loaded one clock early so that it is ready to start down-counting by the time the following instruction executes.

The main disadvantage to loading the REPEAT counter with an #immediate value is that such value is fixed and cannot be changed or made variable, due to it being in program memory, which is to say, loading REPEAT with an #immediate value does not permit “computed” repeats, wherein the REPEAT amount is variable. Loading the REPEAT counter using the contents of one of the Auxiliary Registers enables computed repeat amounts. This works because the contents of Auxiliary Registers are immediately available for this purpose. To implement a computed/variable repeat amount, perform the computation and push the

result into one of Auxiliary Registers (ARn), then push the contents of the same ARn into the REPEAT counter. For example:

demoN:	—	_2:REPEAT = _2:#31
—	—	_8:*AR1++[8] = _8:*AR2++[8]
demoP:	—	_8:ADD.0 = (_8:work.2, _1:work.7)
—	—	_8:AR3 = _8:ADD.0
—	—	_2:REPEAT = _2:AR3
—	—	_8:*AR1++[8] = _8:*AR2++[8]

In the “demoN” example above, REPEAT is loaded with the #immediate value #31, causing the instruction following it to execute a total of thirty-two times.

In the “demoP” example above, an amount is computed by adding the contents of work.7 to the contents of work.2, with the result of the ADD operation being pushed into AR3. The contents of AR3 is then pushed into the REPEAT counter, with the following instruction executing the previously computed number of times+1.

It should be understood that the REPEAT counter will only decrement if DEST **120**, srcA **130** or srcB **140** in the following instruction employ the indirect (auto-post modify) addressing mode. If none of the addressing mode fields in the instruction that follows the REPEAT load are indirect (auto-post modify) addressing mode, then the value previously loaded into the REPEAT counter will not decrement. Stated another way, the REPEAT counter is intended to be used only for repeating indirect, auto-post-modify instructions.

FIG. 20 is a schematic diagram illustrating an exemplary embodiment of the disclosed universal floating-point ISA’s memory-mapped hardware loop-counter operators **1950**. In the disclosed embodiment, there are two hardware loop-counters, LPCNT1 and LPCNT0 that employ identical hardware loop-counter logic **1950**. LPCNT1 and LPCNT0 may be sixteen bits wide and may be implemented as down-counters using adder **1970** that adds a -1 to the current contents of the LPCNT **1960**. Each loop counter circuit **1950** has logic that detects when the current value of its respective counter is not zero and outputs a signal LPCNT_nz **1990** that is read simultaneously with the current count value, with LPCNT_nz being the most significant bit. Stated another way, when a loop counter is pulled, LPCNT[15:0] value **1980** occupies the first sixteen LSBs and LPCNT_nz **1990** occupies the MSB position of the 17-bit word being read.

Either loop counter may be read/pulled anytime and can be written anytime, but only with the direct addressing mode. Once loaded with a non-zero value, the loop counter automatically decrements, but only as a result of a read/pull, wherein the direct destination address (PCS_) is the PCS operator, and the direct srcA address is the desired loop counter address, effectuating a conditional load of the PC when the tested bit is set (i.e., “1”).

Employment of the hardware loop counters is much more efficient for tight looping scenarios than conventional methods that involve initializing a memory location with a start value, then manually subtracting one from it, pulling the result out of the SUBtract operator to register its condition signals in the STATUS Register, and then manually testing the Z flag in the STATUS Register to determine if a branch should be taken, each of these steps being done for each iteration. Such methods work fine for routines that are average or longer length.

However, for tight looping scenarios where the sequence of instructions in the loop is only three or four instructions, a hardware loop counter is much more efficient and desirable.

The following example shows a conventional memory-based loop counter in software.

```
DemoQ:
loop3:    _2:work.0 = _2:#24
          _2:SUB.0 = (_2:work.0, _2:#1)
          :
          <hypothetical code>
          :
          _2:work.0 = _2:SUB.0
          _4:PCS = (_4:STATUS, 0, loop3)
```

The following is an example using one of the instant invention's memory-mapped hardware loop counters.

```
DemoR:
loop3:    _2:LPCNT0 = _2:#24
          :
          <hypothetical code>
          :
          _4:PCS = (_4:LPCNT0, 16, loop3)
```

In the “DemoR” example above, each time the last line is executed, the **LPCNT_nz 1990** is tested to see if it is set (“1”). If set, i.e., the **LPCNT** value is not zero, the PC will branch to “loop3”. If clear, it will exit the loop. Each time **LPCNT0** is tested, such test automatically decrements that counter by 1 if the pre-registered value immediately on the output of 1970 is not already 0.

FIG. 21 is a schematic diagram of an exemplary embodiment of the disclosed universal floating-point ISA's optional floating-point exception capture module **1700** that can be used for capturing diagnostic information related to a particular floating-point result and any exception that may be signaled when pulled from a floating-point result buffer when alternate immediate exception handling is enabled for that exception. The disclosed embodiment of the exception capture module comprises four, 64-bit registers **1702** that simultaneously capture the operator's A-side result, B-side result (if any), srcA address of the A-side operator being pulled, srcB address of the B-side operator being pulled, the PC value of the instruction that caused the pull, and the destination address of where the pulled results were to be pushed as a result of the instruction.

These 64-bit capture registers are mapped into the processor's programming model block **2000** shown in FIG. 11 and can be read in software by the exception service routine.

FIG. 21 shows, among other things, that whenever an alternate exception handler is enabled via the corresponding bits in **5090** of the STATUS Register **5000** (FIGS. 18A-18B) and a floating-point exception interrupt occurs as a result, the exception capture logic of the exception capture module kills (by way of “writeAbort”) the processor's write cycle to prevent the excepted result from being written to the destination address specified in the instruction that did the pull. Instead, the pulled result is automatically captured by the capture register. This is so the exception service routine can perform a substitution, if desired, in software and then perform the write to the destination from within the service routine using the captured destination address.

FIG. 22A is a block diagram illustrating an exemplary embodiment of the disclosed universal floating-point ISA's hardware implementation of IEEE 754-2008 mandated computational operator module **3010** showing dual operand inputs, their exception signals, dual result outputs and ready semaphore output. Module **3010** implements in hardware, all the IEEE 754-2008 mandated computational operators. These operators can input and output in binary64, binary32,

and binary16 formats, in any combination. The core operators can be double, single, or half precision or any floating-point operator or combination of operators that together can simultaneously accept up to 1024 bits on the A side and up to 1024 bits on the B side. Some operators can directly input decimal character sequences up to 28 decimal digits in length on both operandA and operandB simultaneously and such inputs can be mixed and matched with binary64, binary32, and binary16 formats in any combination, all without having to explicitly convert from one format to another.

In some implementations, operand A and operandB universal input converters **3020** and **3030**, respectively, can convert binary16, binary32, and binary64 format numbers to FP1262 format numbers having a sign bit, a 12-bit exponent, and a 62-bit fraction. Some implementations may not need double-precision and in such cases, converters **3020** and **3030** would convert to a smaller FPxxxx format. For instance, instead of the base operators of module **3010** being FP1262 format (sufficient for computing with double-precision subnormal numbers), an application may only require the ability to compute with single-precision numbers, including their subnormal representations, thus FP928 (9-bit exponent, 28-bit fraction) operators in module **3010** would be sufficient.

In other applications, converters **3020** and **3030** not only automatically convert binary16, binary32, and binary64 numbers to the target hardware binary format of the installed hardware operators of module **3010**, but also automatically convert decimal character sequences to the binary format of the installed hardware operators of module **3010**.

In still other applications, there may not be any binary conversion or decimal character sequence conversions by converters **3020** and **3030** at all.

Still further, there may be some implementations where only a select few of the installed operators require the ability to directly accept mixed formats, such as, for example operandA being a binary32 representation and operandB being a H=20 decimal character sequence. In such implementations, it may be more advantageous for that operator to have its own dedicated conversion circuits embedded in that operator module.

From FIG. 22A, it can be seen that there is a rddataA bus for reading the A-side of an operator's result buffer and a rddataB bus for simultaneously reading the B-side of an operator's result buffer. Such simultaneous reads for A-side and B-side can be from the same or different operators or RAM. It can also be seen that both A-side and B-side buses also comprise a 5-bit exception bus that contains the exception signals (divide-by-zero, invalid, overflow, underflow, and inexact) that were stored along with the result in the result buffer for the operator results being pulled.

Also shown is a “ready” semaphore that signals to the main processor pipeline that the result being pulled is ready. Both the A-side result and the B-side result must be complete, otherwise “ready” will be low (“0”), indicating a not-ready state, which will cause the processor's PC to rewind and make another attempt at pulling the result(s) by re-fetching and executing the instruction at the original fetch address, which will continue until a ready state is signaled as being ready.

FIG. 22B illustrates exemplary memory-mapped hardware implementations of IEEE 754-2008 convertToDecimalCharacter **9400**, Addition **9420**, Fused-Multiply-Add **9700**, and convertFromDecimalCharacter **9460** operator module inputs and outputs implemented in module **3010** in an exemplary embodiment of the present disclosure. The

memory-mapped floating-point operator hardware implementations are in the form of modules that can be easily and conveniently included or excluded in a particular implementation, depending on that implementation's specific requirements.

For example, one implementation may require that the parent CPU contain the full and complete repertoire of IEEE 754-2008 mandated computational and non-computational operators while any child XCU attached to it contain only a small subset of such operators. Thus, if a given XCU will never employ a hardware "remainder" operation, for example, then logic or routing resources on the chip are not wasted implementing the remainder operation in that particular XCU.

Another example is the case where the parent CPU implements a full or partial repertoire of floating-point operators, while a first child XCU attached to it is implemented as a general-purpose micro-controller to handle a given communications protocol with the host system and, consequently, includes no floating-point hardware at all, while a second child XCU is implemented as a general-purpose fixed-point DSP for speech recognition and, consequently, includes no floating-point operators, while third and fourth child XCUs are implemented with only the floating-point hardware necessary to carry out deep learning algorithms for web-based artificial intelligence applications.

FIG. 23 is a block diagram illustrating an exemplary embodiment of the disclosed universal floating-point ISA's logical and integer arithmetic operator module **4000** illustrating dual operand inputs, their signals, dual result outputs and ready semaphore output. Like the memory-mapped IEEE 754-2008 floating-point module **3010** described in conjunction with FIG. 22A, module **4000** has an A-side rddataA bus and a B-side rddataB bus and corresponding signalA and signalB bus, respectively. Note that the signal bus carries the four condition code signals (C, V, N, and Z) that were stored simultaneously with the result of the computation in the selected result buffer for that operator. Just like the floating-point module **3010**, module **4000** has a "ready" semaphore indicating that the selected results are ready. Like the result buffers in the floating-point operator module **3010**, the integer arithmetic and logic operator module **4000** result buffers are also fully restorable.

ConvertFromDecimalCharacter Circuit

FIG. 24 is a block diagram illustrating an exemplary embodiment of the disclosed universal floating-point ISA's hardware implementation of a double-precision IEEE 754-2008 H=20+ convertFromDecimalCharacter circuit **8000** illustrating a virtually identical dual half-system approach. The circuit converts decimal character sequence inputs **8320** up to 28 decimal digits in length to IEEE 754 binary64, binary32, or binary16 formats **8300**. On the front-end is a universal decimal character sequence-to-default format translator **8100**, which translates character sequence **400** (see FIG. 5) inputs to the default character sequence format **300** (see FIG. 4), the translated sequence still being in character form.

The function of the translator **8100** is to separate the integer part and the fraction part of the input character sequence **8320**, placing the characters comprising such parts into their assigned default positions. The translator **8100** also translates the exponent and places it in its assigned default position **310** (FIG. 4). If the character sequence input **400** (FIG. 5) has no explicit exponent or if the input has only a token exponent, the translator **8100** creates the correct character sequence exponent for it and places it in its assigned default position within sequence format **300**. In the

process, the translator **8100** produces a 21-character integer part **350**, a 20-character fraction part **340**, a 5-character exponent (e.g., "e+009") **310**, NaN input payload **510** (if any) and some exception and class-like signals (e.g., isInfinite, isZero, isValid, isInvalid, isExact, etc.) that are propagated alongside the conversion pipeline via a series of delays for later use in determining a final result, if needed.

The circuit **8000** includes half-systems **8400** and **8500**, which are virtually identical in terms of the steps and process 10 the decimal characters comprising the integer part and the characters comprising the fraction part go thru to arrive at an encoded 53-bit floating-point mantissa, one for the integer part and one for the fraction part. In cases where the significand of the intermediate result comprises both an 15 integer part and a fraction part, the appropriate upper bits of the integer part mantissa are concatenated with the appropriate upper bits of fraction part mantissa, for a total of fifty-three bits, with the fifty-third bit (MSB), i.e., the hidden bit, being later discarded, which is why in IEEE 754 20 parlance it is commonly referred to as the "hidden" bit. The 20 process for determining and encoding the two intermediate mantissas (one for integer part and one for fraction part) are 25 virtually identical, except for subnormal fractions, which is explained later. Block **8190** of integer half-system **8400** carries out the integer part intermediate mantissa encoding and block **8330** of the fraction part half-system **8500** carries out the fraction part intermediate mantissa encoding.

It should be understood that the exemplary circuit **8000** shown in FIG. 24 is for producing an IEEE 754-2008 30 "H=20" result that satisfies that standard's minimum requirement for double-precision representations, i.e., a binary64 format result with a minimum 20-decimal character representation for twenty significant digits on the character sequence input. This same process and method can be 35 used for converting IEEE 754-2008 "H=12" (i.e., 12-significant-digit decimal character sequences) for producing a single-precision binary32 format result; IEEE 754-2008 "H=8" (i.e., 8-significant-digit decimal character sequences) for producing a half-precision binary16 format result; or 40 virtually any "H=n" value required. This is done by simply scaling the components, such as ROMs, look-up tables, adders, etc., described herein, sufficiently to produce the desired result format. In implementations that do not require double-precision capability, it is much cheaper to implement 45 this converter as an H=12 (single-precision) circuit, because it is less than half the size of an H=20 circuit. By comparison, an H=8 (half-precision) circuit would be very small in comparison to an H=20 circuit.

Converter **8110** of the integer part half-system **8400** 50 converts each 8-bit character comprising the 21-character integer part to a 4-bit hex digit representation. Likewise, converter **8110** of the fraction part half-system **8500** converts each 8-bit character comprising the 20-character fraction part to a 4-bit hex digit representation. The three 8-bit 55 decimal characters representing the exponent value each get converted by converter **8120** to three 4-bit hex digit sequence that represents the exponent at this stage.

Next, the 21-hex-digit sequence of the integer part is 60 computed by decimal multipliers **8130** of the integer part half-system to produce an equivalent binary value, according to each digit's position. For example, the least significant hex digit is multiplied by 1, the second hex digit is multiplied by 10, the third hex digit is multiplied by 100, and so on, until all twenty-one hex digits of the integer part are 65 computed. Thus, at this stage, for the integer part, there will be twenty-one intermediate results that will eventually be added together by that half-system's adder block **8140** to

compute a total binary value for the integer part. The same is done by the fraction part half-system for the twenty hex digit intermediate values of the fraction part. Likewise, for the three hex-digit intermediate representation for the exponent, decimal multipliers **8135** and adder block **8150** compute the total value of the exponent of the default format input.

At this stage, the total value of the exponent is adjusted in a decimal exponent computation block **8170** so that the adjusted exponent result can be used as a table look-up index into integer part ROMs **8160** and fraction part ROMs **8180** of their respective half-systems. Rather than show a schematic for the decimal exponent computation block **8170**, the actual Verilog RTL source code used for performing this adjustment in the instant implementation is included here:

```
always @(*) //for adjusting integer part exponent for look-up
  if (intIsZero) decExpForLookUp = 0;
  else if (decExp_del_2 > 38) decExpForLookUp =
    decExp_del_2 + 39;
  else if (!intLeadZeroDigits_q2 && expIsMinus_q2)
    decExpForLookUp =
      (decExp_del_2 - 20) + (20 - intLeadZeroDigits_q2);
  else decExpForLookUp = (decExp_del_2 + 20) +
    (20 - intLeadZeroDigits_q2);
//for adjusting fraction part exponent for look-up
wire [8:0] fractExp_q2;
assign fractExp_q2 = (intLeadZeroDigits_q2 != 21) ? 9'b0 :
  (decExp_del_2 - 19);
```

The original exponent value input must be adjusted because the default input format exponent is referenced relative to the last decimal digit character position of the default character sequence input format and the exponent value needed for ROM look-up needs to be referenced relative to the least significant integer part decimal character digit. Thus, in the above example RTL source code, two adjusted exponents are computed: one for the integer part weight ROM look-up and one for the fraction part weight ROM look-up.

The integer part half-system **8400** also includes two look-up table ROMs **8160**. One ROM is a 67-bit by 309-entry ROM wherein each entry corresponds to the greatest weight for D52 (the hidden bit) of the previously computed integer part significand indexed by the previously-adjusted exponent value. The other look-up table ROM in **8160** comprises quantity 309-entry 11-bit unsigned biased binary exponent values corresponding to the same previously adjusted decimal exponent value used for the integer part look-up. These biased binary exponent values correspond to the greatest weight simultaneously indexed from the integer part greatest weight ROM. It should be noted here that, at this stage, twenty-one decimal digits require 67 bits to represent its value and that the 21st decimal digit of the integer part will never be greater than 1.

Block **8160** also comprises interpolation and other logic described in more detail below.

The fraction part half-system **8500** also includes four look-up table ROMs **8180** including various interpolation logic. Like the integer part ROMs **8160**, a previously-adjusted decimal exponent value is used as an index into the ROM **8180** for looking up the greatest weight corresponding to that adjusted decimal exponent as well as interpolation logic, described below. Also like the integer part **8160** block, the fraction part **8180** block includes a quantity 309-entry 11-bit unsigned biased binary exponent look-up table indexed by the same adjusted decimal exponent value used to index the fraction part weights ROM.

Finally, to enable support of IEEE 754-2008 mandated subnormal numbers, ROM block **8180** also includes a separate 67-bit, 17-entry weight ROM and a 6-bit, 17-entry “shift amount” ROM for the greatest weight corresponding to, and indexed by, the adjusted fraction part decimal exponent input. The subnormal weight has its own interpolation logic, explained later. In the case of subnormal fractions, since the binary exponent is always zero, the subnormal shift amount specifies the number of places the subnormal result is to be shifted right to denormalize the final result according to the adjusted decimal exponent input.

From ROM block **8160**, the interpolated 67-bit integer part weight for D52 (the hidden bit) of the integer part mantissa then enters the integer part quantizer/encoder block **8190**, while the integer part unsigned and unbiased binary exponent value from the integer part look-up table enters a first input of result exponent selector **8280**.

At the same time, from ROM block **8180**, the interpolated 67-bit fraction part weight for the MSB of the fraction part mantissa enters the fraction part quantizer/encoder block **8330**, while the fraction part unsigned and unbiased binary exponent value from the fraction part look-up table enters a second input of the exponent selector **8280**. It can be seen that if the original character sequence is fraction-only (i.e., the integer part is zero), then **8280** selects the fraction part exponent as the binary exponent for the final result. If the original character sequence comprises a non-zero integer part, then the integer part binary exponent is selected as the final result binary exponent.

Logic in circuit **8000** determines whether the original character sequence input is subnormal. If so, **isSubnormal 8200** is asserted to logic “1” to indicate that the input is subnormal. If the integer part of the character sequence input is zero then **integerIsZero 8210** is asserted to logic “1” to indicate that the integer part is zero. If input **8200** indicates a subnormal input or **8210** indicates the integer part is zero, the integer part weight and exponent ROM block **8160** outputs 0 for each. If the fraction part is zero, **8180** outputs zero for the fraction part weight and exponent.

To handle NaNs, converter **8250** converts the ASCII hex characters (“0” thru “F”) of the NaN’s payload to 4-bit equivalent hex values. These are, in turn, delayed by use of pipeline registers **8260** to maintain coherency for later use by a Universal IEEE 754 final formatter **8290**.

As the final stage, Universal IEEE 754 binary final formatter **8290** receives the computed intermediate 53-bit mantissa for the integer part from integer part half-system **8400**, which also supplies a IntegerInexact signal and three integer GRS bits (IntegerGRS[2:0]) to the final formatter **8290**. Likewise, fraction part half-system **8500** supplies a computed intermediate 53-bit mantissa, fraction inexact signal (FractionInexact) and three fraction GRS bits (FractionGRS[2:0]) to the final formatter **8290**.

Here it should be understood that if the integer part is not zero and the binary exponent from ROM block **8160** is within a range so that there could also be a fraction part, the integer part will always be exact. Stated another way, the integer part inexact signal can only be active in the range of exponents that have no fraction part. Similarly, if the integer part is zero, then the final result will be fraction-only (assuming the result is not zero, NaN, etc.). Thus, if the decimal character sequence contains a non-zero fraction, then the fraction inexact signal will be used; otherwise, the integer inexact signal will be used for producing the final inexact exception signal in signal group **8310**. DivideByZero in signal group **8310** will always be 0. The final output states of exception signals Invalid, Overflow,

and Underflow depend on what happened in the translation process of translator **8100**. Additionally, Overflow also depends on whether or not the directed rounding step in the final formatter **8290** caused the intermediate result of the final formatter to overflow.

The integer part GRS bits are used for directed rounding of the integer part. If the final result has a non-zero fraction part, the integer part GRS bits will always be {0, 0, 0}, meaning the integer part in such cases will always be exact and the fraction part GRS bits are used for directed rounding of the final result.

The final formatter **8290** uses inputs roundMode[1:0] **8220**, Away **8230** (if active), Size_Dest[1:0] and all the other inputs shown to do the final formatting to the IEEE 754 binary format specified by the size input **8240**. If size="01", then a binary16 format representation is output with forty-eight zeros padded in the most significant bits to create a 64-bit result **8300**. If size is ="10", then a binary32 format representation is output with thirty-two zeros padded in the most significant bits. If size="11" then a 64-bit binary64 format representation is output.

If the decimal character sequence input is integer-only, then IntegerPartOut[51:0] is used as the mantissa for the unrounded intermediate result. If the decimal character sequence input is fraction-only and not subnormal, then Fraction PartOut[51:0] is used as the mantissa for the unrounded intermediate result.

If the decimal character sequence input comprises both a non-zero integer part and a non-zero fraction part, i.e., the interpolated binary exponent of ROM block **8160** is less than 52, then n most-significant bits (not counting the hidden bit) of the integer part intermediate mantissa and 52-n most significant bits (counting the hidden bit) of the fraction part intermediate mantissa are concatenated to form the mantissa of the unrounded result. Here, n is the interpolated binary exponent value from ROM block **8160**.

For subnormal fraction-only numbers, there is no "hidden" bit, meaning bit **52** of the mantissa becomes visible and will occupy bit position **51** of the unrounded result for subnormal numbers having a decimal exponent of -308, meaning the fraction part mantissa is shifted right one bit position to appropriately denormalize it. For subnormal numbers having a decimal exponent of less than -308, e.g. -309, -310, etc., more shifts are required for proper denormalization of the unrounded result mantissa. The subnormal shift amount is supplied by the subnormal shift amount look-up ROM **8180**, which is indexed by the previously adjusted decimal exponent of the decimal exponent computation block **8170**.

FIG. 25 is a schematic diagram illustrating an exemplary embodiment of a circuit employed by the integer part quantizer/encoder **8190** of FIG. 24 to compute/encode the integer part intermediate mantissa using a first input **8490** and a second input **8390**. The circuit includes subtractors, comparators, and shifters arranged in series with the level that immediately follows the previous level.

With reference to FIGS. 24 and 25, the first level receives the first input (integerPartBin[67:0]) **8490**, which carries the binary value computed by the adder block **8140**. The second input (Weight_D52[67:0]) **8390**, is the weight received from the ROM block **8160**. This interpolated weight corresponds to the weight value for the most significant bit for the interpolated binary exponent computed by the decimal exponent computation block **8170**.

Next, the Weight_D52[67:0] is subtracted from the integerPartBin[67:0] with the result of such subtraction entering the B input of data selector **8340** while the integerPartBin

[67:0] enters the A input of the data selector **8340**. If it is determined at subtractor/comparator **8380** that the integerPartBin[67:0] is greater than or equal to the Weight_D52 [67:0], then the subtracted value propagates to the next level subtractor/comparator A input. Otherwise, the unsubtracted value of the previous subtractor/comparator A input propagates to the next level subtractor/comparator A input. First level comparator **8350** tests to see if the first level input binary value is not equal to 0. AND gate **8360** will yield a logic one for integer part intermediate mantissa bit D52 **8370** if the integerPartBin[67:0] is greater than or equal to the looked-up and interpolated Weight_D52[67:0], but only if the integerPartBin[67:0] is not zero.

This same process continues for the next sequential level 15 for determining the logic state of integer part intermediate mantissa bit D51 **8420**, except for the B input of that level's subtractor, the looked-up and interpolated weight **8390** is divided by two using a logic right shift of one bit. Thus, except for the divide by two of the B input for that level, the logic for determining the logic state of the corresponding intermediate mantissa bit is identical to the that of the level 20 that preceded it.

Note that for each next level B input to that level's subtractor, the number of shifts increases by one. Thus for 25 the level that determines the state of D50 **8430**, the B input to that level's subtractor, the weight **8390** must be divided by 4, which is two shifts to the right as shown in **8440**. This same logic, except for the number of shifts that increases by 30 1 for each subsequent level, continues down past DO **8450** and includes the same logic for determining the state of the integer part guard, round, and sticky bits **8460**, **8470**, and **8480**, respectively. The integer part guard, round, and sticky bits are used by the final formatter **8290** (FIG. 24) for 35 directed rounding. Note that the logic for determining the state of the sticky bit **8480** is slightly different than the logic in the levels that precede it, in that it includes an OR gate at the end that drives the sticky bit **8480**. The purpose of this OR gate is to ensure that if there is a remainder after all the subtractions that precede it, the sticky bit will reflect this 40 condition. Thus, if guard bit **8460**, round bit **8470**, or sticky bit **8480** are ever a logic "1", the integer part intermediate result is inexact.

FIG. 26 is a schematic diagram of an exemplary embodiment of a circuit employed by the fraction part quantizer/encoder **8330** of FIG. 24 to compute/encode the fraction part intermediate mantissa using a first input **8510** and a second input **8790**. The circuit includes subtractors, comparators, and shifters arranged in series with the level that immediately follows the previous level.

The illustrated circuit is identical to the integer part circuit shown in FIG. 25, except for the inputs. The first level receives the first input (fractPartBin[67:0]) **8510**, which carries the binary value computed by the fraction part adder block **8140**. The second input (Weight_D52[67:0]) **8790**, is the weight received from the ROM block **8180**. This interpolated weight corresponds to the weight value for the most significant bit for the interpolated binary exponent computed by the decimal exponent computation block **8170**. Like the integer part **8190** circuit (FIG. 25), the fraction part (FIG. 60 26) computes its own guard, round, and sticky bits.

FIGS. 25 and 26 and certain others in this disclosure do not explicitly show pipeline registers for storing intermediate results, but they are there in the circuit. For example, in both FIGS. 25 and 26, there is a rather large register that temporarily holds the intermediate results of the D52, D51, and D50 bits computation, the input weight, the fraction part binary input value, and the like. With respect to these first

three levels of logic, this could be called “stage 1” for that part of the circuit. A second rather large register would then temporarily hold the intermediate results of D52 (stage 1), D51 (stage 1), D50 (stage 1), D49, D48, D47 bits computation, etc., which could be called “stage 2”, and so on. Thus the pipelines of FIGS. 25 and 26 would be roughly 18 stages deep if intermediate results are registered every three levels/bits of computation. Registers are not explicitly shown here because an implementor will more likely than not want to re-time the operator to clock faster (by registering every two levels instead of three) or re-time the operator to clock slower by registering every four levels instead of three. Re-timing to clock faster will be at the expense of more registers. Re-timing to run slower, frees up registers. It should be understood that the main processor is just a push-pull shell whose pipeline is completely decoupled from the pipelines of any installed operators in its memory map. Thus the implementer is free to decide whether to re-time (and by how much) a given operator.

FIG. 27 is a schematic diagram illustrating an exemplary embodiment of the disclosed universal floating-point ISA’s convertFromDecimalCharacter operator’s look-up ROMs 8160 for the integer part, showing the interpolation method for determining the weights and binary exponents derived from a decimal exponent input obtained from the original decimal character sequence input. FIG. 27 shows, among other things, an excerpted 67-bit x 309-entry look-up ROM 8570 holding the greatest weight for D52 (hidden bit) of integer part intermediate mantissa that corresponds to the adjusted decimal exponent input, the unresolved weight 8580 output of which is fed into an interpolation circuit 8590 that interpolates a 67-bit integer part D52 weight based on a integer part binary value 8490 and a decimal exponent 8560.

Input 8540 is the previously adjusted 9-bit decimal exponent that is fed into a selector so that when the integer part being input is zero or the fraction part is subnormal as detected by OR gate 8550, the decimal value 9'D308 (8530) is substituted for decimalExponentIn 8540. In the weight look-up ROM 8570, the weight residing at address “[308]” is 0 (67'D000000000000000000000000000000).

If the decimal exponent input value 8560 is 9'D307, then the weight residing at location “[307]” in ROM 8570 (i.e., 67'D89884656743115795386) will be output on 8580. The weights output on 8580 are the greatest weight for integer part intermediate mantissa D52 and are unresolved at this point. These weights are unresolved at this point because the unbiased exponent 8560 used to look up the unresolved weight 8580 is decimal and, as such, there could be three or four binary exponents (and thus weights) that correspond to that same decimal exponent. To compute a correct weight corresponding to a decimal exponent input using ROM look-up table 8570, the instant invention employs an interpolation circuit 8590, which resolves a correct weight for the integer part by first comparing the integer part binary value 8490 with unresolved weight 8580. If the integer part binary value 8490 is greater than or equal to the unresolved weight 8580, then the unresolved weight 8580 becomes resolved and is selected as the integer part D52 weight 8390.

If the integer part binary value 8490 compares less than unresolved weight 8580, then unresolved weight 8580 is divided by two using a single right shift as shown in the interpolation circuit 8590. This first scaled weight is then compared with the integer part binary value 8490 to determine if the integer part binary value is greater than or equal to it. If so, the first scaled weight is selected to be output as 8390. If not, the integer part binary value 8490 is compared

with a second scaled weight, the second scaled weight being the unresolved weight 8580 divided by four using two right shifts. Another comparison is made using the second scaled weight and if the integer part binary value 8490 is greater than or equal to the second scaled weight, then the second scaled weight is selected as the 8390 output. Otherwise, a third scaled weight is selected as the 8390 output, this third scaled weight being unresolved weight 8580 divided by eight using three right shifts.

Look-up ROMs 8160 also include an 11-bit x 309 entry ROM 8630 containing the largest value biased binary exponent corresponding to the previously adjusted decimal exponent 8560, along with interpolation logic 8620 to determine the correct binary exponent 8610 according to the decimal exponent 8560 and the integer part binary value 8490. As shown, unresolved biased binary exponent 8640 is looked-up from entry ROM 8630 using the previously adjusted decimal exponent 8560 as an index simultaneously and in same manner as the unresolved weight 8580. Exponent interpolation circuit 8620 computes a resolved binary exponent by simultaneously subtracting 0, 1, 2, and 3 from the unresolved biased binary exponent 8640, the result of each subtraction going to a selector that selects the correct biased binary exponent 8610 using the results of comparisons from interpolation circuit 8590 as the input selector, such that unresolved exponent 8640 is resolved as the correct binary exponent 8610 and correctly corresponds to the correct integer part D52 weight 8390.

FIG. 28A is a schematic diagram of the disclosed processor’s convertFromDecimalCharacter operator’s look-up ROMs 8180 for the fraction part 8500, illustrating the interpolation method for determining the weights and binary exponents derived from a decimal exponent input obtained from the original decimal character sequence input. The fraction part ROM block 8180 is essentially the same in function as the integer part ROM block 8160, except for the actual contents of their respective ROMs and additional logic 8680 in the fraction part ROM block, which substitutes a 0 for the binary exponent resolved by interpolation logic 8710, such substitution occurring when the fraction part character sequence is zero or subnormal. Fraction part D52 weight output 8520 is resolved from unresolved weight 8650 by interpolation circuit 8700 in the same manner the integer part D52 weight is resolved. The fraction part’s greatest weight look-up ROM begins with the 67-bit (20-decimal-digit) value “50000000000000000000000000000000” corresponding to the biased decimal exponent of 0 used to index into ROM 8640. “50000000000000000000000000000000” is the greatest weight for the fraction part intermediate mantissa D52 when decimal exponent 8540 is 0 or 1. If the decimal exponent 8540 is 2, then the unresolved weight 8650 is “62500000000000000000000000000000” and so on. Likewise, the fraction part binary exponent ROM 8660 outputs an unresolved exponent 8670 based on the decimal exponent 8540 as index into ROM 8660. Fraction part exponent interpolation circuit 8710 resolves the correct binary exponent based on the fraction part binary value 8510 and the decimal exponent 8540 as inputs, with such being substituted with a zero if the fraction part character sequence input is zero or subnormal.

FIG. 28B is a schematic diagram illustrating an exemplary embodiment of the disclosed universal floating-point ISA’s convertFromDecimalCharacter operator’s look-up ROMs 8180 for the fraction part subnormal exponent inputs, illustrating the interpolation method for determining the weights and binary exponents derived from a decimal exponent input obtained from the original decimal character sequence input. IEEE 754-2008 mandates operability with

subnormal numbers. To satisfy this requirement, circuit **8800** includes two additional but much smaller ROMs (**8720** and **8740**), along with an interpolation circuit.

As shown, when the fraction part character sequence input is zero as indicated at **8770**, the fraction part weight output **8790** is also zero. If the fraction part character sequence input is subnormal as shown at **8760**, the interpolated subnormal weight **8750** is output as the fraction part weight **8790**, otherwise fraction part D52 weight **8520** is output at **8790**.

Because the “binary” exponent for all subnormal numbers is 0, a shift amount must be looked-up and interpolated in the same way as the unresolved exponent of 8670, i.e., using the decimal exponent **8540** as an index into shift amount ROM **8740**. In the case of subnormal numbers, the computed intermediate fraction part result must be denormalized by right-shifting the entire 53-bit fraction part mantissa by the number of bit positions determined by the decimal exponent **8540**. For example, for a decimal exponent of -308, the computed intermediate fraction part mantissa must be right-shifted one bit position to the right to denormalize it. For a decimal exponent of -309, it must be shifted right two bit positions, etc. If the fraction part character sequence input is not subnormal the subnormal shift amount **8780** will be 0, otherwise the interpolated shift amount from the shift amount ROM **8740** will be output at **8780**.

At this point, the formation of the final result should be straightforward: concatenate the sign of the decimal character sequence sign with the 11-bit binary exponent and final mantissa formed as previously described. For normal numbers, if the decimal character sequence input is fraction-only, then the interpolated fraction part binary exponent **8690** is used as the exponent, otherwise the integer part binary exponent **8610** is used as the exponent. If the decimal character sequence input is subnormal, then the binary exponent is zero and the final mantissa is denormalized by shifting it right the number of bits specified in the interpolated shift amount **8780**, then correctly rounded using the fraction part’s GRS bits described earlier.

ConvertToDecimalCharacter

FIG. 29 is a block diagram illustrating an exemplary embodiment of the disclosed universal floating-point ISA’s hardware implementation of a stand-alone, fully pipelined, memory-mapped H=20+ convertToDecimalCharacter component/module **9400** that performs the “convertToDecimalCharacter” operation mandated by IEEE 754-2008 (the standard), including 32-entry by 381-bit SRAM result buffer, restore capability, and ready semaphore. Features of the disclosed embodiment not mandated by the standard include universal binary format input and quantity (32) fully restorable result buffers into which results automatically spill.

A binary64 to H=20 decimal character sequence converter **9200** accepts as input, binary16, binary32, or binary64 formatted numbers and converts them into the default decimal character sequence format **300** (see FIG. 4) before being automatically stored into result buffer **9250**. As shown, result buffer **9250** may be a 32-entry by 381-bit, dual-port SRAM, which is wide enough to accept and/or read out a 47-character sequence and the five exception bits every clock cycle.

Like all the other operators whose pipeline depth is greater than one clock, it has a semaphore memory **9252** that produces a ready signal when the selected result is selected for reading. If the result being selected is not yet complete, then the ready signal for it will remain low “0” when being referenced, otherwise it will be high “1” if ready/complete.

FIG. 30 is a block diagram illustrating an exemplary embodiment of the binary to H=20 decimal character sequence converter **9200** of FIG. 29. The converter includes a universal IEEE 754 to binary64 converter **9210**, a fully pipelined binary-to-decimal-character conversion engine **9220**, NaN payload generator **9230**, and final formatter **9240** that formats the decimal character sequence intermediate result into the default format **300** and, if appropriate, asserts exceptions invalid, overflow, underflow, and/or inexact, with divX0 always remaining “0”.

The binary-to-decimal-character conversion engine **9220** produces a 21-character integer part, a 20-character fraction part, a 9-bit base-10 exponent, a fraction-only signal, a integer inexact signal, and a fraction inexact signal. The NaN payload generator **9230** simply propagates the payload of any NaNs entering the operator for later use by the final formatter **9240**. Note that the “input_is_ . . . ” signals from the universal IEEE 754 to binary64 converter **9210** are delayed so as to walk alongside the pipeline of the converter’s intermediate results so that they are presented to the final formatter **9240** simultaneously with the output of the binary-to-decimal-character conversion engine **9220**.

FIG. 31 is a block diagram illustrating an exemplary embodiment of the binary-to-decimal-character conversion engine **9220** of FIG. 30 illustrating virtually identical dual half-systems, one for the integer part **9340** and one for the fraction part **9350**. The binary-to-decimal-character conversion engine accepts as input, a 64-bit binary64 format number. The binary-to-decimal-character conversion engine also includes a binary-to-decimal exponent look-up ROM block **9330** that converts a 11-bit binary exponent to a 9-bit decimal exponent. To do the binary exponent to decimal exponent conversion, two ROMs are employed in the ROM block **9330**: one 2048-entry by 9-bit ROM for normal numbers and one 52-entry by 9-bit ROM for subnormal numbers. The value output from ROM block **9330** is a 9-bit base10 representation of the exponent. More detailed operation of the ROM block **9330** is described below.

Integer part half-system **9340** includes a binary-to-decimal-character summing circuit **9310** and a integer part 68-bit binary value to 20.2 digit BCD converter **9360** that converts the 68-bit binary value computed by the integer part summing circuit **9310** to a 4-bit by 20.2 digit BCD representation.

Converters **9360**, used by **9340** and **9350**, are identical. Bits

45 to 80 of bus “BCDintegerDigits[83:0]” are the most significant BCD digit, which will always be either a 4-bit “0000” or “0001”. For the fraction part converter **9360**, these bits are not used, hence bus “BCDfractionDigits[79:0]” is only 80 bits.

As shown, the fraction part binary-to-decimal-character summing circuit **9320** supplies the integer part summing circuit **9310** with the number of leading zeros by way of bus “subnLeadZeros[5:0], which the look-up ROM block of the integer part summing circuit **9310** uses to form a 52-bit fraction mask that it provides to the fraction part summing circuit **9320**. The function of this 52-bit fraction mask is described below.

Fraction part summing circuit **9320** also supplies a 6-bit “subnAddress[5:0]” to the binary-to-decimal exponent look-

up ROM block **9330**, which it uses to look up a decimal exponent value in the case of a subnormal number being converted.

The integer part binary-to-decimal-character summing circuit **9310** feeds a computed and correctly rounded 68-bit binary value “IntegerRoundedBinaryOut[67:0]” to the input of the integer part 68-bit binary value to 20.2 digit BCD converter **9360** for conversion to a 4-bit by 20.2-BCD-digit representation for the integer part.

The fraction part binary-to-decimal-character summing circuit **9320** feeds a computed and correctly rounded 68-bit binary value “FractRoundedBinaryOut[67:0]” to the input of the fraction part 68-bit binary value to 20.2 digit BCD converter **9360** for conversion to a 4-bit by 20-BCD-digit representation for the fraction part.

FIG. 32 is a block diagram illustrating an exemplary embodiment of the integer part binary-to-decimal-character summing circuit **9310** of FIG. 31, including integer part weights look-up ROM block **9314**, conditional summing circuit **9000**, and rounding circuit **9316**. ROM block **9314** includes a 1024-entry by 141-bit integer part weight look-up ROM containing the weight for D52 of the mantissa according to the binary exponent value for normal non-zero input numbers. ROM block **9314** also includes logic for forming an integer part mask and a fraction part mask. These masks are provided to each half-system’s conditional summing circuit **9000**. ROM block **9314** also provides a fraction-only signal to the final formatter **9240** (FIG. 30).

The 141-bit integer weight looked up by ROM block **9314** is split into two parts: a 67-bit value representing the 20 most significant decimal digits of the looked-up weight “IntWeight[66:0]” for integer part mantissa D52 (the hidden bit) and a 74-bit value representing the 22 least significant digits (truncated part) of such weight “IntWeightTrunc[73:0]”. These split weights and mask are used by the integer part half-system summing circuit **9000** to compute a sum “IntSumBinary[67:0]” and GRS bits “IntGRS[2:0]” representing the integer part that are in turn correctly rounded by the integer part rounding circuit **9316** according to the direction indicated by inputs “roundMode[1:0]” and “Away”.

Integer part rounding circuit **9316** produces an intermediate binary value “IntegerBinaryOut[67:0]” that gets converted by the integer part 68-bit binary value to 20.2 digit BCD converter **9360** to a 4-bit by 20.2 digit BCD representation. The integer part rounding circuit **9316** also supplies a integer inexact signal to final formatter **9240** (FIG. 30). Integer inexact signal is set “1” when, in the conditional summing circuit **9000**, any of the chaff GRS bits **9090** are “1” or if any of the truncated GRS bits **9149** are “1” (see FIGS. 34A and 34B).

FIG. 33 is a partial detail of an exemplary embodiment of the contents of the integer part binary-to-decimal-character mantissa D52 weight look-up ROM **9314** illustrating the first 20 decimal digits and the second 22 digits (truncated part of the weight), along with the actual Verilog RTL source code employed to obtain a mantissa mask used during the hardware computation. As can be seen, each of the 1024 entries comprise a total 141 bits, which are further broken down into the a 67-bit field representing the first (most significant) 20 decimal digits of that weight, and a 74-bit truncated field representing the second (least significant) 22 decimal digits of that weight. The unbiased and unsigned binary exponent of the binary64 format input is used as an index into the look-up integer part ROM.

FIGS. 34A and 34B together show the algorithm used by the disclosed processor to compute the integer part interme-

diate binary value according to the bits that are set in the integer part mask supplied by the integer part weights look-up ROM block **9314**. It should be noted that in the case of an integer-only number, the integer part mask will be identical to the mantissa of the binary64 format number and for the integer part, D52 (the hidden bit) will always be “1”. Thus, mantissa D52 does not have a corresponding mask bit, because it is always a “1” when the integer part is non-zero.

FIG. 34A is a diagram illustrating an exemplary embodiment of the method/algorithm **9000** used for computing both the integer part **9310** intermediate value and fraction part **9340** intermediate value that are submitted to their respective BCD converter circuits, including the method for obtaining a Guard, Round, and Sticky bit for each part, of the disclosed universal floating-point ISA’s double-precision IEEE 754-2008 H=20+ convertToDecimalCharacter operator.

It should be understood that for the computation in the conditional summing circuit **9000**, “mantissaD51” in block **9030**, “mantissaD50”, etc., means “masked-mantissaD51” and so on. In the case of numbers that have both a non-zero integer part and a non-zero fraction part, at least one of the LSBs of the mantissa will be masked off (as determined by the exponent input) by the integer part weights look-up ROM block **9314** so that their corresponding scaled weights are not added to the total sum of the integer part. This is because those bits that are masked off will be used by the fraction part computation.

Block **9030** shows the intermediate computation used for each of the bits “mantissaD51” thru “mantissaD0” to compute a scaled weight for that bit based on the looked-up input weight **9020** for mantissa D52. The looked-up input weight **9020** is the same as the ROM look-up for the first 20 digits **9010** showing that the 67-bit representation is the first (most significant) 20 decimal digits. Block **9030** specifies that if “mantissaD51” is a “1”, then divide the ROM look-up for the first 20 decimal digits **9010** by 2 with a single right shift to arrive at the scaled weight for that bit, otherwise use “0” as the weight for that bit. For “mantissaD50”, if that bit is “1”, then divide the ROM look-up for the first 20 decimal digits **9010** by 4 with two shifts to the right to arrive at the scaled weight for that bit, otherwise use “0” as the weight for that bit. The foregoing process continues until a scaled weight is computed for each bit down to “mantissaD0”.

Observe that with each shift for each intermediate computation in determining a scaled weight for each of the mantissa bits D51 thru DO (thereby effectuating a divide by 2 for each shifted bit position to determine the scaled weight corresponding that mantissa bit number), the least significant bits fall off the end. These are referred to as “chaff” bits **9040**. The chaff bits must be preserved with registers so they can be used later in computing an overall intermediate sum and intermediate chaff GRS bits **9090**.

Once all the corresponding weights for mantissa bits D51 thru DO have been computed, they are all added to the looked-up input weight **9020** to arrive at an intermediate sum **9050** for the integer part. At the same time, all the chaff from bits D51 thru DO intermediate computation are added together and the chaff carries **9060** out of the chaff sum D51 are added to the intermediate sum **9050**. Also at the same time, a single carry bit **9142** out of D73 of the sum of all truncated weight computations shown in **9130** is added to the intermediate sum **9050**. As shown, the three chaff GRS bits **9090** are added to the three second-22-digits GRS bits **9149**, with the single-bit GRS carry **9080** also being added to the intermediate sum **9050**, all the foregoing arriving at a second intermediate sum **9120** for the integer part. Finally,

GRS **9110** is used to correctly round the intermediate integer part result according to “round_mode[1:0]” and “Away” as shown in the rounding circuit **9316** (FIG. 32). Note that, although not explicitly shown, the LSB of GRS **9110** is “sticky”, meaning that if the sticky bit of the chaff GRS **9090** or the sticky bit of the three second-22-digits GRS bits **9149** is “1”, even after the two have been added together to arrive at GRS **9110**, the sticky bit of GRS **9110** should be a “1”. This can be accomplished by simply logically ORing the sticky bit of the chaff GRS **9090** with the sticky bit of the three second-22-digits GRS bits **9149** and using that output as the GRS sticky bit (LSB) for GRS **9110**. For some implementations requiring less precision, it may be more advantageous to omit the computation of the truncated part value and its single carry bit for each of the integer and fraction parts altogether. By doing so, the 22-decimal-digit (74-bit) truncated part weights can be omitted, making the weight look-up ROMs for the integer part and fraction parts much smaller. In such cases, only the chaff GRS bits are used for directed rounding.

FIG. 34B is a diagram illustrating an exemplary embodiment of the method/algorithm **9130** used for computing the sum of the truncated part (i.e., second 22 digits) used in the computation of both the integer part **9310** of half-system **9340** and fraction part **9350** of half-system **9350** of the disclosed universal floating-point ISA’s convertToDecimal-Character operator, including a method for deriving a truncated part GRS used in the final sum. FIG. 34B shows the intermediate computations **9130** involving the second 22 digits (truncated part) of the looked-up ROM weight. Note that the chaff bits from the shifts of the truncated weight are discarded. The sum of all the computed weights for each bit of the 74-bit truncated mantissa input is shown at **9140**. Note that the weight for mantissa D52 has a shift of zero, meaning it is not shifted, but it also is included in the sum **9140**. A carry bit **9142** is the carry bit out of bit 73 of the sum **9140**.

To determine the GRS bits of the truncated part intermediate sum, magnitude comparisons are made for each of them. Truncated part guard bit “Trunc_G” **9144** is determined by comparing the truncated part sum “TruncSum[73:0]” to twenty-two decimal digit value “500000000000000000000000”. If “TruncSum[73:0]” is greater than or equal to “500000000000000000000000”, then Trunc_G is set to “1”, otherwise it is cleared to “0”.

Truncated part round bit “Trunc_R” **9146** is determined by first subtracting twenty-two decimal digit value “500000000000000000000000” from “TruncSum[73:0]”, then comparing the result of that subtraction to twenty-two decimal digit value “050000000000000000000000”. If the result of that subtraction is greater than or equal to “050000000000000000000000”, then Trunc_R is set to “1”, otherwise it is cleared to “0”.

Truncated part sticky bit “Trunc_S” **9148** is determined by first subtracting twenty-two decimal digit value “550000000000000000000000” from “TruncSum[73:0]”, then comparing the result of that subtraction to twenty-two decimal digit value “000000000000000000000001”. If the result of that subtraction is greater than or equal to “000000000000000000000001”, then Trunc_S is set to “1”, otherwise it is cleared to “0”.

Finally, the 3-bit representation “Trunc_GRS” **9149** is formed by concatenating “Trunc_G”, “Trunc_R”, and “Trunc_S” as shown.

FIG. 35 is a schematic diagram **9316** illustrating an exemplary embodiment of an integer part rounding circuit that correctly rounds the integer part intermediate result **9120** according to the integer part GRS bits **9110**, sign,

Away, and roundMode[1:0] prior to submission of the intermediate result to the BCD conversion circuit **9360** (FIG. 31).

FIG. 36 is a block diagram illustrating an exemplary embodiment of the fraction part binary-to-decimal-character summing circuit **9320** of half-system **9350**, comprising fraction part weights look-up ROM block **9324**, conditional summing circuit **9000**, and rounding circuit **9326**.

ROM block **9324** includes a 1024-entry by 141-bit fraction part weight look-up ROM containing the weight for D52 of the mantissa according to the binary exponent value for normal input numbers that have a non-zero fraction part. Since the fraction part must also be able to handle subnormal numbers, ROM block **9324** also includes a 52-entry by 141 bit subnormal fraction part look-up ROM containing the weight for the most significant “set” bit of the mantissa, since subnormal numbers are denormal, i.e., D52 (the hidden bit) of the mantissa is never “1” for subnormal numbers.

FIG. 37 is a partial detail illustrating an exemplary embodiment of the convertToDecimalCharacter fraction part ROM weight look-up contents **9324** illustrating the first 20 decimal digits and the second 22 decimal digits (truncated part of the weight), along with the actual Verilog RTL source code employed to obtain a mantissa mask used during the hardware computation. The casex statement at the center of FIG. 37 loads 6-bit registers “subnAddrs” and “subnLeadZeros” with a value between “0” and “51” inclusive, based on the first non-zero bit in the mantissa. This is for special handling in the case of subnormal numbers. In the case of subnormal numbers, “subnAddrs” is used as an index into the subnormal look-up ROM (ROMB). If the input exponent is normal, i.e., not zero, then the weight from fraction part ROMA is output from the fraction part weights look-up ROM block **9324**; otherwise, if the input exponent is zero (indicating a subnormal number), then the weight from fraction part ROMB is output from the fraction part weights look-up ROM block **9324**. Note that in the case of subnormal numbers, “subnLeadZeros” is used by the integer part weights look-up ROM block **9314** of FIG. 32 to produce the fraction part’s fraction mask “fractMask[52:0]”, which in turn is used by the fraction part half-system conditional summing circuit **9000** to compute the fraction part binary sum/value.

The fraction part summing circuit **9000** is identical to the integer part summing circuit **9000**, except that the weight for “mantissaD52” **9020** is “0” if mantissaD52 is “0”, otherwise the weight for “mantissaD52” **9020** is the same as the ROM look-up for the first 20 digits **9010**. Like the integer part half-system **9340**, the intermediate result **9120** computed by fraction part half-system **9350** summing circuit **9000** is correctly rounded by fraction part rounding circuit **9326** to arrive at “FractRoundedBinaryOut[67:0]”, which is then input into fraction part 20-digit BCD converter **9360** of the fraction part half-system **9350** to produce a 4-bit by 20-BCD-digit sequence representing the fraction part as shown in FIG. 31. Like the integer part computational block **9310** of integer part half-system **9340**, the fraction part computational block **9320** of fraction part half-system **9350** produces a fraction inexact signal that is asserted to “1” when the fraction part is inexact.

FIG. 38 is a partial detail illustrating an exemplary embodiment of the look-up ROM block **9330** and actual Verilog RTL source code used by the convertToDecimal-Character operator for converting the adjusted binary exponent input to an adjusted decimal exponent for both normal and subnormal numbers. Included in the RTL are normal exponent look-up table ROM (RAMA) **9334** and subnormal

exponent look-up table ROM (RAMB) 9336. Note that RAMA 9334 uses the unbiased binary exponent input as the index to the adjusted decimal exponent being looked up. Also note that the decimal exponents of RAMA 9334 corresponding to binary exponents with binary values in the range 2046 to 1087 are pre-adjusted by subtracting 19 from them. This is to satisfy the default binary-to-decimal-character format on the output, in that the default format references the decimal place relative to the last digit position as shown in FIG. 4. Also observe in RAMA 9334 that for binary exponents in the range 1086 to 1023, the adjusted decimal exponent is 0, for the same reason. Further, binary exponents in the range 1022 to 1, their looked-up decimal values are negative, meaning that the final formatter 9240 will place a “e-” instead of “e+” character string in front of the decimal exponent characters in the final result.

FIG. 39 is a schematic diagram illustrating an exemplary embodiment of the fraction part rounding circuit 9326 that correctly rounds the fraction part intermediate result 9120 according to the fraction part GRS bits 9110, sign, Away and roundMode[1:0] prior to submission of the intermediate result to the BCD conversion circuit used in the disclosed processor's convertToDecimalCharacter hardware operator 9360 (FIG. 31).

FIGS. 40A, 40B, and 40C are block diagrams that together show, respectively, the upper left-most, lower right-most and lower left-most sections of the fully pipelined binary-to-binary-coded-decimal (BCD) conversion block 9360 used by integer part half-system 9340 and fraction part half-system 9350 to convert their respective rounded 68-bit binary outputs to BCD. The BCD conversion block 9360 accepts as input, a computed 68-bit binary value and converts it to a 4-bit by 20.2-digit BCD representation. Note that for the integer part, the most significant 4-bit BCD digit of the output BCD representation will always be either a “0001” or “0000” binary. For the fraction part, the most significant 4-bit BCD digit of the output BCD representation will always be “0000” binary.

In the disclosed exemplary embodiment, the BCD conversion block 9360 consists entirely of 4-bit-in-4-bit-out look-up tables 9362 that emulate a 4-bit BCD shift-add operation. FIGS. 40A, 40B, and 40C show how to arrange the 4-bit-in-4-bit-out look-up tables to create the overall 68-bit binary-to-20.2-decimal-character-digit BCD conversion. The disclosed BCD circuit includes 65 rows of look-up tables, meaning it is 65 logic levels deep. Thus, every eight logic levels (rows) the outputs are registered, such that it requires eight clocks to perform the BCD conversion. FIG. 40B indicates by a dotted line 9364 that inputs D17 thru D67 continue across and/or repetitions of 9362 continue across the required number of positions so as to accommodate a 68-bit input to BCD converter 9360 starting from D0 and continuing to D67.

Once the BCD conversion for both integer part and fraction part half-systems are complete, those intermediate results, along with the looked-up decimal exponent, respective inexact signals, and fraction-only signal are passed the final formatter 9240, as shown in FIG. 30. The final formatter 9240 converts the 4-bit BCD digits to characters by simply adding to each 4-bit BCD digit, the 8-bit ASCII hex value for character “0”, which is 30 hex. Thus, BCD 0 becomes 30 h, 1 becomes 31 h, and so forth. The looked-up decimal value exponent is converted to BCD using a much smaller version of the BCD conversion block 9360 comprising just a few look-up tables 9362 sufficient for a 9-bit input and 3 decimal digit BCD (12-bit) output, which in turn are also converted to characters.

Finally, logic within the final formatter 9240 evaluates all the inputs into it and concatenates the character sign of the binary input, the integer part decimal character sequence, the fraction part decimal character sequence, an “e-” or “e+” character sequence (depending on sign of the exponent) and the converted decimal character sequence exponent, forming the 376-bit, correctly rounded, final decimal character sequence result in the default format shown in 300 (FIG. 4). At the same time, the five exception signals, divX0, invalid, overflow, underflow and inexact are set or cleared depending on exceptions that may have arisen during the conversion process.

FIG. 41 is a block diagram illustrating an exemplary embodiment of a memory-mapped, fully restoreable, hardware-implemented, double-precision floating-point “addition” operator module 9420, including 16-entry by 69-bit SRAM result buffer 9426, restore capability, and ready semaphore 9428. The addition operator module 9420 includes a floating-point adder 9422 that accepts 76-bit inputs for both operands A and B. The 76-bit operand A and operand B inputs are the results of their respective upstream “universal input converter” shown in FIG. 22A, which converts binary16, binary32 or binary64 format numbers to binary76 format used by the floating-point adder 9422 and provides interim exception signals 9432 if any occur during the process of the upstream conversion. Because binary16, binary32 and binary64 formats convert “up” to the larger binary76 format, such conversions will always be exact.

This binary76 format comprises a sign bit as the MSB, a 12-bit exponent and a 62-bit fraction (FP1262). The use of larger binary format number in this and various other hardware floating-point operators is to support computations using subnormal numbers, in that the larger format floating-point adder 9422 will never overflow or underflow as a result of the computation. If overflow or underflow occur at all, it occurs as a result of conversion from FP1262, the output format of the floating-point adder 9422.

The result that gets stored in the SRAM result buffer 9426 is 69 bits. This is because the five MSBs are the five exception signals divX0, invalid, overflow, underflow, and inexact. If the target result size specified by the Size_Dest [1:0] input is binary16 or binary32, the most significant bits (not including the exception bits) are padded with zeros by a binary format converter 9424. The number of levels of delay in multi-tap delay registers 9430 is determined by the combined number of clocks the floating-point adder 9422 and the binary format converter 9424 require to complete. Like all the hardware floating-point operators of the floating-point operator module 3010 shown in FIG. 22A, the addition operator module 9420 includes a semaphore block 9428 to signal that the computation is complete when being accessed.

It should be understood that some applications may not need or want a “universal” converter to automatically convert the FP1262 intermediate result of the floating-point adder 9422 to binary16 or binary32 formats, but may prefer instead to do so with a separate and explicit conversion operation once a final result is obtained and stored in the SRAM result buffer 9426. In such cases, the binary format converter 9424 can be substituted with a FP1262 to binary64 (FP1152) converter. Furthermore, some implementors may prefer employing a multiplier that is FP1252 (12-bit exponent with 52-bit fraction) format as opposed to the FP1262 format shown for the floating-point adder 9422, in that an FP1252 will be less expensive to implement than the FP1262 version. In so doing, implementors should take care

to provide faithful GRS bits to the binary format converter **9424** for correct rounding purposes.

FIG. 42 is a block diagram illustrating an exemplary embodiment of a memory-mapped, stand-alone, fully restoreable, hardware-implemented, double-precision floating-point “multiplication” operator module **9440** including 16-entry by 69-bit SRAM result buffer **9446**, restore capability, and ready semaphore **9448**. A floating-point multiplier **9442** accepts 76-bit inputs for both operands A and B. The 76-bit operand A and operand B inputs are the results of their respective upstream “universal input converter” shown in FIG. 22A, which converts binary16, binary32 or binary64 format numbers to binary76 format used by the floating-point multiplier **9442** and provides interim exception signals **9445** if any occur during the process of the upstream conversion. Because binary16, binary32 and binary64 formats convert “up” to the larger binary76 format, such conversions will always be exact.

This binary76 format comprises a sign bit as the MSB, a 12-bit exponent and a 62-bit fraction (FP1262). The use of this larger binary format number in this and various other hardware floating-point operators is to support computations using subnormal numbers, in that the larger format floating point multiplier **9442** will never overflow or underflow as a result of the computation due to the extra bit in the exponent. If overflow or underflow occur at all, it occurs as a result of conversion from FP1262, the output format of the floating-point multiplier **9422**.

The result that gets stored in the SRAM result buffer **9446** is 69 bits. This is because the five MSBs are the five exception signals divX0, invalid, overflow, underflow and inexact. If the target result size specified by the Size_Dest [1:0] input is binary16 or binary32, the most significant bits (not including the exception bits) are padded with zeros by a binary format converter **9447**. The number of levels of delay in multi-tap delay registers **9444** is determined by the combined number of clocks the floating-point multiplier **9442** and the binary format converter **9447** require to complete. Like all the hardware floating-point operators of the floating-point operator module **3010** shown in FIG. 22A, the multiplication operator module **9440** includes a semaphore block **9448** to signal that the computation is complete when being accessed.

It should be understood that some applications may not need or want a “universal” converter to automatically convert the FP1262 intermediate result of the floating-point multiplier **9442** to binary16 or binary32 formats, but may prefer instead to do so with a separate and explicit conversion operation once a final result is obtained and stored in the SRAM result buffer **9446**. In such cases, the binary format converter **9447** can be substituted with a FP1262 to binary64 (FP1152) converter. Furthermore, some implementers may prefer employing a multiplier that is FP1252 (12-bit exponent with 52-bit fraction) format as opposed to the FP1262 format shown for the floating-point multiplier **9442**, in that an FP1252 will be a little bit cheaper to implement than the FP1262 version. In so doing, implementors should take care to provide faithful GRS bits to the binary format converter **9447** for correct rounding purposes.

FIG. 43 is a block diagram of an exemplary embodiment of a memory-mapped, fully restoreable, stand-alone, double-precision (H=20), hardware-implemented double-precision floating-point “H=20” convertFromDecimalCharacter operator module **9460**, including 32-entry by 69-bit SRAM result buffer **9462**, restore capability, and ready semaphore **9464**. The convertFromDecimalCharacter operator module **9460** includes the convertFromDecimalCharacter circuit

8000 (FIG. 24), 32-entry SRAM result buffer **9462**, semaphore circuit **9464**, multi-tap delay block **9466** for delaying the write enable and write address the required number of clocks, and optional binary64 to universal binary output block **9468**.

The convertFromDecimalCharacter circuit **8000** accepts a 376-bit (47-character sequence) input that is fed into optional Universal DCS Translator **500** (FIG. 5) that translates non-default character sequence inputs **400** to the 47-character default character sequence shown **300** (FIG. 4). The Universal DCS Translator may be omitted if the end application will never need to compute directly with non-default character sequences. In other words, if the application will always perform translation in software, then the Universal DCS Translator circuit can be omitted such that the input is fed directly into the decimalCharToBinary64 circuit of the convertFromDecimalCharacter circuit **8000**. The output of the decimalCharToBinary64 circuit is then fed into the binary64 to binary16, binary32 or binary64 converter **9468** that converts the output to binary16, binary32, or binary64 format, depending on the two-bit Size_Dest[1:0] input. Note that if size is specified for 8 bytes, then there is no actual conversion. Furthermore, if the end application only requires binary64 as the target binary format, then the binary64 to binary16, binary32 or binary64 converter **9468** may be omitted. Finally, the output of the binary64 to binary16, binary32 or binary64 converter **9468** is supplied to one input of a selector whose output enters the SRAM result buffer **9462** for storage at the same address it was input on the front end, meaning that when the character sequence was originally pushed into the convertFromDecimalCharacter operator module **9460** for conversion, the address it was pushed into, including its wren state, is delayed in the multi-tap delay block **9466** the same number of clocks as the pipeline of the convertFromDecimalCharacter circuit **8000**, such that by the time the character string is converted, the delayed write address and write enable are presented to the SRAM result buffer **9462** simultaneously with the output of the selector.

Note that the convertFromDecimalCharacter operator module **9460**, like all the other floating-point operators shown in block **3010** (FIG. 22A), is fully restoreable. This means that results and corresponding exception signals of conversions can be pulled out and pushed into the fat SRAM **9520** of dual asymmetric stack **9500** shown in FIG. 44 for preservation during subroutine calls and interrupt service routines and then pulled from the fat SRAM and pushed directly back into their originating result buffers. This is so the convertToDecimalCharacter operator can be used by the called routine or interrupt service routine without corrupting previously computed results or having to re-compute them upon return. Also note that the five exception signals that were simultaneously pulled from the result buffer are pushed into the signal stack **9510** simultaneously with the push of the converted result into fat SRAM **9520** and pulled simultaneously with same during the restore. All floating-point operators of block **3010** and integer arithmetic and logic operators of block **4000** (FIG. 23), including their respective exception/condition signals, may be preserved and restored in this manner, so long as the destination address is specified by the SP for a stack “push” and SP specifies the source address for a stack “pop” within the address range of the fat dual asymmetric stack **9500**.

Note further, that the SRAM result buffer **9462** is 32 entries deep. The main reason why 32-entry memory was chosen over a 16-entry operator is that the pipeline of the convertFromDecimalCharacter circuit **8000** is roughly 28

levels deep. Thus, when using the REPEAT instruction to convert a block of 32 decimal character sequences to binary format, by the time the 32nd decimal character sequence is pushed into the convertFromDecimalCharacter operator module 9460 for conversion, the results can immediately be pulled out with the REPEAT instruction without incurring a single stall in the process, thereby hiding latency, “as if” the conversion for each decimal character sequence actually only consumed 0 clocks to complete, making it virtually free.

FIG. 44 is a block diagram of an exemplary embodiment of a multi-function dual asymmetric “fat” stack and “fat” SRAM block 9500 used for operator context save-restore operations and other general-purpose functions. Among the problems associated with implementing a opcodeless processor (wherein its memory-mapped floating-point, integer arithmetic and logical operators are completely decoupled from the processor’s instruction pipeline, and wherein these operators have their own block of SRAM that results automatically spill into), is that there needs to be a way to save and restore results that may already be stored in one or more of a given operator’s memory-mapped result buffers. The main reason for this is so that interrupt service routines and called subroutines can use these operators without corrupting these results that may already be residing inside the result buffer the service routine or subroutine needs to use.

One problem is that the five exception signals of a given floating-point operator and the four condition signals of integer arithmetic and logical operators need to be preserved simultaneously with its actual computed result. If the computed result is 32 bits wide, there is no room to store these signals within the standard-sized word. If the result is 64 bits, there is still no room. This is because the processor can only read or write in groups of 8, 16, 32, 64, 128, 256, 512 or 1024 bits at a time. It would be possible to store a 32-bit result, along with its 5-bit exception signals, in a 64-bit location, but this is exceedingly inefficient in that almost half the bits are wasted. For example, when a 512-bit result, along with its 5-bit exception, needs to be saved, this would imply that 507 bits of the 1024-bit write would be wasted just to preserve the exception signals. Furthermore, the logic to handle all possibilities of sizes would be exceedingly complicated.

The circuit shown in FIG. 44 solves this problem. When a result is pulled from its result buffer and pushed into the address range of where this block is mapped in the processor’s data RAM using the stack pointer (SP) as the destination address, the computed result is stored in fat SRAM 9520 at the byte address specified in SP and the exception signals are simultaneously stored in the separate signal stack SRAM 9510 at the address specified in SP.

It should be understood that for 8-bit results, such may be pushed anywhere within the fat SRAM 9520 because it is one-byte aligned, meaning it can be written anywhere within the SRAM. A 16-bit result being pushed into the fat SRAM 9520 must be aligned on 2-byte address boundaries; a 32-bit result must be aligned on 4-byte address boundaries; and so forth.

Note that the signal stack SRAM 9510 is only five bits wide, which is just wide enough to accommodate the 5-bit exception signals of floating-point operators and 4-bit condition signals of integer arithmetic and logical operators. If multiple single 8-bit bytes are pushed into the fat SRAM 9520 in consecutive locations, so will their respective signals be pushed into the signal stack SRAM 9510 into consecutive locations. But suppose there are multiple pushes of 8-byte (64-bit) results into consecutive locations on

proper boundaries of the fat SRAM 9520. In such case, the exception signals will be simultaneously pushed into the signal stack SRAM 9510 every eight locations, leaving the locations between them undisturbed. This is because, when performing the push of the 64-bit result into the fat SRAM 9520, the address for the next push must be incremented by eight, because the fat SRAM 9520 is byte-addressable.

When performing the push of a result into the fat SRAM 9520, the user need not be concerned with the signal stack SRAM 9510, in that the simultaneous push of the exception or condition code into the signal stack SRAM 9510 is automatic.

It should be understood that the fat SRAM 9520 is randomly accessible and can be used as ordinary SRAM if desired. This block of SRAM is useful for storage of decimal character sequences because it can be written to 128 characters at a time. Because it has two read-side ports, it can be read up to 256 characters at a time, 128 characters from side A and 128 characters from side B. This comes in handy when computing directly with decimal character sequences where the target operator, such as the instant invention’s Universal Fused-Multiply-Add operator allows two decimal character sequence operands as input.

While the circuit 9500 can be used as general-purpose SRAM for storage decimal character sequences as described above, it’s other main purpose is for use as a stack that is wide enough to save and restore results of operators such as convertToDecimalCharacter, convertToHexCharacter, the instant invention’s Universal FMA, and certain future tensor operators that have “gobs” of data as a result comprising multiple parallel results from multiple parallel operators that can be read and written in a single clock.

As described earlier, contents of result buffers, along with their exception/condition code signals, are pushed into the dual asymmetric stack made up of the signal stack SRAM 9510 and the fat SRAM 9520 using the SP as the destination, such that when the result is written into the fat SRAM 9520, its corresponding exception/condition code signals are automatically and simultaneously written into the signal stack SRAM 9510 at the location specified in SP.

To restore the saved result and its corresponding exception/condition code back into the same or different operator result buffer, the SP is used as the source address, such address abiding in the aligned boundary requirement mentioned previously. When the SP is used as the source address and such address fall within the range of where the circuit 9500 is mapped, not only will the result buffer be restored, but also its corresponding exception/condition code signals simultaneously along with it.

At the top of FIG. 44 is shown the Verilog source code used to create the special logic that detects when the SP is used as either the source or destination address, wherein the address falls within the range of where the circuit 9500 is mapped in the processor’s data memory.

As mentioned previously, the fat SRAM 9520 can be used as data 1, 2, 4, 8, 16, 32, 64, 128-byte-readable/writeable data RAM, provided that such reads and writes occur on boundaries evenly divisible by the number of bytes being read or written. As such, debuggers can easily modify the contents of the computed results of a given result buffer by first pushing such results into the dual asymmetric stack memory of circuit 9500.

Since the signal stack SRAM 9510 shadows the fat SRAM 9520 within the same address range, one problem arises from such configuration is, how to modify the exception/condition code signals saved in the signal stack SRAM 9510 after an operator result buffer save operation. Stated

another way, it is desirable for debuggers to easily access and modify the contents of the signal stack SRAM **9510** during debugging operations. The circuit **9500** solves this problem with the use of the SigD bit **170** of the instruction. When overwriting the contents of a given location within the signal stack SRAM **9510** with a new value, simply set the SigD bit **170** by use of the “s” character for the destination field **120** of the instruction **100**. Thus, with SigD=“1”, any write to an address within the address range of circuit **9500** will be written to the signal stack SRAM **9510** instead of the fat SRAM **9520**. Likewise, if it is desired to randomly read the contents of a location within the signal stack SRAM **9510**, simply precede the size field **180** of destination field **120** with a “s” character, thereby setting destination **120** SigD bit to a “1”, such that when that address is accessed, the contents of the signal stack SRAM **9510** at that address is driven onto the five LSBs of the A-side read bus instead of the fat SRAM **9520**. Thus, the SigD bit **170** of the destination **120** field is used to select from which memory (the fat SRAM **9520** or the signal stack SRAM **9510**) data will be read or written outside of context save/restore operations.

It should be further understood that the size (address range) of circuit **9500** can be increased as much as the application requires by substituting the 64 k-byte version of circuit **9500** for the optional data SRAM block **950** in FIG. **9**. If an even larger version is required, such can be accomplished by mapping a 128 k-byte version beginning at data memory location **0x00020000**, and so forth.

FIG. **45** is a block diagram of an exemplary embodiment of a memory-mapped, fully restorable, stand-alone, double-precision floating-point “fusedMultiplyAdd” (FMA) operator module **9700**, which is designed to also operate as a sum-of-products operator. The FMA operator module **9700** includes a 16-entry by 69-bit SRAM result buffer **9710**, restore capability, and ready semaphore, presented here without universal decimal character front end. The FMA operator module **9700** also includes a floating-point multiplier **9750** that accepts 76-bit inputs for both operands A and B. The 76-bit operand A and operand B inputs are the results of their respective upstream “universal input converter” shown in FIG. **22A**, which converts binary16, binary32, or binary64 format numbers to binary76 format used by a floating-point multiplier **9750** and provides interim exception signals seen entering FP1262-to-universal-IEEE-754 converter **9780** during the process of the upstream conversion. Because binary16, binary32, and binary64 formats convert “up” to the larger FP1262 format, such conversions will always be exact.

The binary76 format comprises a sign bit as the MSB, a 12-bit exponent and a 62-bit fraction (FP1262). The purpose of the 12-bit exponent (as opposed to 11-bit exponent of binary64 format numbers) is to support computation with subnormal numbers, as mandated by IEEE 754-2008, when the FMA operator module **9700** is employed as a pure FMA. The extra 10 bits in the fraction part is to help absorb underflows when the FMA operator module **9700** is employed as a sum-of-products operator in extended series computations involving small numbers.

The FMA operator module **9700** also includes quantity (16) 81-bit C-register/accumulators **9720**, one for each of this operator’s inputs. Note that the five MSBs of a given C-register/accumulator are the exception signal bits divX0, invalid, overflow, underflow, and inexact. The lower 76 bits are the FP1262 representation that was automatically converted upstream by the same universal IEEE 754-to-FP1262 converter (see FIG. **22A**) used for converting operandA. For

initialization to positive zero, a single one-byte write may be used, since the upper bits of a zero value, no matter the size, is automatically zero-extended during the write operation.

The C-registers’ addresses are offset from their respective 5 FMA inputs by 16 locations. Refer to table **3010** of FIG. **12A** for the input addresses of the FMA input buffers and their respective C-register/accumulators. Before computing with the FMA operator module **9700** as a pure FMA, the C-register corresponding to the desired FMA input buffer number 10 must be initialized with a value for operandC of the R=(A*B)+C equation. Once the C-register(s) corresponding to the FMA input buffers that are to be used have been initialized, the FMA operation can be performed by simply pushing operandA and operandB (using a dual-operand push) into the corresponding FMA input buffer location.

Once operandA and operandB are simultaneously pushed into the FMA operator module **9700**, the FP1262 floating-point multiplier **9750** multiplies them together, producing an un-rounded FP1262 result comprising a sign bit, 12-bit 20 exponent, 62-bit fraction and the three GRS bits. For a FP1262 multiplier, this intermediate result can usually be accomplished in a single clock, thus after the first clock, the intermediate result is fed into a FP1262 floating-point adder **9760**, which in turn computes the sum of the product of the floating-point multiplier **9750** just computed and the current 25 value of its corresponding C-register/accumulator in C-register bank **9720**. Depending on the design of the floating-point adder **9760**, the addition operation can be achieved in one to three or four clocks for an FP1260 adder. The output 30 of the adder then enters data selector **9730** and universal IEEE 754 converter **9780**. At this point, the intermediate result from the floating-point adder **9760** remains un-rounded. The other input to data selector **9730** is only used for initializing and restoring a given C-register in 35 C-register bank **9720**. Thus, at all other times, the intermediate result from the floating-point adder **9760** enters data selector **9730** so that it will automatically update the selected C-register/accumulator of C-register bank **9720** for use in the next FMA computation involving it, hence the C-register 40 now becomes an accumulator.

Stated another way, for pure FMA operations, the C-register is always initialized as operandC before operandA and operandB are simultaneously pushed into the corresponding FMA input buffer and, hence, it remains a C-register for holding a operandC value for use in a single R=(A*B)+C 45 computation. For use as a sum-of-products operator, the corresponding C-register is initialized only once, usually with a 0, followed by a series of pushes of operandA and operandB. It can be seen that every time the floating-point adder **9760** produces an intermediate, un-rounded result, the pertinent C-register in C-register bank **9720** is automatically updated with such results and thus becomes an accumulator, the contents of which are used in the next sum-of-products iteration.

Simultaneous with the C-register/accumulator being updated, the intermediate FP1262 results of the floating-point adder **9760** are automatically converted to either binary16, binary32, or binary64 format by the universal IEEE 754 converter **9780** as specified by the Dest_Size[1:0] 50 input, such also being correctly rounded according to the rounding direction specified by inputs roundMode[1:0] and Away. The converted and correctly rounded final result then enters data selector **9740** before being automatically written, along with its five exception bits, into 16-entry SRAM result buffer **9710**. Note that for binary16 and binary32 final result 55 formats, the MSBs are automatically zero extended to 64 bits, with bits D68 thru D64 being occupied by the five

exception signals, for a total of 69 bits. It should be noted here that the other input of the data selector **9740** is for use in result buffer restore operations so that results in the SRAM result buffer **9710** do not have to be re-computed during context save/restore operations typical of subroutine calls or interrupt service routines. Like the other floating-point operators shown in block **3010** of FIG. 22A, the FMA operator module **9700** also comprises a SRAM-based semaphore circuit for generating a ready signal to indicate when a given result is ready when being read by the processor.

Also like the other operators in block **3010** (FIG. 22A), it should be understood that some applications may not need or want a “universal” converter to automatically convert the FP1262 intermediate result of the universal IEEE 754 converter **9780** to binary16 or binary32 formats, but may prefer instead to do so with a separate and explicit conversion operation once a final result is obtained and stored in the SRAM result buffer **9710**. In such cases, the universal IEEE 754 converter **9780** can be substituted with a FP1262 to binary64 (FP1152) converter.

An important consideration regarding the stand-alone FMA operator module **9700** is the fact that if the floating-point adder **9760** takes more than one clock to complete, the time between successive pushes to the same FMA input buffer must be delayed by the number of extra clocks beyond the first. This is because floating-point adder **9760** will not have otherwise completed the computation of the previous result in time to be available for use as an operandC. Thus, if only one FMA input is to be used for successive FMA computations, as might be the case when the FMA operator module **9700** is used as sum-of-products of products operator, NOPs or some other instruction must be executed before subsequent pushes into the same FMA input buffer location. Stated another way, do not use the REPEAT instruction to perform multiple successive pushes to the same input buffer location, because the result of the floating-point adder **9760** will not have time to complete before the next push.

For sum-of-products operations involving multiple vectors, this should not be a problem, because, when employing the REPEAT instruction, the destination address in an Auxiliary Register can be made to automatically index into the adjacent FMA input buffer location. Thus, up to 16 vectors of any length can be computed rapidly using the REPEAT instruction inside a hardware loop counter loop to index into C-register bank **9720** using indirect addressing mode with auto-post-modification, such that, when the 16th buffer is reached (i.e., the REPEAT counter has reached 0), a branch back to the beginning of the loop is taken and the REPEAT counter re-initialized to 15 to begin the next iteration. Stated another way, by the time the next sum-of-products iteration begins, intermediate results from the floating-point adder **9760** will be available for use as operandC in that iteration. Some examples are provided below:

```

demoS:   _ 8:C.0 = _8:PI
         _ 8:FMA.0 = (_8:workA, _8:workB)
         _ 8:C.0 = _1:#0 //this will NOT work for above reasons
         _ 4:AR1 = _4:#vectA
         _ 4:AR2 = _4:#vectA + 8
         _ 4:AR3 = _4:#FMA.0
         _ 2:LPCNT0 = _2:#46
         _ 8:*AR3++[1] = (_8:*AR1++[16], _8:*AR2++[16])
         _ 4:PCS = (_2:LPCNT0, 16, loopU)

```

The first example above, demoS, is a simple FMA operation using direct addressing mode. Location PI contains a binary64 representation for the value of PI. WorkA contains the binary64 value for operandA, and workB contains the

binary64 value for operandB. As can be seen, C-register at location C.0 is initialized with the contents of location PI, followed by a dual-operand push of the contents of location workA and workB.

- 5 The second example above, demoT, shows what NOT to do using the FMA operator module **9700** as a sum-of-products operator. First, C.0 is initialized with an immediate 0. Then AR1 is initialized with the immediate pointer to the first binary64 format operandA in the vector. AR2 is initialized with the immediate pointer to the first binary64 format operandB in the vector. AR3 is initialized with the immediate pointer to the first FMA input buffer location, FMA.0. Next, the REPEAT counter is loaded with the immediate value **45**, causing the next instruction to be executed 46 times. Note that the destination address in AR3 never changes because it is being incremented by zero, because the original intent was to do a sum-of-products on a single vector. The problem with this instruction sequence is that after the second fetch of the instruction following the
- 10 REPEAT instruction, the execution of the first fetch has not yet completed, because the floating-point adder **9760** requires more than one clock to complete and, consequently, its addition result is not yet available for use as the next operandC in the equation, R=(A*B)+C, on the very next clock.
- 15 REPEAT instruction, the execution of the first fetch has not yet completed, because the floating-point adder **9760** requires more than one clock to complete and, consequently, its addition result is not yet available for use as the next operandC in the equation, R=(A*B)+C, on the very next clock.
- 20 REPEAT instruction, the execution of the first fetch has not yet completed, because the floating-point adder **9760** requires more than one clock to complete and, consequently, its addition result is not yet available for use as the next operandC in the equation, R=(A*B)+C, on the very next clock.
- 25 REPEAT instruction, the execution of the first fetch has not yet completed, because the floating-point adder **9760** requires more than one clock to complete and, consequently, its addition result is not yet available for use as the next operandC in the equation, R=(A*B)+C, on the very next clock.

To employ the present disclosure’s stand-alone, hardware FMA operator module **9700** as a sum-of-products operator, where the push of the two operands is always to the same input buffer location, use one of the processor’s hardware loop counters in a software loop instead of the REPEAT instruction. Doing so gives time for the floating-point adder **9760** to complete before the next pair of operands are pushed into that operator. Here is an example routine that uses the FMA operator module **9700** as a sum-of-products operator where the push of the two operands is always to the same input buffer location:

```

demoU:   _ 8:C.0 = _1:#0
         _ 4:AR1 = _4:#vectA
         _ 4:AR2 = _4:#vectA + 8
         _ 4:AR3 = _4:#FMA.0
         _ 2:LPCNT0 = _2:#46
loopU:   _ 8:*AR3++[1] = (_8:*AR1++[16], _8:*AR2++[16])
         _ 4:PCS = (_2:LPCNT0, 16, loopU)

```

- 40 In demoU above, instead of loading the REPEAT counter with an immediate **45**, the hardware loop counter LPCNT0 is loaded with an immediate **46**. The loop comprising just two instructions is executed 46 times and intermediate results are accumulated in FMA.0 C-register/accumulator (C.0) after each iteration, resulting in a final, correctly rounded result being automatically stored in FMA result buffer, FMA.0, which can be pulled out of either the A side or B side of the SRAM result buffer **9710**.
- 45 The example loopU works for use with a single FMA input buffer because the fetch of the PCS instruction will not execute until two clocks later, for a total of three clocks, due to the processor’s three-stage instruction pipeline. The above assumes that the floating-point adder **9760** pipeline is no more than two clocks deep because it takes one more clock to actually write the 81-bit intermediate result into the specified C-register/accumulator of C-register bank **9720**. If a given implementation’s floating-point adder **9760** requires more than two clocks to complete, then NOPs should be inserted just before the branch as shown in the following example, which is for a 4-clock adder circuit. Note that a “_” character standing alone (except for optional label) on the

- 50 character standing alone (except for optional label) on the
- 55 character standing alone (except for optional label) on the
- 60 character standing alone (except for optional label) on the
- 65 character standing alone (except for optional label) on the

assembly line effectuates a NOP because it assembles into machine code 0x0000000000000000, which, when executed, reads data RAM location 0x00000000 from both side A and side B and writes it back into location 0x00000000, further noting that location 0x00000000 in data RAM is always read as 0, no matter what value was previously stored there.

```
loopUa:    _8:*AR3++[1] = (_8:*AR1++[16], _8:*AR2++[16])
           ...
           ...
           _4:PCS = (_2:LPCNT0, 16, loopUa)
```

It should be understood that a correctly rounded intermediate result is automatically stored in the SRAM result buffer **9710** at the completion of each iteration and can be read anytime. It should be further understood that simultaneous with the automatic storage of a correctly rounded result in the SRAM result buffer **9710** at completion of each iteration, so too is a un-rounded, 81-bit intermediate result stored in the corresponding C-register/accumulator (C.0 in this case) in C-register bank **9720** upon completion of each iteration, which can be pulled out of the A side any time.

One application for using the hardware FMA operator module **9700** as a sum-of-products operator is a very efficient computation of multi-vector tensors. In the following example, both the hardware loop counter and REPEAT counter are used. In this example, it is ok to use the REPEAT counter because the destination FMA input buffer changes with each push/clock. In the hardware FMA operator module **9700**, logic and delay registers correctly and automatically distribute for storage in C-register bank **9720**, the intermediate result of the floating-point adder **9760**, as well as a correctly rounded final result into result buffer SRAM **9710**, all corresponding to their original push locations. In the following example, there are quantity (16) double-precision, dual operand vectors of quantity (32) entries each. Each vector has its own FMA input location. For instance, vector0 operand pairs are pushed into FMA.0, vector1 operand pairs are pushed into FMA.1 and so on, all the way up to FMA.15 for vector16.

```
demoV:    _4:AR5 = _4:#C.0
           _2:REPEAT = _2:#15
           _8:*AR5++[1] = _2:#0
           ...
           _4:AR3 = _4:#vectB
           _4:AR2 = _4:#vectB + 8
           _2:LPCNT0 = _2:#32
loopV:    _4:AR0 = _4:#FMA.0
           _2:REPEAT = _2:#15
           p _8:*AR0++[1] = (_8:*AR3++[16], _8:*AR2++[16])
           _4:PCS = (_2:LPCNT0, 16, loopV)
```

The first three instructions of demoV above initialize all 16 C-registers in C-register bank **9720** to 0. Next, Auxiliary Registers AR3 and AR2 are initialized with the pointers to the first operand A and first operandB of the first vector, wherein, vector data is arranged sequentially in SRAM such that the first operandA and the first operandB of the second vector follow the first, and continues in a round-robin fashion, such that the second operandA of the first vector immediately follows the first operandB of the 16th vector in SRAM. The data is arranged in SRAM this way so that the pointers into the vectors need be initialized only one time.

Next, hardware loop counter LPCNT0 is loaded with the immediate value **32**, which is the number of operand pairs in each of the 16 vectors. LoopV is the entry point for the

software loop, wherein, on each entry, Auxiliary Register is loaded with address of the first FMA input buffer, which corresponds to the first vector sum-of-products. Once AR0 is loaded, then the REPEAT counter is loaded with the number of additional times the next instruction is to execute. In this case, the instruction following the REPEAT instruction will be executed a total of 16 times. Note that the mnemonic is the “p” character, which signals the universal IEEE 754 converter **9780** to correctly round each final result towards positive infinity before storing it into the SRAM result buffer **9710**. The PCS instruction is a conditional load of the PC. In this case, bit **16** (the “not_zero” flag) of LPCNT0 is tested to see if it is not zero. If not zero, the PC performs a relative branch to loopV. If zero, then the PC continues incrementing (exits the loop).

Like all the other operators of blocks **3010** (FIG. 22A) and **4000** (FIG. 23), the result buffers of the FMA operator module **9700** and C-register/accumulators of C-register bank **9720** are fully restoreable. It should be noted that the C-register/accumulators are 81 bits wide, thus when pulling and pushing them from/to the asymmetric memory stack **9500** (FIG. 44) or the fat SRAM **9520** or optional data SRAM **950** (FIG. 9), a size of 16 (128 bits) is to be used for the source and destination sizes. When pushed into the dual asymmetric stack using SP as the destination address, the corresponding exception bits of the C-register/accumulator will automatically be pushed into the signal stack **9510** of the asymmetric memory stack **9500**. When pulled using SP as the source address with the FMA operator module **9700** as the destination, the exception bits will be automatically and simultaneously restored from the signal stack **9510**.

FIG. 46A is a high-level block diagram of an exemplary embodiment of the present disclosure's memory-mapped, stand-alone, fully restoreable, multi-function, Universal Fused-Multiply-Add (FMA) (and accumulate) operator module **9800**, including dual convertFromDecimalCharacter converters on the input.

The multi-function universal FMA operator module **9800** is multi-function because it can be employed as an FMA operator, a sum-of-products operator, or a dual IEEE 754-2008 “H=20” convertFromDecimalCharacter sequence operator. The multi-function universal FMA operator module **9800** is universal because it can directly accept as input binary16, binary32, binary64 format numbers, and/or decimal character sequence numbers up to 28 decimal digits in length in any combination, without first having to explicitly convert them to binary beforehand and can output correctly rounded results in binary16, binary32, or binary64 format according to the destination size specified in the size field **180** of instruction **100** (FIG. 1), again without having to explicitly convert the result with an separate operation.

Operands wrdataA[375:0] (operandA) and wrdataB[375:0] (operandB) can each be up to 376 bits wide and each may propagate to an FMA operator **9816** via one of several paths in the pipeline as determined by inputs {SigD_q2, SigA_q2 and SigB_q2}, collectively, and in that order (“SIGNALS”). Referring to instruction format **100** of FIG. 1, SigD_q2 is the size **180** of the DEST **120** delayed by two clocks, SigA_q2 is the size **180** of the srcA **130** delayed by two clocks and SigB_q2 is the size **180** of the srcB **140** delayed by two clocks. The purpose of the 2-clock delay is to align those inputs with the execution cycle of the processor, which happens during Stage q2 (see FIG. 6).

At this point, it should be understood that inputs such as SigA_q2, SigB_q2, wraddr[5:0], wren, etc., must be delayed internally by the multi-function universal FMA operator module **9800** the appropriate number of clocks to

maintain coherency with its overall pipeline and, as such, some of these will have more than one tap because some are needed at various stages within the pipeline.

When SIGNALS is “000” at the time of the push of operandA and operandB into the multi-function universal FMA operator module **9800**, the operands both follow the route that goes thru the universal IEEE 754-to-binary64 converters **9802a** and **9802b**, respectively, that convert the binary format numbers from the format specified by their respective size input to binary64 format. Since the range of binary16 and binary32 numbers are smaller format than binary64, such conversions are always exact. Thus, except for maybe signaling NaN, no exceptions are produced at this point.

The output of converters **9802a** and **9802b** then enter their respective delay blocks **9804a** and **9804b**, respectively. The number of sequential registers in the delay block must equal the number of stages in the IEEE 754 decimalCharToBinary64 converters of **8000a** and **8000b** so that pipeline coherency is maintained in cases where input formats are a mix of binary format and decimal character sequence formats, or in cases where the previous or subsequent push into **9800** contains a decimal character format number. The outputs of **9804a** and **9804b** then enter the B input of data selector **9806a** and **9806b** then the input of binary64-to-FP1262 converter **9810a** and **9810b**, respectively, in that the FMA circuit in **9816** only operates on FP1262 binary format numbers, which have a sign bit, 12-bit exponent and 62-bit fraction. FP1262 formatted numbers have one more bit in the exponent and ten more bits in the fraction than binary64 format numbers. The purpose of the extra bit in the exponent is to prevent underflow and overflow in single iteration computations. The purpose of ten extra bits in the fraction is to help absorb underflows during relatively long series sum-of-products computations involving small numbers.

The output of binary64-to-FP1262 converters **9810a** and **9810b** then enter the B input of data selector **9812a** and **9812b**, respectively, then the B input of data selector **9814a** and **9814b**, respectively, finally arriving at their respective inputs to FMA block **9816** (FIG. 46B), where they are operandA and operandB of the computation to be performed by the FMA block **9816**.

When SIGNALS is “100” this means that both operands are binary format, as in the case above for when SIGNALS is “000”, except delay blocks **9804a** and **9804b** are to be bypassed, thereby substantially reducing the overall pipeline latency of **9800** by roughly 28 clocks in the instant implementation. Beyond this point in the **9800** pipeline, the operands follow the same path as for when SIGNALS is “000” as described above.

This mode can be used in cases where **9800** is being employed as a pure FMA or sum-of-products operator and no decimal character sequences are involved at any stage in the overall computation, for example, a single FMA computation or a sum-of-products computation involving shallow vectors. Otherwise, when employing **9800** as a FMA for a single computation involving binary format operands, the processor will stall for period of time roughly equal to the pipeline depth of **8000** when attempting to immediately pull a result from **9800**’s result buffer.

When SIGNALS is “010”, “001” or “011”, the operand input with a logic “1” for its respective SigA_q2 and/or SigB_q2 input is treated as a decimal character sequence and thus propagates thru converter **8000a** and/or **8000b** where it is first translated to the default decimal character input format **300** if in a format shown in **400** of FIG. 5. Once translated to the default character sequence input format, it

is then converted to binary64 format by the IEEE 754 decimalCharacterToBinary64 converter **8000a** and/or **8000b**, respectively, and then, from this point forward, follows the same path as described above when SIGNALS is “000”.

When SIGNALS is “111”, “110” or “101”, either one or both operands are decimal character sequences, but no FMA computations are involved. Rather, the multi-function universal FMA operator module **9800** is to be used as solely as a single or dual IEEE 754-2008 H=20 convertFromDecimalCharacter converter. Thus, in this mode, the FMA circuit **9820** of FIG. 46B is bypassed and the result of the conversion by converters **8000a** and/or **8000b** is automatically stored in their respective result buffer **9818** and/or **9830**, respectively. The main idea here is, considering the fact that the conversion circuit **8000** employed by converters **8000a** and **8000b** are rather large, in many implementations it may not be feasible to incorporate yet a third conversion circuit **8000** just to handle routine convertFromDecimalCharacter operations. With just a little bit more logic, the multi-function universal FMA operator module **9800** is now able to be employed not only as a universal pure FMA operator and/or sum-of-products operator, but also as a single or dual IEEE 754-2008 H=20 convertFromDecimalCharacter operator. In this mode, when used in combination with the REPEAT instruction, the multi-function universal FMA operator module **9800** can effectively convert quantity (64) decimal character sequences, each up to 28 decimal digits in length, in roughly 64 clocks, which includes the clocks required to push them into, and pull them out of, the multi-function universal FMA operator module **9800**. Stated another way, the apparent latency to convert quantity (64) decimal character sequences to binary is zero, meaning it is free.

FIG. 46B is a block diagram illustrating an exemplary embodiment of the FMA (and accumulate) circuit **9820** and “split” SRAM block **9818** and **9830** for storage of either dual convertFromDecimalCharacter results, FMA or sum-of-products results, readable on side A and side B. The split SRAM block includes side-A data selector **9824**, side-A SRAM **9818**, side-B data selector **9828**, side-B SRAM **9830**, and a mode-based delay block **9832**, which together form the middle and back-end of the multi-function universal FMA operator module **9800**. The SRAMs **9818** and **9830** form a “split” SRAM block for storage of either dual convertFromDecimalCharacter results, FMA results, or sum-of-products results, readable on side A and side B.

The FMA circuit **9820** is bypassed by operandA and operandB when SIGNALS is “111”, which configures the multi-function universal FMA operator module **9800** as a dual convertFromDecimalCharacter converter. In this configuration, the outputs of data selectors **9814a** and **9814b** (i.e., operandA **9822** and operandB **9826**, respectively) are routed directly to the write data inputs of A-side result buffer SRAM **9818** and B-side result buffer SRAM **9830** via the A inputs of data selectors **9824** and **9828**, respectively. Note that unlike the previously described mandated computational floating-point operators of block **3010** (FIG. 22A), which each employ a simple three-port (one write side and two read side) SRAM as a result buffer, the multi-function universal FMA operator module **9800** employs two true dual-port SRAM blocks arranged as a true four-port SRAM to function as a dual result buffer for simultaneous storage of two results when the multi-function universal FMA operator module **9800** is configured to operate as a dual convertFromDecimalCharacter converter. At all other times, the result buffer of FIG. 46B is configured to function as a

three-port result buffer, in that in those cases, there is only one result coming from FMA circuit **9820** via the B inputs of data selectors **9824** and **9828** that is to be stored.

An advantage of being able to convert two decimal character sequences simultaneously, is that both binary format results can be simultaneously pulled out and then immediately pushed back into any dual-operand operator as operandA and operandB. For example, in the disclosed embodiment, only the Universal FMA operator can directly accept decimal character sequences as operandA and operandB. However, an application may need to divide one vector by another, but the data in the vectors are all decimal character sequence format. With the multi-function universal FMA operator module **9800** configured for use as a dual convertFromDecimalCharacter converter, and assuming there are at least 28 elements in each vector, using the REPEAT instruction the decimal character sequence data of both vectors can be simultaneously pushed into the multi-function universal FMA operator module **9800** and the converted results of both can be immediately pushed into the floating-point division operator, again using the REPEAT instruction. Thus, in this example, the cost or apparent latency for converting two decimal character sequences is only one clock per operand pair, which is the clock needed to push an operator pair into the multi-function universal FMA operator module **9800** in the first place, thus the apparent latency to actually convert them is zero, meaning it is free.

Just like all the floating-point operators of block **3010**, the multi-function universal FMA operator module **9800** is fully restoreable, meaning that result buffers **9818** and **9830**, in addition to all 32 of the FMA C-register/accumulator bank **9844** (shown in FIG. 46C) must be randomly readable and writable to support save/restore operations, such as during a subroutine call or interrupt service routine. Thus, data selector **9834** enables selection between result buffer **9818** data on data selector **9834** input B and the 81-bit C-register/accumulator contents specified by rdaddrs[4:0] when rdaddrs[5] and rden are both “1”, such that, when result buffer **9818** is selected, the selected data is driven onto A-side read data bus rddataA[80:0] (with its data bits [80:64] padded with zeros), and when an element in the FMA C-register/accumulator bank **9844** is selected, all 81 bits of the specified C-register/accumulator are driven to the A-side data bus. The format of the C-register/accumulators are the same as that previously described for the C-register bank **9720** (FIG. 45).

The mode-based, multi-tap delay block **9832** does what its name suggests. Because the multi-function universal FMA operator module **9800** is universal and multi-function, each operational mode has its own number of delays for required inputs such as wren (write enable), wraddrs (write address), etc., which must arrive at various places at the right time in the pipeline, depending on the mode of operation.

FIG. 46C is a block diagram illustrating an exemplary embodiment of the FMA circuit employed by the multi-function universal FMA operator module to perform both pure FMA and sum-of-products computations. The FMA circuit is presented here without universal decimal character front end, and is employed in the disclosed universal floating-point ISA's multi-function, universal FMA (& accumulate) operator.

Note that with the exception of the number of C-register accumulators in C-register bank **9844**, the function of the FMA circuit **9820** is identical to the respective computation portions of the FMA operator module **9700** of FIG. 45. For more information on the operation of the FMA circuit **9820**,

refer to the section of the instant specification that pertains to the FMA operator module **9700**.

FIG. 47A is a block diagram of an exemplary embodiment of the present disclosure's optional hardware JTAG-accessible, breakpoint, trace buffer, and real-time monitor/debug module **7000** that enables on-the-fly, real-time-data-exchange operations between the parent CPU **600** and up to quantity (16) child XCs **800** attached to it within the same device. In addition to real-time-data-exchange operations, the hardware debug module **7000** provides a PC discontinuity trace buffer **7200** and hardware breakpoint module **7300** that accepts as inputs debug control bits and PC comparison values for use as breakpoint triggers, along with their respective enables **7050** (collectively).

PC discontinuity trace buffer **7200** captures all bits of the implemented program counter (PC) anytime there is a discontinuity, meaning, anytime the PC is loaded, such load is detected and the PC value just before and after the discontinuity are captured and stored in the first-in, last-out queue of the PC discontinuity trace buffer **7200**, such that the exit and entry addresses of the PC into program memory are recorded and visible to the host workstation by way of update/capture register block **7500**, with such being shifted out via shift register **7130** to the host JTAG interface connected to IEEE 1149.1 JTAG state machine Test Access Port (TAP) **7600**. The PC discontinuity trace buffer **7200** may record two PC discontinuity events, where each event comprises the PC value just before the discontinuity (the exit point) and the PC value immediately after the discontinuity (the entry point). Thus, the PC trace record comprises PC Trace Oldest (the previous PC exit point in history), PC Trace 2 (the previous PC entry point in history), PC Trace 1 (the most recent PC exit point in history) and PC Trace Newest (the most recent PC entry point in history) **7020** (collectively).

The main idea behind tracing only PC discontinuities instead of every PC value in history is to have at least a minimal processor program execution recordation capability for debugging and monitoring purposes without the expense of implementing a full trace buffer, understanding that the host debug station can reconstruct execution path in the software based on just the PC entry points and exit points and compiled program listing file. Note that **7020** can be sampled by the update/capture register block **7500**, which is then loaded into shift register **7130** and scanned out thru via JTAG state machine TAP **7600** transmit data out (TDO).

Various methods and operations involved with employing an IEEE 1149.1 TAP for basic microprocessor debugging operations is well understood in the industry, including its use with a PC discontinuity trace buffer and certain aspects of implementing hardware breakpoints by way of update/capture registers **7500** and JTAG state machine TAP **7600** that provides an update data register clock (UDRCK), update data shift register enable (UDRSHE), update data register enable (UDRUPD), update register strobe (URSTB) **7140** (collectively) and a 8-bit JTAG instruction register output **7150** that contains the JTAG instruction corresponding to the operation that is to be performed by the update/capture register block **7500** and the shift register **7130**.

Among the functions that the hardware debug module **7000** can carry out include, resetting the parent CPU (which also resets any child XCs attached to it when FORCE_RESET **7100** is active or global external RESET_IN **7080** is active), force hardware breakpoint on the parent CPU, single-stepping the CPU after it has encountered a software or hardware breakpoint, programming up to two PC comparison addresses and the type of comparison to be made and

enabling same, event detection 7320 based on such comparison and event counting. Further, the hardware debug module 7000 provides a means for a host debug station to obtain status information such as whether the CPU is in a reset state, whether the CPU has encountered a software or hardware breakpoint and whether the CPU has completed a single-step command, which are all common and well known debug functions that an IEEE 1149.1 interface can be used for.

What sets the CPU's hardware debug module 7000 apart from prior art debug and monitoring capability is that the hardware debug module 7000 gives a host debug and development workstation the ability to monitor and exchange data in real-time with the target CPU (and any XCU included in the same or other JTAG scan chain) without the use of specialized processor opcodes, interrupts and their service routines, debug micro-kernel or DMA cycles.

As described above, the present disclosure does not employ opcodes of any kind in the instruction set. Instead, all processor registers and operators, including operators that carry out real-time monitoring and data exchange operations, are memory mapped and reside at specific locations designated for them in the processor's data memory space.

FIG. 47B is a block diagram illustrating an exemplary embodiment of a breakpoint, single-step and real-time monitor/debug module 7300. With reference to both FIGS. 47A and 47B, to carry out real-time data exchange operations between a host debug station and the target processor via a JTAG interface, the present invention makes use of dedicated "hot-spots" at specific locations in the processor's memory map. For performing a real-time monitor data READ operation, wherein the host debug station desires to read a specific location out of the target processor, the host debug station first scans into the update/capture register block 7500 via JTAG state machine TAP 7600 the monitor read address 7330, monitor write address 7350, and the size 7410 for the read access cycle and the size 7420 of write access cycle, for example "000" for one byte, "001" for two bytes, "010" for four bytes or "011" for eight bytes, and so on. Next, via the JTAG state machine TAP 7600, the host debug station initiates a real-time data exchange operation in the target processor by scanning a "1" into the monitor request bit (bit 3 of the break control register in the update/capture register block 7500), which brings mon_req 7170 to a logic "1" at completion of the JTAG update data register state. Once set in this manner, logic AND gate 7450 detects it, causing state machine 7460 (FIG. 47B) to register a single-operand monitor instruction 7470 assembled on-the-fly by simply concatenating buses mon_write_size[2:0] 7420, mon_write_address[14:0] 7350, mon_read_size[2:0] 7410, and mon_read_address[14:0] 7330, ind_mon_read 7310 and ind_mon_write 7340 into 64-bit register monitor_instructionq 7465 as shown in the following example:

```
wire [63:0] monitor_instruction={5'b0000000, mon_
write_size[2:0], ind_mon_write, mon_write_address[14:0],
1'b0, mon_read_size[2:0], ind_mon_read, mon_read_ad-
dress[14:0], 20'b000000000000000000000000}.
```

Note that the above example monitor_instruction 7470 (and thus monitor_instructionq 7465) conforms to processor instruction format 100 in FIG. 1. What makes this a monitor instruction is the fact that it is registered in 64-bit register monitor_instructionq 7465 and is substituted, by way of data selector 7370 when break_q0 7304 is active, for program instruction 100. Notice that break_q0 7304 can be triggered to active "1" by inputs "broke", "force_break" or "event_det" 7320 when input "skip" is inactive "0", and is

also triggered by "mon_cycl_det" 7450 as a result of a monitor cycle request "mon_req" 7170 when "mon_state[1:0]" is "00".

Referring to state machine 7460 of FIG. 47B, during state "00", if "mon_cycl_det" goes active "1", the assembled 64-bit monitor_instruction 7470 is clocked into monitor_instructionq 7465 and simultaneously triggers one-shot mon_req_q0 to indicate this happened and is presently only used for simulation purposes. Next, state machine 7460 enters state "01" which registers an instruction equivalent to a NOP instruction in register 7465 on the rising edge of the very next clock, ensuring the processor executes the previous assembled monitor instruction only once. Note that state-machine 7460 remains in state "10" until "mon_req" 7170 is brought to inactive "0" low by the update/capture register block 7500 as a result of it being cleared by way of JTAG scan initiated by the host debug environment.

To perform a monitor read operation, the desired monitor read address is scanned into the mon_read_addrs data register, along with the size of the data to be read, in the update/capture register block 7500 via JTAG by the host debug environment. The debug environment then scans into the mon_write_addr data register of the update/capture register block 7500 the value 0x00007FEB, which is the address of the Real-Time Monitor Data register (see Table 2000 of FIG. 11). The value "011" binary is scanned in as the size of the data to be written into the Real-Time Monitor Data register, even though this might not match the size of the data that is read during the monitor cycle, because the size of the Real-Time Monitor Data register in the present disclosure has a fixed width of 64-bits. The host debug station then makes a monitor request by setting the mon_req bit in the update/capture register block 7500 to a "1".

When break_q0 7304 goes active as a result of mon_req 7170 going active, data selector 7370 drives monitor_instructionq 7465 onto processor instruction bus 7360 and thus enters the processor's instruction pipeline for processing as if it had been an actual instruction fetch from program memory. Upon execution, the processor's pipeline will have read the specified number of bytes from the srcA address specified in instruction_q0_del 7360 and written it to the Real-Time Monitor Data register (FIG. 11) at location 0x0007FEB, as specified in the assembled monitor instruction. The host workstation then performs a JTAG capture of the Real-Time Monitor Data register (mon_read_reg[63:0] 7120) and scans out the captured data via JTAG state machine TAP 7600. As a final step, the host debug station then scans a "0" into the mon_req bit of the update/capture register block 7500 to remove the monitor request, causing state machine 7460 go back to state "00" and wait for another monitor request.

As can be seen from state machine 7460, monitor_instructionq 7465 is driven onto Instruction_q0_del bus 7360 for only one clock cycle. During this one clock cycle the PC is frozen and the pre-PC is rewound to the previous fetch address so it can re-fetch the instruction that was being fetched when the monitor request was made (refer to FIG. 15 and FIG. 16).

To perform a monitor write operation, the data to be written is scanned into the 64-bit monitor_write_reg in the update/capture register block 7500 via JTAG state machine TAP 7600 by the host debug station. Next, the monitor write address along with the size of the data to be written is scanned into their respective registers in the update/capture register block 7500. Next, the value 0x0007FEB (which is the address in data memory where the contents of monitor_write_register can be read by the processor), along

with the size “011” binary, is scanned into the monitor read address and monitor read size registers in the update/capture register block **7500**. Once all the foregoing elements of the monitor write instruction are scanned in, the host workstation scans a “1” into the mon_req bit of the update/capture register block **7500**. When the processor executes the assembled monitor instruction, the processor will have read the contents of the monitor_write_reg in **7110** (which is visible to the processor at location 0x0007FEB in its data memory space) and written this data to the location specified in the assembled monitor instruction. Again, as a final step, the host debug station then scans a “0” into the mon_req bit of the update/capture register block **7500** to remove the monitor request, causing state machine **7460** go back to state “00” and wait for another monitor request.

Observe that the read and write address fields in the instruction **100** format can accommodate only 15 bits of direct address and thus only the lower 15 bits of the monitor read address **7330** and monitor write address **7350** scanned into the update/capture register block **7500** are used in the assembled monitor instruction, mon_write_addrs[14:0] for DEST **120** direct address field **210** and mon_read_addrs[14:0] for srcA **140** direct address field **210** (instruction **100** bits **34** thru **20**). Also notice that ind_mon_write **7340** and ind_mon_read **7310** are set to “1” if any bit in the upper 17 bits of mon_write_addrs or mon_read_addrs (respectively) is a “1”. These two outputs are used as the IND **190** bits for their respective fields in **100** during assembly of the monitor instruction **7470**.

Further notice in FIG. 14C, that if ind_mon_read **7310** is a “1”, meaning that current read cycle is indirect address mode, then mon_read_addrs[31:0] **7330** will be driven onto SrcA_addrs_q0 bus **6200**. Likewise, if ind_mon_write **7340** is a “1”, then mon_write_addrs[31:0] **7350** will be driven onto Dest_addrs_q0 bus **6230**. Thus, monitor read or monitor write cycles with addresses less than 0x00008000, are direct access mode, and monitor read or monitor write cycles with addresses greater than 0x00007FFF are indirect access mode. Furthermore, notice that due to the arrangement of their respective address bus selectors in FIG. 14C, monitor read and monitor write operations have priority over ordinary processor accesses.

State-machine **7380** in FIG. 47B handles processor hardware breakpoints and single-stepping. If a pre-programmed event is detected, as indicated by event_det output **7320**, or if a “1” is scanned into the force break bit (bit 1) of the break control register in the update/capture register block **7500**, the “broke” bit will be registered as a “1”, indicating that the processor is now in a break state and all instruction fetching has been suspended, as a result of the PC now being frozen. Once “broke” has been set to one (1), state-machine **7380** enters state “01” where it waits for the host debug station to scan a “1” into the single-step (“sstep”) bit (bit 2) of the break control register in the update/capture register block **7500**. When a “1” is scanned into the sstep bit, register “skip” is set to “1” for one clock cycle, causing the break state to be released for exactly one clock cycle, as shown in logic block **7302**. This allows the processor to fetch and execute just one instruction. Note that the REPEAT counter circuit **1900** in FIG. 19, pre-PC **610** and PC **1620** remain frozen while break_q0 is active. Thus, when “skip” goes to “1” for exactly one clock cycle, break_q0 **7304** goes inactive for exactly one clock cycle, releasing the break state for exactly one clock cycle, effectuating a single-step of the processor. When the single-step takes place, pre-PC **610** and PC **1620** are allowed to increment, just once, thereby

advancing the fetch address to the next instruction, and REPEAT, if not zero, is allowed to decrement just once.

When state-machine **7380** enters state “10”, register “skip_cmplt” is set to “1”, indicating that the processor has completed its single-step as commanded by the host. The state of “skip_cmplt” can be captured by the break status register (bit 5) in the update/capture register block **7500** and scanned out by the host debug station. State-machine **7380** remains in state “10” until the host debug station scans a “0” into the “sstep” bit of the break control register in the update/capture register block **7500** to indicate that it acknowledges skip complete and to advance to the next state, state “11”.

While in state “11”, state-machine **7380** checks to see if a force break has been asserted. If not, state-machine **7380** returns to state “00” to begin searching for another break event or another force break. If force break (“frc_brk”) is active “1”, then state-machine **7380** returns to state “01” to begin looking for another single-step command issued by the host debug station.

It should be understood that the CPU need not be at a breakpoint to perform real-time-data-exchange operations, as these can be performed on-the-fly while the CPU is executing a program. It should be further understood that this real-time-data-exchange capability requires zero software overhead, meaning it works straight out of the box without any firmware in program memory.

FIG. 47C is a block diagram of a conventional industry standard IEEE 1149.1 JTAG state machine Test Access Port (TAP) **7600**. The JTAG state machine TAP includes TAP state machine **7630**, 8-bit JTAG instruction shift register **7610**, and instruction register **7620**. Outputs UDRCAP, UDRSH, and UDRUPD from state-machine **7630**, as well as 8-bit output UIREG[7:0] from instruction register **7620**, are used by the update/capture register block **7500** to load its registers with data scanned into and out of 64-bit data shift register **7130**. There is nothing novel about the JTAG state machine TAP **7600**, as it is well understood in the industry, but is included here for completeness and as an aid in describing the present invention’s JTAG debug module **7000** (FIG. 47A).

FIG. 48A is a simplified schematic and pertinent Verilog RTL source code describing behavior of a child XCU **800** breakpoint module **7700** in an exemplary embodiment of the present disclosure.

FIG. 48B illustrates exemplary snippets of Verilog RTL showing memory mapping and behavioral description of the parent CPU’s XCU hardware breakpoint control and status registers in an exemplary embodiment of the present disclosure.

With reference to FIGS. 48A and 48B, the theory of operation of the child XCU breakpoint module **7700** will be explained. The child XCU breakpoint module **7700** operates similarly to the hardware breakpoint module **7300** (FIG. 47A), except, instead of monitor request, mon_read_addrs, mon_write_addrs, and their respective sizes coming from the JTAG update/capture registers **7500**, they come directly from a memory-mapped XCU control register owned by the parent CPU.

Stated another way, in the present implementation, there is no IEEE 1149.1 accessible debug module **7000** connected directly to any XCUs. Instead, child XCUs each have a relatively simple breakpoint module **7700** attached to them, wherein the breakpoint and single-step control signals come from a 64-bit XCU control register **7742** mapped in the parent CPU’s data space and status inputs from each XCU

are visible to the CPU via a 64-bit XCU status register **7756** input, also mapped in the CPU's data memory space (see **2000** in FIG. 11).

For example, to force a hardware breakpoint on any one or any combination of XCUs, the parent CPU simply sets corresponding force_break bits **7750** in its XCU control register. To single-step an XCU, the CPU simply sets corresponding sstep bits **7752** in its XCU control register. To force a hard reset on any one XCU or any combination of up to 16 XCUs, the CPU simply sets the corresponding force reset bit **7748** in its XCU control register. In like manner, the parent CPU can force a preemption on any one or combination of XCUs attached to it by simply setting the corresponding preemption bits **7754** in its XCU control register **7742**.

The parent CPU has a XCU status register **7756** mapped in its data memory space that provides it with visibility as to the 64-bit status of up to 16 child XCUs attached to it. Thus, referring to XCU_STATUS_REG **7756** in FIG. 48B, the CPU's XCU status register is made up of 16 XCU done bits **7758**, 16 XCU software break detected bits **7760**, 16 XCU broke bits **7762**, and 16 XCU skip complete bits **7764**.

The XCU done bits **7758** (bit positions **63** to **48**) of XCU_STATUS_REG **7756** come from each of the XCU's DONE flag in their respective STATUS REGISTERS. On power-up or upon being reset, their DONE flag is set to "1". If there is no XCU attached for that bit position, then that bit will be read as "0", meaning that XCU slot is empty. Thus, on power-up initialization, the CPU can simply test that bit position to see if a XCU corresponding to that bit position is present. This ISA has adopted the policy that when the CPU has spawned a given XCU or group of XCUs to perform a task, the XCU routine will clear its DONE flag to signal that it is now busy performing the task and when complete, bring its DONE flag back to "1" to signal it has completed the task and that results are available.

XCU software break detected bits **7760** provide a means for the parent CPU to determine if one or more of the child XCUs attached to it has encountered a software breakpoint instruction **7708**. If so, the bit(s) in the XCU software break detected bits **7760** are read as "1" for XCUs that have encountered a software breakpoint. These bits also provide a means for the parent CPU to determine what caused the XCU to break in cases where both software breakpoints and hardware breakpoints are employed, in that if a hardware breakpoint is encountered by a XCU, the corresponding XCU_BROKE **7762** will be set. It should be understood that the corresponding bit in XCU_BROKE **7762** is set to "1" if either a software breakpoint or hardware breakpoint is encountered. Therefore, XCU_SWBRKDET **7760** provides a means for determining what caused the break to occur.

Bit group XCU_SKIPCMPLT **7764** provides a means for the parent CPU to determine if a XCU or group of XCUs have performed a single-step command issued by the parent CPU via its XCU_CNTRL_REG **7742** XCU single-step control group **7752**. When a given XCU is in break mode as indicated by its corresponding XCU_BROKE bit in **7762**, the parent CPU can command it to perform a single-step (i.e., execute only the next instruction) by setting that XCU's single-step command bit in **7752** to "1". The parent CPU then tests that XCU's skip-complete bit in **7764** to see if that XCU has completed the single-step operation. If so, the CPU can then issue another single-step command after it brings the previous single-step bit back low again, and so forth. Referring to XCU breakpoint state-machine **7720** of FIG. 48A, it can be seen that this state-machine is very similar to that of the CPU's breakpoint state-machine **7380** (FIG.

47B). As such, it can be seen from either, that once the sstep bit is set to "1" and the target processor has completed the single-step, the sstep bit must be brought back to "0", otherwise the state-machine will not advance to the next state.

As can be seen from the XCU control register **7742** and XCU direct address map **7740**, XCU_CNTRL_REG **7744** is made up of qty (4) 16-bit groups: FORCE_RESET **7748**, FORCE_BREAK **7750**, SSTEP **7752**, and PREEMPT **7754**. Bits **7746** show that all bits of all these groups can be written to simultaneously with a single direct write to XCU_STATUS_REG_ADDRS 15'h7FDE.

XCU_CNTRL_REG **7744** shows that on reset, XCU_PREEMPT[15:0], XCU_SSTEP[15:0], and XCU_FORCE_RESET[15:0] are cleared to 0x0000, while XCU_FORCE_BREAK[15:0] is set to 0xFFFF. This means that once the external reset is released, all XCUs will immediately be forced into a hardware breakpoint. The main reason for this is because it is anticipated that most XCU implementations will employ volatile SRAM for program memory and thus will have nothing to execute on power-up. Forcing a breakpoint immediately will cause each XCU to freeze, allowing an opportunity for the parent CPU to push a program or thread into the XCUs' program memory space using its real-time-data-exchange capability. Once a program is loaded into a given XCU's program memory, the parent CPU then brings the respective FORCE_BREAK bit in **7750** low, followed by bringing the corresponding single-step bit in **7752** to "1", which causes the target XCU to step out of the hardware break state. Implementors must remember to bring the single-step bit back low again, otherwise state-machine **7720** will not advance to subsequent states.

Observe that XCU_CNTRL_REG_ADDRS 15'h7FDE in the XCU direct address map **7740** is a global write address for the parent CPU's XCU control register **7742**, meaning that all 64 bits of the parent CPU's XCU control register **7742** can be written simultaneously by pushing a 64-bit value there. Also notice that each 16-bit group within the parent CPU's XCU control register **7742** can be written exclusively, without disturbing the other groups, in that each 16-bit group has its own address that shadows XCU_CNTRL_REG_ADDRS 15'h7FDE as shown in the XCU direct address map **7740**.

For example, to force a hardware break on XCU2 and XCU3, write a 0x000C to location 15'h7FD2, the XCU_FORCE_BREAK_ADDRS as shown in the XCU direct address map **7740**. Then, to single-step only XCU3, write a 0x0004 to location 15'h7FD3, the XCU_SSTEP_P_ADDRS. To clear the hardware both hardware breakpoints, write a "0" to XCU_FORCE_BREAK_ADDRS then write 0x000C to XCU_SSTEP_ADDRS to cause both XCUs to step out of their respective breakpoints. Finally write a "0" to XCU_SSTEP_ADDRS to bring the respective single-step bits back low again, otherwise the respective XCU state-machine **7720** will not advance to subsequent states.

In FIG. 48A, it can be seen that the disclosed exemplary child XCU implementation also supports software breakpoints. When XCU_instruction_q0[63:0] **100** compares equal to 64'h127FF30000000000 (the software breakpoint instruction **7708**), swbreakDetect **7710** is driven high and enters XCU breakpoint logic block **7726** eventually causing break_q0 **7724** to go high. As similarly explained for the parent CPU, break_q0 **7724** for the XCU is used by its pre-PC, PC, and REPEAT counter, causing them to freeze while **7724** is active "1". "Skip" and "broke" inputs to **7726** come from state-machine **7720**, such that when frc_brk or swbreakDetect are active, broke gets set to "1" and stays that

way until the corresponding frc_brk bit from the parent CPU XCU_CTRL_REG **7742** is brought back low, allowing the target XCU state-machine **7720** to be single-stepped out of break mode as explained above.

Note that XCU (and parent CPU) software break instruction $64'h127FF300000000000$ is actually a conditional load PCC instruction that tests bit **0** of byte location **0x0000** in the processor's data RAM to see if it is "0". Since a read of location "0" always returns a "0" regardless of what value was previously written to it, the test always evaluates as true and a relative branch is always taken using the relative offset specified in the instruction. Since the offset specified is zero, the PC just continuously branches to itself—forever, unless or until a enabled interrupt, non-maskable interrupt occurs. But since interrupt service routines, when complete, ordinarily will return to where it left off, in this instance back to the software break, it will continue to loop here upon return.

When actually employing a software break instruction, a common practice is to save a copy of the original instruction in program memory at the location where a software breakpoint is desired, using a monitor read operation from program memory. Once a copy of the original instruction is saved, the software breakpoint instruction is written into that location with a monitor write operation initiated by the parent CPU. When a software breakpoint is encountered by the target processor, program execution comes to a halt and there it will remain until the software break instruction is replaced with another instruction that does not loop to itself, usually the copy of the original that was previously saved. Note that if the software breakpoint instruction has not been replaced, any attempt to single-step thru it will be futile because it will just continue to loop each step.

Thus, the common practice is, when a software breakpoint is encountered, replace the software breakpoint instruction with the copy of the original at that location, then perform the single-step to step out of the breakpoint and continue running. To remain in single-step mode after encountering a software breakpoint, assert the force hardware break bit in breakpoint control register, then replace the software breakpoint instruction.

CPU-XCU Fused Instruction Register

As mentioned previously and as can be seen from FIG. **48A**, the theory of operation of the XCU breakpoint module **7700** is very similar to that of the parent CPU breakpoint module **7300**. The main difference is that none of the child XCUs in the present implementation have a JTAG port. Instead, the parent CPU employs a body-snatching technique that involves hot-spots at specific locations in its direct data space and logic that substitutes a monitor instruction instantly assembled by such logic when such hot-spots are touched by parent CPU.

As can be seen in the XCU breakpoint module **7700**, there are thirty-three hot-spots in the parent CPU's data memory: SrcA_addrs **0x00007FB0** thru **0x00007FBF**, Dest_addrs **0x00007FB0** thru **0x00007FBF**, and direct Destination address **0x00007FC0**. This is due to the fact that only the upper 28 bits of SrcA_addrs_q0, SrcB_addrs_q0, and Dest_addrs_q0 of address comparators **7702**, **7704**, and **7706** (respectively) are being compared with their corresponding hot-spot (or trigger) address, thereby creating a 16-location window for each. This is explained in greater detail immediately below.

Referring to address comparator **7702**, when the most significant 28 CPU SrcA address bits equal $28'h00007FB$, comparator **7702** output is driven to "1", thereby asserting a monitor read request ("monRDreq") when the parent CPU reads from any SrcA location in the address range

0x00007FB0 thru **0x00007FBF**, which is one location per XCU. For instance, a CPU read from location **0x00007FB0** initiates a monitor read operation from XCU0, a CPU read from location **0x00007FB1** initiates a monitor read operation from XCU1, and so forth.

Referring to address comparator **7704**, when the most significant 28 Destination (write) address bits equal to $28'h00007FB$, comparator **7704** output is driven to "1". When direct (write) address OPdest_q0 is equal to **15'h7FC0**, comparator **7706** is driven to "1". Thus, when either address comparator **7704** is a "1" or address comparator **7706** is a "1", OR gate **7722** is driven to "1", thereby asserting a monitor write request ("monWRreq") when the parent CPU writes to either hot-spot. Note that, like address comparator **7702**, address comparator **7704** will go active "1" anytime the parent CPU writes to data memory in the range **0x00007FB0** thru **0x00007FBF**, which is one location per XCU. For instance, a CPU write to location **0x00007FB0** initiates a monitor write operation to XCU0, a CPU write to location **0x00007FB1** initiates a monitor write operation to XCU1, and so forth.

Address comparator **7706** is used to create a hot-spot in the parent CPU's data memory space that will allow it to initiate monitor write operations to all attached child XCUs—simultaneously. This capability is useful for copying threads, programs, data, parameters, etc. common to all XCUs so that the parent CPU does not have to waste time pushing the same data into the XCUs separately, since some threads and data blocks can be quite large.

Verilog RTL source code block **7718** shows the logic that assembles a monitor instruction **7714** that is substituted for the instruction **100** fetched by the child XCU from its own program memory. When the parent CPU pushes or pulls from the any hot-spot thus defined, a monitor request is asserted on breakpoint logic block **7726**, causing break_q0₇₇₂₄ to go active "1" for exactly one clock. During this time, data selector **7712** drives assembled monitor instruction **7714** onto the target child XCU(s)' instruction bus **7728** during its instruction fetch cycle, which gets decoded and executed by the target XCU(s).

FIG. **48C** is a diagram illustrating, in an exemplary embodiment, fusing of the parent CPU monitor read instruction to the child XCU monitor read instruction assembled by the parent CPU. In other words, logic in the parent CPU instantaneously assembles a monitor read instruction using information from within its own instruction when a parent CPU instruction references the respective hot-spot within the parent CPU's data space.

A monitor read instruction is defined as one in which the instruction **100** that the parent CPU has just fetched has a monitor-read hot-spot address (direct or indirect) in its srcA **130** field. As explained previously, any read from locations **0x00007FB0** thru **0x00007FBF** will cause an XCU monitor read instruction to be assembled and instantaneously substituted for the instruction being simultaneously fetched by the target child XCU from its program memory for exactly one clock cycle. The target child XCU is specified by the last four bits of SrcA in CPU Monitor Read from XCU Instruction **7770**.

The resulting address fields that make up the assembled XCU monitor read instruction fetched and executed by the target XCU are shown in Assembled XCU Monitor Read Instruction **7772**. Notice that the destination address of the Assembled XCU Monitor Read Instruction **7772** is always **0x0000** in the target XCU's data space. SrcA address of the Assembled XCU Monitor Read Instruction **7772** is the address specified by the SrcB address field **140** (FIG. 1) of

the parent CPU Monitor Read from XCU Instruction 7770, such that when the target XCU executes the Assembled XCU Monitor Read Instruction 7772, it will perform a data read from its data memory space at the address specified by the srcA address field in the Assembled XCU Monitor Read Instruction 7772, and then write the results to location 0x0000 in its own direct memory space. During Stage q2 of its execution pipeline, the parent CPU intercepts the data being written by the target XCU, “as if” the CPU had performed the read operation itself, and then the parent CPU writes the intercepted data to the CPU destination data address specified in the DEST 120 field of CPU Monitor Read from XCU Instruction 7770.

FIG. 48D is a diagram illustrating, in an exemplary embodiment, fusing of the parent CPU monitor write instruction to the child XCU monitor write instruction assembled by the parent CPU. In other words, logic in the parent CPU instantaneously assembles a monitor write instruction using information from within its own instruction when a parent CPU instruction references the respective hot-spot within the parent CPU’s data space.

A monitor write instruction is defined as one in which the instruction 100 that the parent CPU has just fetched has a monitor-write hot-spot address (direct or indirect) in its DEST 120 field. As explained previously, any write to locations 0x00007FB0 thru 0x00007FBF will cause a XCU monitor write instruction to be assembled and instantaneously substituted for the instruction being simultaneously fetched by the target child XCU from its program memory for exactly one clock cycle. The target child XCU is specified by the last four bits of the DEST Monitor Write Hot-Spot specified in the CPU Monitor Write to XCU Instruction 7774. In cases where the hot-spot is the “poke-all” address 0x00007FC0 within the parent CPU’s data space, all child XCUs attached to the CPU will be issued the monitor write instruction, rather than just one, resulting in a write into each XCU memory space of the same data provided by the parent CPU, hence the term “poke-all” or “push-all”.

The resulting address fields that make up the assembled XCU monitor write instruction fetched and executed by the target XCU are shown in Assembled XCU Monitor Write Instruction 7776. Notice that the destination address DEST of the Assembled XCU Monitor Write Instruction 7776 is the address specified by the srcB 140 field of the parent CPU Monitor Write to XCU Instruction 7774. Also observe that the SrcA 130 address field and SrcB 140 address field of the Assembled XCU Monitor Write Instruction 7776 are both always 0x0000. Thus, when the parent CPU executes the CPU Monitor Write to XCU Instruction 7774, the parent CPU reads from its own data memory space, the contents of the address specified by the SrcA 130 address of the parent CPU Monitor Write to XCU Instruction 7774 and writes it to the monitor write hot-spot specified by DEST 120 address. During the child XCU Stage q2, the data being written by the parent CPU is intercepted by target child XCU internal logic and is written to the target XCU data memory address specified by the Assembled XCU Monitor Write Instruction 7776 XCU destination address DEST 120.

It should be understood that, fundamentally, the instruction registers of the parent CPU and child XCUs are fused. The parent CPU can reach into any child XCU anytime and grab what ever is needed and can write whatever is needed to any resource mapped into any child XCU’s data space without the use of interrupts, opcodes, or DMA, because the parent-child relationship is one in which they share the same instruction register, albeit fused, anytime a monitor read or

monitor write is fetched and executed by the parent CPU. Thus, in one instance, they are executing their own programs and then in another instance, they together and simultaneously execute an instruction fetched by the parent CPU.

Observe further that either direct or indirect addressing modes may be used, thereby giving the parent CPU full reach into any and child XCUs attached to it. Moreover, the CPU REPEAT instruction, along with indirect addressing (auto-post-modify) mode, can be used to easily and efficiently transfer entire blocks of memory into and out of any child XCU

FIG. 49 is a block diagram illustrating an exemplary embodiment of a double-quad, single-precision (H=12) Universal FMA operator 9900 that can accept in a single push, quantity (8) 16-character decimal character sequence or binary32 format numbers as operandA, and quantity (8) 16-character decimal character sequence or binary32 format numbers as operandB, outputting as quantity (8) correctly rounded binary32 format numbers (including corresponding exceptions) for a total of 512 bits, or quantity (8) binary32 format numbers only, for a total of 256 bits, for each pull. Universal FMA blocks 9800a, 9800b, 9800c, 9800d, 9800e, 9800f, and 9800g are virtually identical to Universal FMA 9800 in FIG. 46A, except the decimal character to binary conversion circuit is H=12 compliant, meaning it can convert decimal character sequences up to twelve decimal digits in length. The Universal FMA operator 9900 is universal in the same way Universal FMA 9800 is, in that it can accept scientific notation, integers, characters sequences with token exponents, or binary16, binary32, or binary64 format numbers as operands.

Since all the FMAs share the same size and signal inputs, the numbers making up 1024-bit operandA gob must be the same format and numbers making up 1024-bit operandB gob must be the same format, but gob A can be a different format than gob B.

Like the Universal FMA in 9800, each Universal FMA in 9900 has a split 32-entry result buffer that is 37 bits wide (instead of 69 bits wide as shown in 9816), which is wide enough to accommodate a binary32 format result plus five exception bits. Results are read out as one 512-bit gob or one 256-bit gob for each result buffer location. The wider gob is for situations where the exception bits for each FMA are read out as part of that FMA’s results in their respective position in the gob, in which case 64-bit words are pulled out in one 512-bit gob. Many applications can do without the exception bits, in such cases results can be read out as quantity (8) binary32 format numbers in one 256-bit gob.

Once pushed into fat memory, results can be pulled out in properly aligned 1, 2, 4, 8, 16, 32, 64, or 128-byte accesses. For example, the results of a single 128-byte push into fat memory can be read out as a sequence of quantity (128) 1-byte pulls, (64) 2-byte pulls, (32) 4-byte pulls, and so on. However, care must be exercised to ensure that the original 128-byte push is on a even 128-byte address boundary. The results of a single 64-byte push into fat memory can be read out as a sequence of quantity (64) 1-byte pulls, (32) 2-byte pulls, (16) 4-byte pulls, and so on. Like the previous example, care must be exercised to ensure that the original 64-byte push is on an even 64-byte address boundary. Thus the same can be done for pushes of 32, 16, 8, etc. bytes in length.

Each of the eight FMAs in the Universal FMA operator 9900 can be used as sum-of-products operators, in that each has quantity 32 C-register/accumulators and operate the same was as those in the Universal FMA 9800.

Finally, for intensive pure convertFromDecimalCharacter to binary conversion operations, the Universal FMA operator **9900** can be used to convert quantity (16) H=12 decimal character sequences every clock cycle.

FIG. 50A thru FIG. 50L comprise a example assembled source listing of a actual working 3D transform program **9950** written in the present disclosure's ISA assembly language that employs up to qty. (16) child XCUs to perform a 3D transform (scale, rotate, translate) on all three axes of a 3D object in .STL file format and write the resulting transformed 3D object back out to external memory when the transformation is complete. This program was simulated using the XILINX® Vivado® development environment (Vivado HLx Edition, v2018.2 (64-bit)) targeted to its Kintex® UltraScale and UltraScale+ brand FPGAs.

Referring to 3D transform program **9950**, the parent CPU first determines how many child XCUs are available to perform the transformation by first counting the active DONE, one corresponding to each XCU, and then divides the number of triangles that make up the 3D object as evenly as possible. The parent then performs a “poke-all” monitor instruction to push the thread code simultaneously into all XCUs. It sets a breakpoint simultaneously on all XCUs at the same location in their program memory and then simultaneously asserts then releases a reset so that all XCUs begin to execute at the same location until they all reach their breakpoint.

The parent CPU then pushes the number of triangles each XCU is to transform, the location of where the triangles are located in XCU memory space, and then the triangles themselves. The parent then releases the breakpoint and single-steps all child XCUs out of the break state simultaneously, at which time the XCUs begin processing the triangles. During this time, each XCU brings its DONE bit inactive low, to signal it is busy. At this time the parent is monitoring its XCU status register for the DONE bits to be brought back to active high, indicating XCU processing is complete. When complete, the parent XCU then pulls the transformed triangles from each XCU result memory and pushes them back out to external memory.

If, at the initial stages described above, the parent CPU determines there are no XCUs available to perform the transform, the parent CPU performs the entire 3D transform solo (without the use of any XCU) and pushes the transformed triangles back out to external memory when complete. It should be understood that the parent CPU and the child XCUs (if any) execute the same instruction set. In this implementation, the 3D transform routine physically resides in the parent CPU program memory space. Thus, in this instance, when the parent performs a “push-all” of the required routine and parameters, it pulls the routine code out of its program memory space and pushes it into all XCUs program memories simultaneously. If the required program is not resident in the CPU's program space, it could alternatively pull it in from any child XCU program memory space or external memory space.

FIG. 51 is an actual wire-frame “Before” and “After” rendering **9955** of a simple “olive” 3D model in .STL file format performed by from 1 to 16 child XCUs or solo parent CPU using the scale, rotate, and translate parameters shown for each axis.

In the drawings and specification, there have been disclosed typical preferred embodiments of the disclosure and, although specific terms are employed, they are used in a generic and descriptive sense only and not for purposes of limitation, the scope of the invention being set forth in the following claims.

What is claimed is:

1. A universal Instruction Set Architecture (ISA) computer having an instruction pipeline that does not include opcode fetch, decode, or execute logic, said universal ISA computer comprising:

a conditionally writable, memory-mapped, self-incrementing Program Counter (PC) residing at a plurality of addresses within a memory map of the ISA and an output of which is used as a read address for fetching instructions from a program memory, wherein memory-mapped conditions for writing to the PC shadow a memory map of the PC, wherein the PC is updated with write data if, and only if, a written-to address corresponds to a memory-mapped condition and the memory-mapped condition is true;

a plurality of memory-mapped Auxiliary Registers configured as automatically post modifiable, indirect address pointers for reading or writing data from or to memory, memory-mapped registers, or memory-mapped computational hardware operators;

an instruction register that does not include an opcode field, said instruction register comprising at least the following fields:

a Rounding Mode field that specifies which rounding mode is to be used when a memory-mapped hardware operation implied by a destination address of an instruction is an IEEE 754-2008 floating-point operation, wherein when the operation is not an IEEE 754 floating-point operation, the Rounding Mode field is available for other purposes;

a Data Access Mode field that specifies whether sourceA and sourceB addressing modes are each either direct or indirect, whether sourceA is direct or indirect and source B is immediate, whether sourceA is table-read and sourceB is direct or indirect, or whether an only source is direct or indirect or whether an only source is immediate; and

destination, sourceA, and sourceB address fields, each address field having its own corresponding general-purpose signal field, size field, indirect (IND) bit, indirect mode (IMOD) bit, and combination direct address and post-modification-amount field, wherein the post-modification-amount field specifies a post-modification amount and which memory-mapped Auxiliary Register is to be used as a pointer register when IND addressing mode is specified; and

at least one memory-mapped computational hardware operator configured to accept one or more operands simultaneously written to the at least one memory-mapped computational hardware operator by the ISA computer and to deliver an available computed result when the ISA computer accesses a memory-mapped, randomly accessible, hardware-operator result buffer at a read address that is the same as an address of the at least one memory-mapped computational hardware operator to which the one or more operands were written.

2. The universal ISA computer as recited in claim 1, further comprising a memory-mapped instruction Repeat Counter for use with instructions specifying that either the destination, sourceA, or sourceB addressing mode are indirect, such that when the Repeat Counter is written with a non-zero value and an instruction specifies that at least one of the destination, sourceA, or sourceB addressing modes are indirect, execution of an immediately following instruc-

tion is repeated a number of times specified in the Repeat Counter during an execution stage within the instruction pipeline.

3. The universal ISA computer as recited in claim 1, further comprising at least one memory-mapped hardware Loop Counter, wherein when the Loop Counter is read in response to a first instruction that has the PC as its destination, the Loop Counter is automatically decremented if not already zero, and if already zero, no decrement takes place and the PC is updated with a value specified by the first instruction.

4. The universal ISA computer as recited in claim 1, wherein fully pipelined hardware is configured in series with an operand input stream, the fully pipelined hardware automatically converting a human-readable decimal character sequence floating-point representation to an IEEE 754 binary floating-point format representation along with any corresponding exceptional signals produced by the fully pipelined hardware during conversion, while in-flight to a memory-mapped IEEE 754 binary floating-point hardware operator input specified by a destination address of an instruction that caused one or more operands to take flight.

5. The universal ISA computer as recited in claim 4, wherein an exponent of the human-readable decimal character sequence floating-point representation includes a token exponent.

6. The universal ISA computer as recited in claim 1, having a memory-mapped, fully pipelined, hardware operator that, with a single instruction, is configured to convert one or more received binary floating-point representations into resulting human readable decimal character sequence floating-point representations and to automatically store the resulting human-readable decimal character sequence floating point representations along with any IEEE 754 exceptional signals produced by the memory-mapped, fully pipelined hardware operator during conversion;

wherein the resulting human-readable decimal character sequence floating point representations and IEEE 754 exceptional signals, if any, are stored in a randomly accessible result buffer dedicated to the memory-mapped, fully pipelined, hardware operator at a direct or indirect destination address specified in the single instruction; and

wherein the memory-mapped, fully pipelined hardware operator is configured to accept new IEEE 754 binary floating-point representations every clock cycle until the dedicated result buffer becomes full.

7. The universal ISA computer as recited in claim 1, having a memory-mapped, fully pipelined, hardware operator that, with a single instruction, is configured to convert one or more received human readable decimal character sequence floating-point representations into resulting IEEE 754 binary floating-point representations and to automatically store the resulting IEEE 754 binary floating-point representations along with any IEEE 754 exceptional signals produced by the memory-mapped, fully pipelined hardware operator during conversion;

wherein the resulting IEEE 754 binary floating-point representations and IEEE 754 exceptional signals, if any, are stored in a randomly accessible result buffer dedicated to the memory-mapped, fully pipelined, hardware operator at a direct or indirect destination address specified in the single instruction; and

wherein the memory-mapped, fully pipelined hardware operator is configured to accept new human-readable

decimal character sequence floating-point representations every clock cycle until the dedicated result buffer becomes full.

8. The universal ISA computer as recited in claim 1, further comprising a dual asymmetric stack.

9. The universal ISA computer as recited in claim 1, wherein at least some of the memory is three-port Random Access Memory (RAM), wherein one port is a write port for writing data, one port is a read port for reading source A data, and one port is a read port for reading source B data.

10. The universal ISA computer as recited in claim 1, wherein the universal ISA computer is configured as a parent opcodeless universal ISA computer connected to at least one child universal ISA computer via a fused instruction register, and is further configured to:

fetch a real-time monitoring and data exchange instruction that does not include an opcode field; and execute the real-time monitoring and data exchange instruction simultaneously with the at least one child universal ISA computer connected to the parent universal ISA computer.

11. The universal ISA computer as recited in claim 10, wherein an IEEE 1149.1-1990 (JTAG) interface from a host debug workstation is connected to the parent universal ISA computer, and the parent universal ISA computer includes a JTAG real-time monitor and debug logic block configured to receive JTAG signals from the host debug workstation, translate the JTAG signals into opcodeless real-time monitoring and data exchange instructions, selectively issue the opcodeless real-time monitoring and data exchange instructions to the parent universal ISA computer or to the at least one child universal ISA computer, and to retrieve results once executed, without the use of hardware breakpoints, interrupts, or Direct Memory Access (DMA) cycles.

12. The universal ISA computer as recited in claim 1, wherein one or more of the at least one memory-mapped computational hardware operators include:

memory-mapped hardware for carrying out an operation of a memory-mapped computational hardware operator specified by a direct or indirect destination address in an instruction; and
wherein the memory-mapped, randomly accessible, hardware-operator result buffer is configured to store results and any exceptional signals or status signals produced by the at least one memory-mapped computational hardware operator, wherein the results and any exceptional signals or status signals automatically spill into the memory-mapped, randomly accessible, hardware-operator result buffer, wherein when the results are subsequently accessed with a read operation, an optional memory-mapped Status Register is automatically updated to reflect the state of any exceptional signals or status signals.

13. The universal ISA computer as recited in claim 1, wherein assembly language thereof does not include mnemonics representing opcodes, and wherein a single non alphanumeric character is used as the first non-space character on an assembly line as a mnemonic to represent a default mode.

14. The universal ISA computer as recited in claim 1, further comprising as one of its memory-mapped hardware operators, a fully pipelined, floating-point Fused-Multiply-Add (FMA) operator configured to accept on its inputs, a new set of human-readable decimal character sequence floating-point representation or IEEE 754 binary floating-point format representation operands every clock cycle until an associated result buffer becomes full.

15. The universal ISA computer as recited in claim 14, wherein the FMA operator further comprises a hardware logic situated in series with the FMA pipeline output and associated result buffer input, wherein the hardware logic is configured to automatically convert the FMA pipeline output to human-readable decimal character sequence floating-point representations before automatic storage into an associated result buffer. 5

16. The universal ISA computer as recited in claim 1, wherein all IEEE 754 floating-point operations that the universal ISA computer carries out in memory-mapped hardware require only one instruction per operation, and such memory-mapped hardware is configured to accept one or more new operands every clock cycle, with results and any corresponding exceptional signals produced by the 15 memory-mapped hardware automatically spilling into dedicated, randomly addressable result buffers.

* * * *